


SQL: The Query Language

Part II

R&G - Chapter 5



Announcement

- **Midterm exam (Oct 26 & 28 tentative)**
- **Assignment 2 is available in Moodle**
 - Relational Algebra and SQL
 - Due time: Monday Oct 31 (11:59pm this time!)

SQL

- The form:

**SELECT A_1, A_2, \dots, A_n
FROM r_1, r_2, \dots, r_m
WHERE P**

- A_i represents an attribute
 - r_i represents a relation
 - P is a predicate
- This query is equivalent to the relational algebra expression:

$$\Pi_{A_1, A_2, \dots, A_n} (\sigma_P (r_1 \times r_2 \times \dots \times r_m))$$

Example Schemas

Sailors(sid, sname, rating, age)

Boats(bid, bname, color)



Reserves(sid, bid, day)

FOREIGN KEY (bid) REFERENCES Boats,


FOREIGN KEY (sid) REFERENCES sailors

)


Set Operators



UNION/UNION ALL



INTERSECT



MINUS/EXCEPT

Set Operations

Union: $R \cup S$

In SQL:

SELECT * FROM R

UNION

SELECT * FROM S;

Intersection: $R \cap S$

In SQL:

SELECT * FROM R

INTERSECT

SELECT * FROM S;

Set difference: $R - S$

In SQL:

SELECT * FROM R

EXCEPT

SELECT * FROM S;

Example of Union

Find sids of sailors who've reserved a red or a green boat

Solution 1 (Without using Set Operations)

$\pi_{sid}(\sigma_{color='red' \vee color='green'} Boats \bowtie Reserves)$

```
SELECT R.sid
FROM Boats B,Reserves R
WHERE R.bid=B.bid AND
(B.color='red' OR B.color='green')
```

Example of Union

Find sids of sailors who've reserved a red or a green boat


Solution 2 (with Set operations)

- Find sailors who've reserved red boats *Tempred*;
- Find sailors who've reserved green boats *Tempgreen*;
- find the Union of *Tempred* and *Tempgreen*

```
SELECT  R.sid
FROM    Boats B, Reserves R
WHERE   R.bid=B.bid AND B.color='red'
UNION
SELECT  R.sid
FROM    Boats B, Reserves R
WHERE   R.bid=B.bid AND B.color='green'
```


Question: how to revise the SQL statement if the query is changed to be “find names of sailors who've reserved a red or a green boat

UNION Operator



The UNION operator returns results from both queries after eliminating duplications.

UNION ALL Operator



The UNION ALL operator returns results from both queries, including all duplications.

Examples of UNION and UNION ALL

Sid	Sname	Rating	Age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.5

S1

Sid	Sname	Rating	Age
28	Yuppy	9	35.0
31	Lubber	8	55.5
44	Guppy	5	35.0
58	Rusty	10	35.5

S2

SELECT sname
FROM S1
WHERE Age >40

UNION

SELECT sname
FROM S2
WHERE Age >40

Sname
Dustin
Lubber

Result

Examples of UNION and UNION ALL

Sid	Sname	Rating	Age
22	Dustin	7	45.0
31	Lubber	8	55.5
58	Rusty	10	35.5

S1

Sid	Sname	Rating	Age
28	Yuppy	9	35.0
31	Lubber	8	55.5
44	Guppy	5	35.0
58	Rusty	10	35.5

S2

```
SELECT sname
FROM S1
WHERE Age >40
UNION ALL
SELECT sname
FROM S2
WHERE Age >40
```

Sname
Dustin
Lubber
Lubber

Result

Intersection

- **Find sids of sailors who've reserved a red and a green boat**

- Identify sailors who've reserved red boats (in *Tempred*),

- Identify sailors who've reserved green boats (in *Tempgreen*),

- find the intersection of *Tempred* and *Tempgreen*

$\rho(\text{Tempred}, \pi_{sid}((\sigma_{color='red'} \text{Boats}) \bowtie \text{Reserves}))$

$\rho(\text{Tempgreen}, \pi_{sid}((\sigma_{color='green'} \text{Boats}) \bowtie \text{Reserves}))$

$\pi_{sname}((\text{Tempred} \cap \text{Tempgreen}) \bowtie \text{Sailors})$

Intersection (Cont.)

- **INTERSECT**: Can be used to compute the intersection of any two *union-compatible* sets of tuples.

```
SELECT sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
INTERSECT
SELECT sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='green';
```

Find sid's of sailors who've reserved a red and a green boat

- Can we write equivalent SQL statement without using set operations?

```
SELECT R.sid  
FROM Boats B,Reserves R  
WHERE R.bid=B.bid AND  
(B.color='red' AND B.color='green')
```

Find sid's of sailors who've reserved a red
and a green boat



- Can we write equivalent SQL statement without using set operations?
 - Hint: join with Boats table TWICE (one for red boats, one for green boats)
 - Solution:

```
SELECT R1.sid
FROM Boats B1, Reserves R1,
      Boats B2, Reserves R2
WHERE R1.sid=R2.sid
      AND R1.bid=B1.bid
      AND R2.bid=B2.bid
      AND B1.color='red'
      AND B2.color='green';
```


Find sid's of sailors who've reserved a red and a green boat



- Can we write equivalent SQL statement without using set operations?
 - Hint: join with Boats table TWICE (one for red boat, one for green boat)
 - **Question:** can we use one single Reserves table in the two joins (as shown below)?

```
SELECT R1.sid
FROM Boats B1, Boats B2, Reserves R
WHERE R.bid=B1.bid
      AND R.bid=B2.bid
      AND B1.color='red'
      AND B2.color='green';
```

Find sid's of sailors who've reserved a red and
a green boat



- Can we write equivalent SQL statement without using set operations?
 - Hint: join with Boats table TWICE (one for red boat, one for green boat)
 - **Question:** can we use one single Reserves table in the two joins (as shown below)?

```
SELECT R1.sid
FROM Boats B1, Boats B2, Reserves R
WHERE R.bid=B1.bid
      AND R.bid=B2.bid
      AND B1.color='red'
      AND B2.color='green';
```

Find sid's of sailors who've reserved a red but did not reserve a green boat

EXCEPT (sometimes called MINUS)

Included in the SQL/92 standard, but many systems don't support them.

```
SELECT sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
EXCEPT
SELECT sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='green'
```