

SQL: The Query Language

Part IV

Aggregate Queries

R&G - Chapter 5

Simple and Nested SQL Queries

- **Simple SQL Queries**

```
SELECT  $A_1, A_2, \dots, A_n$   
FROM  $r_1, r_2, \dots, r_m$   
WHERE  $P$ 
```

- **Nested SQL queries**

```
SELECT  $A_1, A_2, \dots, A_n$   
FROM  $r_1, r_2, \dots, r_m$   
WHERE (NOT)IN|(NOT) EXISTS|Op ANY/ALL  
      ( $\text{SELECT } A_1, A_2, \dots, A_n$   
        FROM  $r_1, r_2, \dots, r_m$   
        WHERE  $P$ )
```

Today's plan

- **Aggregate queries**

Aggregate Function

- **Aggregate functions: significant extension of relational algebra.**
- **Behavior of Aggregate Functions:**
 - Operates on a single column
 - Input: a collection of values.
 - Output: a single value.
 - Used ONLY in the SELECT list and in the HAVING clause.

Types of SQL Aggregate Functions

- **SUM**
- **AVG**
- **MIN**
- **MAX**
- **COUNT**

COUNT (*)
COUNT ([DISTINCT] A)
SUM ([DISTINCT] A)
AVG ([DISTINCT] A)
MAX (A)
MIN (A)

single column

```
SELECT COUNT (*)  
FROM Sailors S
```

```
SELECT AVG (S.age)  
FROM Sailors S  
WHERE S.rating=10
```

```
SELECT COUNT (DISTINCT S.rating)  
FROM Sailors S  
WHERE S.sname='Bob'
```

Input to Aggregate Function

- **SUM and AVG :**
 - ➔ Only accept numerical values as input.
 - **MIN , MAX and COUNT**
 - ➔ Accept both numerical and categorical values at input.
- ❑ Each function eliminates NULL values and operates on non- null values.

Aggregate Function Continued...

- Accepts:
 - ✓ DISTINCT: consider only distinct values of the argument expression.
 - ✓ ALL : consider all values including all duplicates.

Example: `SELECT COUNT(DISTINCT column_name)`

Example of DISTINCT and ALL

<u>sno</u>	salary
SL100	30000
SL101	10000
SL102	10000
SL103	20000

Staff

Query 1:

SELECT AVG(DISTINCT salary) AS avg_sal
FROM Staff;

Result:

avg_sal
20000

= (30K+10K+20K)/3

Query 2:

SELECT AVG(ALL salary) AS avg_sal
FROM Staff;

Result:

avg_sal
17500

=

(30K+10K+10K+20K)/4

Example of Aggregate Queries

Query: Find name and age of the oldest sailor(s)

```
SELECT S.sname, S.age
FROM Sailors S
WHERE S.age = (SELECT MAX (S2.age)
               FROM Sailors S2)
```

Solution 1

Example of Aggregate Queries

Query: Find name and age of the oldest sailor(s)

```
SELECT S.sname, MAX (S.age)
FROM Sailors S;
```

Solution 2

Question: is Solution 2 correct?

Answer

Query: Find name and age of the oldest sailor(s)

```
SELECT S.sname, MAX (S.age)
FROM Sailors S
```

- MySQL: return the name of the 1st tuple, and the maximum age
- SQL server and Oracle: return error message.
- Rule: For aggregate queries without GROUP BY clause, **DO NOT** put non-aggregate attributes and aggregate functions together in SELECT clause.

More examples

Find the names of sailors who are older than the oldest sailor of rating 10

```
SELECT S.sname
FROM Sailors S
WHERE S.age >
      (SELECT MAX (S2.age)
       FROM Sailors S2
       WHERE S2.rating =10)
```

GROUP BY and HAVING

- Sometimes, we want to ask for groups of tuples.
- The query description normally starts with “for each...”.
- Consider: *for each rating level, find the age of the youngest sailor.*
 - IF we know that rating values go from 1 to 10; we can write 10 queries that look like this (!):

For $i = 1, 2, \dots, 10$:

```
SELECT MIN (S.age)
FROM Sailors S
WHERE S.rating = i
```
 - In general, we don't know how many rating levels exist, and what the rating values for these levels are!

GROUP BY Clause

- To generate values for a column based on groups of rows, use aggregate functions in SELECT statements with the GROUP BY clause

```
SELECT  attr1, attr2, ... attrk, aggregate_function
FROM    tables
[WHERE  qualification]
GROUP BY attr1, attr2, ... attrn;
```

*: the attributes attr₁, attr₂, ... attr_k in SELECT clause that do not have any aggregation function must appear in GROUP BY clause ($k \leq n$)

Tips of GROUP BY Queries

- If the query description starts with “for each **X**”.
- Write the GROUP BY queries

```
SELECT  X, ...  
FROM    ...  
[WHERE  ...]  
GROUP BY X;
```

Group By Examples

For each rating, find the average age of the sailors

```
SELECT S.rating, AVG (S.age)
FROM Sailors S
GROUP BY S.rating
```

For each rating, find the age of the youngest sailor with age ≥ 18

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
```


Grouping over a join of two relations.

Find the number of reservations for each **red** boat.

```
SELECT B.bid, COUNT(*) AS numres
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
GROUP BY B.bid
```

COUNT(*) AS numres: name the output count values under a column named “numres”



GROUP BY Examples

For each rating, find the name and age of the oldest sailor

```
SELECT S.name, MAX (S.age)
FROM Sailors S
GROUP BY S.rating
```

Is this query correct?

This query is not correct, as the non-aggregate attribute name does not appear in GROUP BY Class

How to fix?

Queries With GROUP BY and HAVING

```
SELECT    attr1, attr2, ... attrk, aggregate_function
FROM      tables
[WHERE    qualification]
GROUP BY  attr1, attr2, ... attrn
HAVING    group-qualification;
```

- HAVING clause restricts which group-rows are returned in the result set
- NO GROUP BY, NO HAVING;
- Attributes in *group-qualification* of HAVING clause must be either an aggregate op or appear in the *grouping-list* of GROUP BY clause

Query: For each age group that has more than 1 sailor, find the lowest rating of sailors

```
SELECT S.age, MIN (S.rating)
FROM Sailors S
GROUP BY S.age
HAVING COUNT (*) > 1
```

sid	sname	rating	age
22	Dustin	7	45
31	lubber	8	45
71	zorba	10	18
64	horatio	7	35
29	brutus	1	33
58	rusty	10	35

age	rating
18	10
33	1
35	7
35	10
45	7
45	8

age	min-rating	count
18	10	1
33	1	1
35	7	2
45	7	2

age	Min-rating
35	7
45	7

Answer relation

1: projection

2: grouping

Query:

Find the name of sailors who've reserved all boats.

- **Recall our solution of using nested queries**

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS (SELECT B.bid
                  FROM Boats B
                  WHERE NOT EXISTS (SELECT R.bid
                                    FROM Reserves R
                                    WHERE R.bid=B.bid
                                           AND R.sid=S.sid))
```

Query: Find the name of sailors who've reserved all boats.

- **Can you do this using GROUP BY and HAVING?**
 - Hint: for each sailor, check whether the number of distinct boats he/she has reserved equals to the total number of (distinct) boats.

```
SELECT S.name
FROM Sailors S, reserves R
WHERE S.sid = R.sid
GROUP BY S.name, S.sid
HAVING COUNT(DISTINCT R.bid) =
        ( SELECT COUNT (*) FROM Boats)
```

Note: It must have both sid and name in the GROUP BY clause. Why? Can we only use sid?
Can we only use sname?