



**STEVENS**  
INSTITUTE of TECHNOLOGY  
THE INNOVATION UNIVERSITY®

# CS 492: Operating Systems

## *Threads*

*Instructor: Iraklis Tsekourakis*


Email: [itsekour@stevens.edu](mailto:itsekour@stevens.edu)



# Exercise

- A computer has 4 GB of RAM of which the operating system occupies 512 MB. The processes are all 256 MB (for simplicity) and have the same characteristics. If the goal is 99% CPU utilization, what is the maximum I/O wait that can be tolerated?

# Threads



# Goals for Today

- Threads
  - Concept
  - Multiple-threading
  - Implementation of threads in kernel


# Processes Continues...

- A process is created to run a program to perform a duty.
- What if we need to perform two or more similar duties?
- One approach: create multiple processes, each handling one of the duties.

# Why processes are not always ideal...

- Processes are not very efficient
  - Each process has its own PCB and OS resources
  - Typically high overhead for each process: e.g., 1.7 KB per `task_struct` on Linux!
- Processes don't (directly) share memory
  - Each process has its own address space
  - Parallel and concurrent programs often want to directly manipulate the same memory
- Can we do better?

# Threads




- Thread = an independent sequential execution stream within process
- Simple programs use one or multiple threads per process

# Advantages of Threads

- Performance
  - Thread creation is 10-100 times faster than processes.
- Efficiency
  - Allows one process to use multiple CPUs or cores
  - Allows program to overlap I/O and computation

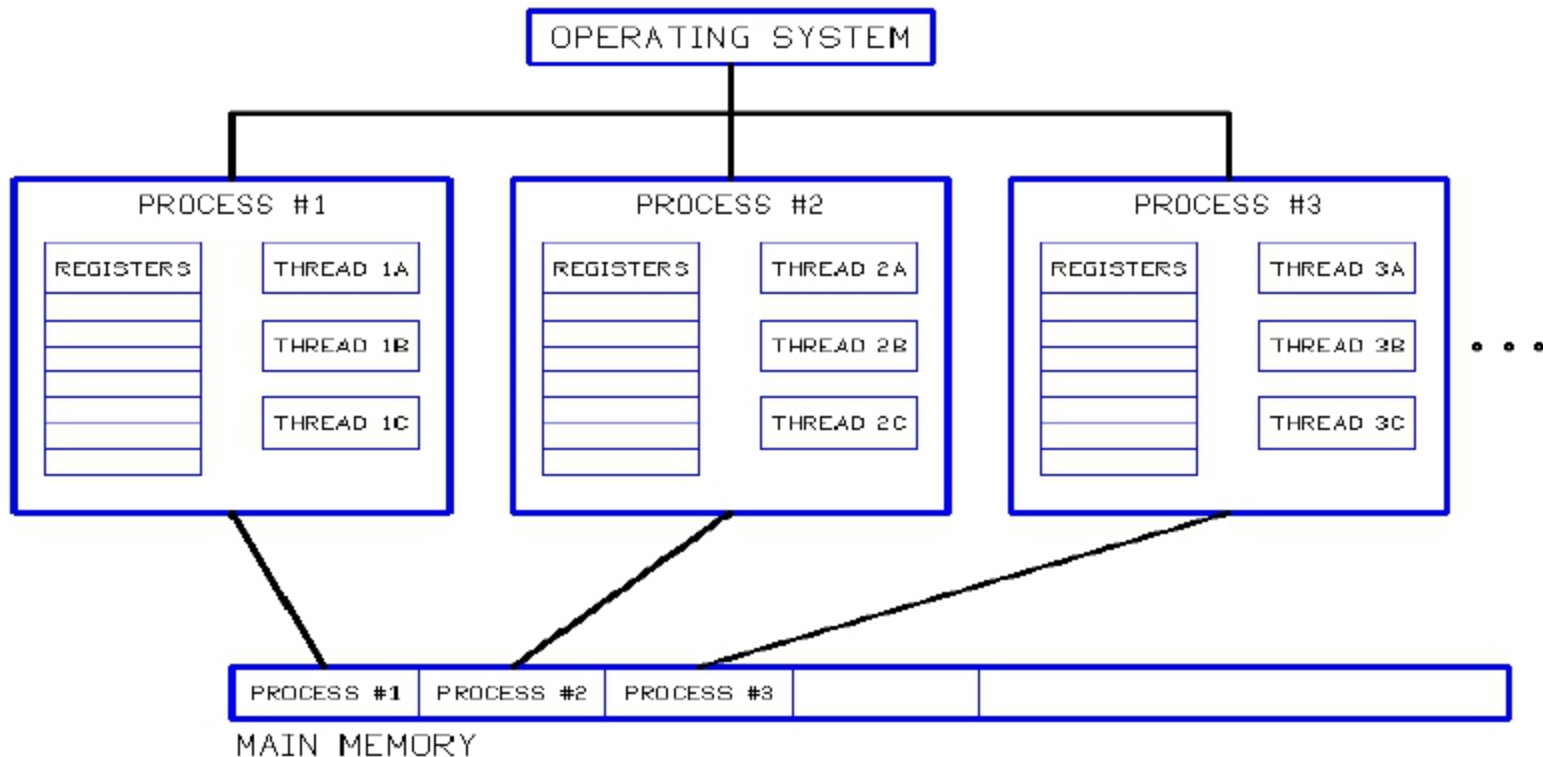


# Thread Usage



- Example: word processor process
  - One thread to read from keyboard
  - One thread to format document
  - One thread to write document to disk
- Would multiple processes work here?

# Process vs Thread





- A *thread* cannot exist without a *process*, thus a process is a “container” for threads
- A process may contain *multiple threads*

# Goals for Today

- Threads
  - Concept
  - Multiple-threading
  - Implementation of threads in kernel

# Multiple Threads in the Same Process

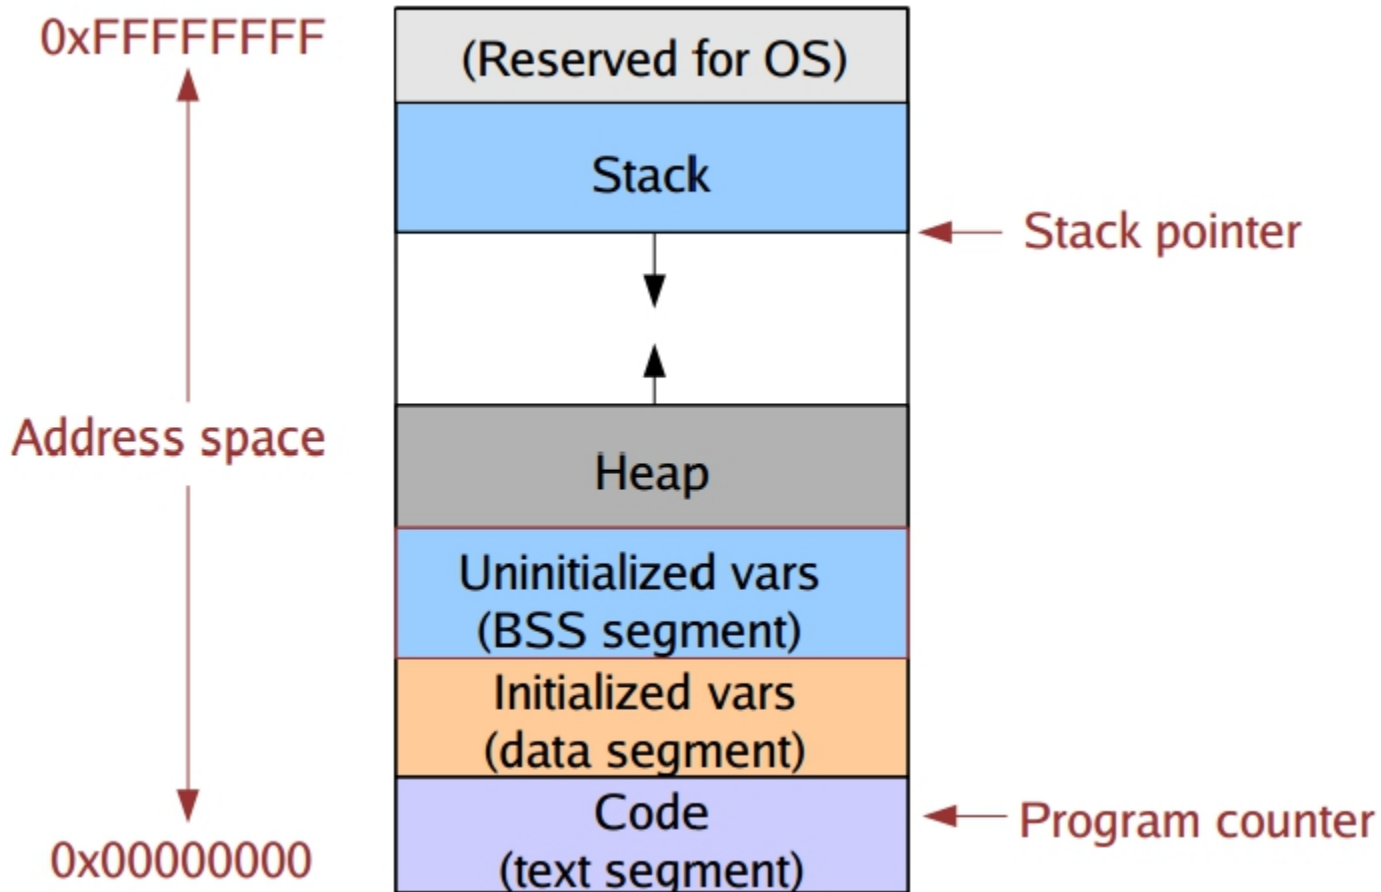


- Share same address space
- Share global variables
- Owned by a single user
- No protection between threads

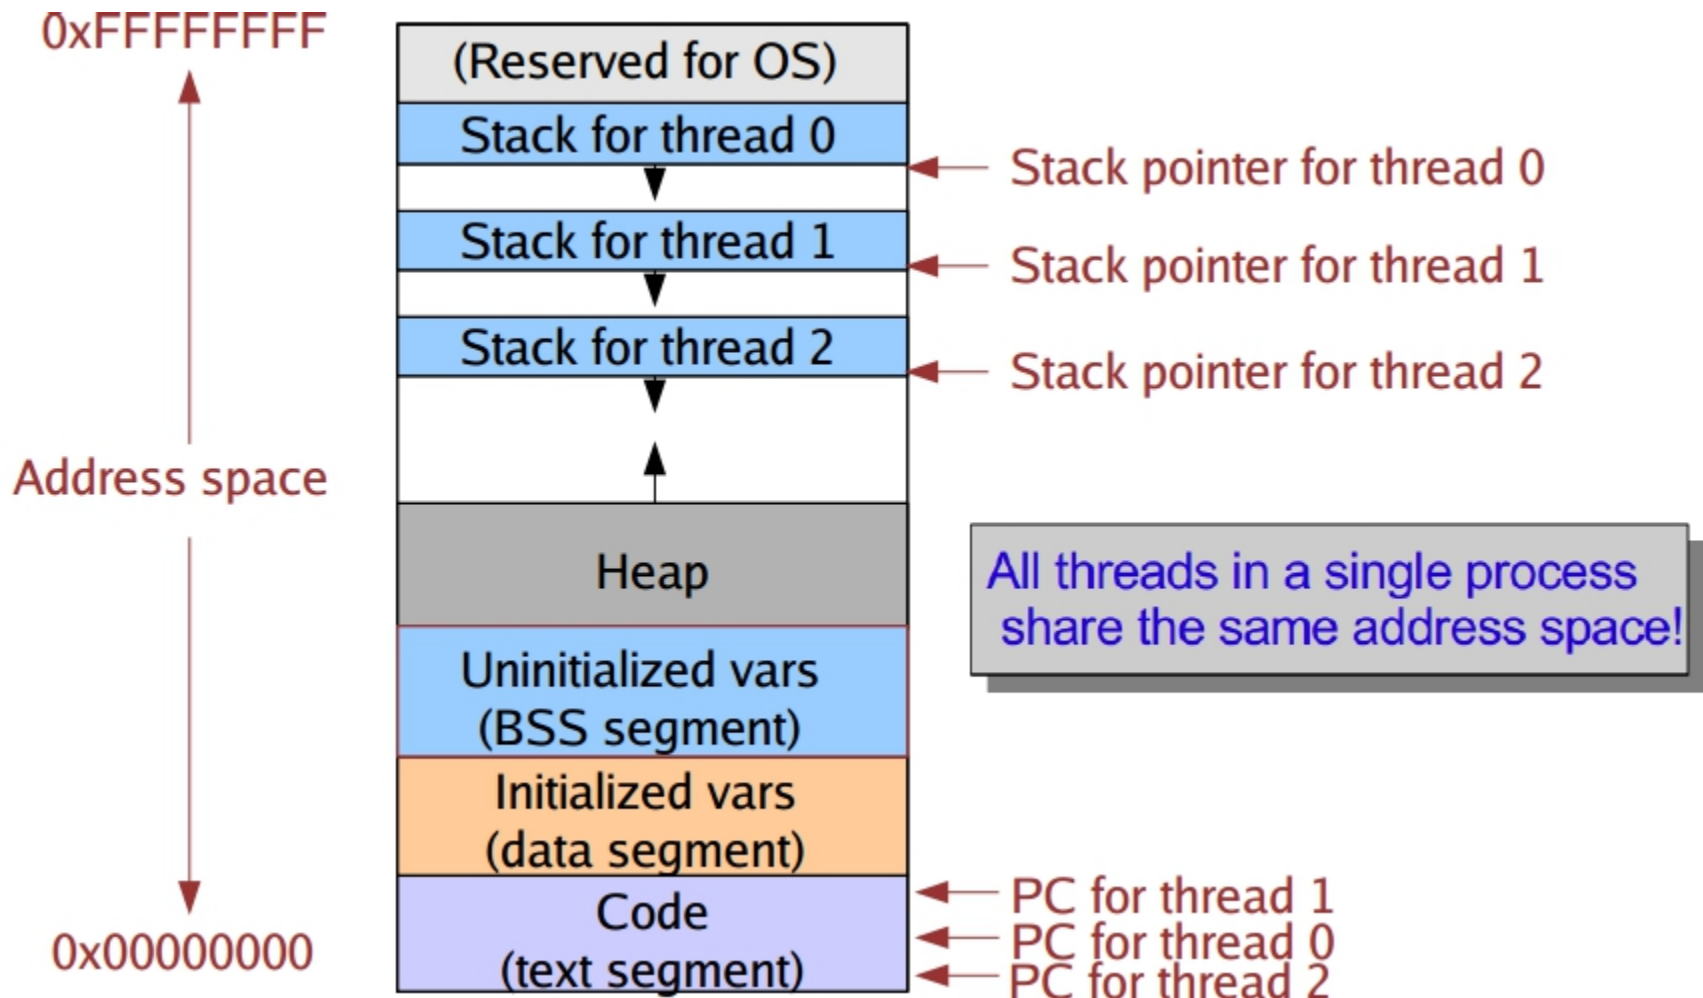
# Question

- The register set is listed as a per-thread rather than a per-process item. Why? After all, the machine has only one set of registers.

# (Old) Process Address Space




# (New) Address Space with Threads



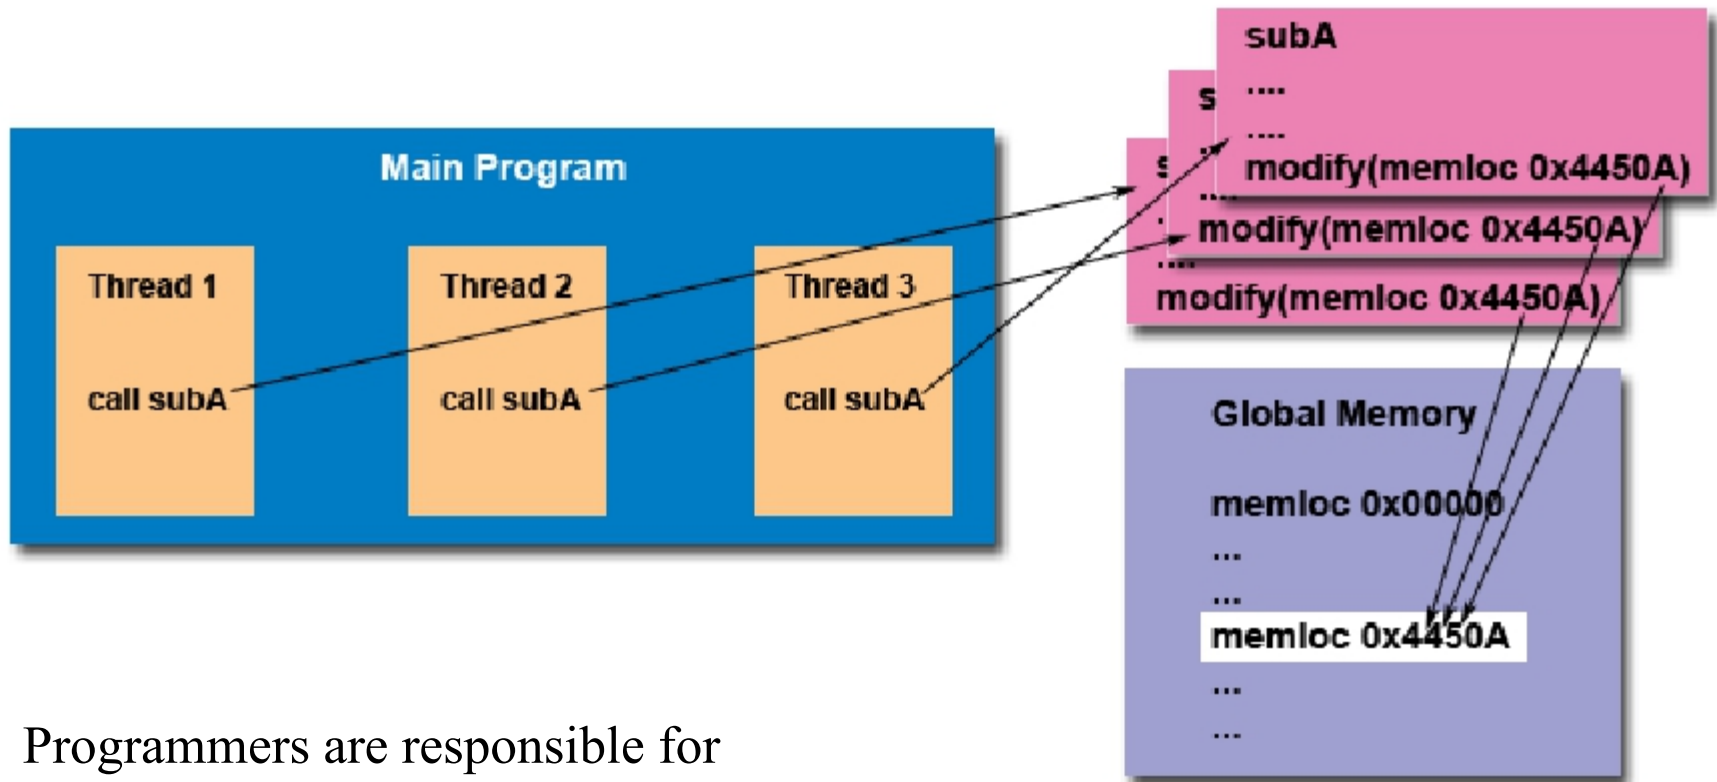
# Local and Global Variables for Threads

- *Global variables*: the same for all threads
- *Local variables*: each thread gets a copy of the variables
- Threads within a process share memory, open files, and I/O streams.
  - File descriptors are always considered as global variables





# Thread Safety



Programmers are responsible for synchronizing access (protecting) globally shared data.

# An Example of Unsafe Multi-Threading

```
void push_back(LIST *l, int v) {  
    NODE * n = malloc(sizeof(NODE)); // line 1  
    n->value = v; // line 2  
    n->next = NULL; // line 3  
    l->tail->next = n; // line 4  
}
```

Assume

- Two threads running the same code on the same list at the same time.
- Thread 1 just finishes line 3 and Thread 2 is at line 4
- Thread 1 continues on to finish push\_back and then Thread 2 runs next.
- What can happen?

# How to Avoid Mess?

- Synchronization, Lock mechanism, etc.
- Details: next lecture

# Goals for Today

- Threads
  - Concept
  - Multiple-threading
  - Implementation of threads

# POSIX Threads

Thread call	Description
Pthread_create	Create a new thread
Pthread_exit	Terminate the calling thread
Pthread_join	Wait for a specific thread to exit
Pthread_yield	Release the CPU to let another thread run
Pthread_attr_init	Create and initialize a thread's attribute structure
Pthread_attr_destroy	Remove a thread's attribute structure

Some of the Pthreads function calls.

# POSIX Threads (2)

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
```

```
#define NUMBER_OF_THREADS 10
```

```
void *print_hello_world(void *tid)
{
    /* This function prints the thread's identifier and then exits. */
    printf("Hello World. Greetings from thread %d\n", tid);
    pthread_exit(NULL);
}
```

```
int main(int argc, char *argv[])
{
    /* The main program creates 10 threads and then exits. */
    pthread_t threads[NUMBER_OF_THREADS];
    int status, i;


    for(i=0; i < NUMBER_OF_THREADS; i++) {
        printf("Main here. Creating thread %d\n", i);
        status = pthread_create(&threads[i], NULL, print_hello_world, (void *)i);
    }
}
```

# POSIX Threads (3)

```
int status, i;  
  
for(i=0; i < NUMBER_OF_THREADS; i++) {  
    printf("Main here. Creating thread %d\n", i);  
    status = pthread_create(&threads[i], NULL, print_hello_world, (void *)i);  
  
    if (status != 0) {  
        printf("Oops. pthread_create returned error code %d\n", status);  
        exit(-1);  
    }  
}  
exit(NULL);  
}
```

An example program using pthreads.

# Implementing Threads in User, or Kernel Space



- (a) A user-level threads package.  
(b) A threads package managed by the kernel.



# Question?

- In a system with threads, is there one stack per thread or one stack per process when user-level threads are used? What about when kernel-level threads are used? Explain.