



STEVENS
INSTITUTE of TECHNOLOGY
THE INNOVATION UNIVERSITY®

CS 492: Operating Systems


Threads (2)

Instructor: Iraklis Tsekourakis

Email: itsekour@stevens.edu



Implementing Threads in User Space



- (a) A user-level threads package.
(b) A threads package managed by the kernel.

Pros of Threads in User Space

- A user-level threads package can be implemented on an OS that does not support threads
- Thread switching is at least an order of magnitude faster, than trapping to the kernel
- Thread scheduling very fast: no context switching, no kernel trap, no flushing of memory cache
- Each process can have its own scheduling algorithm

Cons of Threads in User Space

- Blocking System calls (eg. waiting for keyboard input)
- Page faults (partial load of programs into memory)
- Threads need to voluntarily give up the CPU for multiprogramming
- Programmers generally want threads precisely in applications where threads block often (Web Server)

Pros of Threads in Kernel

- No run-time system needed in each process
- No thread table in each process
- Blocking system calls are not a problem, since the kernel scheduler can schedule another thread in that case


Cons of Threads in Kernel

- If thread operations are common (creation, termination), much more kernel overhead will be incurred
- Fork a multithreaded process?
- Signals sent to processes. Should the kernel assign it to a specific thread to handle?
- Slower than user-space threads

Question?


- What is the biggest advantage of implementing threads in user space?
- What is the biggest disadvantage?

Hybrid Implementations



Multiplexing user-level threads
onto kernel-level threads.

Pop-up Threads



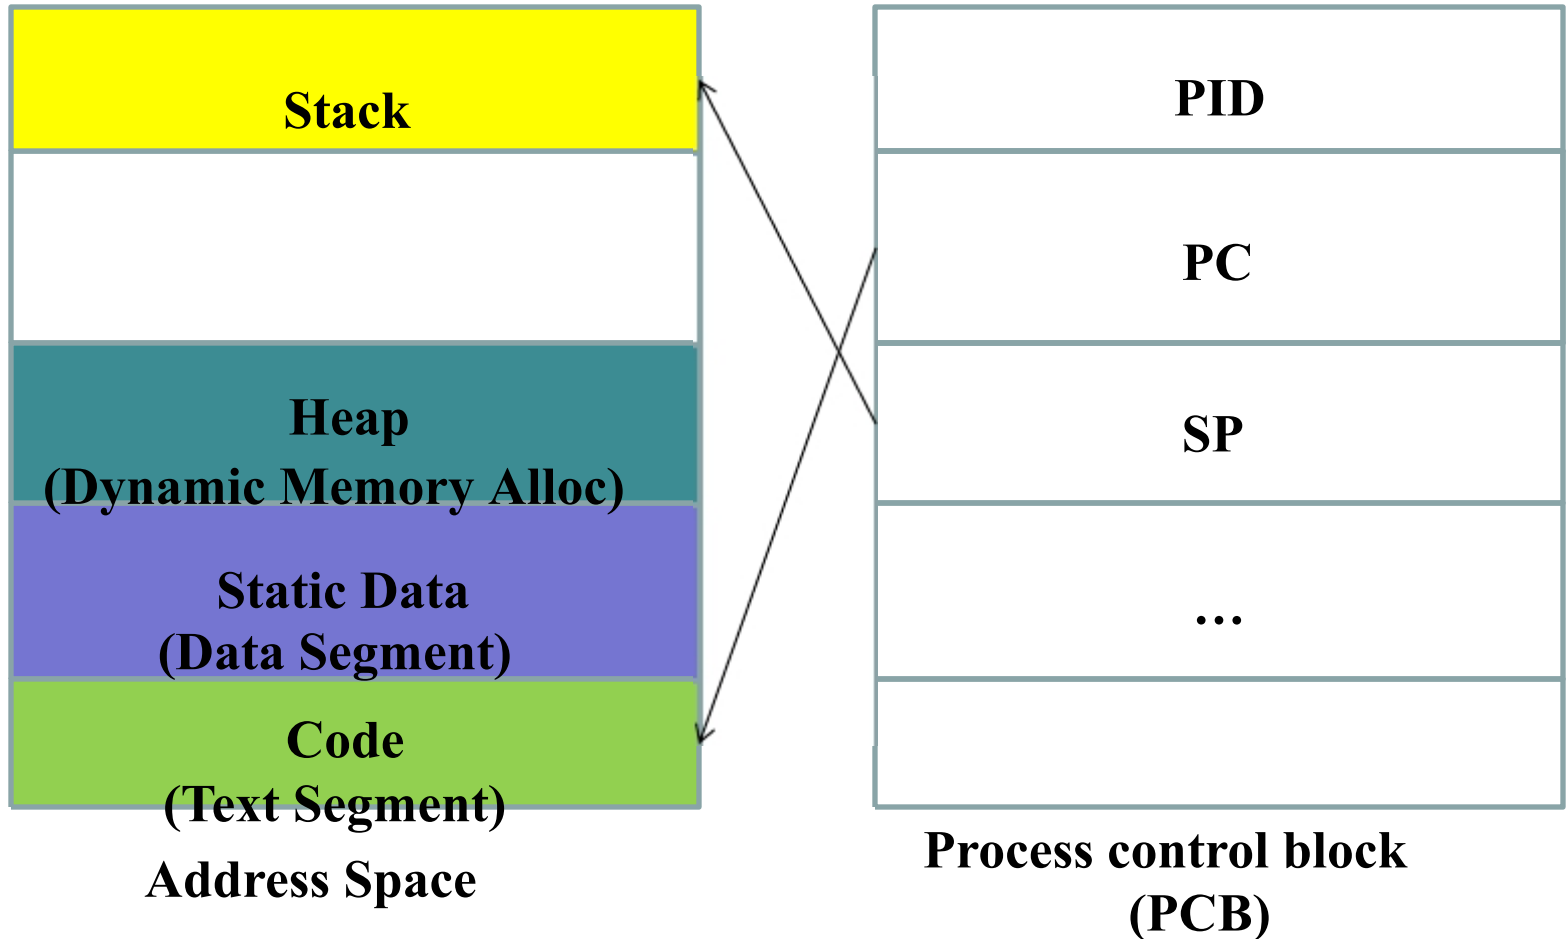
Creation of a new thread when a message arrives. (a) Before the message arrives. (b) After the message arrives.

Recall: Process Control Block (PCB)

- Each PCB contains
 - Process state
 - Process ID
 - Program Counter (PC)
 - Current CPU registers (if not executing)
 - CPU scheduling info (e.g., priority)
 - Memory-management info
 - Resources allocated to it
 - Resources it needs (e.g., I/O status information)

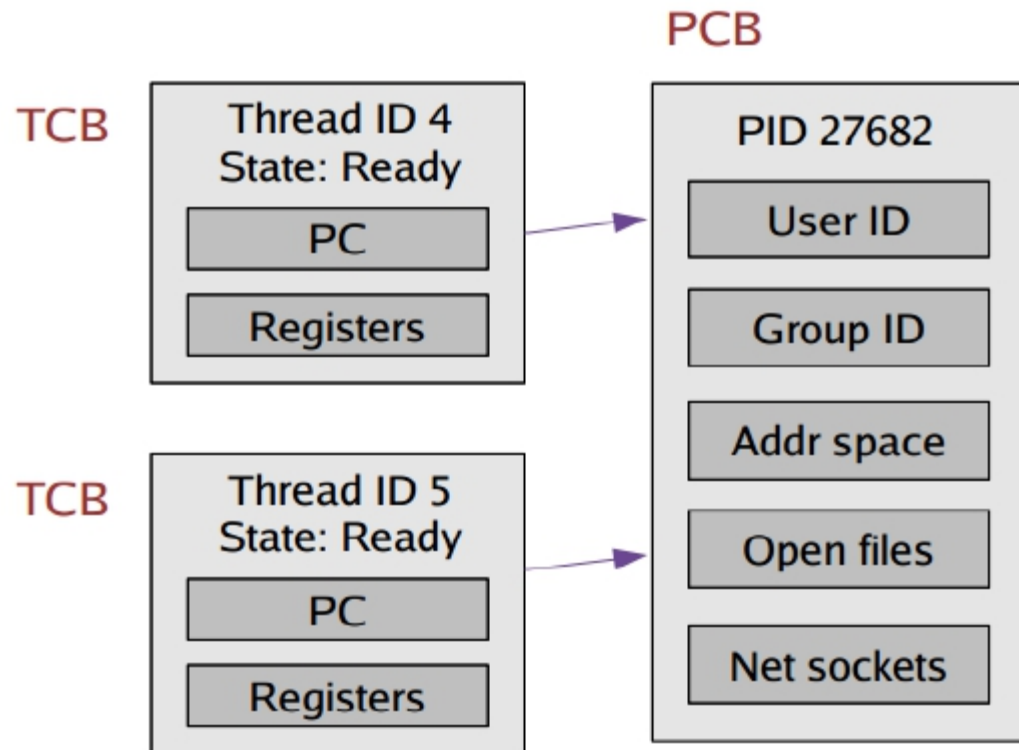
pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
⋮	

PCB Diagram

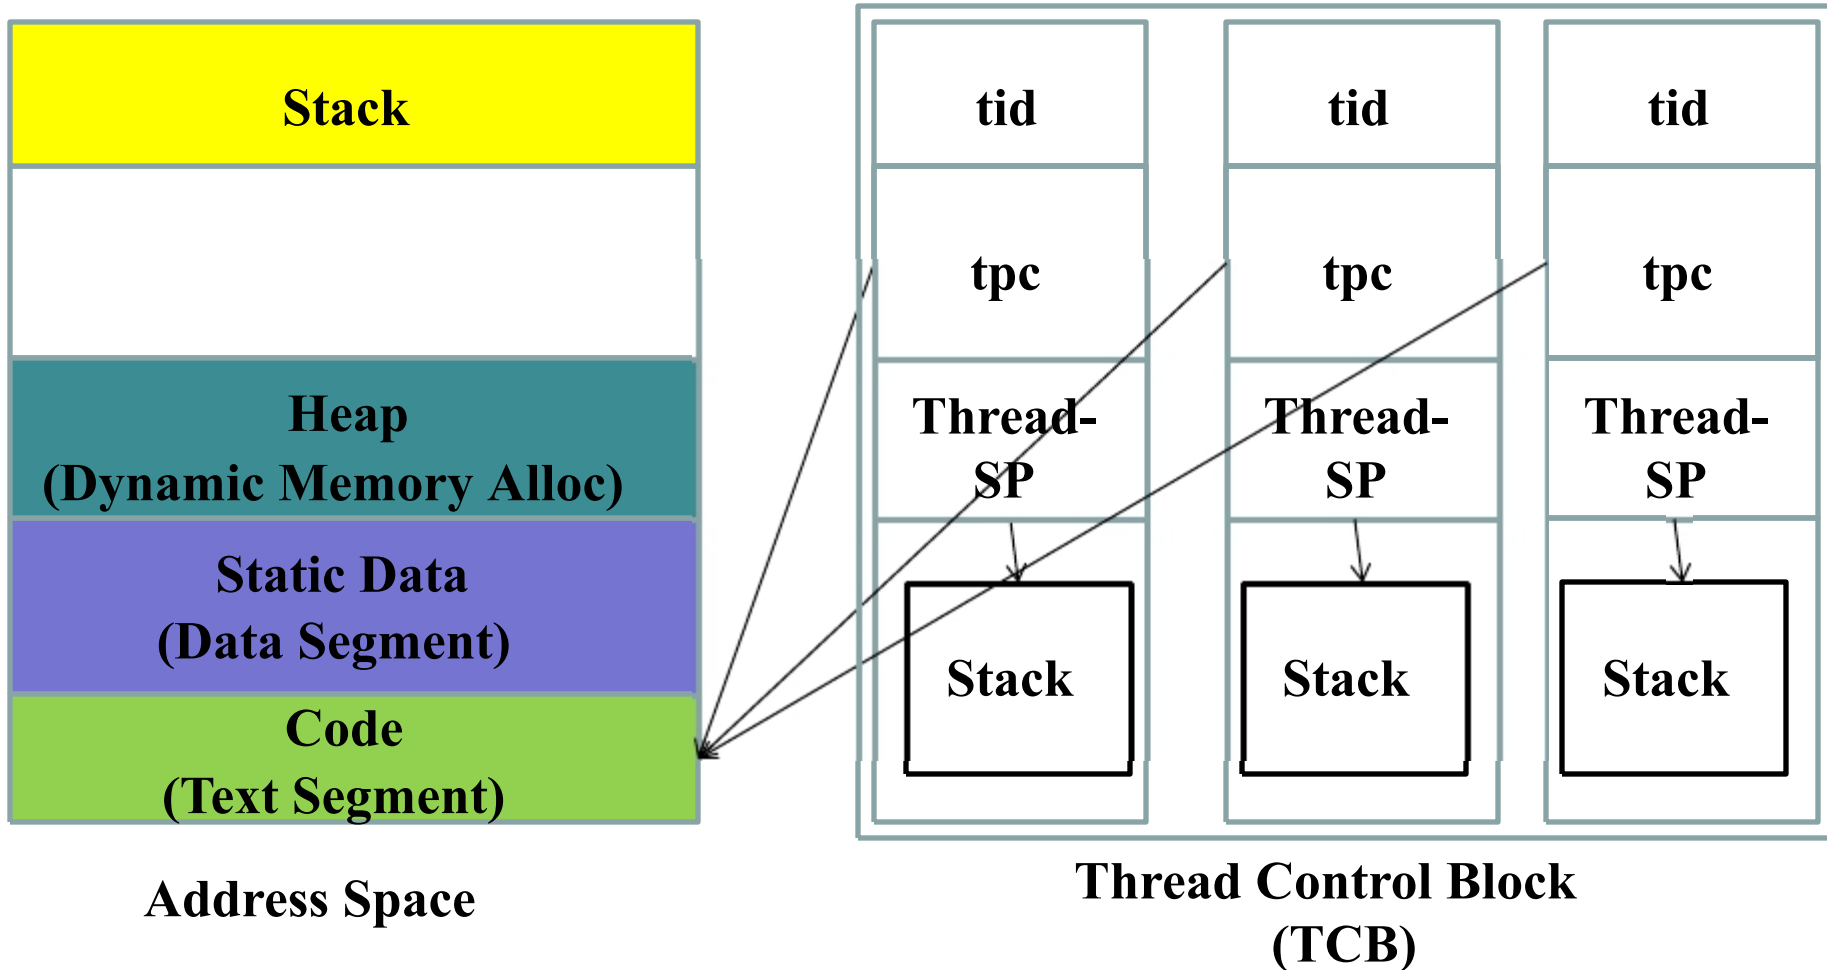


Thread Control Blocks (TCBs)

- Idea: Break the PCB into two pieces:
 - Thread-specific stuff: Process state
 - Process-specific stuff: Address space and OS resources (open files, etc.)



TCB Diagram



Thread Control Block (TCB)

- Information directly related to process execution:
stored in Thread Control Block (TCB)
 - Program counter
 - CPU registers
 - CPU scheduling information
 - Pending I/O information
- Other information associated with processes:
stored in Process Control Block (PCB)
 - Memory management information
 - Accounting information

Thread Control Block (TCB)

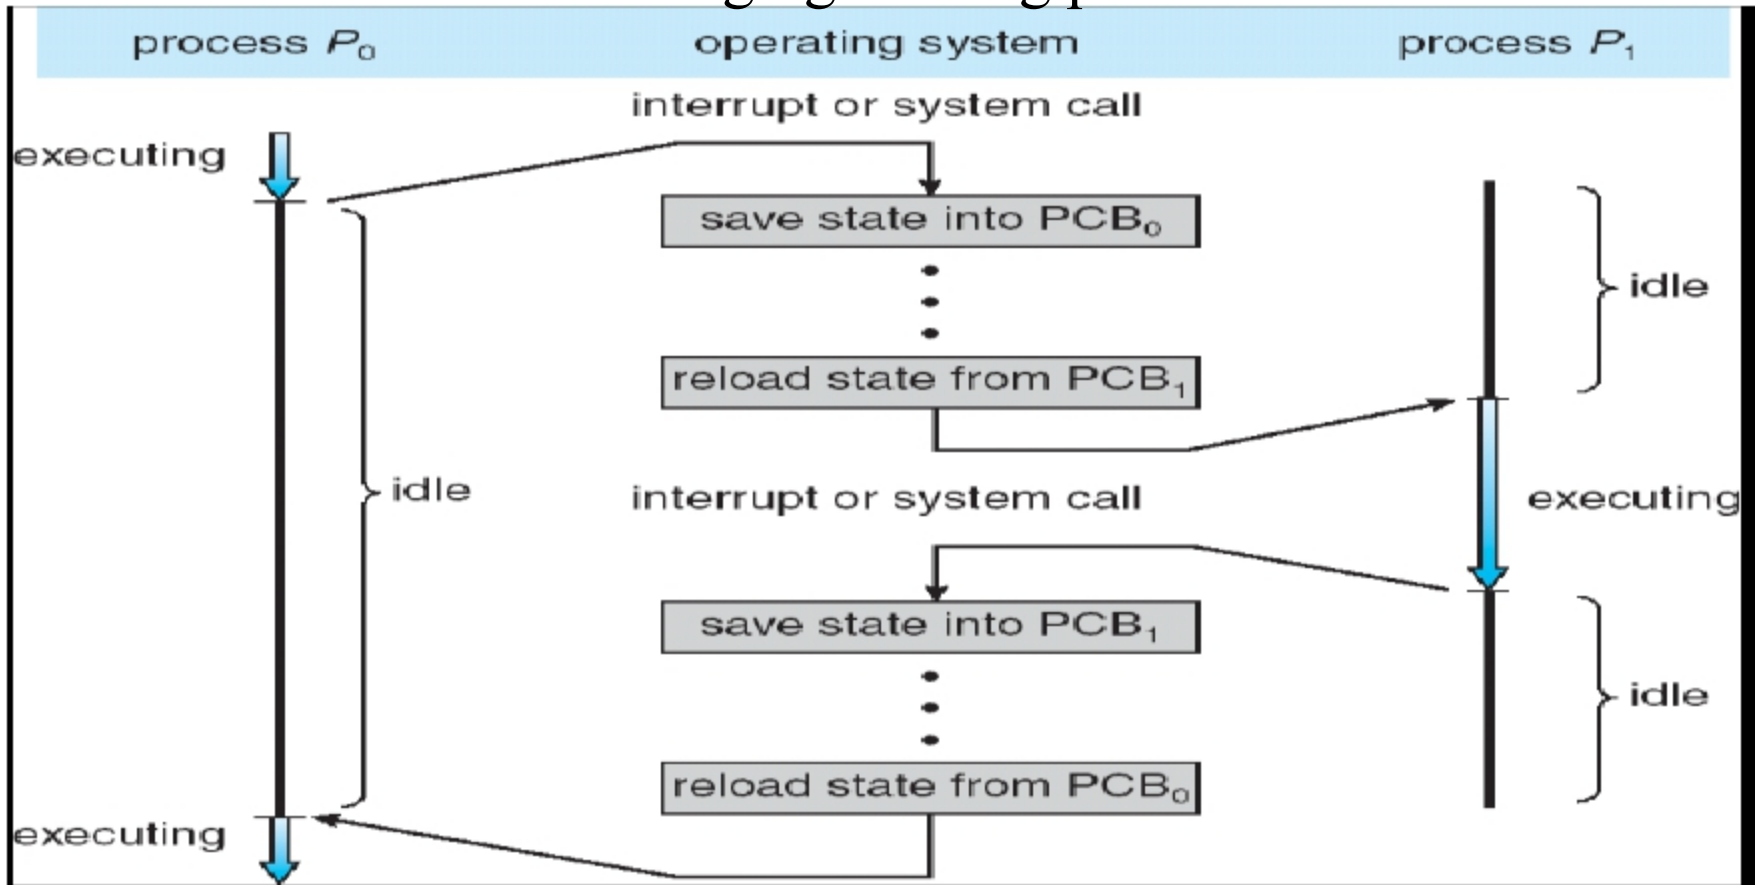
- TCB's are smaller and cheaper than processes
 - Linux TCB (`thread_struct`) has 24 fields
 - Linux PCB (`task_struct`) has 106 fields

Effect of TCBs

- Threads in a process can execute different parts of the program code at the same time.
- Threads can execute the same parts of the code at the same time, but with different execution state:
 - They have independent current instructions.
 - They are working with different data.

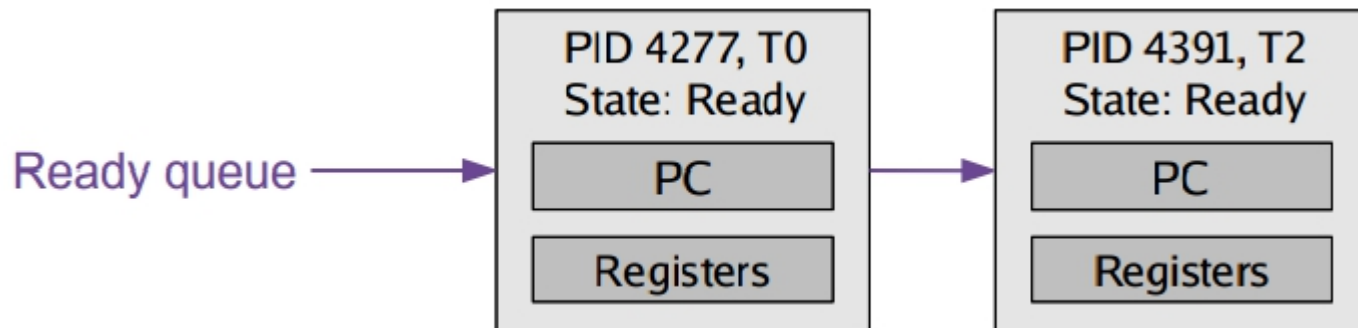
Recall: Context Switch

- *Context switch*: changing running proc



Now Context Switching with TCB

- TCB is now the unit of a context switch
 - Ready queue, wait queues, etc. now contain pointers to TCB's
 - Context switch causes CPU state to be copied to/from the TCB



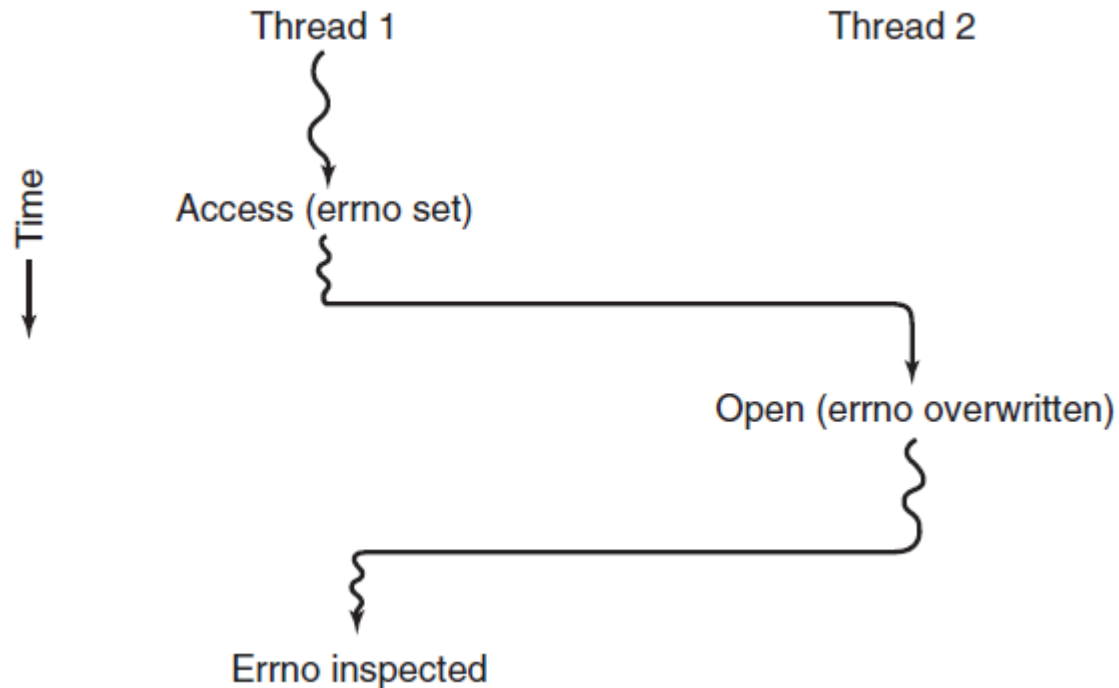
Context Switching (Cont.)

- Context switch between two threads in the same process:
 - No need to change address space
- Context switch between two threads in different processes:
 - Must change address space, sometimes invalidating cache
 - This will become relevant when we talk about virtual memory.

Concept Revisit


- Thread = an **independent** sequential execution stream within process
- The independency is accomplished because a thread maintains its own:
 - Stack pointer
 - Registers
 - Scheduling properties (such as policy or priority)

Making Single-Threaded Code Multi-Threaded



Conflicts between threads over the use of a global variable.

Making Single-Threaded Code Multi-Threaded (2)



Threads can have private global variables.

Making Single-Threaded Code Multi-Threaded (3)

- `create_global("bufptr");` *//allocates storage for a pointer*
- `set_global("bufptr", &buf);`
- `bufptr = read_global("bufptr");`
 - We used a procedure `create global` to allocate storage for a pointer to the variable, rather than the variable itself. Is this essential, or could the procedures work with the values themselves just as well?

Making Single-Threaded Code Multi-Threaded (4)

- Library procedures not reentrant
- Memory allocation procedures
- Signals (especially in user-level threads)
- Stack management