

CS 492: Operating Systems

Inter Process Communication (2)


Software/Hardware Solution

Instructor: Iraklis Tsekourakis

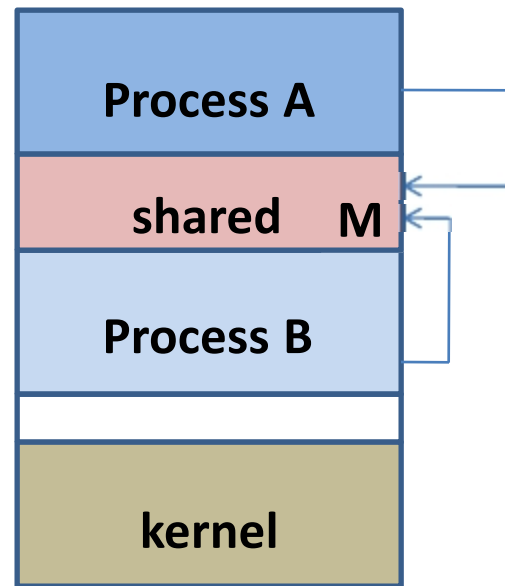
Email: itsekour@stevens.edu

Recap: Inter-process Communication

- How can processes interact in real time?
 - (a) By passing messages through the kernel
 - (b) By sharing a region of physical memory
 - (c) Through asynchronous signals or alerts



(a)



(b)

Recap

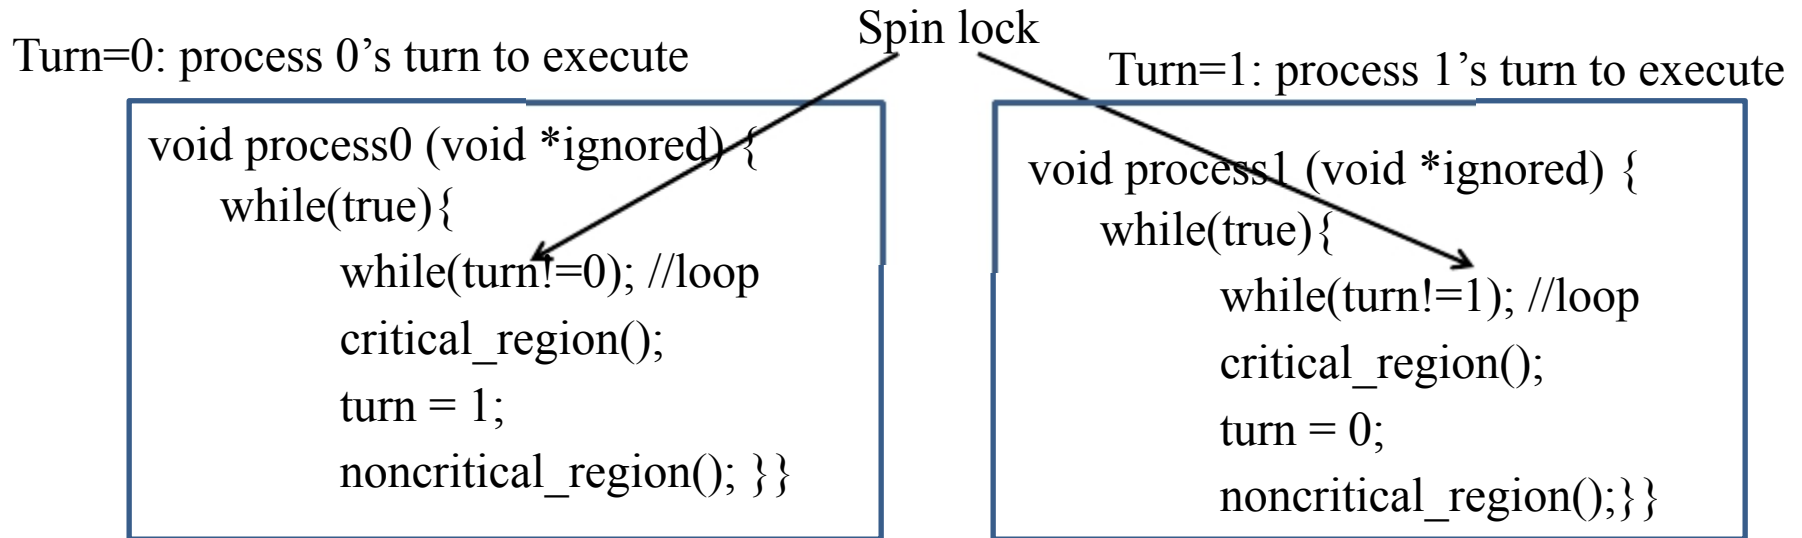
- Sharing a region of physical memory
 - Race conditions
 - Cooperating processes read and write the shared memory
 - Critical section
 - Section of code where the shared memory is accessed
 - Mutual exclusion
 - Only one thread can be in critical section at a time

Mechanisms for Mutual Exclusion

- 1) Software solution
 - 1.1) Strict alternation
 - 1.2) Peterson's solution
- 2) Hardware solution
 - 2.1) Interrupt disabling
 - 2.2) Test-and-Set lock (TSL)
- 3) Higher level solutions

Software Mechanism 1: Strict Alternation

- Use a shared variable **Turn** to strictly alternate between processes
- Waiting process continually reads the variable to see if it can proceed
 - Called **spin-lock** – lock wherein a process **busy waits**



Strict Alternation: More Discussion

- Disadvantages
 - Horribly wasteful!
 - Only allows two processes (extensions to more than 2 processes are expensive)
 - Fast process will get blocked by the slow process

Software Mechanism 2: Peterson's Solution

- Assume two threads, P0 and P1.
- Variables:
 - int turn – whose turn to enter critical section
 - bool interested[2]
 - interested[i]=true means Pi is ready to enter critical section

Peterson's Solution

```
#define FALSE 0
#define TRUE 1
#define N      2                /* number of processes */

int turn;                       /* whose turn is it? */
int interested[N];              /* all values initially 0 (FALSE) */

void enter_region(int process);  /* process is 0 or 1 */
{
    int other;                  /* number of the other process */

    other = 1 - process;        /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;             /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process)   /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```


Peterson's solution: More Discussion

- Mutual exclusion

- Q1: Why Peterson's solution guarantees mutual exclusion?
- Q2: Why Peterson's solution is more efficient than the strict alternation?

```
void enter_region(int process);      /* process is 0 or 1 */
{
    int other;                      /* number of the other process */

    other = 1 - process;            /* the opposite of process */
    interested[process] = TRUE;     /* show that you are interested */
    turn = process;                 /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}
```

- Still only works for 2 processes

- Can generalize to n processes, but for some fixed n

Mechanisms for Mutual Exclusion

- 1) Software solution
 - 1.1) Strict alternation
 - 1.2) Peterson's solution
- 2) Hardware solution
 - 2.1) Interrupt disabling
 - 2.2) Test-and-Set lock (TSL)
- 3) Higher level solutions

Hardware Mechanism 1: Disabling Interrupts

- Idea:
 - Right before entering the critical section: the process disables system interrupts;
 - Before leaving the critical section: the process re-enables the interrupts;
 - Effect: CPU cannot be switched to other processes when a process is in the critical section.
- Disadvantages??

Hardware Mechanism 1: Disabling Interrupts (2)

- On the other hand, it could be convenient for the kernel to disable interrupts for a few instructions
 - Update variables or lists (e.g. ready processes list)
- Actually, not anymore

Hardware Mechanism 2:

Locking Test-and-Set Lock (TSL)

- Test and modify the content of a word atomically
 - The word here is the lock.
 - **TSL REGISTER,LOCK**

TestAndSet function:

- Return the current value of memory word *lock*, then set *lock* to be true.
- Atomic operation (the code cannot be interrupted during execution)

Hardware Mechanism 2: The TSL Instruction (1)

enter_region:

TSL REGISTER,LOCK

CMP REGISTER,#0

JNE enter_region

RET

| copy lock to register and set lock to 1
| was lock zero?
| if it was nonzero, lock was set, so loop
| return to caller; critical region entered

leave_region:

MOVE LOCK,#0

RET

| store a 0 in lock
| return to caller

Entering and leaving a critical region
using the TSL instruction.

Hardware Mechanism 2: Locking Test-and-Set Lock (TSL)

Implementation

Enter_region:

//If the lock is set, wait till it is reset by other processes

//Else enter the critical section and set the lock

Critical section;

Leave_region:

//reset the lock

Hardware Mechanism 2: The TSL Instruction (2)

enter_region:

```
MOVE REGISTER,#1
XCHG REGISTER,LOCK
CMP REGISTER,#0
JNE enter_region
RET
```

```
| put a 1 in the register
| swap the contents of the register and lock variable
| was lock zero?
| if it was non zero, lock was set, so loop
| return to caller; critical region entered
```

leave_region:

```
MOVE LOCK,#0
RET
```

```
| store a 0 in lock
| return to caller
```

Entering and leaving a critical region
using the XCHG instruction

Busy Waiting..

- Disabling interrupts, using TSL, strict alternation, and Peterson's algorithm
 - All -> busy waiting.
 - Not only wasteful, but can also have unexpected effects!! **The priority Inversion Problem**
- Can we do better?

Mechanisms for Mutual Exclusion

- 1) Software solution
 - 1.1) Strict alternation
 - 1.2) Peterson's solution
 - 1.3) Sleep and wakeup**
- 2) Hardware solution
 - 2.1) Interrupt disabling
 - 2.2) Test-and-Set lock (TSL)
- 3) Higher level solutions

Sleep and Wakeup

The producer-consumer problem (1)

```
#define N 100                                     /* number of slots in the buffer */
int count = 0;                                    /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                                /* repeat forever */
        item = produce_item();                    /* generate next item */
        if (count == N) sleep();                  /* if buffer is full, go to sleep */
        insert_item(item);                        /* put item in buffer */
        count = count + 1;                        /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);        /* was buffer empty? */
    }
}
```

```
void consumer(void)
```

```
{
    ...
}
```

Sleep and Wakeup

The producer-consumer problem (2)

```
if (count == 1) wakeup(consumer);
}
}

void consumer(void)
{
    int item;

    while (TRUE) {
        if (count == 0) sleep( );
        item = remove_item( );
        count = count - 1;
        if (count == N - 1) wakeup(producer);
        consume_item(item);
    }
}
```

/ repeat forever */*
/ if buffer is empty, got to sleep */*
/ take item out of buffer */*
/ decrement count of items in buffer */*
/ was buffer full? */*
/ print item */*

Sleep and Wakeup....

```
#define N 100                                     /* number of slots in the buffer */
int count = 0;                                    /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                                /* repeat forever */
        item = produce_item();                    /* generate next item */
        if (count == N) sleep();                  /* if buffer is full, go to sleep */
        insert_item(item);                        /* put item in buffer */
        count = count + 1;                        /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);        /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {                                /* repeat forever */
        if (count == 0) sleep();                  /* if buffer is empty, got to sleep */
        item = remove_item();                     /* take item out of buffer */
        count = count - 1;                        /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer);    /* was buffer full? */
        consume_item(item);                       /* print item */
    }
}
```

Any race conditions here? Any cases where it fails?