



STEVENS
INSTITUTE of TECHNOLOGY
THE INNOVATION UNIVERSITY®

CS 492: Operating Systems

Inter Process Communication (4)

Monitors

Instructor: Iraklis Tsekourakis

Email: itsekour@stevens.edu

Announcement

- Hw1 due tonight!
- Teams have been formed
- Make sure you coordinate with your CAs (compilers etc.)
 - Try with them a hello world application.
- (Prob. next week) I will also introduce GIT.
 - You should create a repo, and add your CA

Announcement 2

- Next Tuesday (Feb 13): I will introduce the lab assignment
- Next Wednesday (Feb 14): Pthreads tutorial (required for Assignment 1, and Lab)
 - Introduction of thread programming
- Next Friday (Feb 16): Lab session (Inter-process communication)

Recap: Semaphores

- Semaphore: a count variable that counts the numbers of wakeups saved
- Operations:
 - *Down (for sleep):*
 - decrements the value by 1 (used 1 saved wakeup).
 - Process sleeps when it equals zero
 - *Up (for wakeup):* increments the value by 1 (save a wakeup)

Recap: Mutex

- A simplified version of semaphores
 - Two states: blocked and unblocked
- Implementation of mutual exclusion by using mutex

Reader-Writer Problem

- Readers: read data
- Writers: write data
- Rules:
 - Multiple readers can read the data simultaneously
 - Only one writer can write the data at any time
 - A reader and a writer cannot be in critical section together.
- Locking table: whether any two processes can be in the critical section simultaneously

	Reader	Writer
Reader	OK	No
Writer	No	No

Readers-Writers Problem

- Allow multiple processes do reading concurrently
- Allow one single process do writing at a time
- Semaphore mutex = 1; /* control critical section */
- Semaphore wrt = 1; /* # of writer process */
- int readcount = 0; /* reader counter */


Reader-Writer Solution


```
reader() {
    while(TRUE) {
        <other stuff>;
        down(mutex);
        readCount++;

        if(readCount == 1)
            down(wrt);
        up(mutex);

        /* Critical section */
        access(resource);


        down(mutex);
        readCount--;
        if(readCount == 0)
            up(wrt);
        up(mutex);
    }
}
```

 First reader blocks writers

 Last reader unblocks writers

```
int readCount = 0;
semaphore mutex = 1;
semaphore wrt = 1;

writer() {
    while(TRUE) {
        <other computing>;
        down(wrt);
        /* Critical section */
        access(resource);
        up(wrt);
    }
}
```

 Writer blocks all readers

Is Semaphores the Final Solution?

- Semaphores accomplish 2 purposes:
 - 1) Mutual exclusion of the shared memory (mutex)
 - 2) Scheduling constraints: *down* and *up*
- But tricky to use
 - if *down()* and *up()* are scattered among several processes it may be difficult to use them correctly
 - We have shown that
 - Changing order of *down* may cause deadlocks
 - Changing order of *up* affects efficiency
- We need something that is simple and easy
 - Answer: monitors

Monitors

- Motivation
 - Would be nice if we had some clean, well-defined **language support** for synchronization...
 - Take complexity out of programmer's hands – let compiler deal with mutual exclusion
 - Cleaner idea
 - Use ***condition variables*** for scheduling constraints
- Rules:
 - 1) **Only one** process can be **active** in a monitor at any time
 - 2) Processes cannot access internal data of monitor
 - 3) Put all critical sections inside monitor

Implementation of Monitors

- A package that consists of a collection of procedures, variables and data structures

Monitor

Integer I

Condition c

Procedure producer()

Procedure consumer()

End Monitor

- A programming language construct
- Compiler (not the programmer!) implements the mutual exclusion on monitor entries

Condition Variables in Monitors

- Mutual Exclusion is not enough. (req 2)
- Operations:
 - `Wait (&condVar)` : Causes the calling process to block. Re-acquire lock later, before returning.
 - `Signal (&condVar)` : Wake up one blocked process, if any
- Rule: Must hold Condition Variable when doing condition variable operations!

Condition Variables in Monitors (2)

- Signal: how do we make sure we don't have 2 active processes in the monitor?
- *Wait, signal*: similar to *sleep* and *wakeup*.
Lost wakeUp calls in case of context switch.
 - Does this hold for Monitor Condition variables too?

```
monitor example
  integer i;
  condition c;

  procedure producer();
  .
  .
end;

procedure consumer();
.
.
end;
end monitor;
```

Condition Variables in Monitors

```
monitor ProducerConsumer
  condition full, empty;
  integer count;

  procedure insert(item: integer);
  begin
    if count = N then wait(full);
    insert_item(item);
    count := count + 1;
    if count = 1 then signal(empty)
  end;

  function remove: integer;
  begin
    if count = 0 then wait(empty);
    remove = remove_item;
    count := count - 1;
    if count = N - 1 then signal(full)
  end;

  count := 0;
end monitor;
```

```
procedure producer;
begin
  while true do
    begin
      item = produce_item;
      ProducerConsumer.insert(item)
    end
  end;

procedure consumer;
begin
  while true do
    begin
      item = ProducerConsumer.remove;
      consume_item(item)
    end
  end;
end;
```

Question

Synchronization within monitors uses condition variables and two special operations, wait and signal. A more general form of synchronization would be to have a single primitive, *waituntil*, that had an arbitrary Boolean predicate as parameter. Thus, one could say, for example, *waituntil* $x < 0$ or $y + z < n$.

The signal primitive would no longer be needed. This scheme is clearly more general than that of Monitor Condition Variables, but it is not used.

Why not?

Monitor: Reader-Writer Problem

- The monitor solution uses two variables:
 - int **readercount**: Number of active readers; initially = 0
 - bool **busy**: indicate whether a writer is active initially = false

And two condition variables:

- Condition variable **OKtoread** = NIL
- Condition variable **OKtowrite** = NIL

Readers, and Writers

```
readers-writers : monitor;
begin
  readercount : integer;
  busy : boolean;
  OKtoread, OKtowrite : condition;

  procedure startread;

    begin
      if busy then OKtoread.wait;
      readercount := readercount + 1;
      OKtoread.signal;
      (* Once one reader can start, they all can *)
    end startread;

  procedure endread;

    begin
      readercount := readercount - 1;
      if readercount = 0 then OKtowrite.signal;
    end endread;

  procedure startwrite;

    begin
      if busy OR readercount != 0 then OKtowrite.wait;
      busy := true;
    end startwrite;

  procedure endwrite;

    begin
      busy := false;
      if OKtoread.queue then OKtoread.signal
        else OKtowrite.signal;
      end endwrite;

begin (* initialization *)
  readercount := 0;
  busy := false;
end;

end readers-writers;
```

Readers, and Writers

```
procedure startread;  
  
    begin  
        if busy then OKtoread.wait;  
        readercount := readercount + 1;  
        OKtoread.signal;  
        (* Once one reader can start, they all can *)  
    end startread;
```

Readers, and Writers

```
procedure endread;  
  
begin  
  readercount := readercount - 1;  
  if readercount = 0 then OKtowrite.signal;  
end endread;
```

Readers, and Writers

```
procedure startwrite;  
  
    begin  
    if busy OR readercount != 0 then OKtowrite.wait;  
    busy := true;  
    end startwrite;
```

Readers, and Writers

```
procedure endwrite;  
  
    begin  
        busy := false;  
        if OKtoread.queue then OKtoread.signal  
            else OKtowrite.signal;  
        end endwrite;
```