



**STEVENS**  
INSTITUTE of TECHNOLOGY  
THE INNOVATION UNIVERSITY®

# CS 492: Operating Systems

## *Scheduling*

*Instructor: Iraklis Tsekourakis*

Email: [itsekour@stevens.edu](mailto:itsekour@stevens.edu)



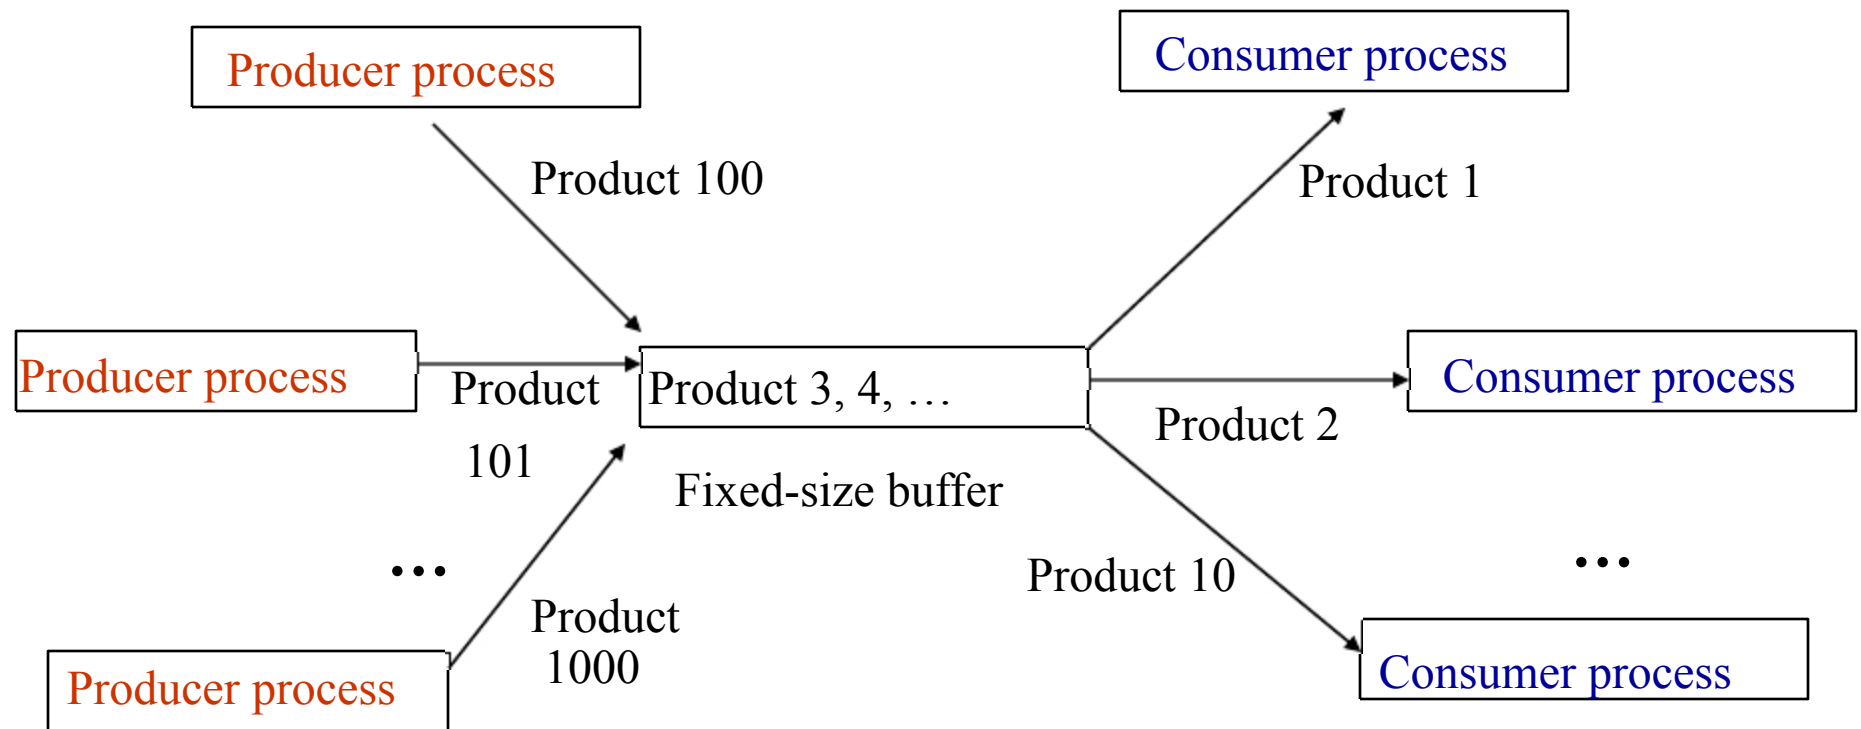
# MIDTERM!

- Friday, March 9<sup>th</sup>
  - Computer Test on Canvas
  - Install lockdown browser

# Assignment 1

(Due: 11:59pm, Monday, March 5<sup>th</sup> )

# Assignment 1



# Input Parameters

- P1: Number of producer threads
- P2: Number of consumer threads
- P3: Total number of products to be generated by all producer threads
- P4: Size of the queue to store products for both producer and consumer threads (0 for unlimited queue size)
- P5: 0 or 1 for type of scheduling algorithm
  - 0: First-Come-First-Serve;
  - 1: Round-Robin
- P6: Value of quantum used for round-robin scheduling
- P7: Seed for random number generator

# Product

- Each product has:
  - A unique product ID;
  - The timestamp when the product is generated
  - The “life” of the product
- After a product is consumed, the consumer thread prints out the product ID, and sleeps 100 milliseconds (by calling `usleep()`).

# Critical Section

- Protect the queue
  - When the queue is full
  - When the queue is empty
- Protect the shared variables (e.g., total number of products generated by all producer threads)

# Scheduling (1/2)

- Producers: first come, first serve
- Consumer threads: two scheduling algorithms
  1. First come, first serve
    - Products are sorted in the order that they are inserted into the queue
    - The **consumption** of a product is simulated by calling the  $fn(10)$  function  $l$  times, where  $l$  is the product's life




# Scheduling (1/2)

- Producers: first come, first serve
- Consumer threads: two scheduling algorithms
  - First come, first serve
  - Round-robin
    - Each product is assigned a quantum  $q$
    - If the product's life  $l > q$ 
      - Call  $fn(10)$  function  $q$  times
      - Update product's life  $l = l - q$ , and keep product in the queue.
    - Otherwise,
      - Call  $fn(10)$  function  $l$  times
      - Print product ID, and remove product from queue.

# Experiments and Reports

- Experiments: change the setting of the following parameters
  - Number of producers/consumers
  - Value of quantum used for round-robin algorithm
- Report: collect statistics of
  - Total time for processing all products
  - Min, max, and average *turn-around* times (the time between the product is produced and it is removed from queue)
  - Min, max, and average *wait* times (the time that products wait in the queue)
  - Producer throughput (number of products inserted into queue per minute)
  - Consumer throughput (number of products printed out per minute)
- More details/requirements are in the assignment description

# Scheduling



# Recap

- Inter-process communication (IPC)
  - Mechanisms for shared memory
    - Software-based techniques
    - Hardware-based techniques
    - Higher level (Semaphores, Mutex, Monitors)

# What's Next?

- Multiple processes/threads are ready to run
- We assume single CPU scenario
  - So only one process/thread can be picked
- Which one?
- But, we have to worry about efficient use of CPU as well

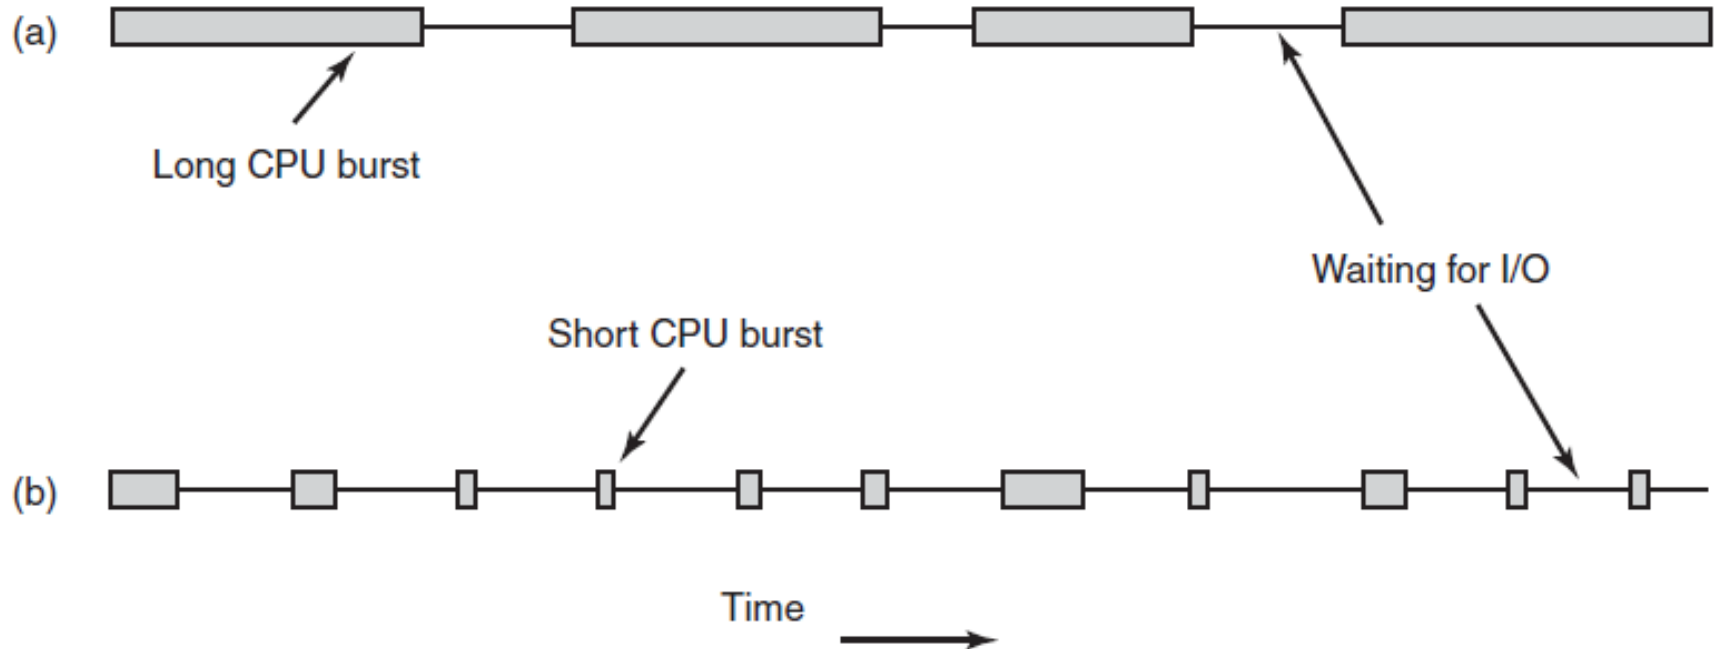
# Goals for Today

- Scheduling
  - Issues in scheduling
  - Basic scheduling algorithms
    - Batch system
    - Interactive system
    - Real-time system

# Scheduling

- The Problem
  - $K > 1$  processes/threads that are ready
  - $N \geq 1$  CPUs (  $N$  may be  $< K$  )
  - Which processes/threads to assign to which CPU(s)?

# Introduction to Scheduling Behavior



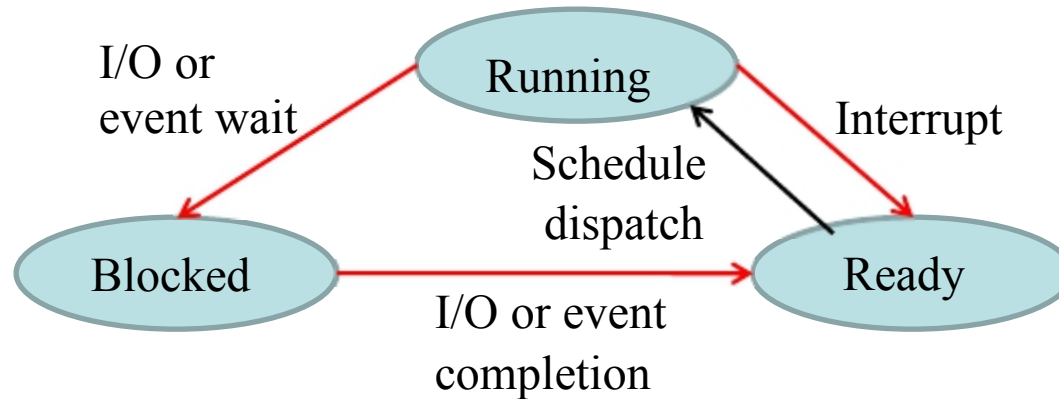
CPU-bound vs I/O-bound



# Question?

- Can a measure of whether a process is likely to be CPU bound or I/O bound be determined by analyzing source code? How can this be determined at run time?

# When to Schedule?



- Scheduling decisions may take place when
  1. A process is created
  2. A process in the running state exits
  3. A process switches from *Running* to *Blocked* state (e.g., by semaphore)
  4. A process switches from *Running* to *Ready* state
  5. A process switches from *Blocked* to *Ready*

# Clock Interrupts

- Scheduling takes place at clock interrupts
- Clock interrupts: how system keeps track of time
  - Interrupt at periodic intervals (clock *tick*)
  - Typically 1msec
  - Implemented using a hardware clock interrupt
  - High priority, second to power-failure interrupt.

\*However, Ivan Sutherland,

[http://www.cs.virginia.edu/~robins/Computing\\_Without\\_Clocks.pdf](http://www.cs.virginia.edu/~robins/Computing_Without_Clocks.pdf)

# Non-preemptive vs. Preemptive Scheduling

- *Non-preemptive* scheduling
  - For every clock interrupt, running process keeps going
- *Preemptive* scheduling
  - For every clock interrupt, running process is suspended and switched with another process (if there is any)
  - Maximum or fixed time execution of processes (timeslice)

# Scheduling Criteria (1/2)

- User-oriented
  - Minimize *response* time - time from issuing request to get the first response
    - Time to echo a keystroke in editor
    - Time to compile a program
    - Real-time Tasks: Must meet deadlines
  - Minimize *turnaround* time – time between submission and termination
  - Maximize *throughput* – # of jobs completed per unit time
  - Freedom from starvation
  - Fairness

## Scheduling Criteria (2/2)

- System-oriented
  - Maximize CPU *utilization* - keep the CPU as busy as possible
  - Minimize *overhead* – keep CPU doing useful work as opposed to context switching

# Goals for Today

- Scheduling
  - Issues in scheduling
  - Basic scheduling algorithms
    - Batch system
    - Interactive system
    - Real-time

# Scheduling in Batch Systems

- Batch systems - off-line, jobs come with resource estimates
- Assumption: know how much time is needed to complete job ahead of time
- Scheduling goals
  1. Throughput: maximize jobs per unit time (hour)
  2. Turnaround time: minimize time users wait for jobs
  3. CPU utilization: keep the CPU as busy as possible
- Algorithms:
  - First-Come First-Served (FCFS)
  - Shortest job first (SJF)




# First-Come, First-Served (FCFS) Scheduling

- Example

Process	CPUTime
P1	24
P2	3
P3	3

**Gantt Chart for Schedule**



*turnaround time* – average time takes for job to be completed after submission


- Suppose the arrival order for the processes is P1, P2, P3
- Turnaround time
  - P1 = 24;
  - P2 = 27;
  - P3 = 30;
- Average turnaround time  $(24+27+30)/3 = 27$
- Short process delayed by long process: *Convoy effect*

# First-Come, First-Served (FCFS) Scheduling (Cont.)

- Example

Process	CPUTime
P1	24
P2	3
P3	3

**Gantt Chart for Schedule**



- Suppose now the arrival order for the processes changes: P2, P3, P1
- Turnaround time
  - P1 = 30
  - P2 = 3;
  - P3 = 6;
- Average turnaround time
  - $(30+3+6)/3 = 13$
- Much better than the previous case

# Shortest-Job-First (SJR) Scheduling

- Associate with each process the length of its CPU time.
- Use the CPU time length to schedule the process with the shortest CPU time.
- Two variations:
  - *Non-preemptive* – once CPU given to the process it cannot be taken away until it completes.
  - *preemptive* – if a new process arrives with CPU time less than the remaining time of current executing process, preempt. (*Shortest Remaining Time Next*)

# Non-Preemptive SJF Example

- Example

Process	ArrivalTime	CPUTime
P1	0	7
P2	2	4
P3	4	1
P4	5	4


What's the turnaround time for each process?  
What's the average turnaround?

# Non-Preemptive SJF Solution

- Example

Process	Arrival Time	CPU Time
P1	0	7
P2	2	4
P3	4	1
P4	5	4

**Gantt Chart for Schedule**



- Turnaround time

- $P1 = 7$ ;
- $P2 = 12 - 2 = 10$ ;
- $P3 = 8 - 4 = 4$ ;
- $P4 = 16 - 5 = 11$

- Average turnaround time

$$(7+10+4+11)/4 = 8$$

# Scheduling in Batch Systems: More Discussions

- Preemptive SJF is optimal – gives minimum average turnaround time for a given set of processes.
- Can use non-preemptive algorithms since this reduces process switching

# Goals for Today

- Scheduling
  - Issues in scheduling
  - Basic scheduling algorithms
    - Batch system
    - Interactive system
    - Real-time system

# Scheduling in Interactive Systems

- Scheduling goals
  - Response time: respond quickly to users' requests
  - Proportionality: meet users' expectations
- Algorithms:
  - Round-Robin Scheduling
  - Priority Scheduling
  - Multiple Queues



# Round-Robin Scheduling

- Round-robin
  - Each process is allowed to run specified time quantum.
  - After this time has elapsed, the process is preempted and added to the end of the ready queue.
- Quantum – a time interval, e.g., 100 msec
- Advantages:
  - Solution to fairness and starvation
  - Fair allocation of CPU across jobs
  - Low average waiting time when job lengths vary
  - Good for responsiveness if small number of jobs

# Round Robin Example

- Time Quantum = 20


Process	CPUTime
P1	53
P2	17
P3	68
P4	24

**33 left    13 left**

**48 left    28 left    8 left**

**4 left**

Gantt Chart for Schedule



**Typically, higher average turnaround time than SJF,  
but better response time**