



**STEVENS**  
INSTITUTE of TECHNOLOGY  
THE INNOVATION UNIVERSITY®

# CS 492: Operating Systems

## *Deadlocks (2)*

*Instructor: Iraklis Tsekourakis*


Email: [itsekour@stevens.edu](mailto:itsekour@stevens.edu)

# Recap


- Deadlock
  - Concepts
  - How to deal with deadlocks
    - Ignoring them: ostrich algorithm
    - Detecting & recovering from deadlock
    - Preventing deadlock
    - Avoiding deadlock

# Conditions for Resource Deadlocks

- Mutual exclusion
  - Only one thread at a time can use a resource
- Hold and wait
  - Thread holding at least one resource is waiting to acquire additional resources held by other threads
- No preemption
  - Resources are released only voluntarily by the thread holding the resource, after thread is finished with it
- Circular wait
  - There exists a set  $\{T_1, \dots, T_n\}$  of waiting threads
    - $T_1$  is waiting for a resource that is held by  $T_2$
    - $T_2$  is waiting for a resource that is held by  $T_3$
    - ...
    - $T_n$  is waiting for a resource that is held by  $T_1$



Is there a deadlock?



Is there a deadlock?

# Question

Suppose that there is a resource deadlock in a system. Give an example to show that the set of processes deadlocked can include processes that are not in the circular chain in the corresponding resource allocation graph.

# Solution

Consider three processes, A, B and C and three resources R, I and S. Suppose A is waiting for I that is held by B, B is waiting for S held by A, and C is waiting for R held by A. All three processes, A, B and C are deadlocked. However, only A and B belong to the circular chain.

# Goals for Today

- Deadlock
  - How to deal with deadlocks
    - Ignoring them: ostrich algorithm
    - Detecting & recovering from deadlock
    - Preventing deadlock
    - Avoiding deadlock



# Recap: Graph Analysis VS Deadlock Detection

- Facts:
  - No cycle in resource allocation graph  $\Rightarrow$  NO deadlock.
  - Cycle in resource allocation graph  $\Rightarrow$ 
    - if only one instance per resource type, then deadlock.
    - if several instances per resource type, possibility of deadlock.
- Detect deadlocks (for single instance per resource type)  
= detect the cycles in the resource allocation graph.

# Detection with One Resource of Each Type

- We need a formal algorithm for detecting deadlocks
  - i.e., to detect the cycles in the resource graph
- A simple one to detect cycles:
  - Use resource graph
  - Perform DFS (depth first search) starting at **each** node, looking for cycles (recurrence of starting node)
  - If it comes to a node it has encountered in this run, then there exists a cycle.

# Pseudocode for Deadlock Detection

```
func detectCycle()
  for node in graph:
    visited = bool[N]
    set all visited to false
    detectCycle(n, n, visited)

func detectCycle(n, origin, visited)
  for neighbour in graph[n]
    if neighbour == origin
      cycle detected
    if not visited[neighbour]
      visited[neighbour] = true
      detectCycle(neighbour, visited)
      visited[neighbour] = false
```

# Detection: Time Complexity

- $N$  nodes,  $E$  edges
- Time complexity for DFS is  $O(N + E)$
- Need to run DFS  $N$  times (once for each node)
  - $O(N(N + E))$
- Complexity
  - $O(N(N + E)) = O(N^2 + NE)$
  - $N \ll E$ , so  $N^2 \ll NE$ , total complexity is  $O(NE)$

# Recovery from Deadlock (1)

- Recovery through preemption
  - take a resource from some other process
  - depends on nature of the resource
- Recovery through rollback
  - checkpoint a process periodically
  - use this saved state
  - restart the process if it is found deadlocked

# Recovery from Deadlock (2)

- Recovery through killing processes
  - crudest but simplest way to break a deadlock
  - kill one of the processes in the deadlock cycle
  - the other processes get its resources
  - choose process that can be rerun from the beginning

# What's Next?

Question: How to detect deadlocks with multiple resources of each type?

# Detection with Multiple Resources of Each Type

- Basic idea: Allocate resources to a process that can be executed to completion
- Data structure

Resources in existence  
( $E_1, E_2, E_3, \dots, E_m$ )

Resources available  
( $A_1, A_2, A_3, \dots, A_m$ )

Current allocation matrix

Request matrix

$$\begin{bmatrix} C_{11} & C_{12} & C_{13} & \cdots & C_{1m} \\ C_{21} & C_{22} & C_{23} & \cdots & C_{2m} \\ \vdots & \vdots & \vdots & & \vdots \\ C_{n1} & C_{n2} & C_{n3} & \cdots & C_{nm} \end{bmatrix}$$

$$\begin{bmatrix} R_{11} & R_{12} & R_{13} & \cdots & R_{1m} \\ R_{21} & R_{22} & R_{23} & \cdots & R_{2m} \\ \vdots & \vdots & \vdots & & \vdots \\ R_{n1} & R_{n2} & R_{n3} & \cdots & R_{nm} \end{bmatrix}$$

Row n is current allocation  
to process n

Row 2 is what process 2 needs



# Detection with Multiple Resources of Each Type

Resources in existence  
( $E_1, E_2, E_3, \dots, E_m$ )

Resources available  
( $A_1, A_2, A_3, \dots, A_m$ )

Current allocation matrix

Request matrix

$$\begin{bmatrix} C_{11} & C_{12} & C_{13} & \cdots & C_{1m} \\ C_{21} & C_{22} & C_{23} & \cdots & C_{2m} \\ \vdots & \vdots & \vdots & \cdots & \vdots \\ C_{n1} & C_{n2} & C_{n3} & \cdots & C_{nm} \end{bmatrix}$$

$$\begin{bmatrix} R_{11} & R_{12} & R_{13} & \cdots & R_{1m} \\ R_{21} & R_{22} & R_{23} & \cdots & R_{2m} \\ \vdots & \vdots & \vdots & \cdots & \vdots \\ R_{n1} & R_{n2} & R_{n3} & \cdots & R_{nm} \end{bmatrix}$$

Row n is current allocation  
to process n

Row 2 is what process 2 needs

# Detection with Multiple Resources of Each Type: Algorithm

- Deadlock detection is based on comparing vectors
- Algorithm
  - Look for an unmarked process,  $P_i$ , for which  $R[i]$  (i.e. the  $i$ -th row of  $R$ )  $\leq A$  (i.e.,  $P_i$  's request can be granted by available resources)
  - If such a process is found, update  $A = A + C[i]$  (i.e. add the  $i$ -th row of  $C$  to  $A$ ), mark the process and go back to step 1 (i.e.,  $P_i$  releases its resources)
  - If no such process exists, the algorithm terminates

# Output of Algorithm

- If the algorithm outputs a sequence of processes, then no deadlock!
- Otherwise, deadlock exists

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ;
- Three resource types: A (7 instances), B (2 instances), and C (6 instances). Existence matrix:  $E = (7, 2, 6)$
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>C</u>	<u>Request</u>	<u>R</u>	<u>Available</u>	<u>A</u>
	<i>ABC</i>		<i>ABC</i>		<i>ABC</i>	
$P_0$	010		000		0 0 0	$P_0$ gets resource,
$P_1$	200		202			
$P_2$	303		000			
$P_3$	211		100			
$P_4$	002		002			


# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ;
- Three resource types:  $A$  (7 instances),  $B$  (2 instances), and  $C$  (6 instances).  $E = (7, 2, 6)$

	<u>Allocation</u>	<u>C</u>	<u>Request</u>	<u>R</u>	<u>Available</u>	<u>A</u>
	<i>ABC</i>		<i>ABC</i>		<i>ABC</i>	
$P_0$	010		000		0 1 0	P <sub>0</sub> finishes and releases resources
$P_1$	200		202			
$P_2$	303		000			P <sub>2</sub> gets resource,
$P_3$	211		100			
$P_4$	002		002			

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ;
- Three resource types: A (7 instances), B (2 instances), and C (6 instances).  $E = (7, 2, 6)$

	<u>Allocation</u>	<u>C</u>	<u>Request</u>	<u>R</u>	<u>Available</u>	<u>A</u>	
	ABC		ABC		ABC		
$P_0$	0	1	0	0	0	1	  P <sub>2</sub> finishes and releases resources P <sub>3</sub> gets resource
$P_1$	2	0	0	2			
$P_2$	3	0	3	0	3	1	
$P_3$	2	1	1	0			
$P_4$	0	0	2	0			

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ;
- Three resource types:  $A$  (7 instances),  $B$  (2 instances), and  $C$  (6 instances).  $E = (7, 2, 6)$

	<u>Allocation</u>	<u>C Request</u>	<u>R Available</u>	<u>A</u>
	<i>ABC</i>	<i>ABC</i>	<i>ABC</i>	
$P_0$	010	000		
$P_1$	200	202		✓ P <sub>1</sub> gets resource
$P_2$	303	000	3 1 3	✓ P <sub>3</sub> finishes and releases resources
$P_3$	211	100	5 2 4	
$P_4$	002	002		

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ;
- Three resource types:  $A$  (7 instances),  $B$  (2 instances), and  $C$  (6 instances).  $E = (7, 2, 6)$

	<u>Allocation</u> <u><math>C</math></u>	<u>Request</u> <u><math>R</math></u>	<u>Available</u> <u><math>A</math></u>	
	$ABC$	$ABC$	$ABC$	
$P_0$	010	000		✓
$P_1$	200	202	7 2 4	$P_1$ finishes and releases resources
$P_2$	303	000		✓
$P_3$	211	100	5 2 4	✓
$P_4$	002	002		$P_4$ gets resource



# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ;
- Three resource types:  $A$  (7 instances),  $B$  (2 instances), and  $C$  (6 instances).  $E = (7, 2, 6)$

	<u>Allocation</u>	<u>C</u>	<u>Request</u>	<u>R</u>	<u>Available</u>	<u>A</u>
	<i>ABC</i>		<i>ABC</i>		<i>ABC</i>	
$P_0$	010		000			✓
$P_1$	200		202		7 2 4	✓
$P_2$	303		000			✓
$P_3$	211		100			✓
$P_4$	002		002		7 2 6	P <sub>4</sub> finishes and releases resources

Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in *all* processes finished  
No deadlock!

# Example (Cont.)

- $P_2$  requests an additional instance of type  $C$ .

	<u>Allocation</u> <u><math>C</math></u>	<u>Request</u> <u><math>R</math></u>	<u>Available</u> <u><math>A</math></u>
	$ABC$	$ABC$	$ABC$
$P_0$	010	000	0 0 0
$P_1$	200	202	
$P_2$	303	00 1	
$P_3$	211	100	
$P_4$	002	002	

# Example (Cont.)

- $P_2$  requests an additional instance of type C.

	<u>Allocation</u>	<u>C</u>	<u>Request</u>	<u>R</u>	<u>Available</u>	<u>A</u>
	ABC		ABC		ABC	
$P_0$	010		000		0 1 0	
$P_1$	200		202			
$P_2$	303		00 1			
$P_3$	211		100			
$P_4$	002		002			

- State of system?
  - Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes requests.
  - Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$ .

# Detection: How often to run?

- Possibilities:
  - After every allocation of resource
    - Too expensive
  - Every few minutes
  - When CPU is idle – indication of blocked processes waiting for resources

# Detection: How often to run? (Cont.)

- When, and how often, to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - One for each disjoint cycle