



**STEVENS**  
INSTITUTE of TECHNOLOGY  
THE INNOVATION UNIVERSITY®

# CS 492: Operating Systems

## *Deadlocks (3)*

*Instructor: Iraklis Tsekourakis*

Email: [itsekour@stevens.edu](mailto:itsekour@stevens.edu)

# Goals for Today

- Deadlock
  - How to deal with deadlocks
    - Ignoring them: ostrich algorithm
    - Detecting & recovering from deadlock
    - Preventing deadlock
    - Avoiding deadlock

# Four Necessary Conditions for Deadlock

1. Mutual exclusion condition
    - Each resource assigned to at most 1 process
  2. Hold and wait condition
    - Process holding resources can request for additional resources
  3. No preemption condition
    - Assigned resources cannot be claimed involuntarily
  4. Circular wait condition
    - Must be a circular chain of 2 or more processes
- ❖ All four conditions must be present when a deadlock occurs
  - ❖ If one condition is absent, the deadlock is impossible to happen

# Deadlock Prevention

## **Principle:**


If any one of the four conditions for deadlock (with reusable resources) is denied, deadlock is impossible.

# Deadlock Prevention

## 1. Attacking the Mutual Exclusion Condition

- Principle:
  - Allow resources to be shared, if possible
- Not all devices can be spooled
  - However, the mutual-exclusion condition must hold for non-sharable resources

# Example of Deadlock: Dining Lawyers Problem




- Five chopsticks/Five lawyers (really cheap restaurant)
  - Free-for all: Lawyer will grab any one they can
  - Need two chopsticks to eat
  - May cause deadlock!
- Can we attack the mutual exclusion condition? If so, how?

# Deadlock Prevention

## 2. Attacking the Hold and Wait Condition

- Goal: Prevent processes that hold resources from waiting for more resources

# Breaking Deadlocks for Dining Lawyers Problem



- Can we attack the hold and wait condition? If so, how?



# Deadlock Prevention

## 3. Attacking the No Preemption Condition

- If a process is holding some resources and request another resource that is not available, then all resources the process is currently holding are preempted (released).
- Applied to resources whose state can be easily saved and restored later
  - CPU registers, memory space, etc.

# Breaking Deadlocks for Dining Lawyers Problem




- Can we attack the no Preemption condition? If so, how?

# Deadlock Prevention

## 4. Attacking the Circular Wait Condition

1. Imagesetter
2. Scanner
3. Plotter
4. Tape drive
5. CD Rom drive

(a)



(b)


- Normally ordered resources
- A resource graph

# Deadlock Prevention

## 4. Attacking the Circular Wait Condition

- Rule: All requests of a process must be made in numerical order
- Problem: impossible to find an ordering to satisfy everyone

# Breaking Deadlocks for Dining Lawyers Problem



- Can we attack the circular wait condition? If so, how?

# Summary of Deadlock Prevention

Condition	Approach
Mutual Exclusion	Spool everything
Hold and wait	Request all resources initially
No preemption	Take resources away
Circular wait	Order resources numerically

# Next..

- Deadlock
  - Concepts
  - How to deal with deadlocks
    - Ignoring them: ostrich algorithm
    - Detecting & recovering from deadlock
    - Preventing deadlock
    - Avoiding deadlock

# Deadlock - What to do?

- Different strategies
  - Detection/Recovery
    - Clean up only when trouble really happens
  - Prevention
    - Break one of four necessary conditions
  - Avoidance
    - Processes *pre-declare usage patterns*
    - Request manager avoids “unsafe states”



# Assumptions

- Processes pre-declare usage
  - declare *maximal resource usage*
    - E.g., P1 will never need more than 7 tape drives and 1 printer
- Processes proceed to completion
  - Don't hold onto resources forever
  - Complete in reasonable time
    - Worst-case: ok to stall P2 until P1 finishes

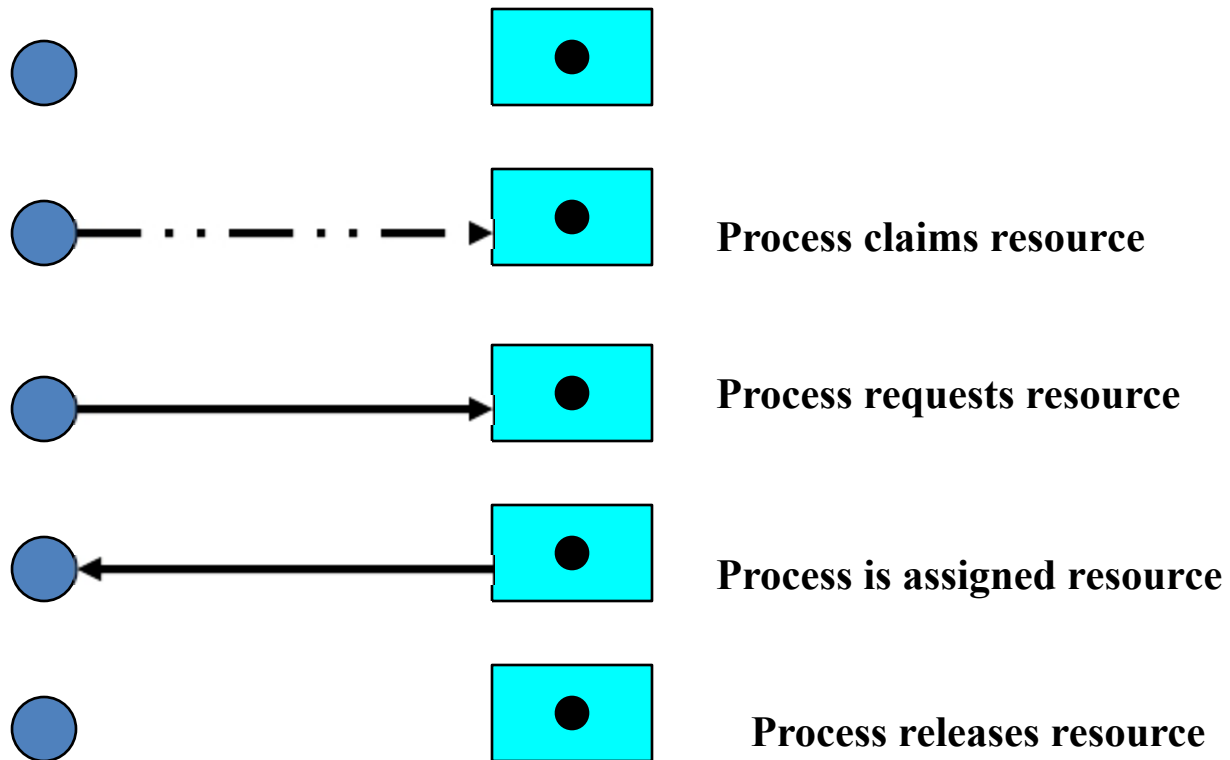
# Deadlock Avoidance Algorithm

- Basic idea
  - Single instance for each resource: resource allocation graph
  - Multiple instances for each resource: Banker's algorithm

# Resource Allocation Graph

- Used for deadlock avoidance when there is *only one instance* of each resource type.
- Three kinds of edges:
  - **Claim** edge  $P_i \dashrightarrow R_j$  indicates that process  $P_i$  *may* request resource  $R_j$
  - **Request** edge:  $P_i \rightarrow R_j$
  - **Assignment** edge:  $R_j \rightarrow P_i$


# Resource Allocation Graph



# Resource Allocation Graph (Cont.)

- Resources must be claimed a priori in the system
  - Add claim edges at beginning.
- When a process requests a resource
  - The claim edge converts to the *request* edge (change dotted line to solid line).
  - Check cycles in the graph
    - If resource assignment does not result in the formation of a cycle in the resource allocation graph, safe state.
    - Otherwise, unsafe state.

# Resource Allocation Graph



# Deadlock Avoidance Algorithm

- Basic idea
  - Single instance for each resource: resource allocation graph
  - Multiple instances for each resource: Banker's algorithm (can also be applied for single instance per resource type)

# Safe Execution Sequence

- $P_1, P_2, P_3, \dots P_n$  is a *safe sequence* if
  - Every process  $P_i$  can be satisfied using
    - currently-free resources; PLUS
    - resources held by  $P_1, P_2, \dots P_{i-1}$
- Bound  $P_i$ 's waiting
  - $P_1$  will run to completion, release resources
  - $P_2$  can complete with now-free +  $P_1$ 's +  $P_2$ 's
  - $P_3$  can complete with now-free +  $P_1$ 's +  $P_2$ 's +  $P_3$ 's
- $P_i$  won't wait forever, no wait cycle, no deadlock



# Safe State

- System in a *safe state* if
  - there exists one safe sequence
- Worst-case behavior
  - Everybody asks for everything at once
  - Follow the safe sequence
- Serial execution is *worst-case*, not typical
  - Usually execute in parallel

# Safe and Unsafe States (1)

Has Max			Has Max			Has Max		
A	3	9	A	3	9	A	3	9
B	2	4	B	4	4	B	0	—
C	2	7	C	2	7	C	2	7
Free: 3			Free: 1			Free: 5		
(a)			(b)			(c)		
Has Max			Has Max					
A	3	9	A	3	9			
B	0	—	B	0	—			
C	7	7	C	0	—			
Free: 0			Free: 7					
(d)			(e)					

Demonstration that the state in (a) is safe.

# Safe and Unsafe States (2)

Has Max		
A	3	9
B	2	4
C	2	7

Free: 3  
(a)

Has Max		
A	4	9
B	2	4
C	2	7

Free: 2  
(b)

Has Max		
A	4	9
B	4	4
C	2	7

Free: 0  
(c)


Has Max		
A	4	9
B	—	—
C	2	7

Free: 4  
(d)

Demonstration that the state in (b) is not safe.

# Basic Facts

- If a system is in a safe state  $\Rightarrow$  no deadlocks.
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock.
- Avoidance  $\Rightarrow$  ensure that a system will never reach an unsafe state.



# Avoidance - Key Ideas

- Safe state
  - Some safe sequence exists
  - Prove it by *finding one*
- Unsafe state: No safe sequence exists

# Deadlock Avoidance

## *Multiple Instances* of resource type

- Banker's algorithm
- Similar to deadlock detection algorithm
  - Consider each request as it occurs and see if granting it leads to a safe state
  - If a resource granting leads to a deadlock, then the state is unsafe and does not grant
- Key assumptions:
  - All processes need to declare their max resources
  - # of processes and resources are fixed.

# Banker's Algorithm for Multiple Resources (1)

1. Look for a row,  $R$ , whose unmet resource needs are all smaller than or equal to  $A$ . If no such row exists, system will eventually deadlock.
2. Assume the process of row chosen requests all resources needed and finishes. Mark that process as terminated, add its resources to the  $A$  vector.
3. Repeat steps 1 and 2 until either all processes are marked terminated (safe state) or no process is left whose resource needs can be met (deadlock)

# Banker's Algorithm for Multiple Resources (2)

	Process	Tape drives	Plotters	Printers	CD ROMs
A	3	0	1	1	
B	0	1	0	0	
C	1	1	1	0	
D	1	1	0	1	
E	0	0	0	0	

Resources assigned

	Process	Tape drives	Plotters	Printers	CD ROMs
A	1	1	0	0	
B	0	1	1	2	
C	3	1	0	0	
D	0	0	1	0	
E	2	1	1	0	

Resources still needed

$E = (6342)$   
 $P = (5322)$   
 $A = (1020)$

The banker's algorithm with multiple resources.



```
void process_A(void) {  
    enter_region(&resource_1);  
    enter_region(&resource_2);  
    use_both_resources( );  
    leave_region(&resource_2);  
    leave_region(&resource_1);  
}
```

# Livelock

Figure 6-16. Busy waiting  
that can lead to livelock.

```
void process_B(void) {  
    enter_region(&resource_2);  
    enter_region(&resource_1);  
    use_both_resources( );  
    leave_region(&resource_1);  
    leave_region(&resource_2);  
}
```