

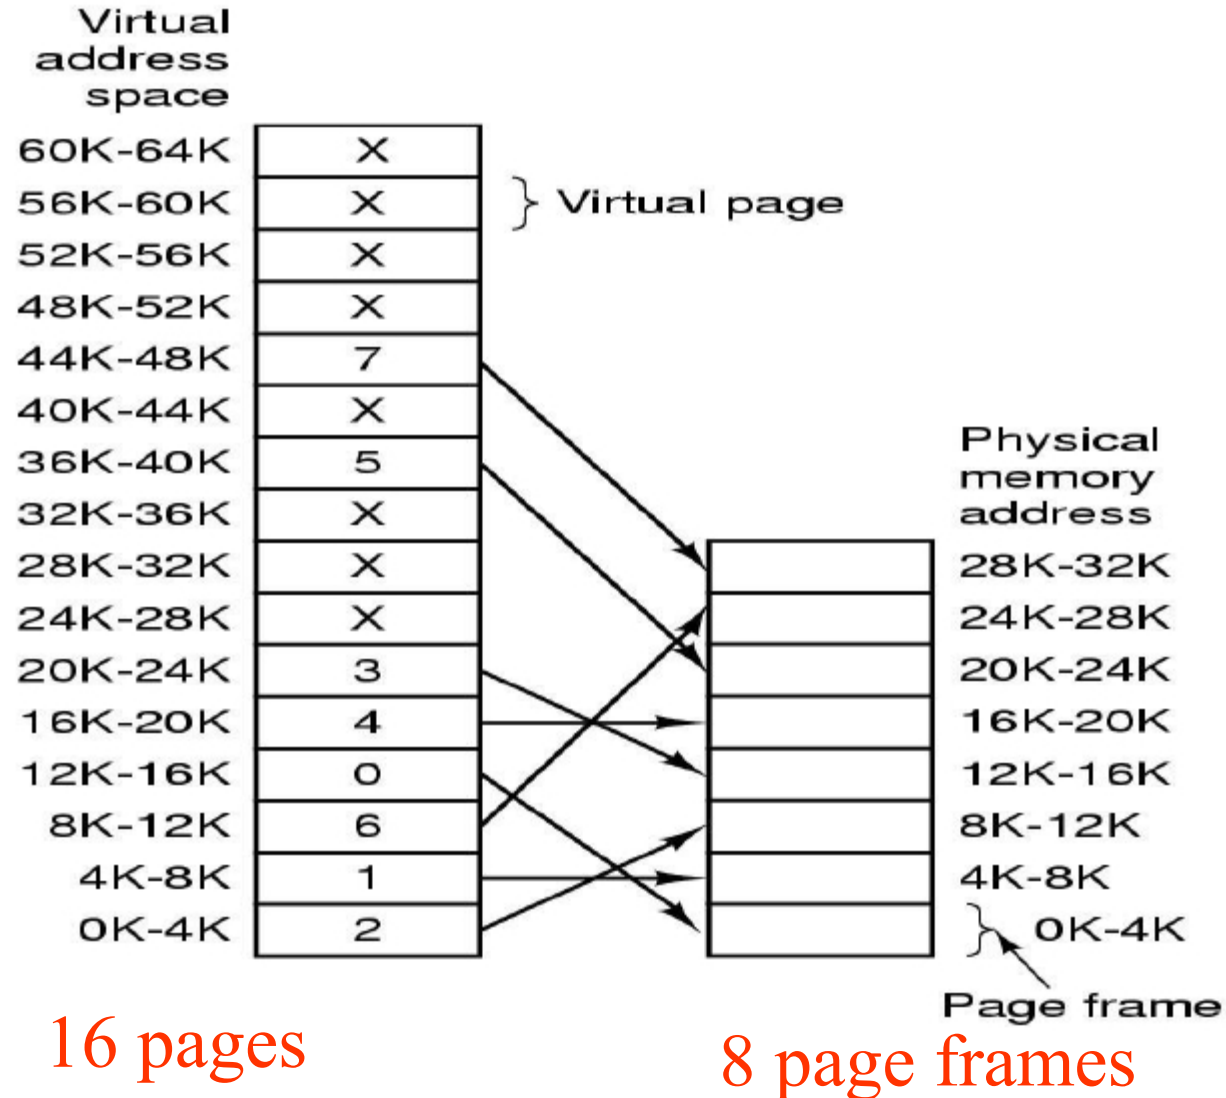
CS 492: Operating Systems

Page Replacement

Instructor: Iraklis Tsekourakis

Email: itsekour@stevens.edu

Virtual Addresses V.S. Physical Memory



Paging and Swapping

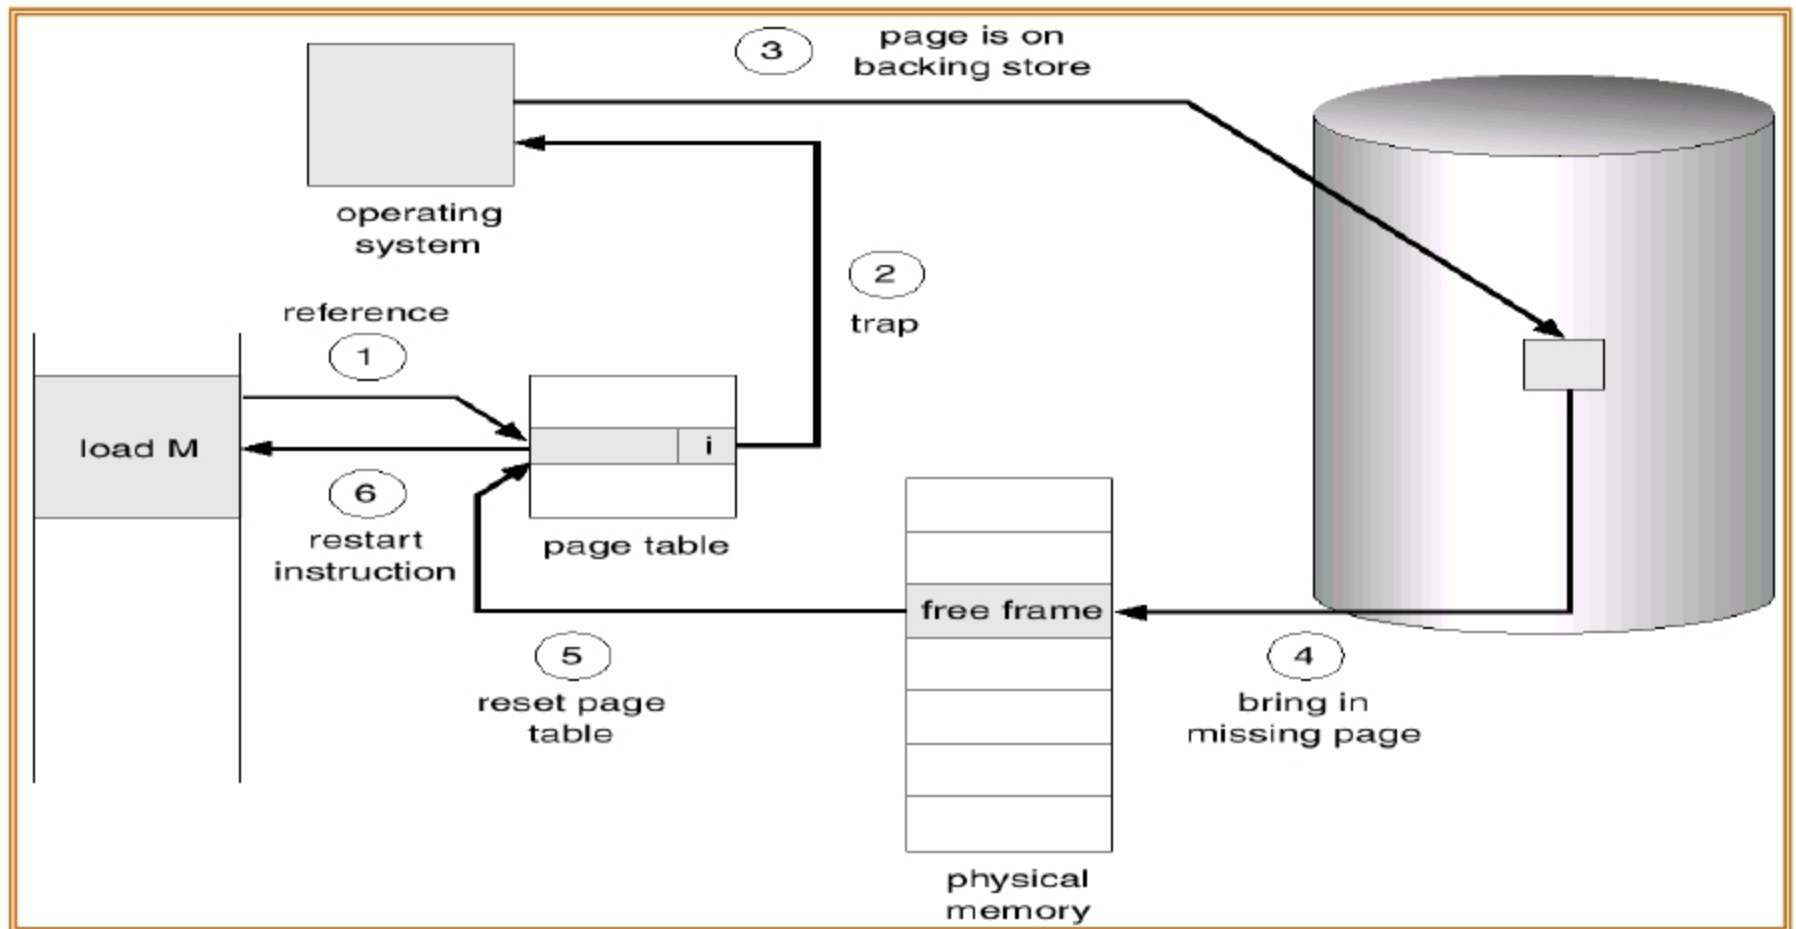
- On heavily-loaded systems, memory can fill up
- To achieve good performance, must move “inactive” pages out to disk
- Two solutions:
 1. Swapping
 - Usually refers to moving the memory for an entire process out to disk
 2. Paging
 - Refers to moving individual pages out to disk (and back)

Outline

- How to move “inactive” pages out to disk to make space?
- Two solutions:
 1. Swapping
 2. Paging
 - Refers to moving individual pages out to disk (and back)

Page Fault

- If there is ever a reference to a page that is not in memory \Rightarrow **page fault**



Theoretical Foundation: Performance of Page Faults

- Page fault rate p : $0 \leq p \leq 1$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault

- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p \times (\text{page fault overhead} \\ & \quad + \text{swap page out overhead} \\ & \quad + \text{swap page in overhead} \\ & \quad + \text{restart overhead}) \end{aligned}$$

Example

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $EAT = (1 - p) \times 200 + p (8 \text{ milliseconds})$
 $= (1 - p) \times 200 + p \times 8,000,000$
 $= 200 + p \times 7,999,800$
- If one access out of 1,000 causes a page fault (i.e. $p=1/1000$), then

$$EAT = 8.2 \text{ microseconds.}$$

This is a slowdown by a factor of 40 compared with memory access time (200 nanoseconds)!!!

Our goal: keep page fault as infrequent as possible!

What happens if there is no free frame in page table?

- Page replacement – find some page in memory and swap it out
- Page replacement forces choice
 - which page(s) must be removed

Page Replacement

- Question: which page to evict when memory is full and a new page is demanded?
- Goal: reduce the number of page faults

Page Replacement Algorithms

- Goal: **lowest** page-fault rate.
- *Input of algorithm*: a particular string of memory references (reference string)
 - Assume reference string in examples to follow is
(page) 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5.
- *Output of algorithm*: the number of page faults on that string.

Algorithm 1: Optimal (OPT) algorithm

- What's the best we can possibly do?
 - Assume OS knows the perfect knowledge of the future
- Algorithm: replace the page that will be used *furthest* in the future
- Estimate by ...
 - logging page use on previous runs of process
 - although this is impractical

OPT Algorithm: Example

Assume there are three page frames in the memory.
How many page faults will there be for the following
memory reference string.?

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

OPT Algorithm: Example

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2			2							7		
	0	0	0		0		4			0				0			0		
		1	1		3		3			3				1			1		

page frames

9 page faults

Infeasibility of Optimal Algorithm

- Nice, but not achievable in real systems!
 - Only works if we know the whole sequence of page references!
 - Not realizable in practice (usually)
- However,
 - Can be approximated by running the program twice
 - Once to generate the reference trace
 - Once (or more) to apply the optimal algorithm