



STEVENS
INSTITUTE of TECHNOLOGY
THE INNOVATION UNIVERSITY®

CS 492: Operating Systems

Page Replacement (2)

Instructor: Iraklis Tsekourakis

Email: itsekour@stevens.edu



Question?

- A machine has 48-bit virtual addresses and 32-bit physical addresses. Pages are 8 KB. How many entries are needed for a single-level linear page table?

Algorithm 2: First In First Out (FIFO)

- FIFO: First in First Out
- Maintain a linked list of all pages
 - In the order that they came into memory
- Page at the beginning of list is replaced

FIFO Example

- Assume 3 frames
- Reference string: 0, 1, 2, 3, 0, 1, 4, 0, 1, 2, 3, 4

3 frames

0	1	2	3	0	1	4	0	1	2	3	4	Tail of List
0	1	2	3	0	1	4	4	4	2	3	3	
	0	1	2	3	0	1	1	1	4	2	2	
		0	1	2	3	0	0	0	1	4	4	
P	P	P	P	P	P	P			P	P		Beginning of List

9 Page faults

Discussions: FIFO

- Advantages
 - Simple to implement
- Disadvantages
 - Page in memory the longest may be often replaced
 - FIFO suffers from “Belady’s Anomaly”
 - The rate fault might increase when the algorithm is given more memory.

FIFO & Belady's Anomaly

- Given a reference string: 0, 1, 2, 3, 0, 1, 4, 0, 1, 2, 3, 4
- We have shown that there are 9 page faults for the given reference string when there are 3 page frames in memory?

0	1	2	3	0	1	4	0	1	2	3	4
0	1	2	3	0	1	4	4	4	2	3	3
	0	1	2	3	0	1	1	1	4	2	2
		0	1	2	3	0	0	0	1	4	4
P	P	P	P	P	P	P			P	P	

- How many page faults are there if now it has **4 page frames** in memory?

FIFO & Belady's Anomaly

All pages frames initially empty

		0	1	2	3	0	1	4	0	1	2	3	4
Youngest page		0	1	2	3	0	1	4	4	4	2	3	3
			0	1	2	3	0	1	1	1	4	2	2
Oldest page				0	1	2	3	0	0	0	1	4	4
		P	P	P	P	P	P	P			P	P	

9 Page faults

(a)

		0	1	2	3	0	1	4	0	1	2	3	4
Youngest page		0	1	2	3	3	3	4	0	1	2	3	4
			0	1	2	2	2	3	4	0	1	2	3
Oldest page				0	1	1	1	2	3	4	0	1	2
					0	0	0	1	2	3	4	0	1
		P	P	P	P			P	P	P	P	P	P

10 Page faults

(b)


- Belady's Anomaly: more frames \Rightarrow more page faults

FIFO Variant (1): Second Chance


- When free page frames are needed, examine pages in FIFO order starting from beginning
 - Consider “reference bit” R
 - Reference: the page is either read or written
 - If $R=0$, replace page
 - If $R=1$, set $R=0$ and place at the end of FIFO list (hence, the second chance)
 - If not enough replaces on first pass, revert to pure FIFO on second pass

Second Chance Example

Page loaded first




(a)



(b)

FIFO Variant (2): Clock

- Problem of the *second chance* mechanism: moving pages around on list is not efficient
- Clock: variant of FIFO, better implementation of Second Chance
- Details




When a page fault occurs, the page the hand is pointing to is inspected. The action taken depends on the R bit:


R = 0: Evict the page

R = 1: Clear R and advance hand

Clock Algorithm: Example



(a) State of buffer just prior to a page replacement



(b) State of buffer just after the next page replacement

“Use” = “Reference”

Some Observations

- Locality of reference – next reference is more likely to be near the last one
- This motivates the algorithm named Least Recently Used (LRU).

Algorithm 3: Least Recently Used (LRU)

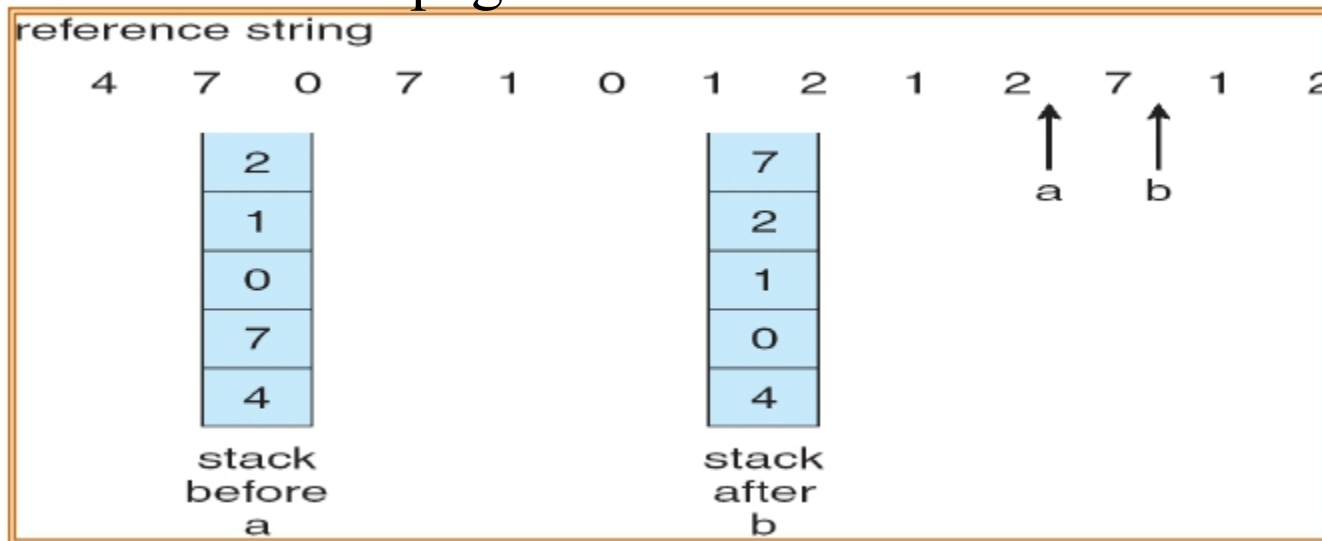
- Make use of locality property: pages used recently will be used again soon
 - Throw out page that has been unused for longest time
- Example: Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

- Optimal VS LRU
 - Optimal: Throw out pages the furthest in the *future*
 - LRU: Throw out pages the furthest in the *past*

How to Implement LRU?

- Stack solution:
 - keep a stack of page numbers in a double link form:
 - Most recently used at front, least at rear
 - Always remove the page at rear
 - Referenced page is moved to the front



Implementing LRU using Stack

- Update this list at every memory reference!
 - Update includes: find a page in the list, move it at front
 - Time consuming!!!
 - So we need other ways to implement LRU

Implementing LRU using Counter

- Counter implementation
 - Equip the hardware with a 64-bit counter
 - At every instruction increment the counter
 - Every page table entry must also have a field to save the counter. When a page is reference copy the counter to that field
 - The more recent the page is used, the larger the counter value is
 - For replacement, evict page with the **lowest** counter value (i.e., the oldest page)

Implementing LRU using Matrix

- Matrix implementation – n pages, a matrix of $n \times n$ bits
- Initially all bits are 0
- When *page* k is referenced,
 - All the bits of row k are set to 1 (increase the binary value of page k)
 - All the bits of column k are set to 0 (decrease the binary value of other pages)
- Page replacement: replace page k , if row k corresponds to lowest binary value

	Page			
	0	1	2	3
0	0	1	1	1
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0

(a)

	Page			
	0	1	2	3
0	0	0	1	1
1	1	0	1	1
2	0	0	0	0
3	0	0	0	0

(b)

	Page			
	0	1	2	3
0	0	0	0	1
1	1	0	0	1
2	1	1	0	1
3	0	0	0	0

(c)

	Page			
	0	1	2	3
0	0	0	0	0
1	1	0	0	0
2	1	1	0	0
3	1	1	1	0

(d)

	Page			
	0	1	2	3
0	0	0	0	0
1	1	0	0	0
2	1	1	0	1
3	1	1	0	0

(e)

	0	0	0	0
1	1	0	1	1
2	1	0	0	1
3	1	0	0	0

(f)

	0	1	1	1
0	0	0	1	1
1	0	0	0	1
2	0	0	0	0

(g)

	0	1	1	0
0	0	0	1	0
1	0	0	0	0
2	1	1	1	0

(h)

	0	1	0	0
0	0	0	0	0
1	1	1	0	1
2	1	1	0	0

(i)

	0	1	0	0
0	0	0	0	0
1	1	1	0	0
2	1	1	1	0

(j)

pages referenced in order 0,1,2,3,2,1,0,3,2,3

Implementing LRU – Not Frequently Used (NFU)

- Most machines do not have the hardware to perform true LRU, but it may be simulated.
- At each clock interrupt, for each page, the R bit (valued 0 or 1) is added to its software counter.
- Page with the lowest R-bit is evicted.
 - What is the problem with this type of history keeping?

Solution: Aging

- 1st step: shift counter to the right
 - Reduce counter values over time by dividing it by two
- 2nd step: add R to the leftmost bit
 - Recent R bit is added as most significant bit, therefore more weight given to more recent references!
- Page replacement: replace page with the lowest counter

Aging Example

	R bits for pages 0-5, clock tick 0	R bits for pages 0-5, clock tick 1	R bits for pages 0-5, clock tick 2	R bits for pages 0-5, clock tick 3	R bits for pages 0-5, clock tick 4
	1 0 1 0 1 1	1 1 0 0 1 0	1 1 0 1 0 1	1 0 0 0 1 0	0 1 1 0 0 0
Page					
0	10000000	11000000	11100000	11110000	01111000
1	00000000	10000000	11000000	01100000	10110000
2	10000000	01000000	00100000	00100000	10001000
3	00000000	00000000	10000000	01000000	00100000
4	10000000	11000000	01100000	10110000	01011000
5	10000000	01000000	10100000	01010000	00101000
	(a)	(b)	(c)	(d)	(e)

- Disadvantage
 - With 8-bit counter, have only a memory of 8 clock ticks back

Question

- A small computer on a smart card has four page frames. At the first clock tick, the R bits are 0111 (page 0 is 0, the rest are 1). At subsequent clock ticks, the values are 1011, 1010, 1101, 0010, 1010, 1100, and 0001. If the aging algorithm is used with an 8-bit counter, give the values of the four counters after the last tick.

Not Recently Used (NRU) Algorithm

- Each page has a Reference bit and a Modified bit
 - Reference bit: the page is accessed (read or written).
 - Modified bit: the page is written.
 - bits are set when page is referenced/modified
- Pages are classified into 4 types
 - Class 0: not referenced, not modified
 - Class 1: not referenced, modified (**Not possible**)
 - Class 2: referenced, not modified
 - Class 3: referenced, modified
- At a page fault, NRU removes a page from the lowest numbered non-empty class

Discussion: NRU Algorithm

- Easy to understand and efficient to implement
- However, pages that are frequently referenced might be removed