



STEVENS
INSTITUTE of TECHNOLOGY
THE INNOVATION UNIVERSITY®

CS 492: Operating Systems

File Systems (2)

Instructor: Iraklis Tsekourakis

Email: itsekour@stevens.edu



Allocation Methods

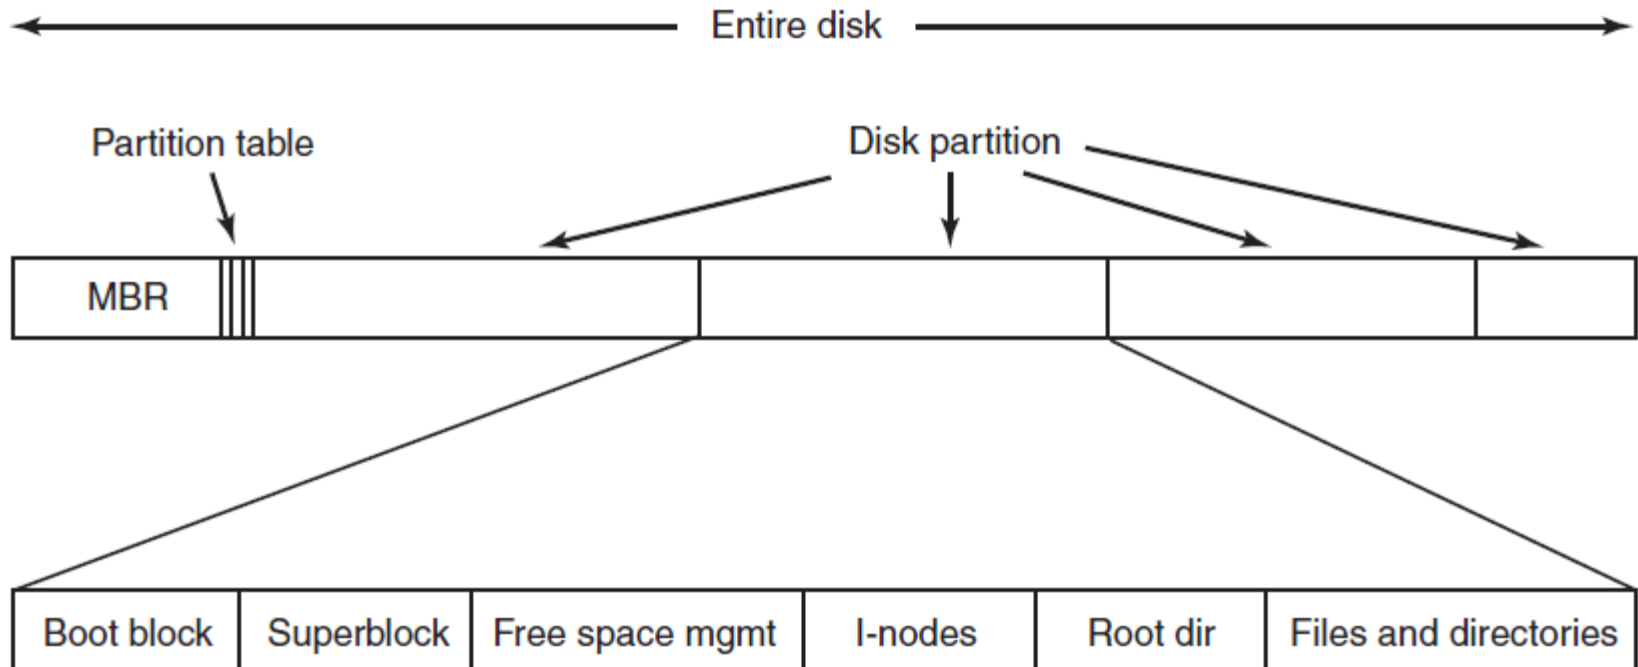
Three allocation methods:

1. Contiguous allocation

2. Linked allocation


3. Indexed allocation

File System Layout



A possible file system layout.

Contiguous Allocation



- Each file occupies a set of contiguous blocks on the disk

Pros:


- Simple – only starting location (block #) and length (number of blocks) are required
- Random access

Cons:

- Wasteful of space (dynamic storage-allocation problem)
 - External fragmentation: may need to compact space
- Files cannot grow

Implementing Files

Contiguous Layout



(a) Contiguous allocation of disk space for seven files. (b) The state of the disk after files *D* and *F* have been removed.

Contiguous Allocation (cont'd)

- Question: Given a logical address LA of file A , how to map LA to its physical address (B,D) (B : block number; D : block offset)?
- Suppose the block size is 512 bytes:


$$\begin{array}{l} \nearrow \text{Quotient } Q \\ LA/512 \\ \searrow \text{Remainder } R \end{array}$$

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2


directory

- *Block number $B = Q + \text{starting address of file } A \text{ in directory}$*
- *Block offset $D = R$*
- *Number of memory access to get the data at address LA :*
 - *1 access for reading directory to get starting address of file A*
 - *1 access for reading data from block B at offset R .*

Linked Allocation



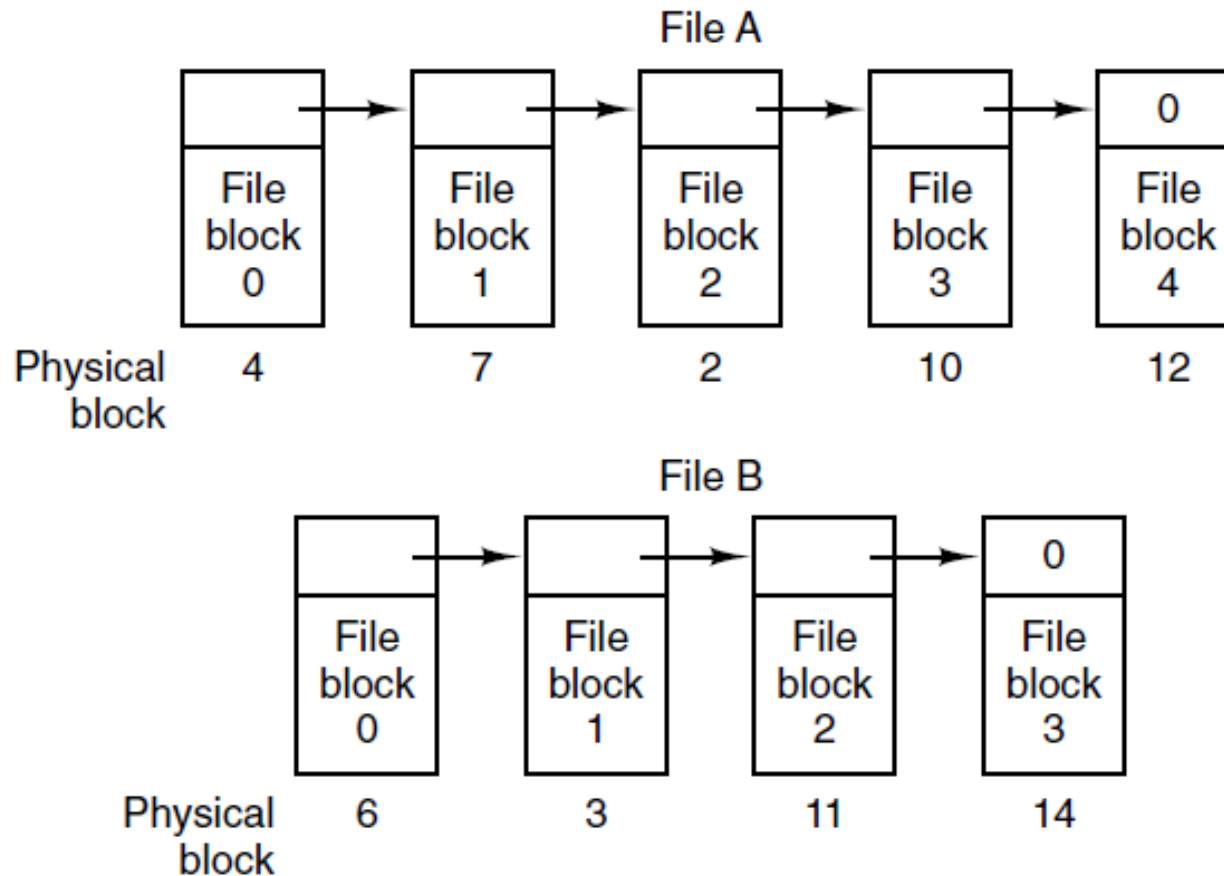
- Each file is a linked list of disk blocks
- Blocks may be scattered anywhere on the disk



- Pros:
 - Simple – need only starting address
 - Free-space management system – no waste of space
- Cons:
 - No efficient random access

Implementing Files

Linked List Allocation




Storing a file as a linked list of disk blocks.


Linked Allocation (Cont.)

- Question: given a logical address LA of file A , how to map LA to its physical address (B,D) (B : block number; D : block offset)?
- Suppose block size is 512 bytes and each block contains 4 bytes reserved for pointer to next block:


$$LA/(512-4) \begin{cases} \text{Quotient } Q \\ \text{Remainder } R \end{cases}$$



- *Block number $B = Q$ th block in the linked chain of blocks, starting from the start block in directory*
- *Block offset $D = R + 4$*
- *Number of memory access to get the data at address LA :*
 - 1 access for reading directory to get starting block of file A
 - Q accesses for traversing Q blocks



File-Allocation Table (FAT)

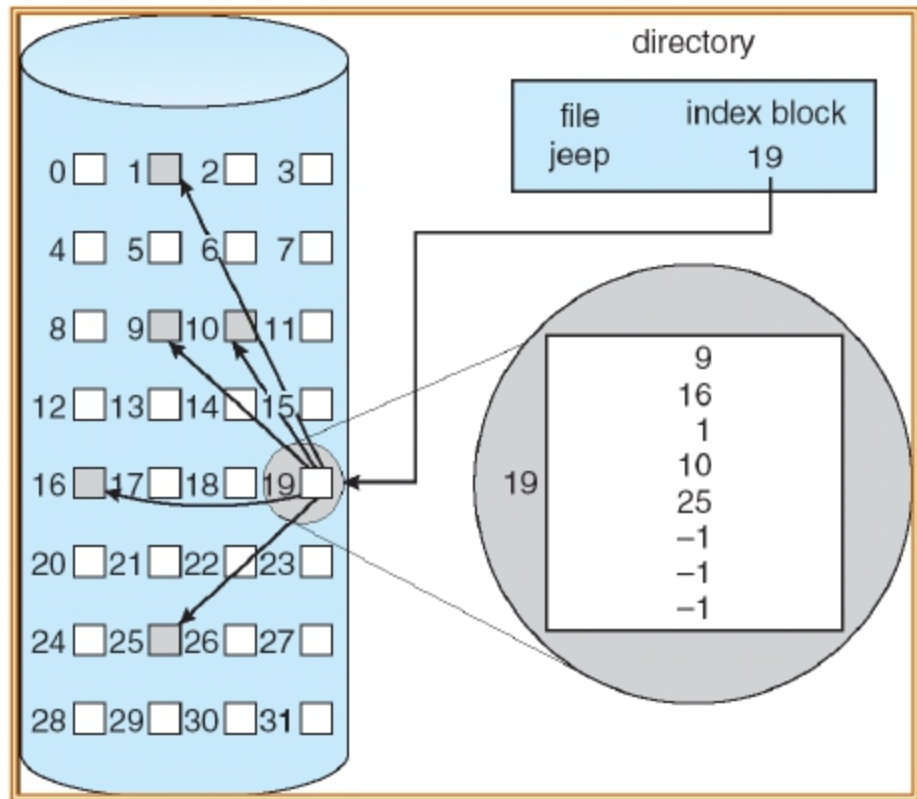


Variant of linked list allocation

Each FAT entry corresponds to disk block number

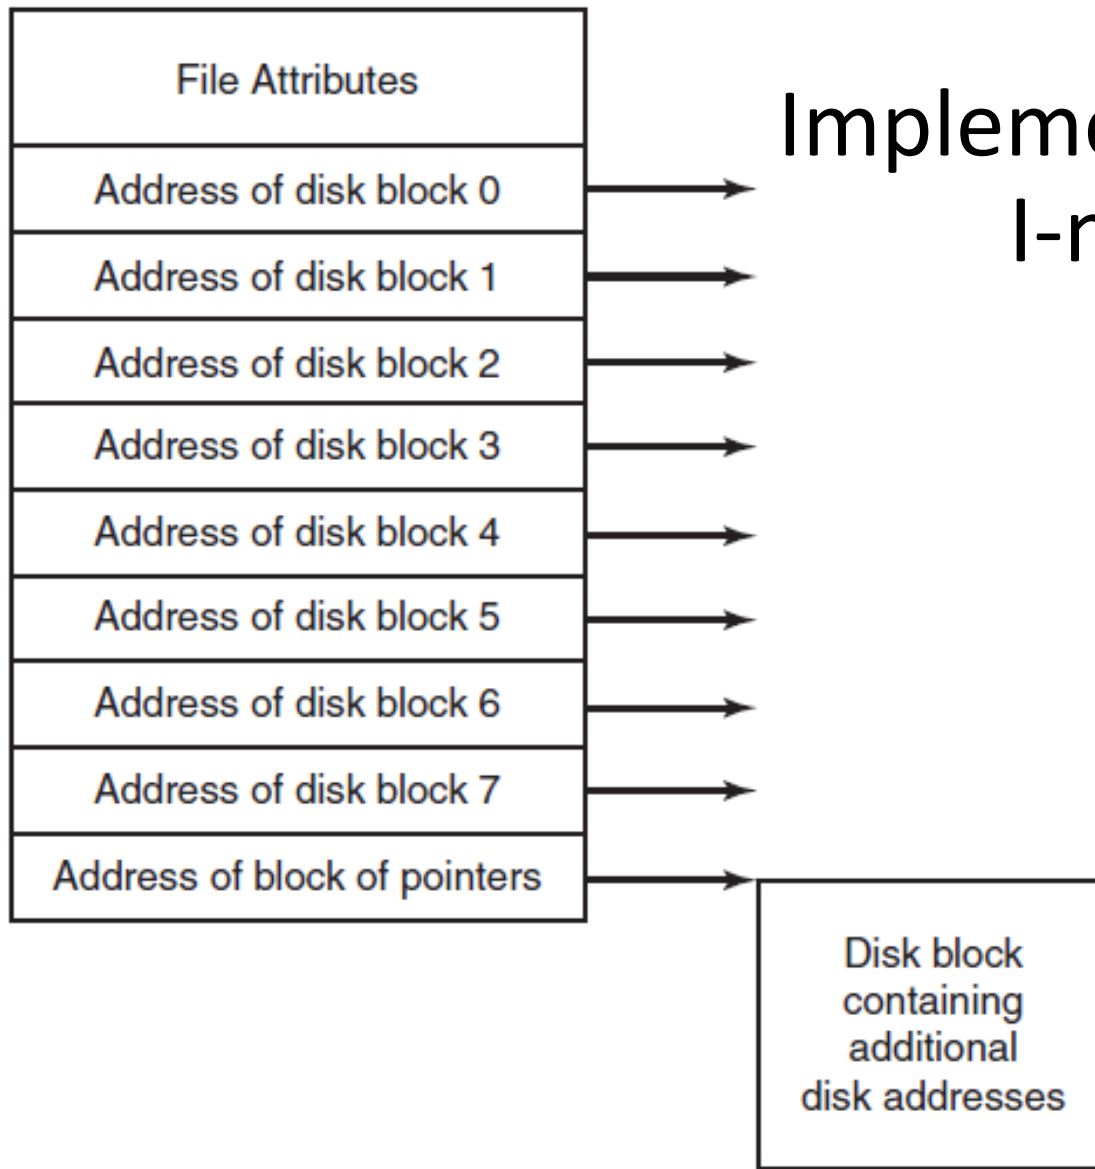
Each FAT entry contains a pointer to the next block or 0

Indexed Allocation



- Brings all pointers together into the **index block** (also called i-node)
- Pros:
 - Efficient random access
 - Dynamic access without external fragmentation
 - Index table storage overhead (way better than FAT)
- Cons:
 - What if i-node doesn't fit all pointers?

Implementing Files I-nodes



Question?


- If the i-node of the previous slide contains 10 direct addresses and these were 8 bytes each and all disk blocks were 1024B, what would the largest possible file be?

Indexed Allocation (Cont.)

- Question: given a logical address LA of file A, how to map LA to its physical address (B,D) (B: block number; D: block offset)?
- Assume the block size is 512 bytes

$$LA/512 \begin{cases} \text{Quotient } Q \\ \text{Remainder } R \end{cases}$$

- Block number B: Look up the Q -th entry of the index table to obtain B
- Block offset $D = R$
- Number of memory access to get the data at address LA :
 - 1 access for reading the index table
 - 1 access to get data at block offset D .



Exercise

Consider a hypothetical file system. A disk block contains 512 bytes. Assume:

- (1) it takes one disk access to find file X's entry in the directory, and
- (2) nothing is cached in memory.

Question:

To read the data at logical address 5116 - 5119 in file X, how many disk block accesses are required (including the initial one to get the directory entry) . Let's consider two file allocation schemes:

- (1) physically contiguous?
- (2) linked in a list, with the "next pointer" from one block to the next embedded in the block and occupying the first 2 bytes of the block?

Question

- Consider a file whose size varies between 4 KB and 4 MB during its lifetime. Which of the three allocation schemes (contiguous, linked and table/indexed) will be most appropriate?