



# CS 558:

## Computer Vision

### 2nd Set of Notes

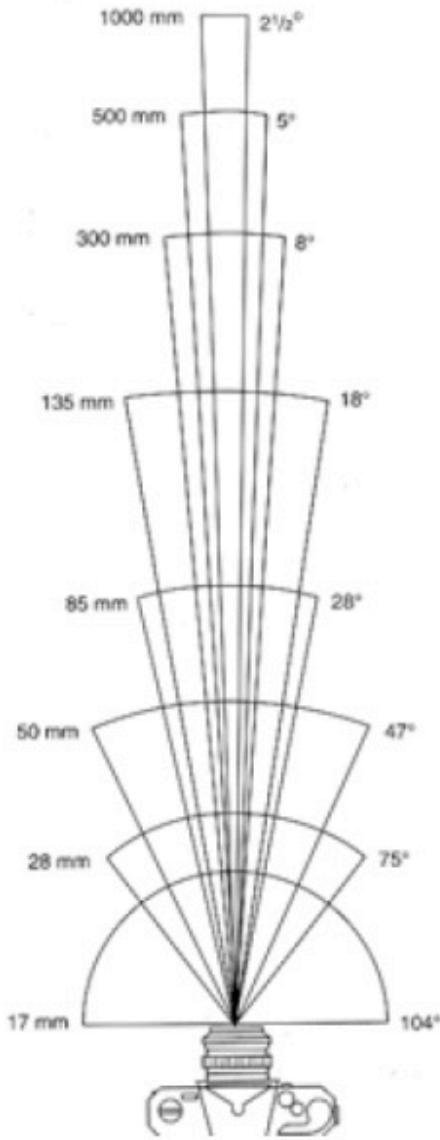
Instructor: Enrique Dunn

Webpage: [www.cs.stevens.edu/~edunn](http://www.cs.stevens.edu/~edunn)

E-mail: [edunn@stevens.edu](mailto:edunn@stevens.edu)

Office: North Bldg 219

# Field of View (Zoom)



17mm



28mm



50mm

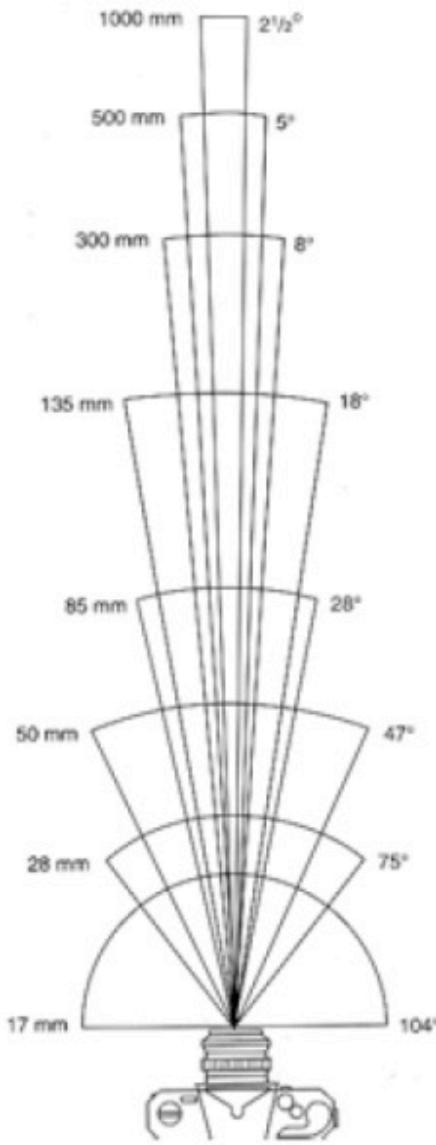


85mm

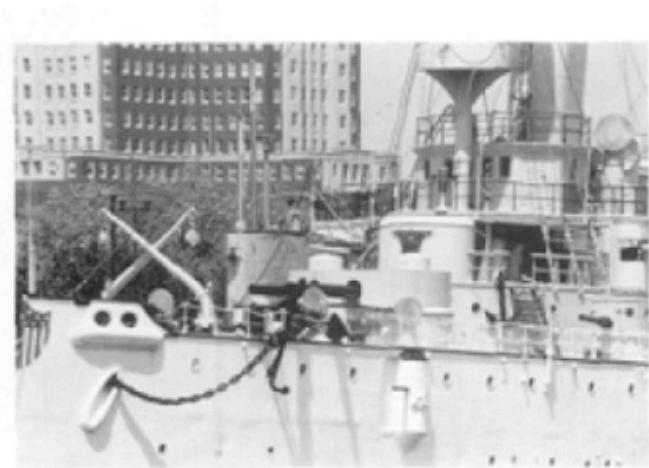
FOV measured diagonally on a 35mm full-frame camera (24 × 36mm)

**From London and Upton**

# Field of View (Zoom)



135mm



300mm



E. P. T. London

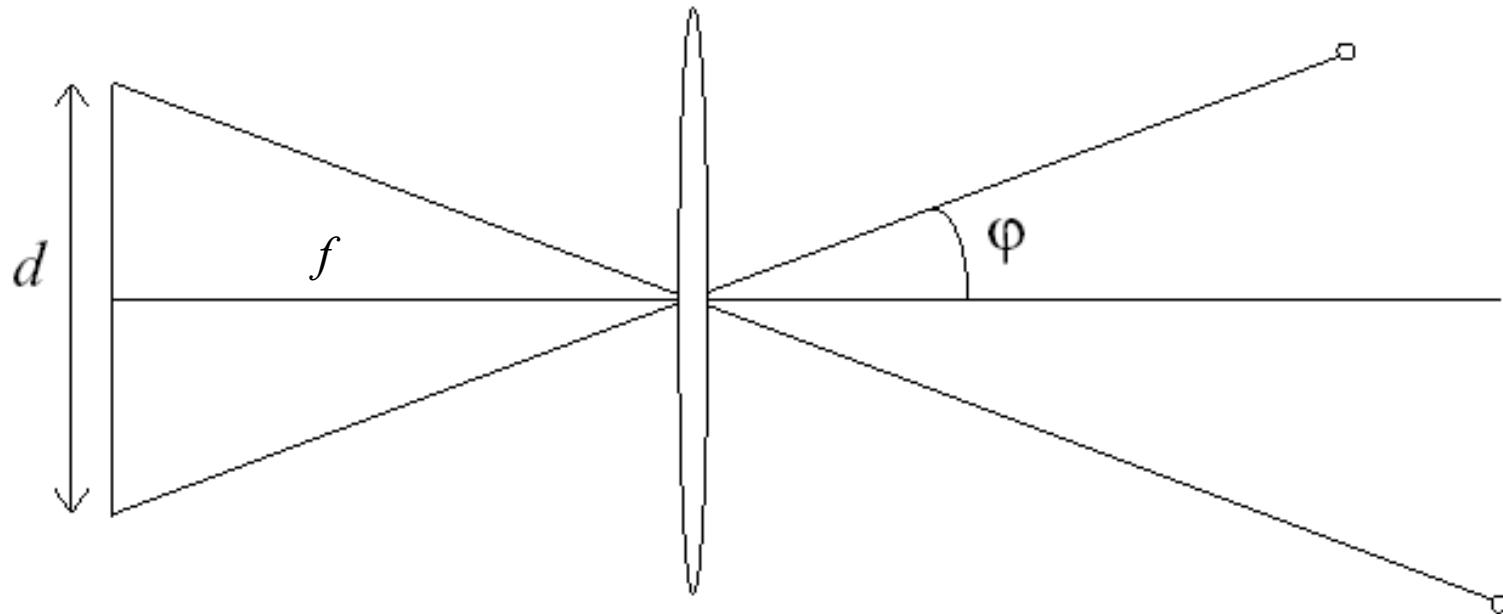


E. P. T. London

**From London and Upton**

FOV measured diagonally on a 35mm full-frame camera (24 × 36mm)

# FOV depends on Focal Length



Size of field of view governed by size of the camera retina:

$$\varphi = \tan^{-1}\left(\frac{d}{2f}\right)$$

FOV =  $2 \times \phi$  (in degrees)

Smaller FOV = larger Focal Length

# Field of View vs. Focal Length

- Portrait: distortion with wide angle
- Why?



Wide angle



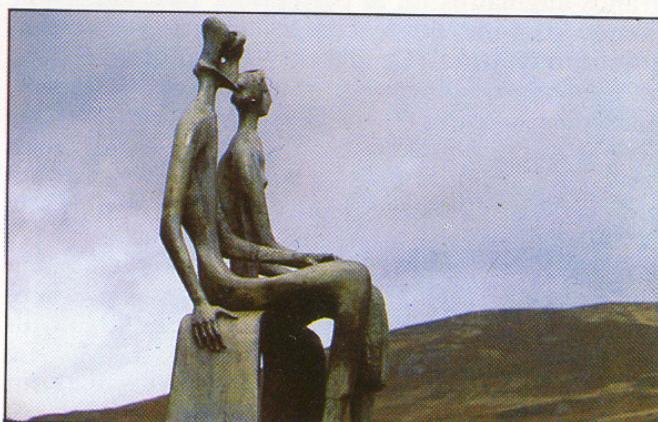
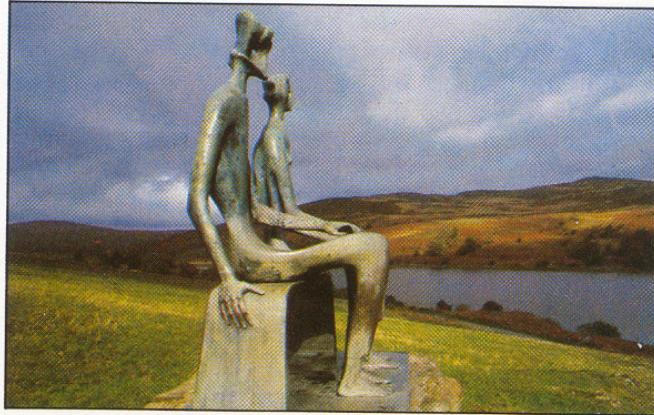
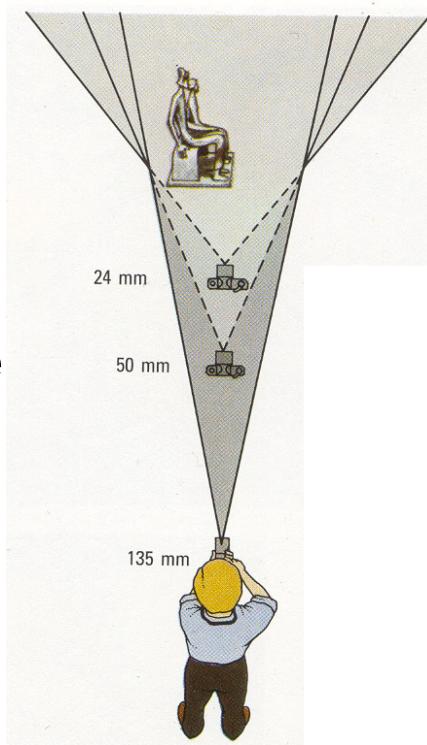
Standard



Telephoto

# Changing the focal length vs. changing the viewpoint

- Telephoto makes it easier to select background (a small change in viewpoint is a big change in background)
  - changing the focal length lets us move back from a subject, while maintaining its size on the image
  - but moving back changes perspective relationships

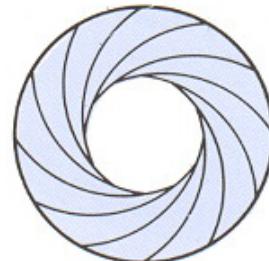


# Aperture

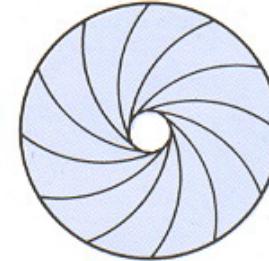
- Diameter of the lens opening (controlled by diaphragm)
- Expressed as a fraction of focal length, in ***f-number N***
  - f/2.0 on a 50mm lens means that the aperture is 25mm
  - f/2.0 on a 100mm lens means that the aperture is 50mm
- Confusing: small f-number = big aperture
- What happens to the area of the aperture when going from f/2.0 to f/4.0?
- Typical f-numbers are (each of them counts as one f/stop)  
f/2.0, f/2.8, f/4, f/5.6, f/8, f/11, f/16, f/22, f/32
  - See the pattern?



Full aperture



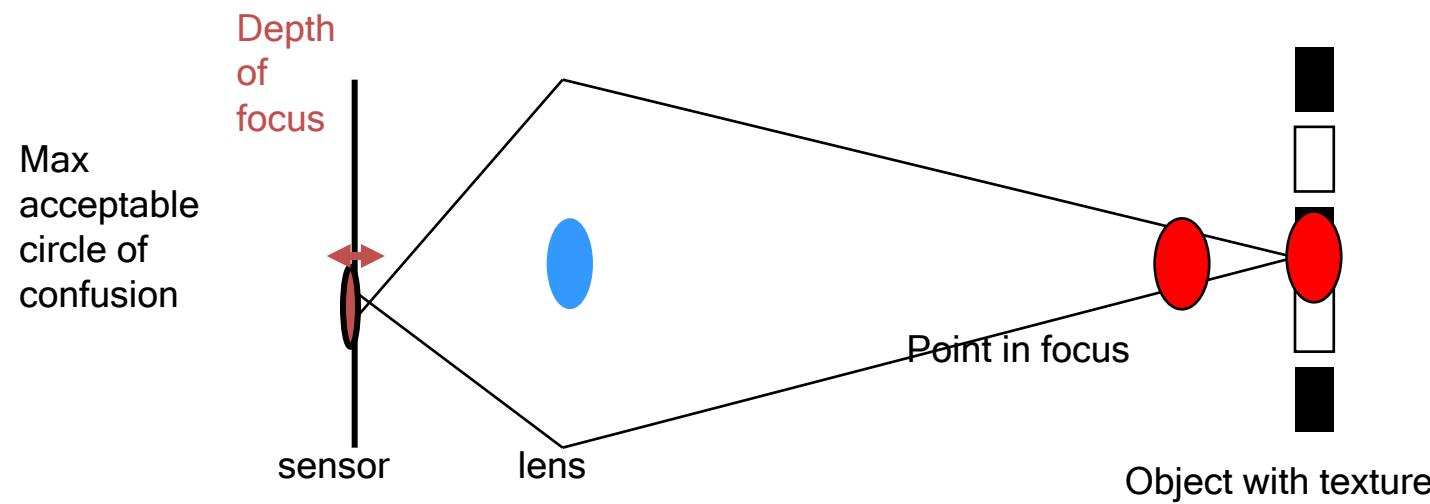
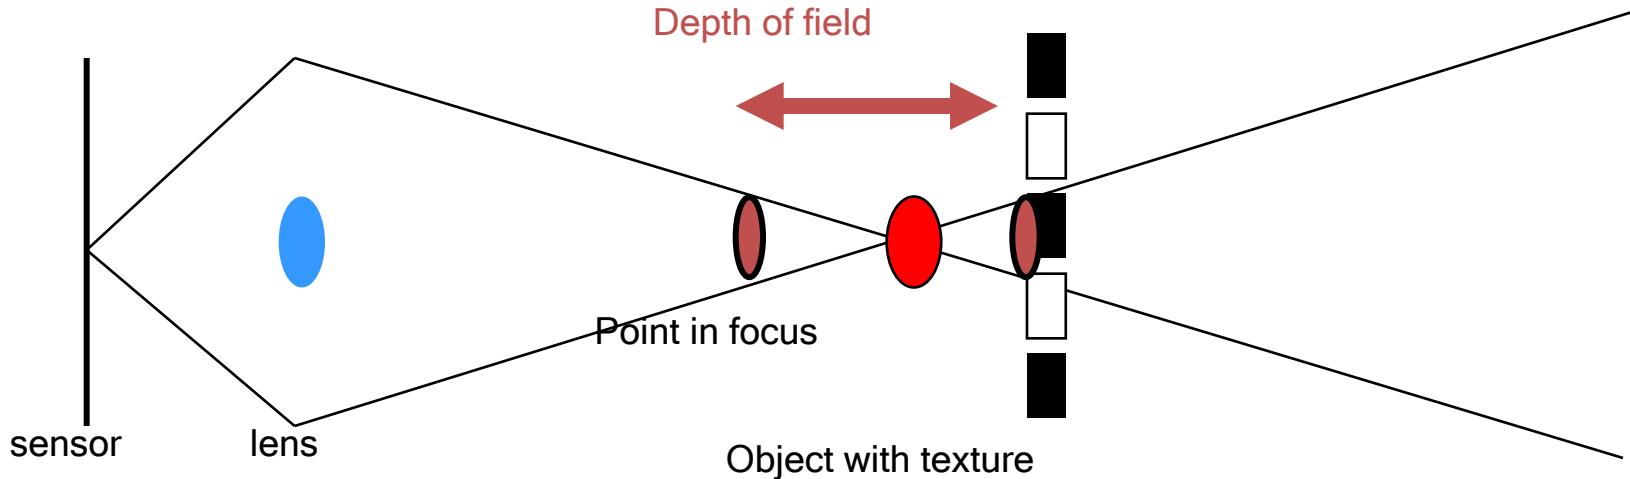
Medium aperture



Stopped down

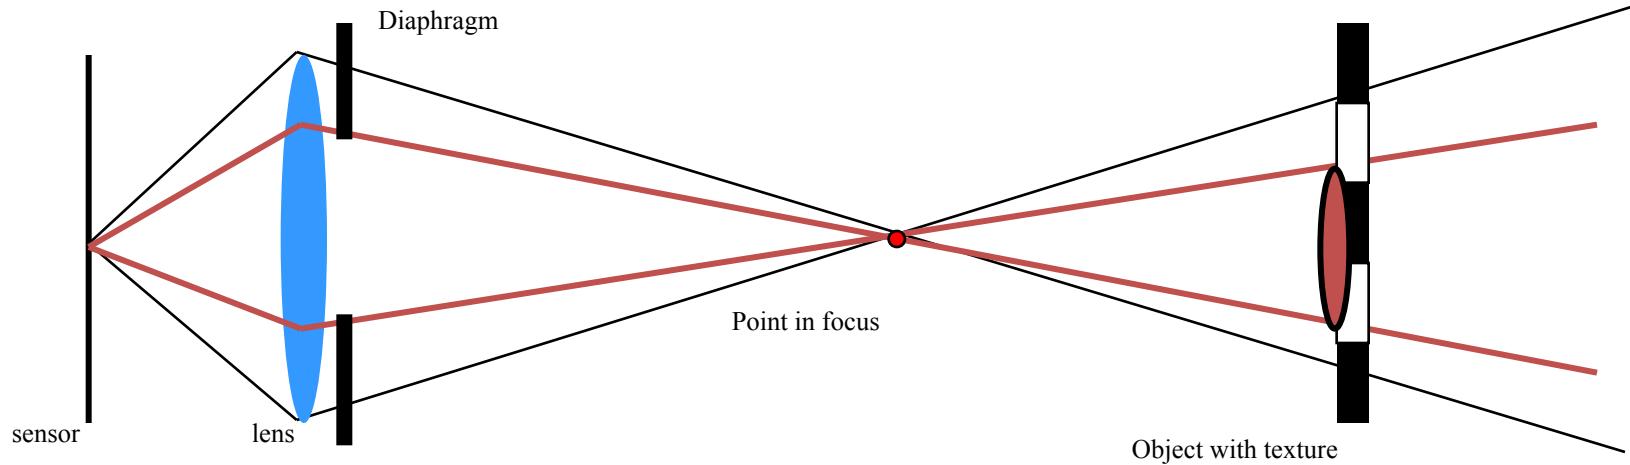
# Depth of field

- We allow for some tolerance

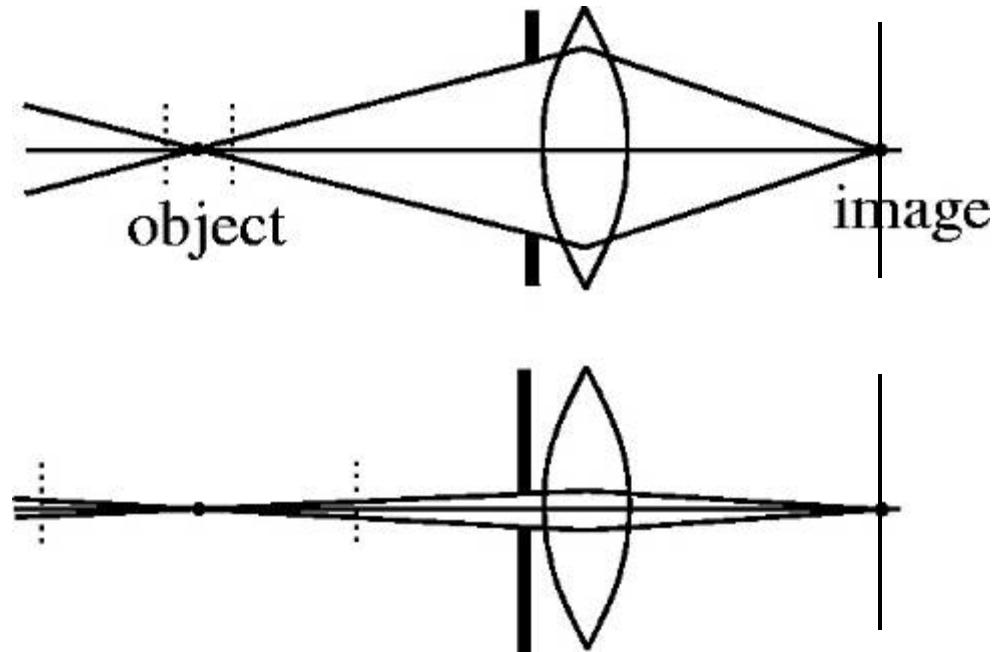


# Aperture controls Depth of Field

- What happens when we close the aperture by two stops?
  - Aperture diameter is divided by two
  - Depth of field is doubled

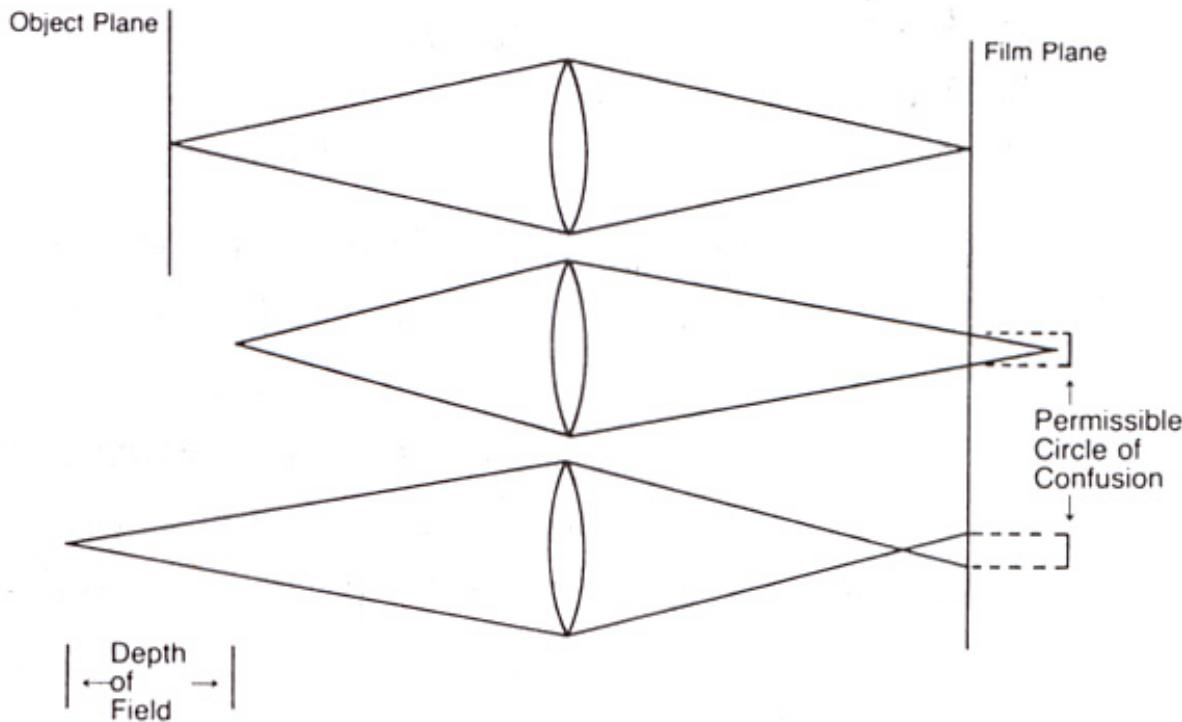


# Aperture controls Depth of Field



- A smaller aperture increases the range in which the object is approximately in focus
- But small aperture reduces amount of light
  - need to increase exposure

# Circle of confusion (C)



- C depends on sensing medium, reproduction medium, viewing distance, human vision,...
  - for print from 35mm film, 0.02mm is typical
  - for high-end DSLR,  $6\mu$  is typical (1 pixel)
  - larger if downsizing for web, or lens is poor

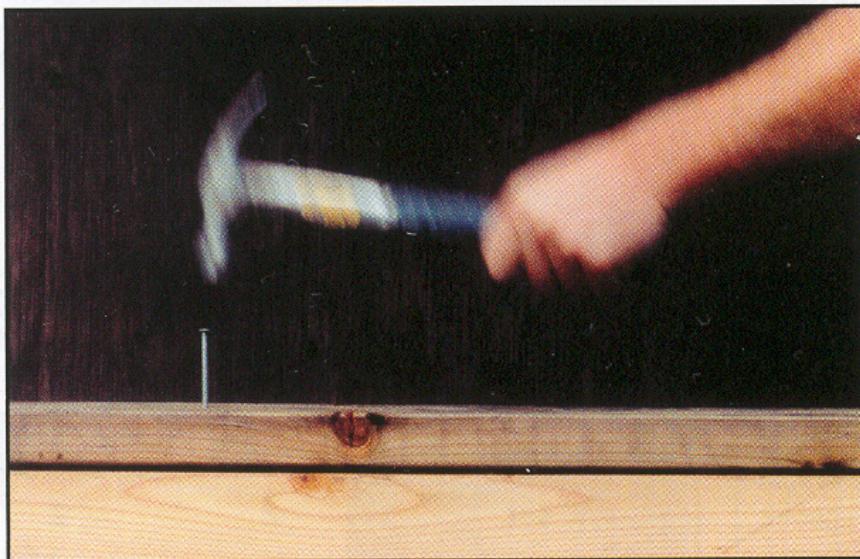
# Shutter speed

- Controls how long the film/sensor is exposed
- Pretty much linear effect on exposure
- Usually in fraction of a second:
  - 1/30, 1/60, 1/125, 1/250, 1/500
  - Get the pattern ?
- On a normal lens, normal humans can hand-hold down to 1/60

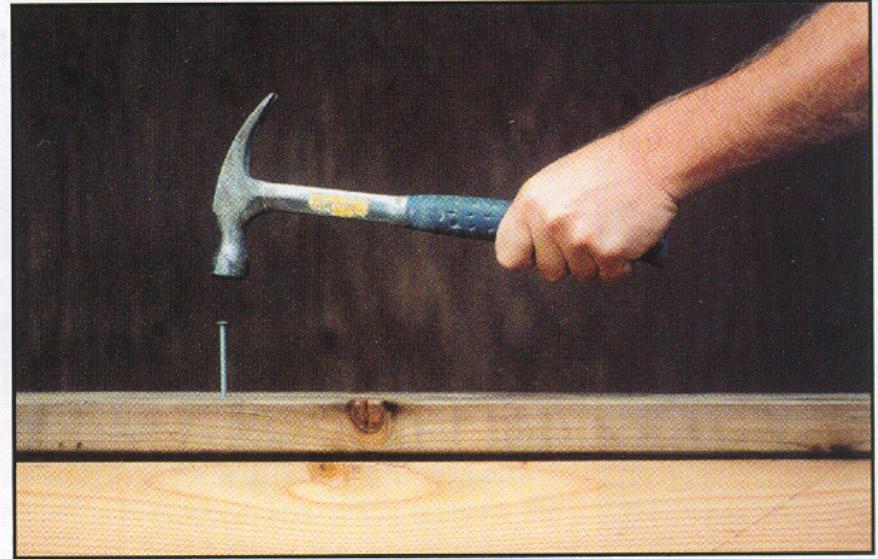
# Main effect of shutter speed

- Motion blur
- Halving shutter speed doubles motion blur

Slow shutter speed



Fast shutter speed



From Photography, London et al.

# Effect of shutter speed

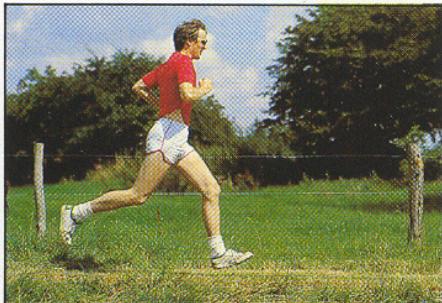
- Freezing motion

Walking people



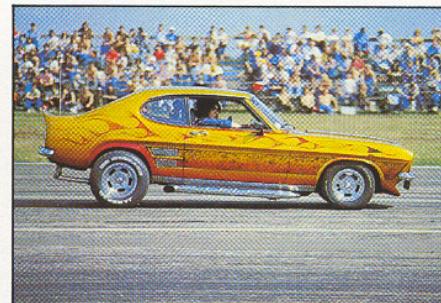
1/125

Running people



1/250

Car



1/500

Fast train



1/1000

# Exposure

- Two main parameters:
  - Aperture (in f number)
  - Shutter speed (in fraction of a second)
- Exposure = irradiance x time

$$H = ExT$$

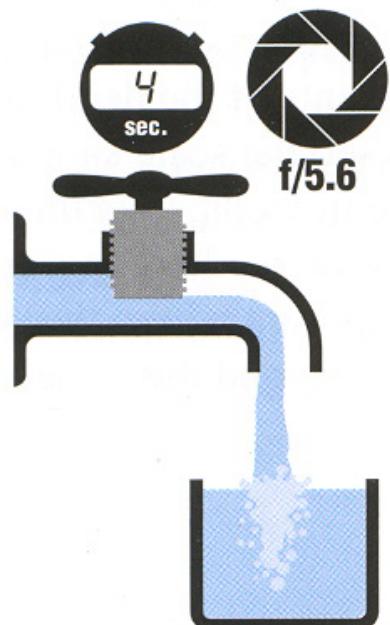
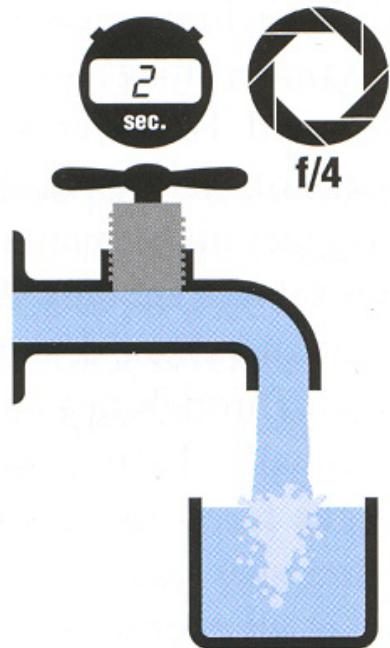
- Irradiance (E)
  - controlled by aperture
- Exposure time (T)
  - controlled by shutter

# Reciprocity

- Reciprocity

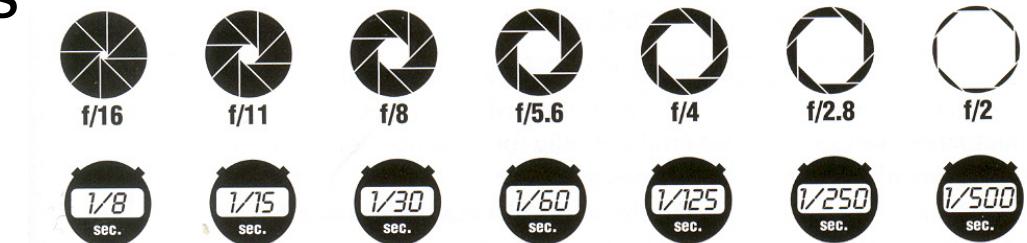
The same exposure is obtained with an exposure twice as long and an aperture *area* half as big

- Hence square root of two progression of f stops vs. power of two progression of shutter speed



# Reciprocity

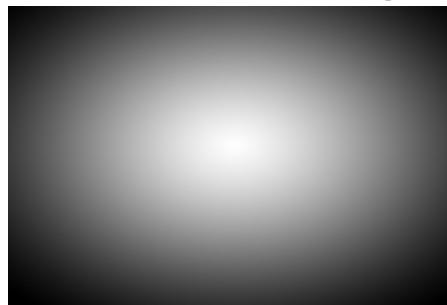
- Assume we know how much light we need
- We have the choice of an infinity of shutter speed/aperture pairs



- What will guide our choice of a shutter speed?
  - Freeze motion vs. motion blur, camera shake
- What will guide our choice of an aperture?
  - Depth of field, distortion reduction, diffraction limit
- Often we must compromise
  - Open more to enable faster speed (but shallow DoF)

# Metering

- Photosensitive sensors measure scene luminance
- Usually TTL (through the lens)
- Simple version: center-weighted average



- Assumption? Failure cases?
  - Usually assumes that a scene is 18% gray
  - Problem with dark and bright scenes

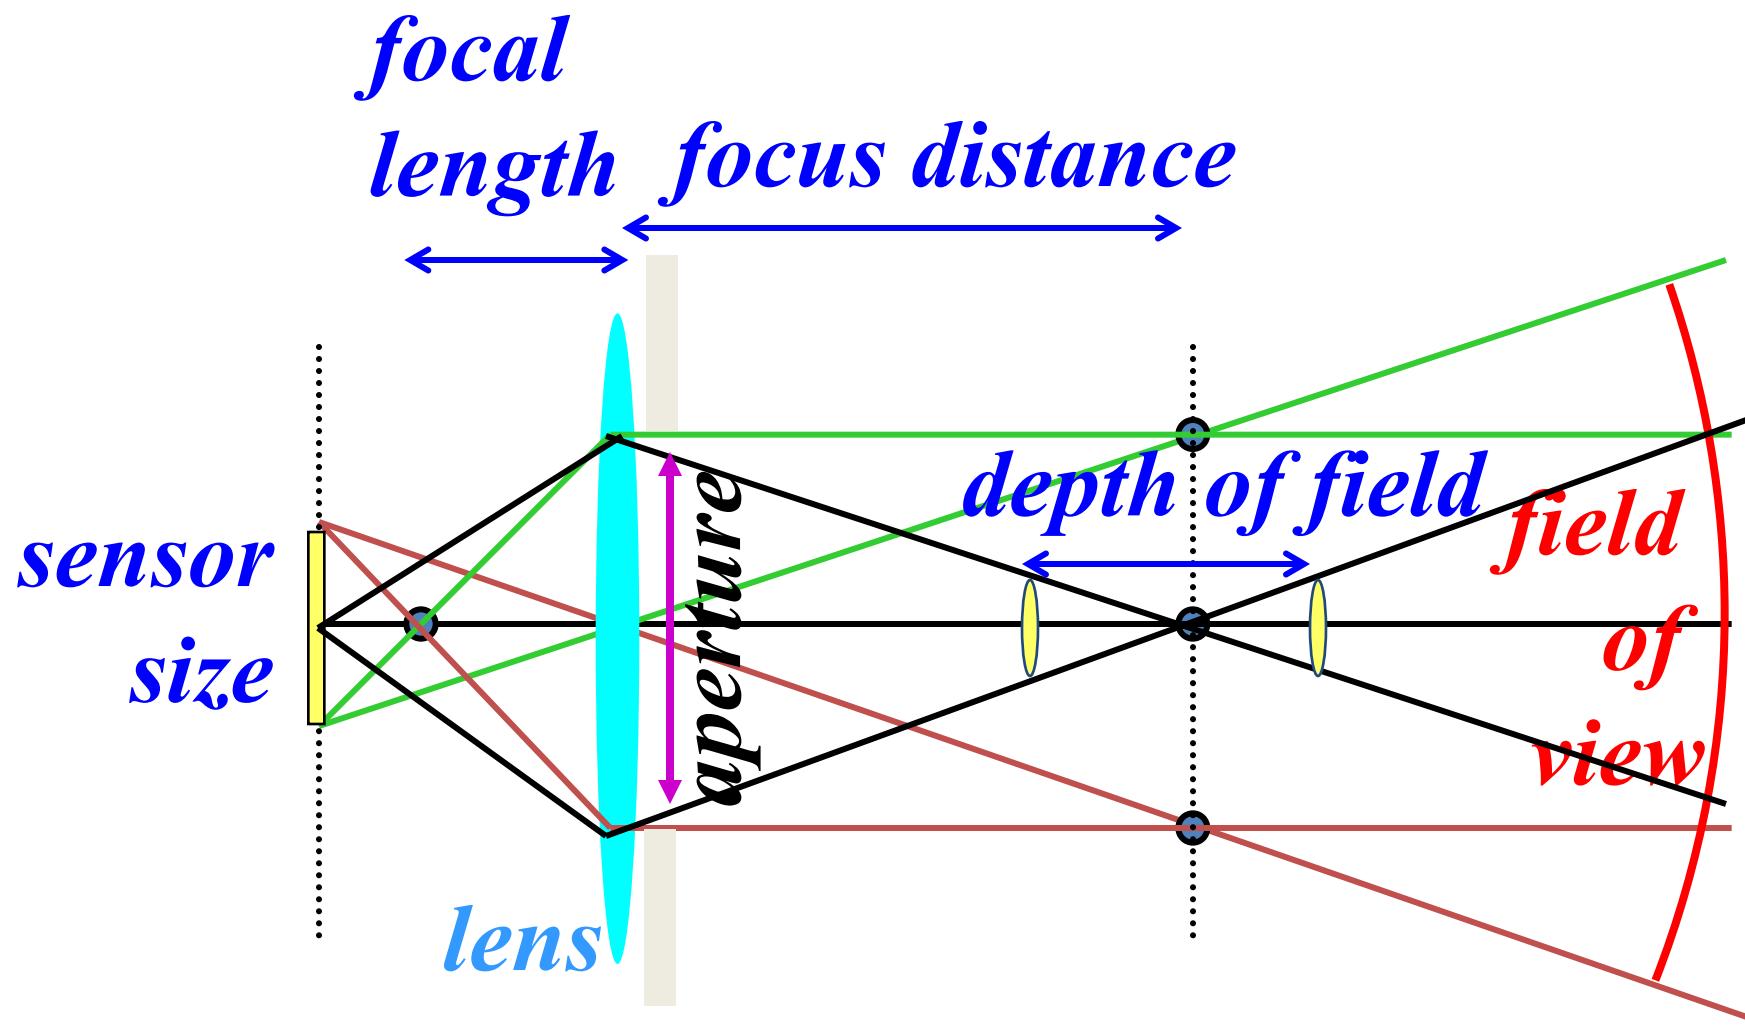
# Exposure & Metering

- The camera metering system measures how bright the scene is
- In Aperture priority mode, the photographer sets the aperture, the camera sets the shutter speed
- In Shutter-speed priority mode, the photographers sets the shutter speed and the camera deduces the aperture
  - In both cases, reciprocity is exploited
- In Program mode, the camera decides both exposure and shutter speed (middle value more or less)
- In Manual, the user decides everything (but can get feedback)

# Pros and cons of various modes

- Aperture priority
  - Direct depth of field control
  - Cons: can require impossible shutter speed (e.g. with f/1.4 for a bright scene)
- Shutter speed priority
  - Direct motion blur control
  - Cons: can require impossible aperture (e.g. when requesting a 1/1000 speed for a dark scene)
    - Note that aperture is somewhat more restricted
- Program
  - Almost no control, but no need for neurons
- Manual
  - Full control, but takes more time and thinking

# Recap



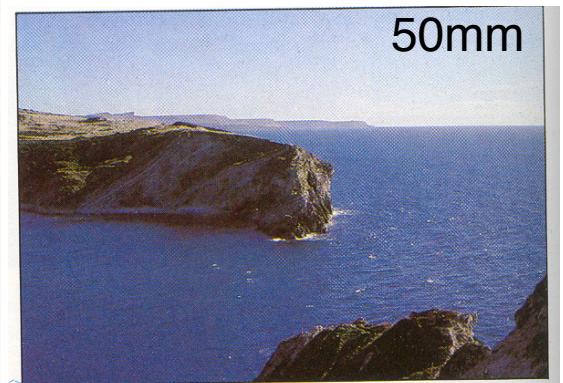
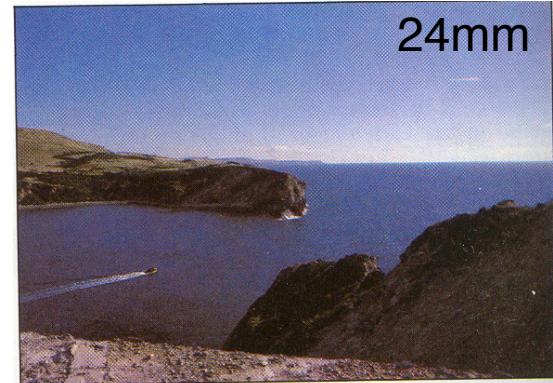
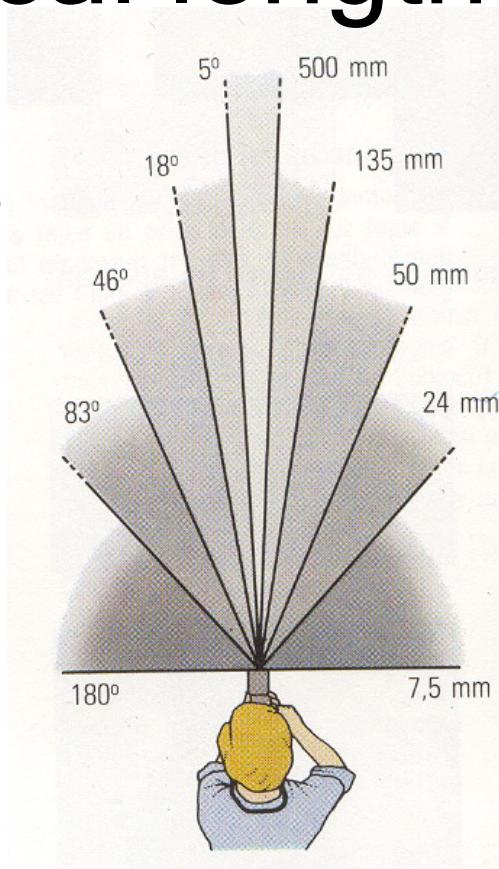
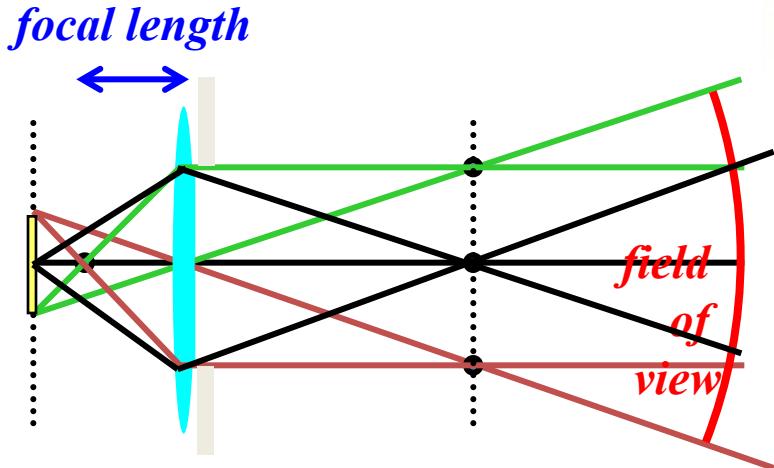
# Focal length

<30mm: wide angle

50mm: standard

>100mm telephoto

Affected by sensor size (crop factor)

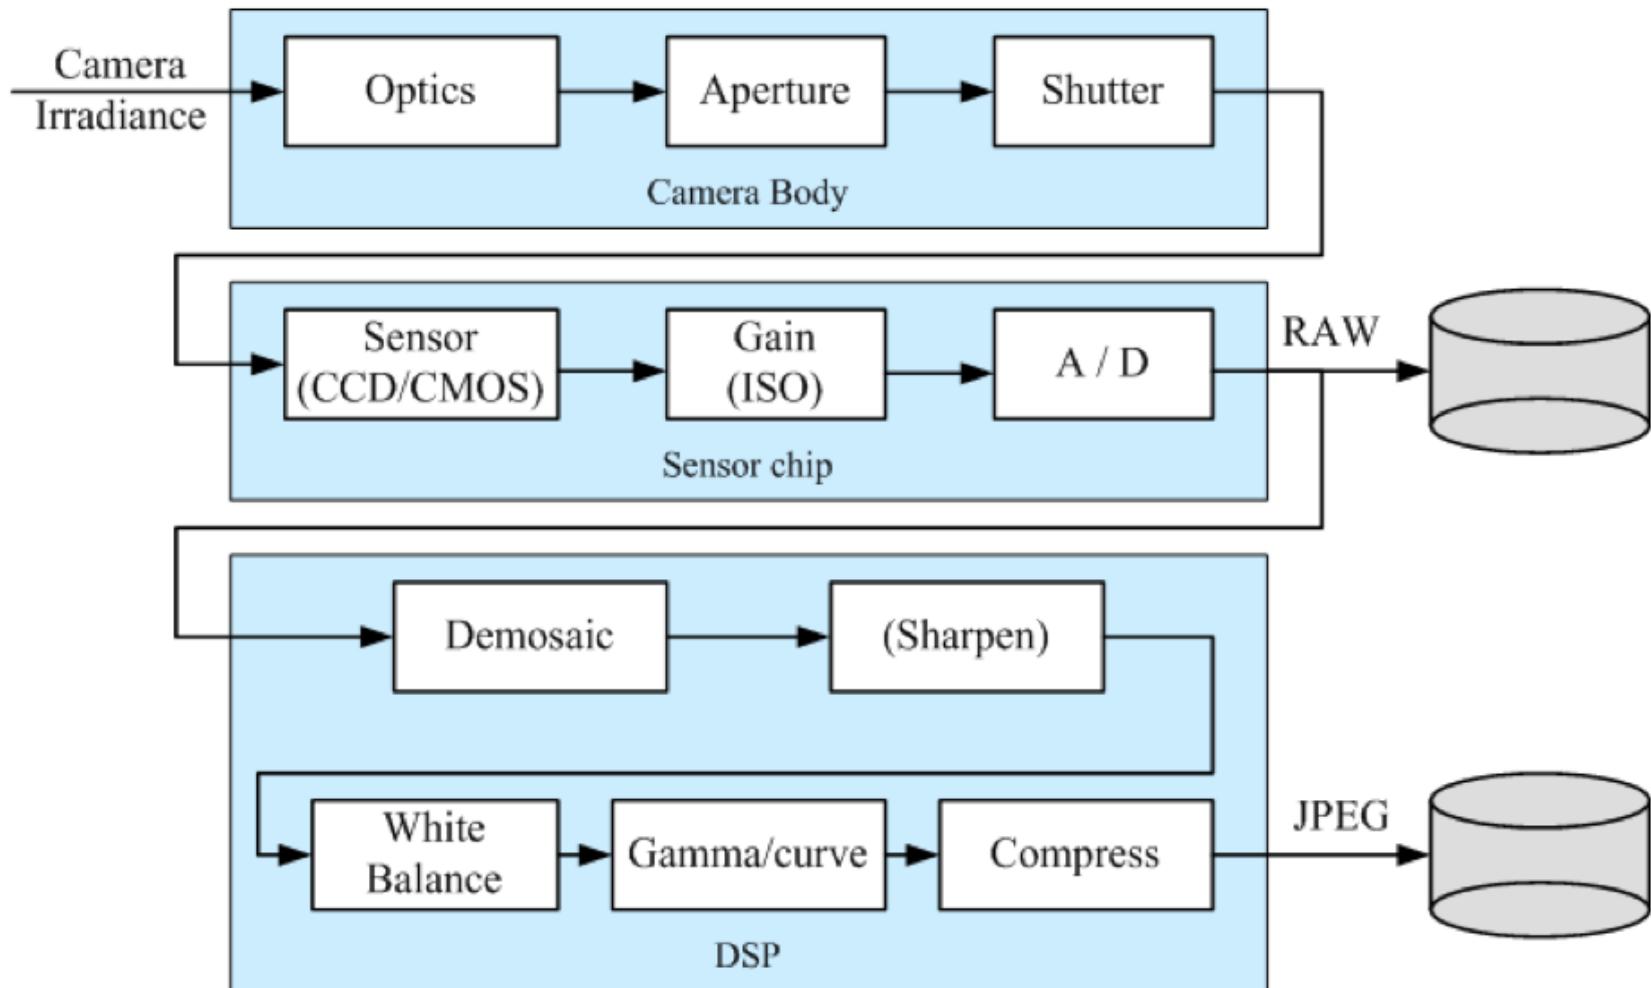


# Exposure

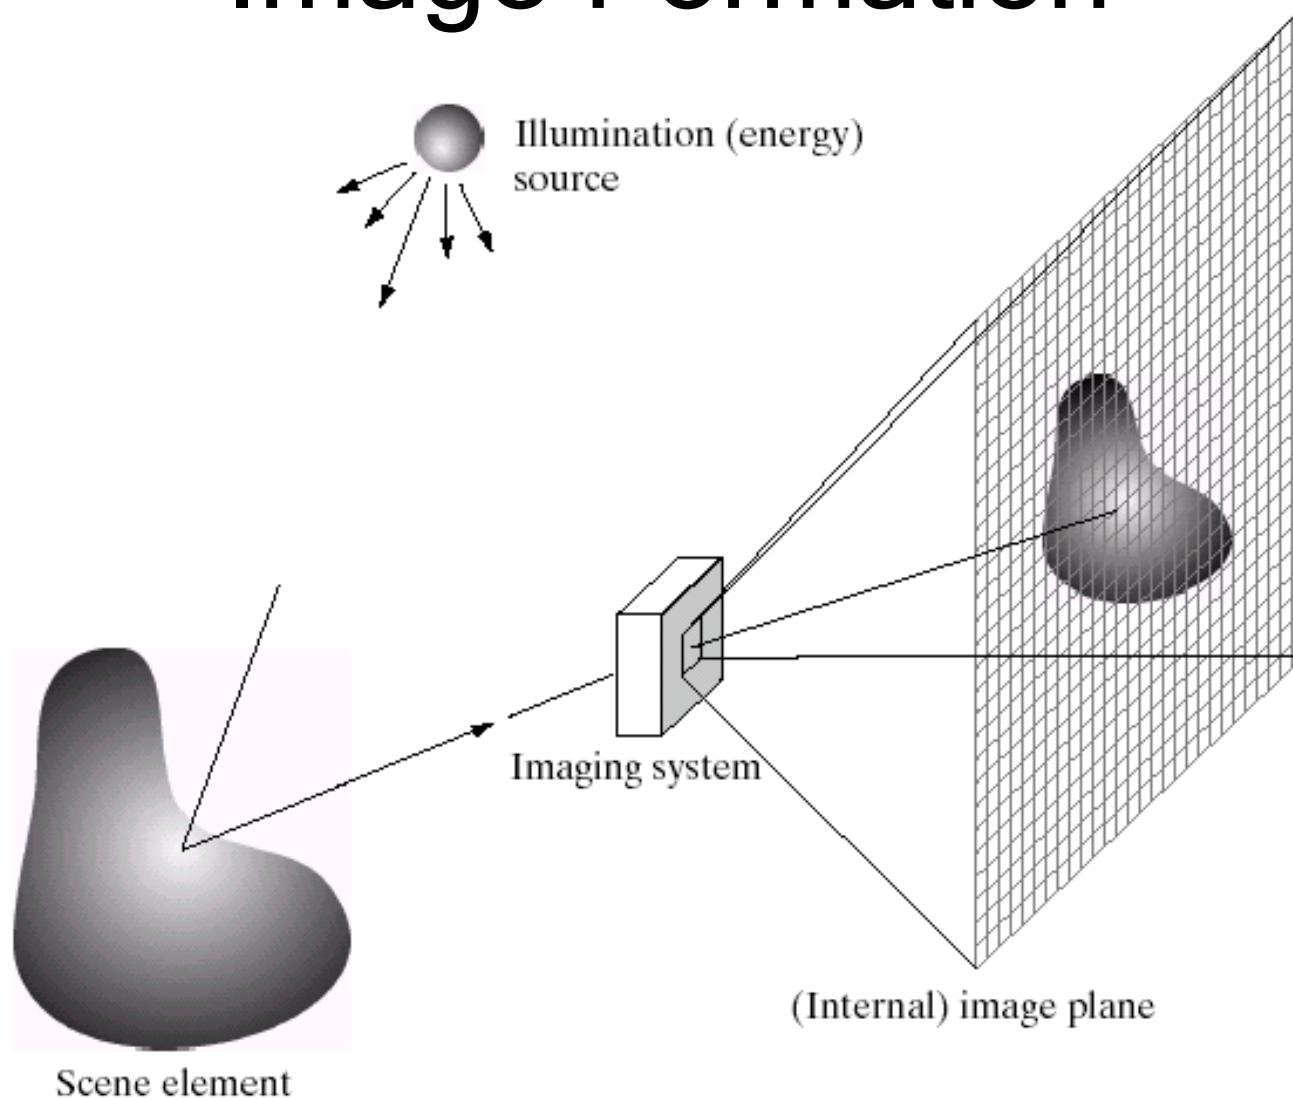
- Aperture (f number)
  - Expressed as ratio between focal length and aperture diameter:  
diameter =  $f / \text{f number}$
  - f/2.0, f/2.8, f/4.0, f/5.6, f/8.0, f/11, f/16 (factor of  $\sqrt{2}$ )
  - Small f number means large aperture
  - Main effect: depth of field
  - A good standard lens has max aperture f/1.8.  
A cheap zoom has max aperture f/3.5
- Shutter speed
  - In fraction of a second
  - 1/30, 1/60, 1/125, 1/250, 1/500 (factor of 2)
  - Main effect: motion blur
- Sensitivity
  - Gain applied to sensor
  - In ISO, bigger number, more sensitive (100, 200, 400, 800, 1600)
  - Main effect: sensor noise

Reciprocity between these three numbers:  
for a given exposure, one has two degrees of freedom.

# Sensor Chip



# Image Formation



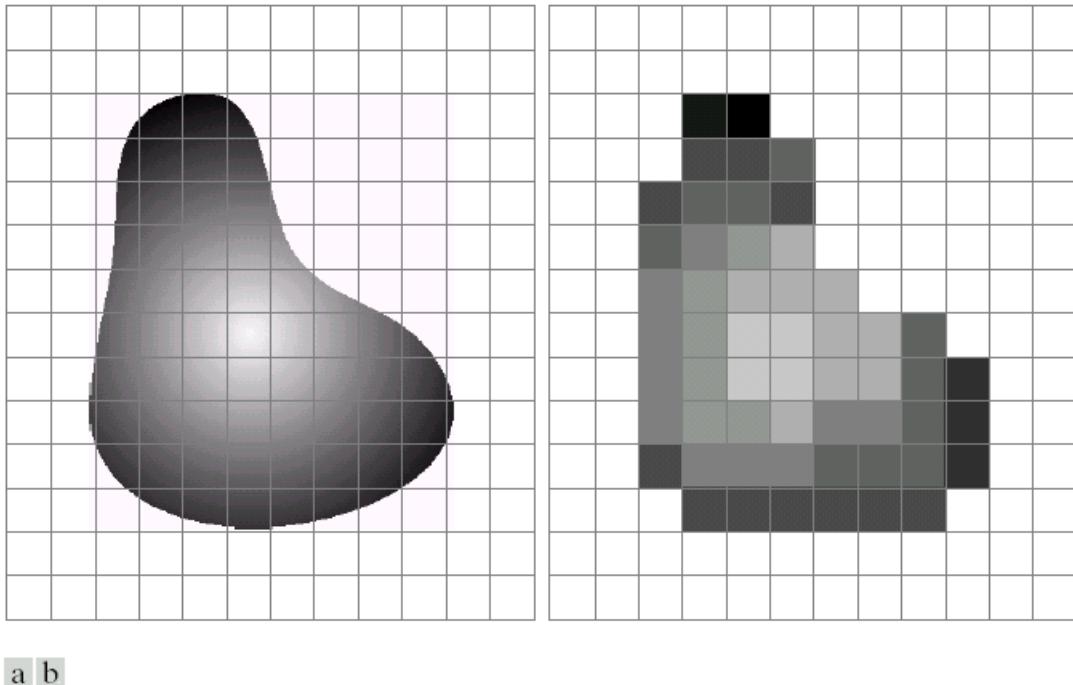
# Digital camera



A digital camera replaces film with a sensor array

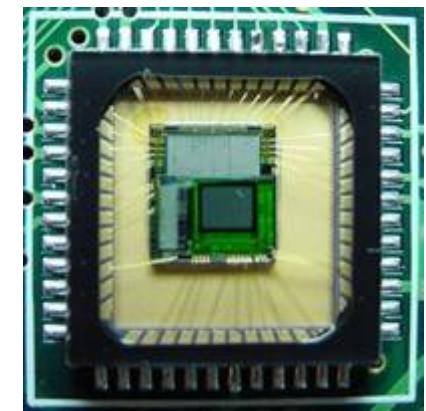
- Each cell in the array is light-sensitive diode that converts photons to electrons
- Two common types: Charge Coupled Device (CCD) and Complementary Metal Oxide Semiconductor (CMOS)

# Sensor Array



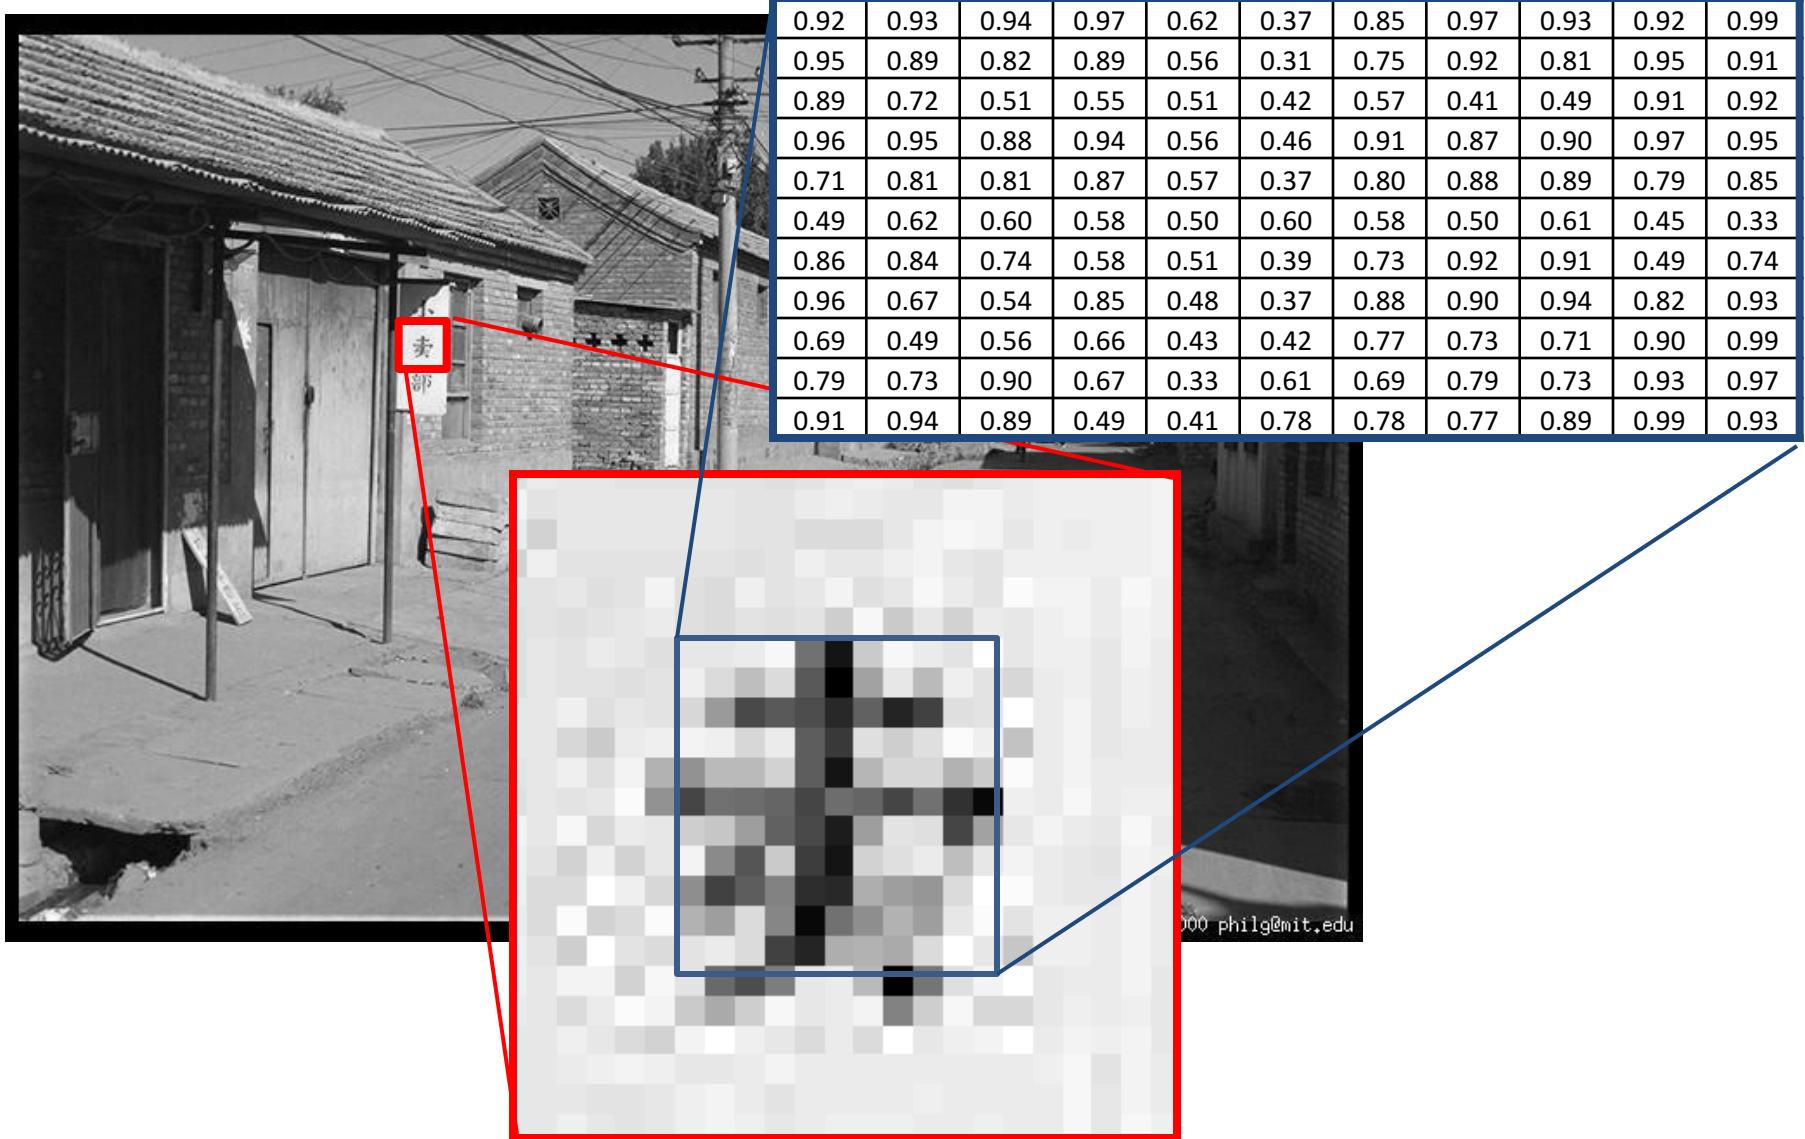
a b

**FIGURE 2.17** (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.

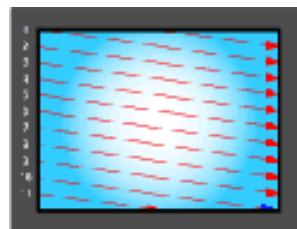
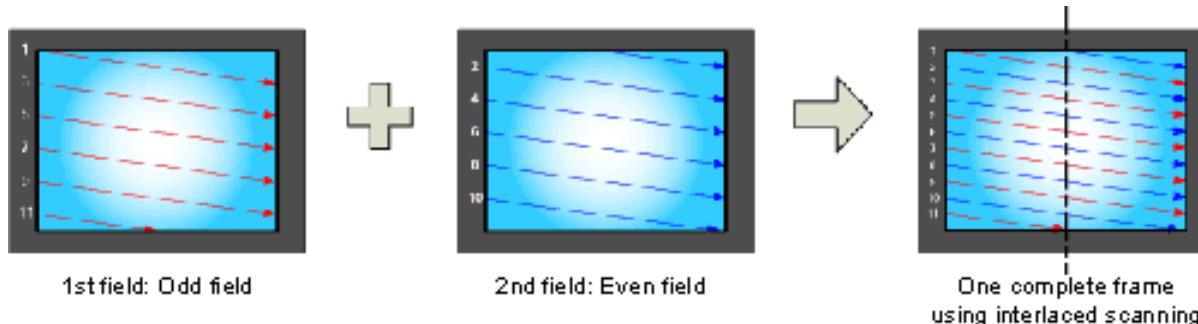


CMOS sensor

# The raster image (pixel matrix)



# Interlace vs. progressive scan



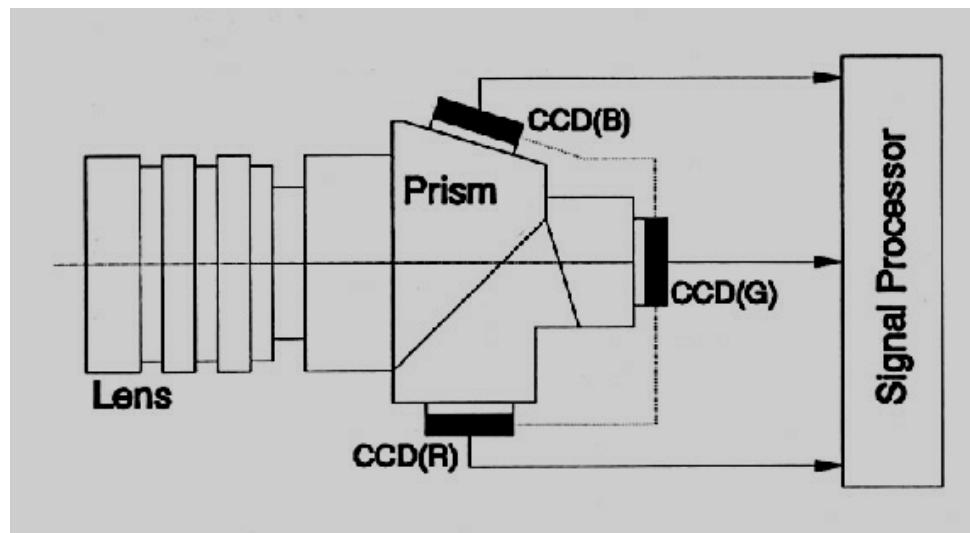
One complete frame  
using progressive scanning

# Some approaches to color sensing

- Scan 3 times (temporal multiplexing)
  - Drum scanners
  - Flat-bed scanners
  - Russian photographs from 1800's
- Use 3 detectors
  - High-end 3-tube or 3-ccd video cameras
- Use spatially offset color samples (spatial multiplexing)
  - Single-chip CCD color cameras
  - Human eye

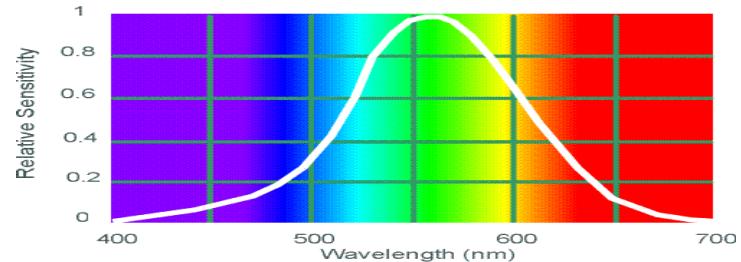
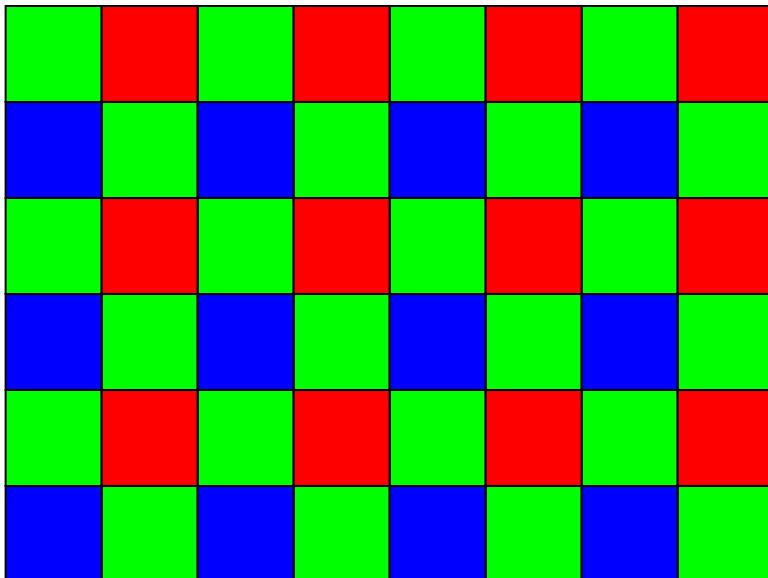
# 3 CCD Sensor

- 3-chip vs. 1-chip: quality vs. cost

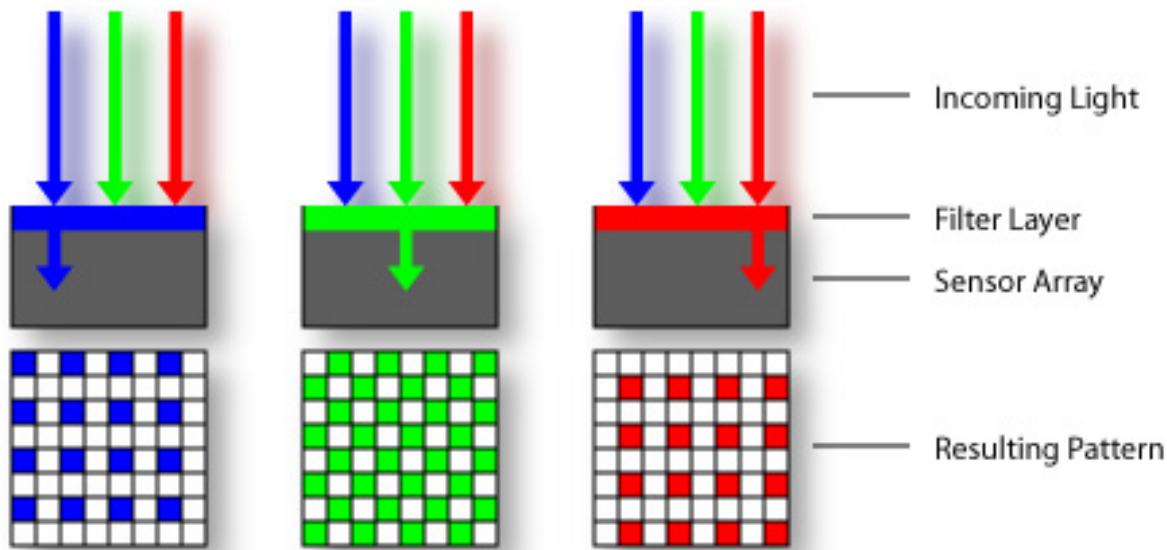
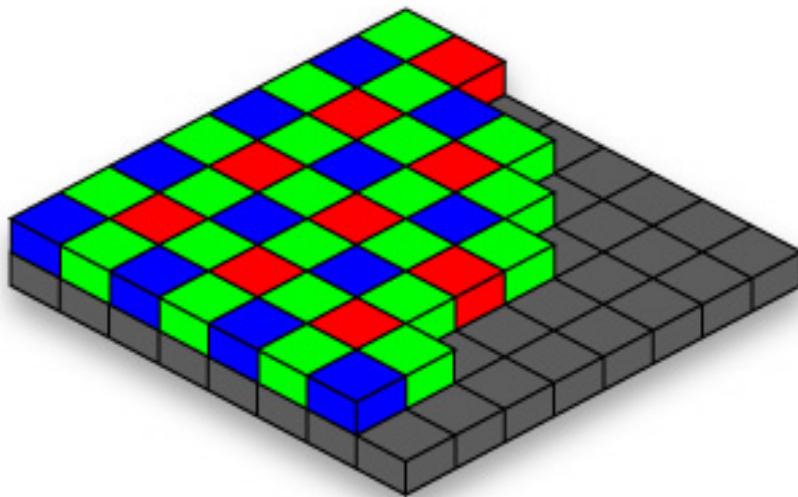


# Spatial Multiplexing: Bayer Grid

- Why more green?
  - We have 3 channels and square lattice doesn't like odd numbers
  - It's the spectrum “in the middle”
  - More important to human perception of brightness

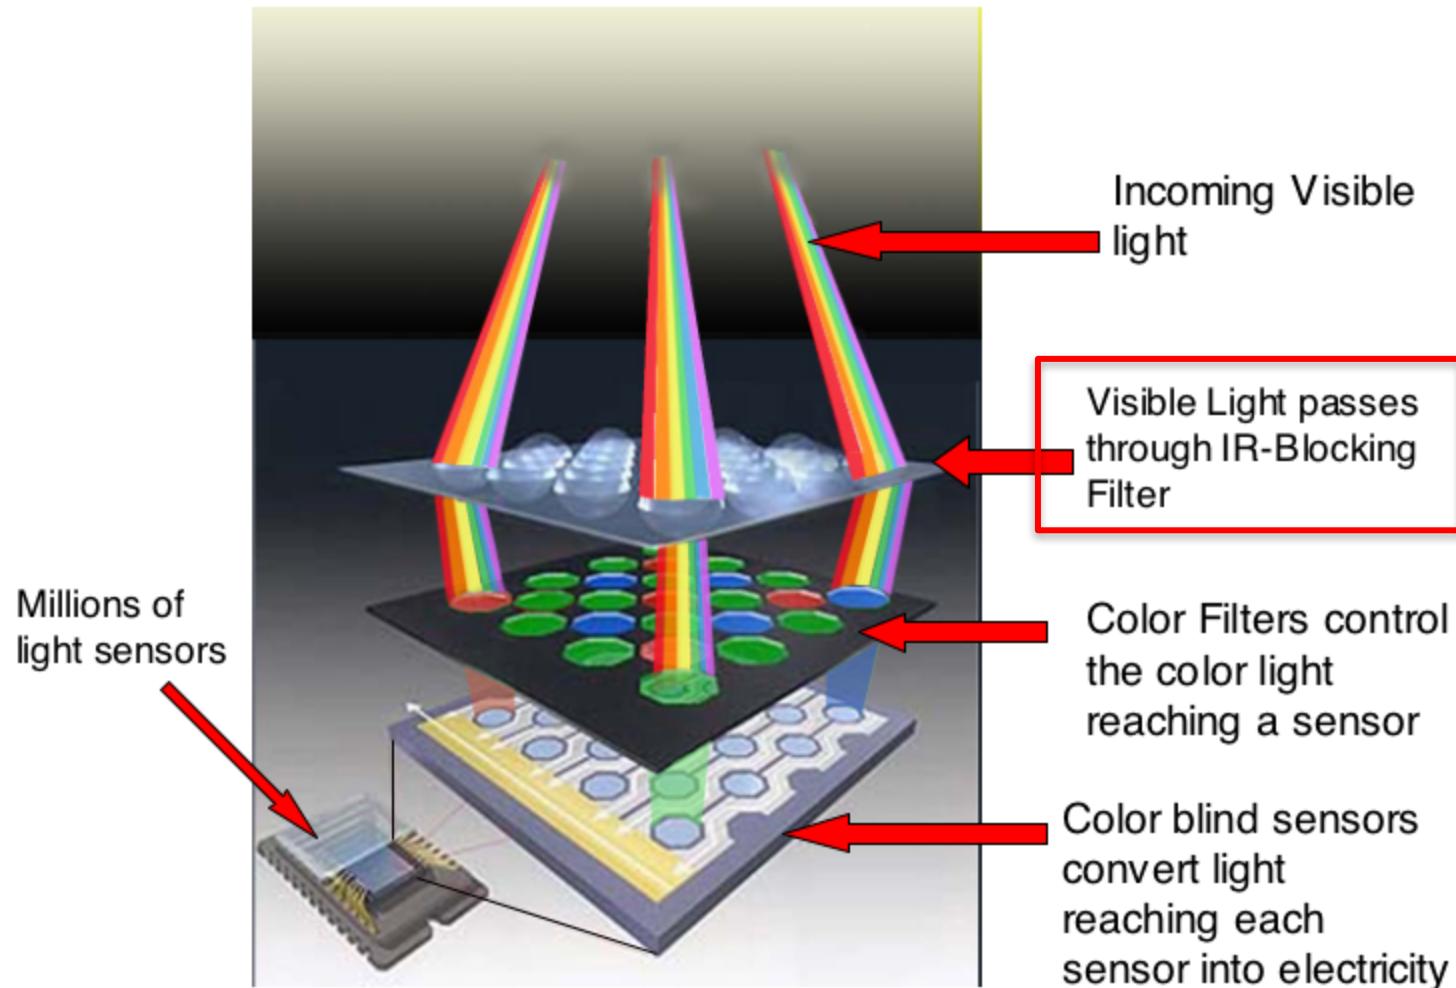


# Practical Color Sensing: Bayer Grid

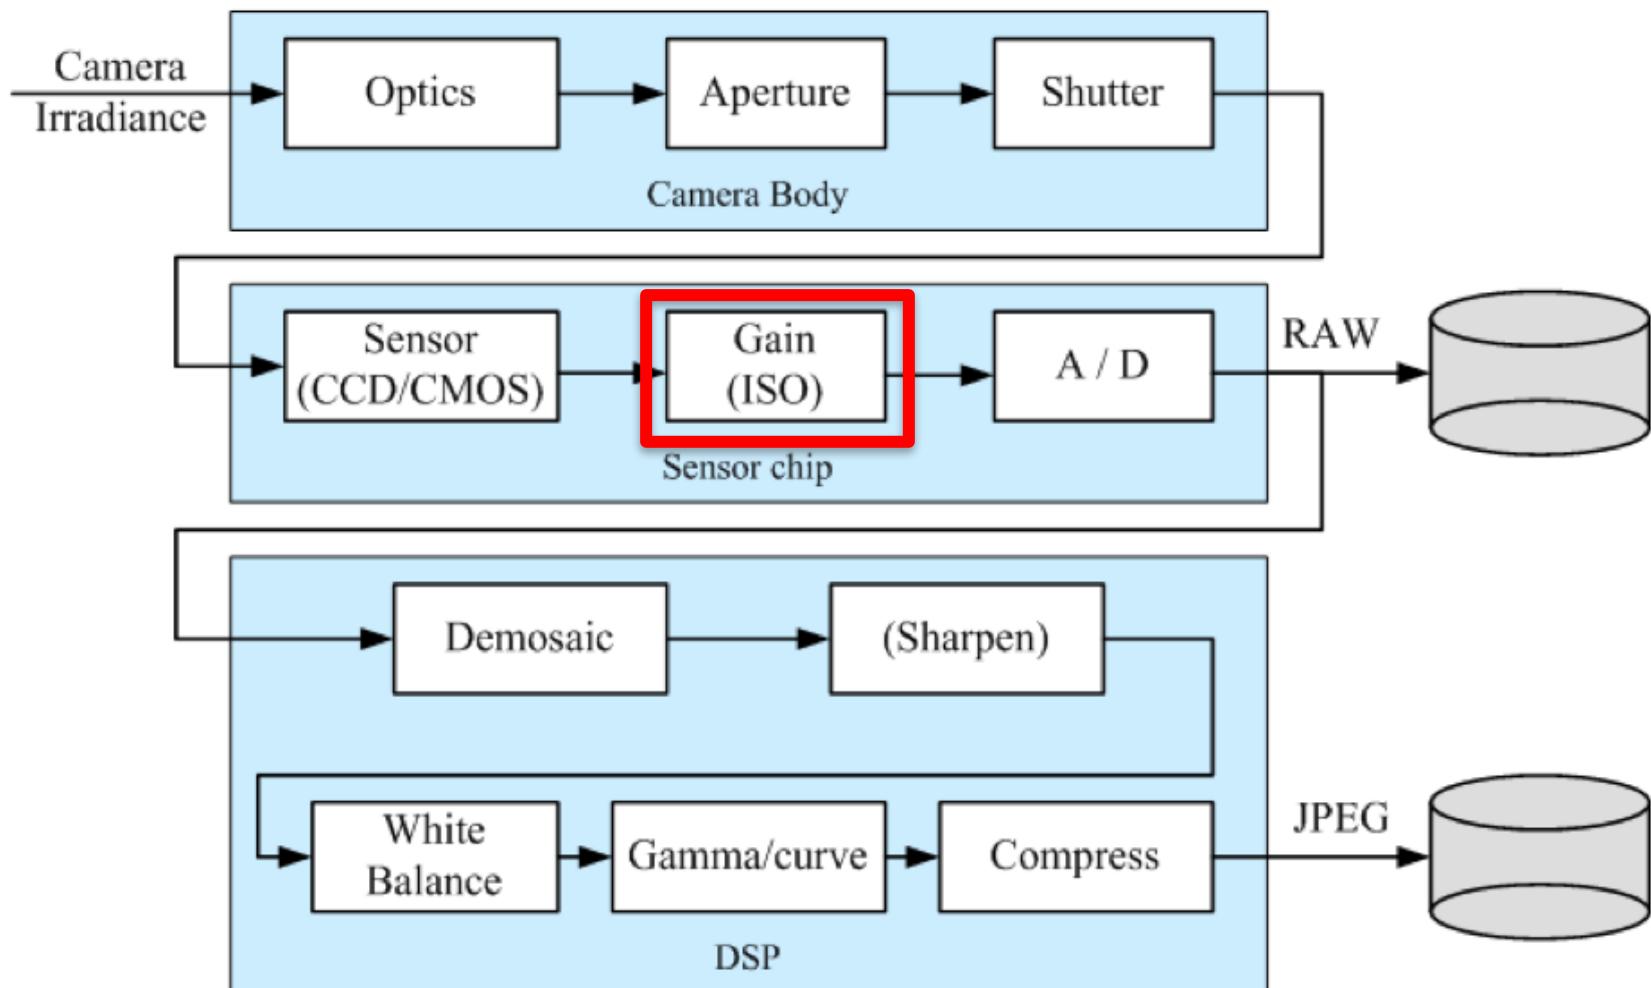


# Recap: Camera sensor

## RGB Inside the Camera



# Sensor Chip: Gain



# Sensitivity (ISO)

- Third variable for exposure: gain applied to sensor
- Linear effect (200 ISO needs half the light as 100 ISO)
- Film photography: trade sensitivity for grain



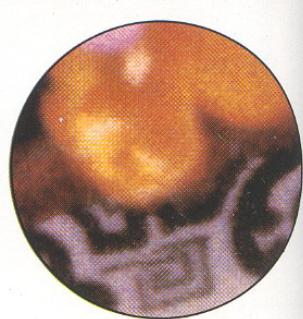
Kodachrome 25 ASA



Ektachrome 64 ASA



Fujichrome 100 ASA



Ektachrome 200 ASA

- Digital photography: trade sensitivity for noise

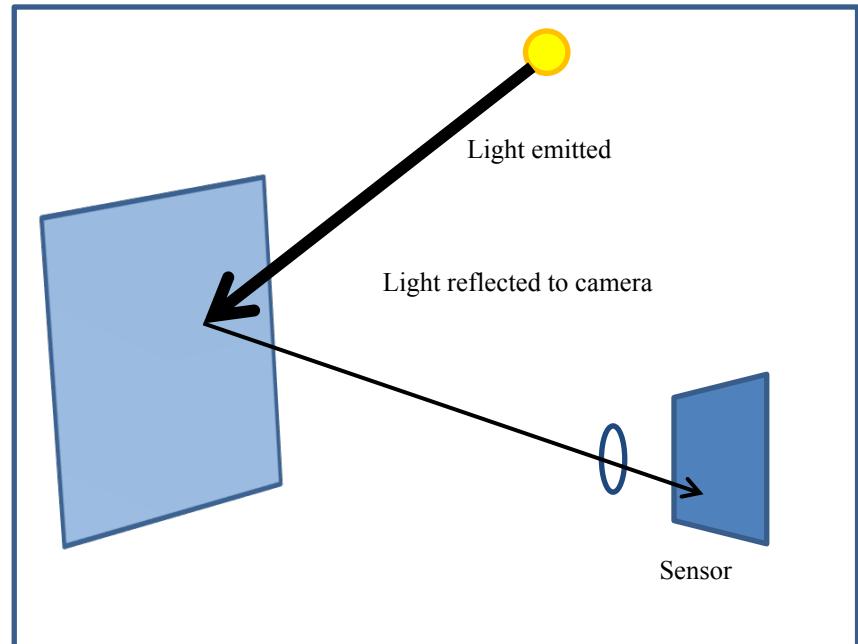
Nikon D2X ISO 100	Nikon D2X ISO 200	Nikon D2X ISO 400	Nikon D2X ISO 800	Nikon D2X ISO 1600	Nikon D2X ISO 3200

# Light and Shading

Slides by D. Hoiem

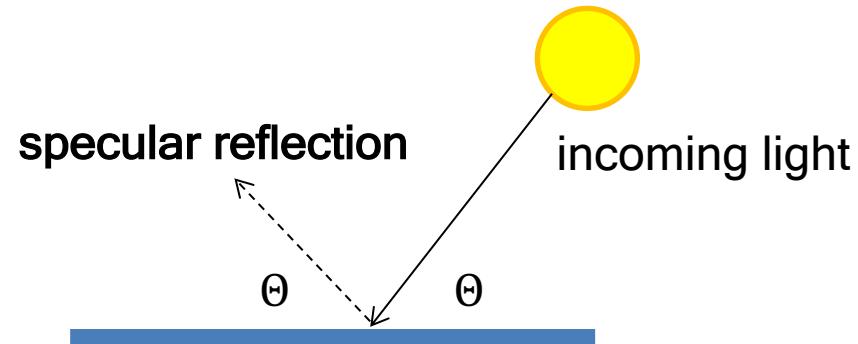
# How does a pixel get its value?

- Major factors
  - Illumination strength and direction
  - Surface geometry
  - Surface material
  - Nearby surfaces
  - Camera gain/exposure

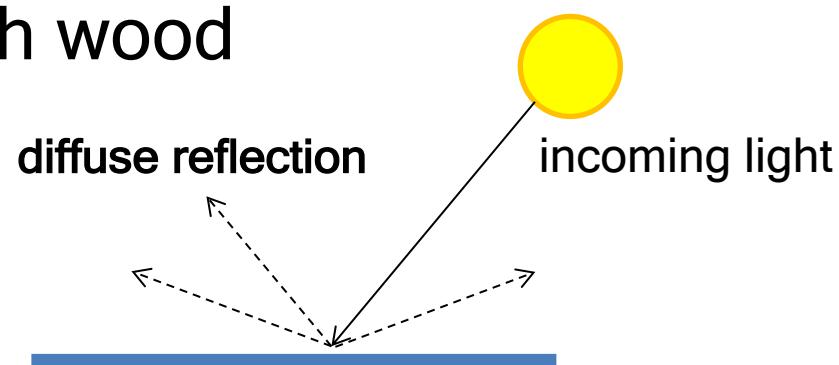


# Basic models of reflection

- Specular: light bounces off at the incident angle
  - E.g., mirror

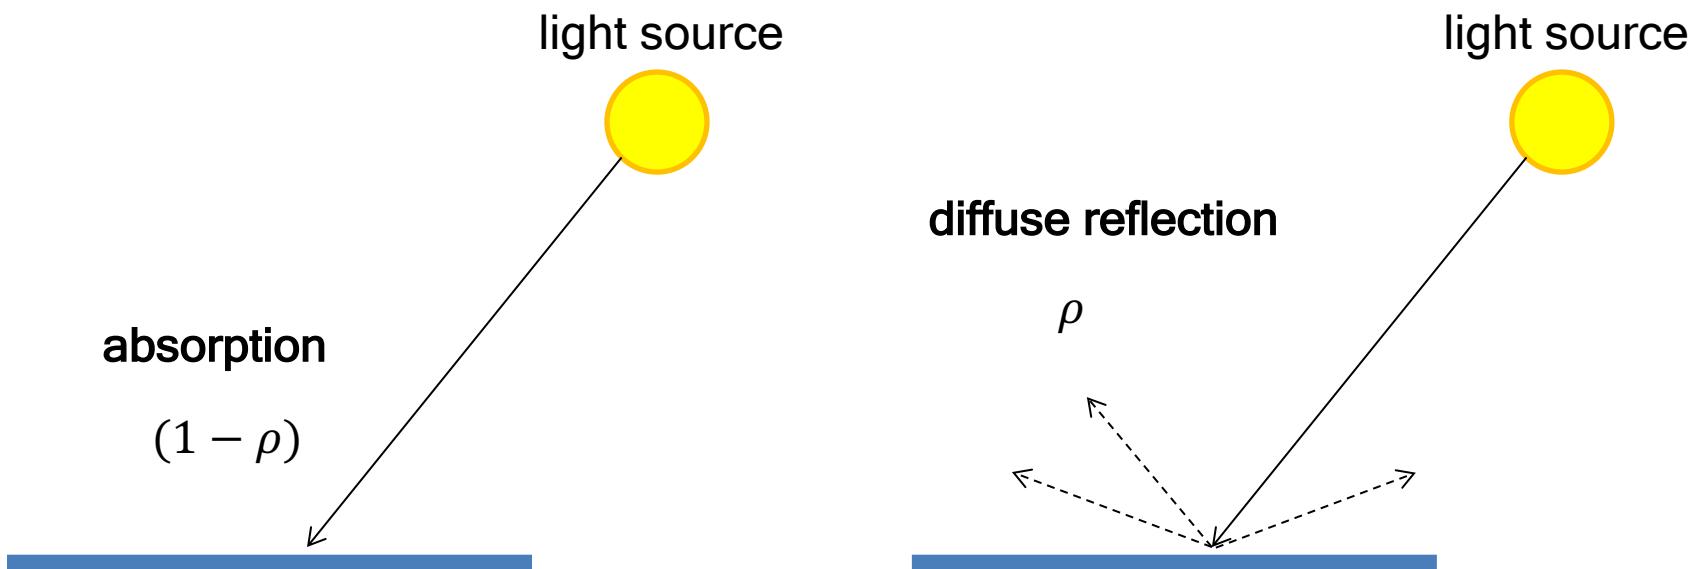


- Diffuse: light scatters in all directions
  - E.g., brick, cloth, rough wood



# Lambertian reflectance model

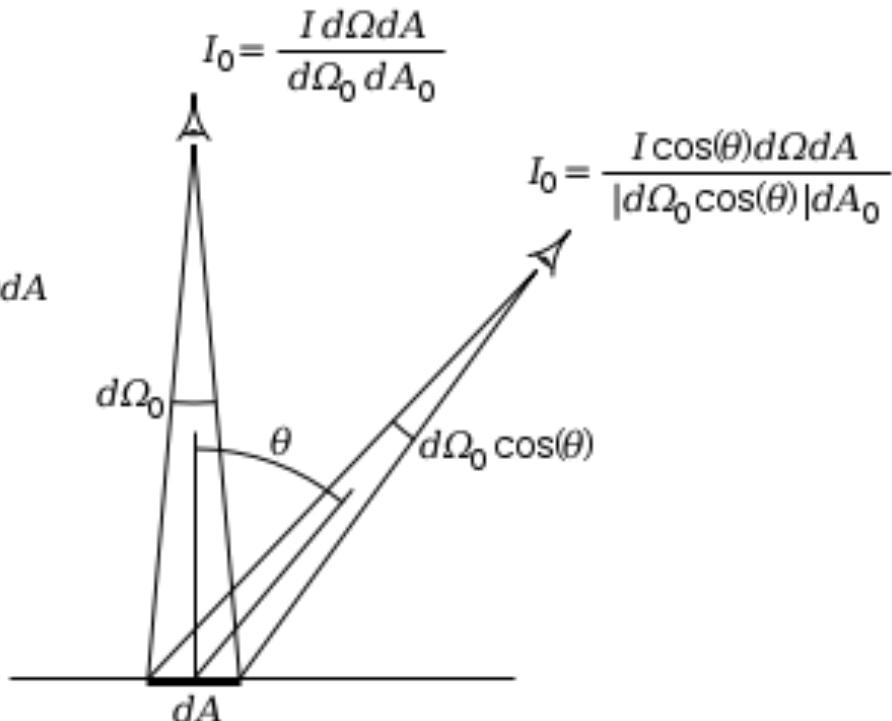
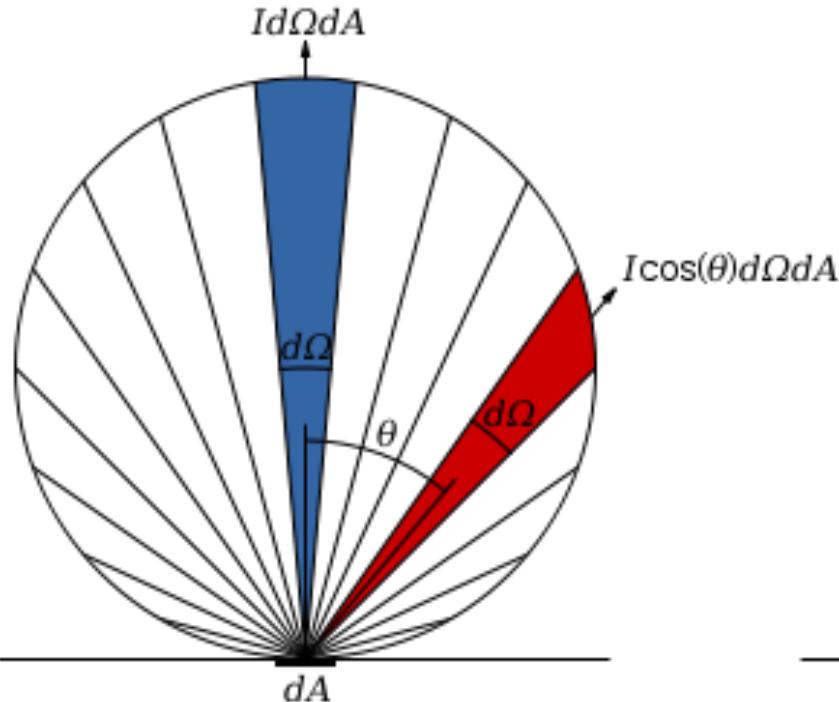
- Some light is absorbed (function of albedo  $\rho$ )
- Remaining light is scattered (diffuse reflection)
- Examples: soft cloth, concrete, matte paints



# Diffuse reflection: Lambert's cosine law

Intensity does *not* depend on viewer angle.

- Amount of reflected light proportional to  $\cos(\theta)$
- Visible solid angle also proportional to  $\cos(\theta)$



# Most surfaces have both specular and diffuse components

- Specularity = spot where specular reflection dominates (typically reflects light source)



Typically, specular component is small

# Intensity and Surface Orientation

Intensity depends on illumination angle because less light comes in at oblique angles.

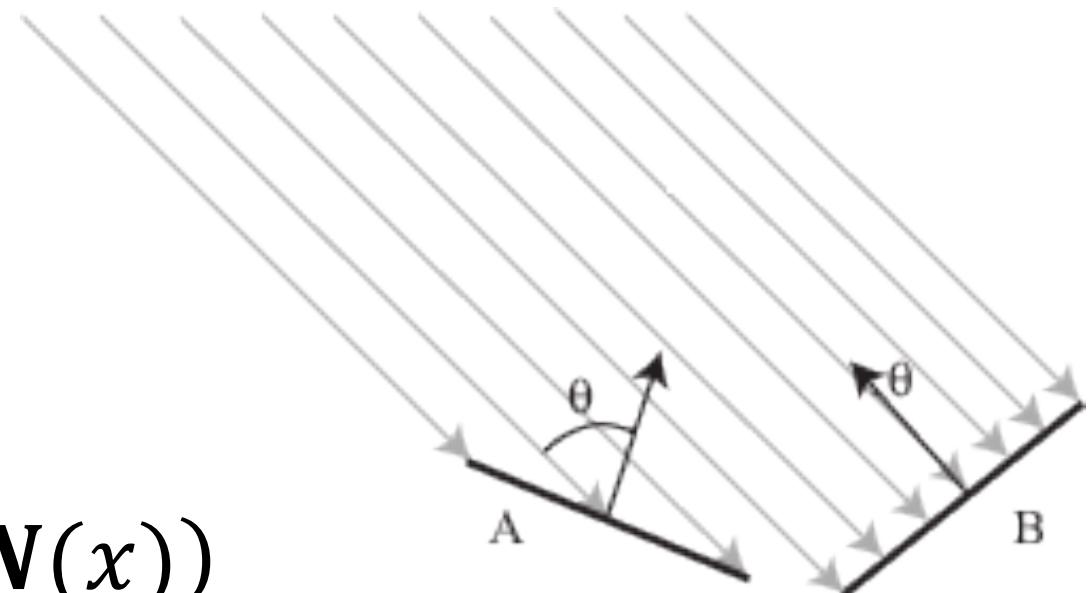
$\rho$  = albedo

$S$  = directional source

$N$  = surface normal

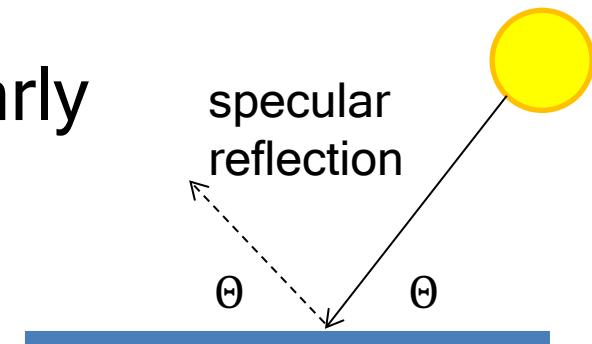
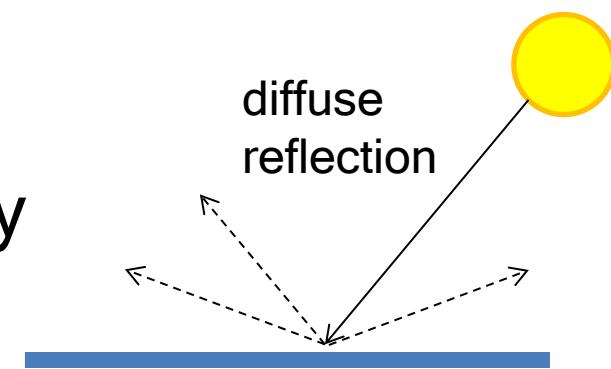
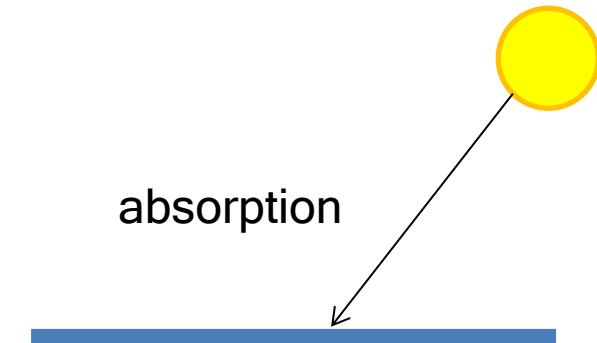
$I$  = reflected intensity

$$I(x) = \rho(x)(S \cdot N(x))$$



# Recap

- When light hits a typical surface
  - Some light is absorbed ( $1-\rho$ )
    - More absorbed for low albedos
  - Some light is reflected diffusely
    - Independent of viewing direction
  - Some light is reflected specularly
    - Light bounces off (like a mirror), depends on viewing direction

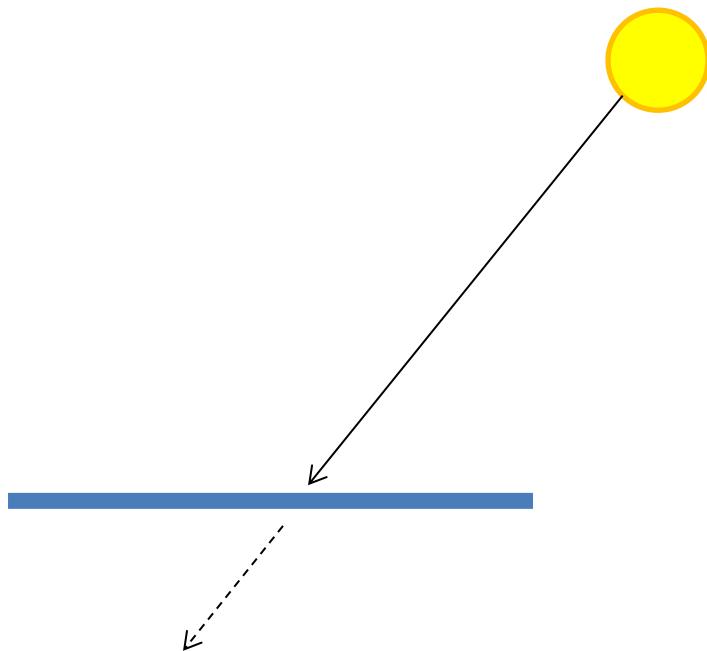


# Other possible effects



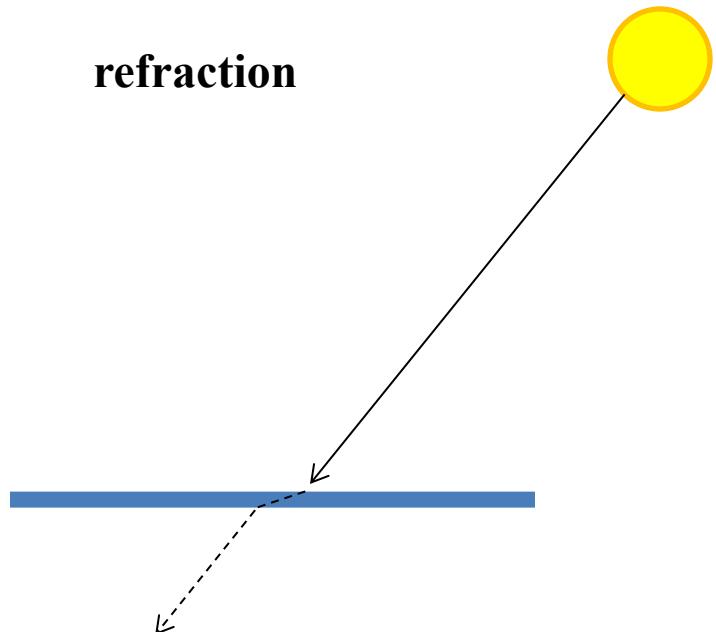
**transparency**

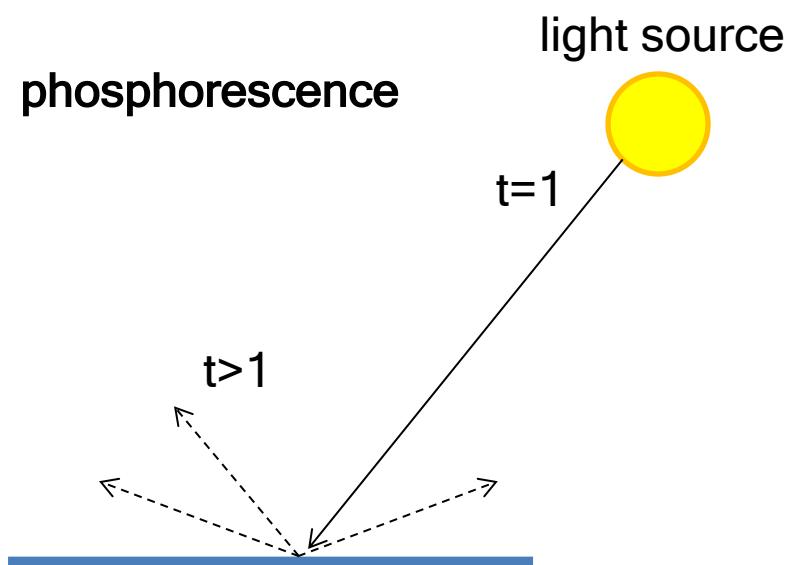
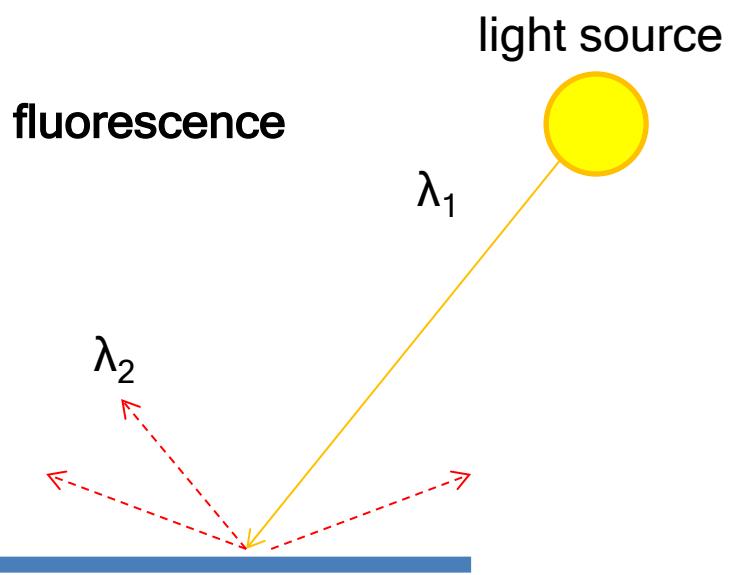
light source



**refraction**

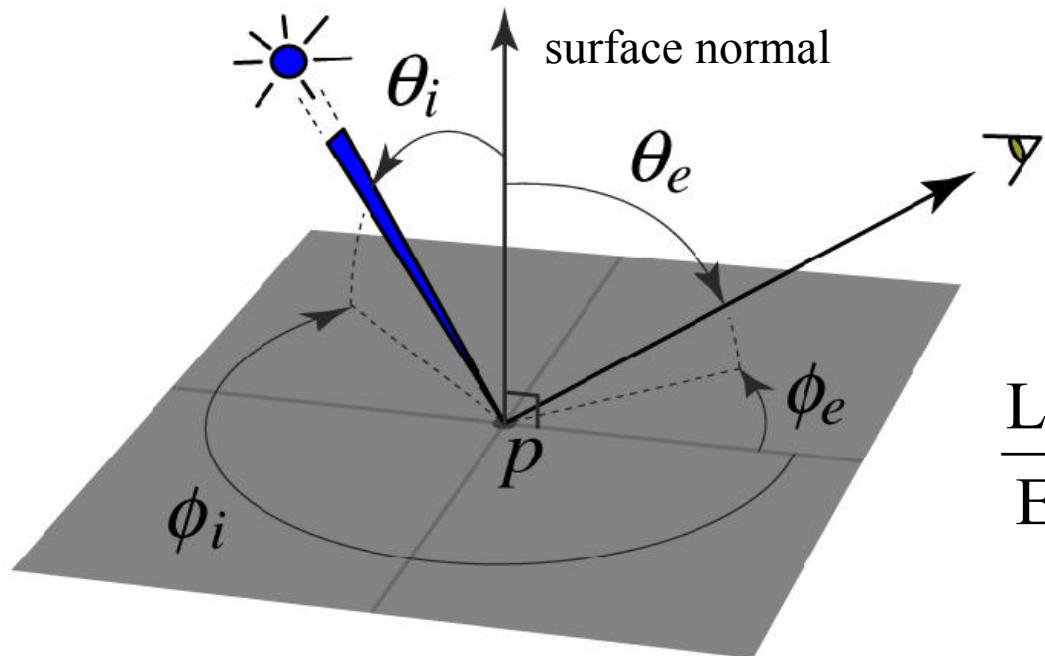
light source





# BRDF: Bidirectional Reflectance Distribution Function

- Model of local reflection that tells how bright a surface appears when viewed from one direction when light falls on it from another
  - Ratio of measured outgoing radiance in direction  $(\theta_e, \phi_e)$  to irradiance from direction  $(\theta_i, \phi_i)$
  - Reciprocal

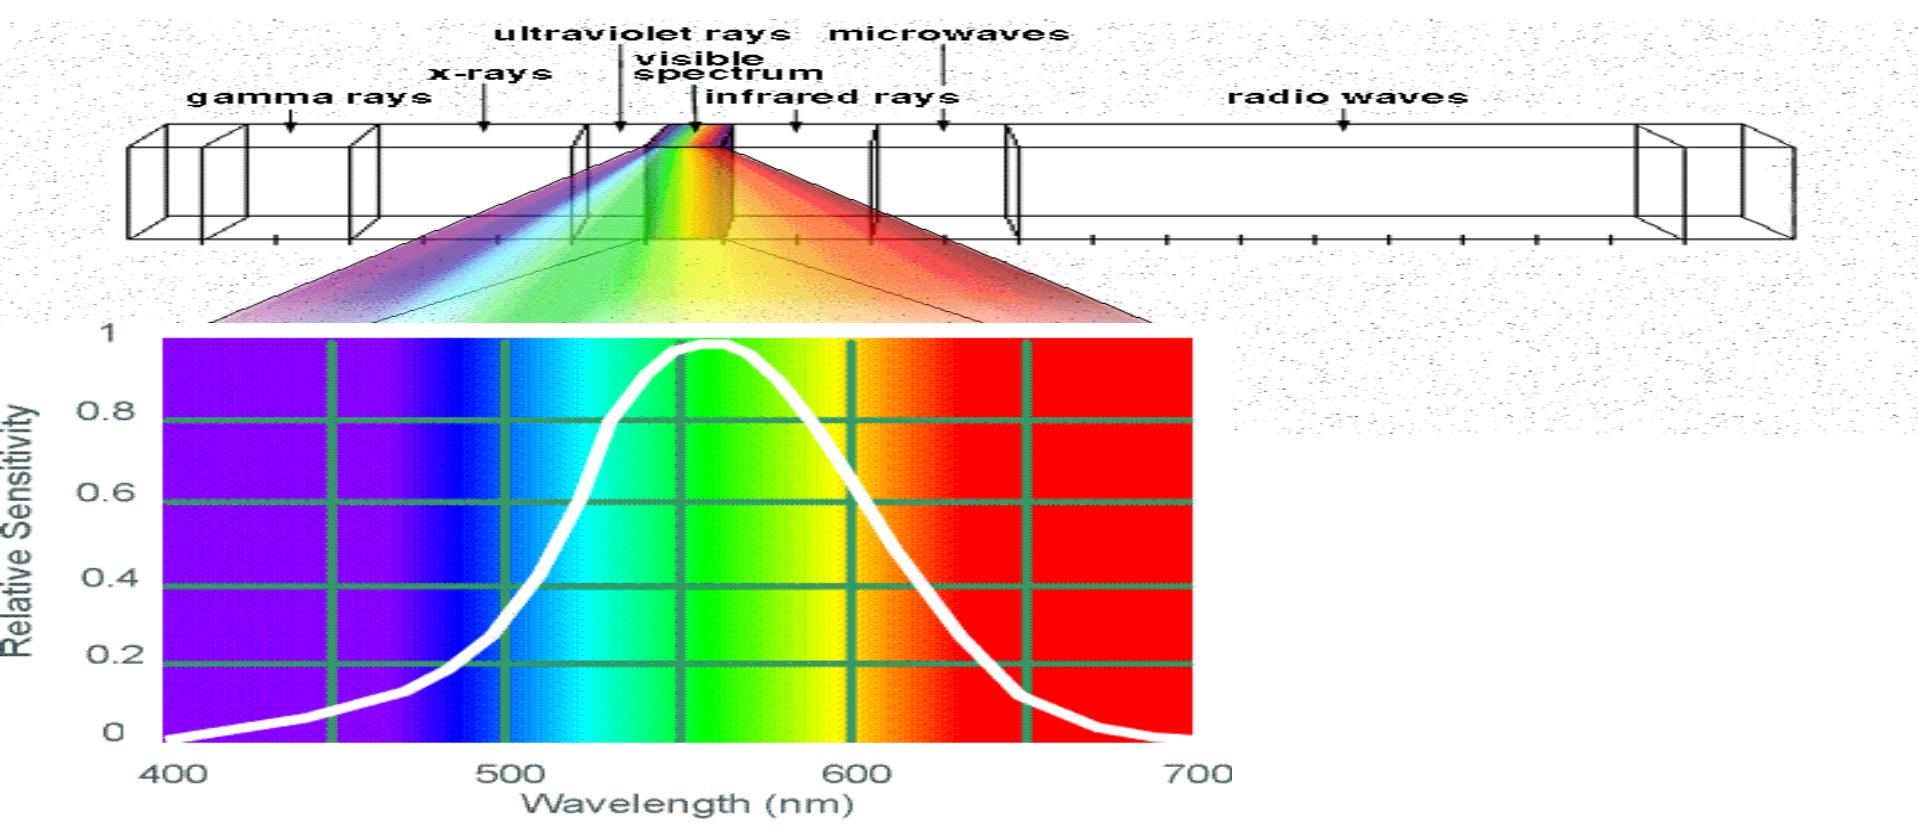


$$\rho(\theta_i, \phi_i, \theta_e, \phi_e; \lambda) =$$

$$\frac{L_e(\theta_e, \phi_e)}{E_i(\theta_i, \phi_i)} = \frac{L_e(\theta_e, \phi_e)}{L_i(\theta_i, \phi_i) \cos \theta_i d\omega}$$

# Color

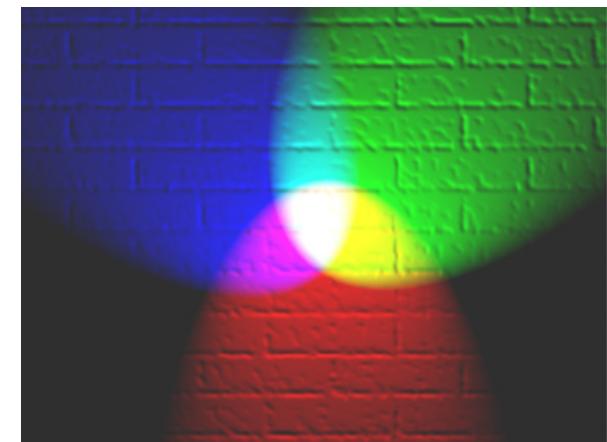
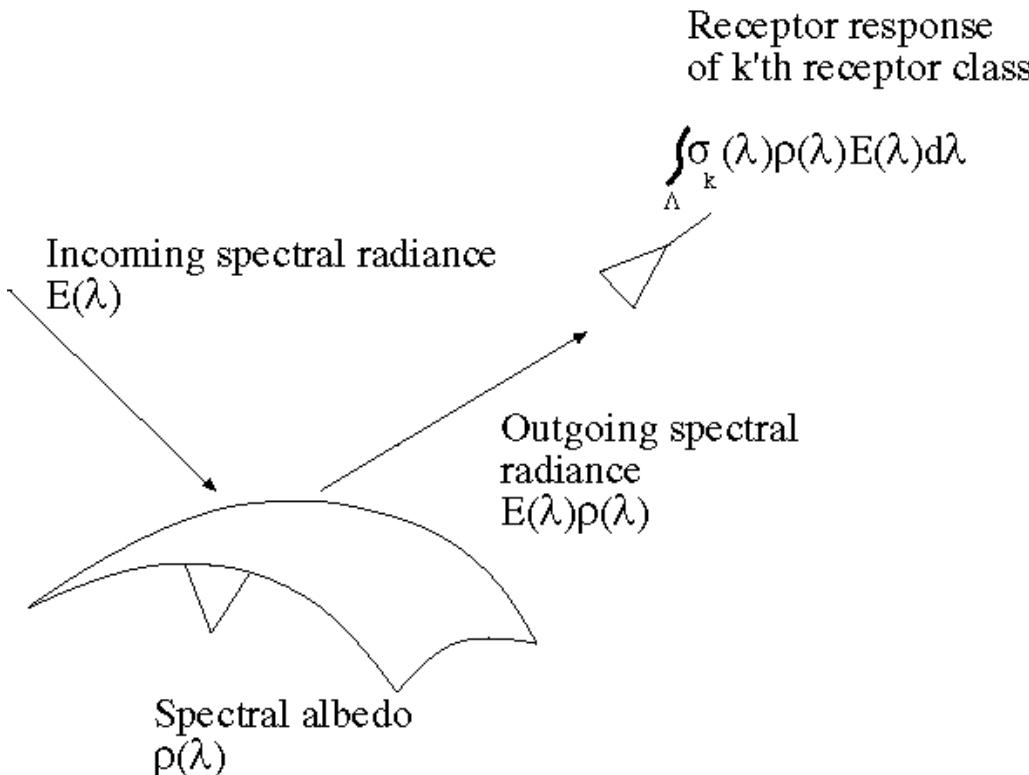
Light is composed of a spectrum of wavelengths



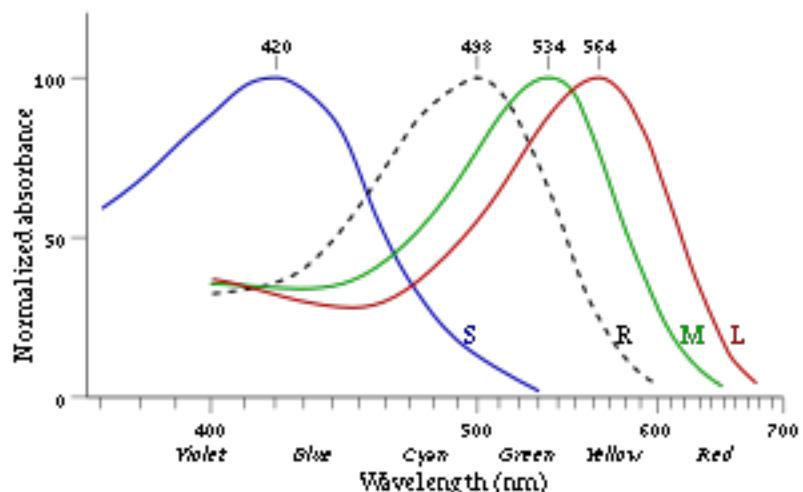
Human Luminance Sensitivity Function

# The color of objects

- Colored light arriving at the camera involves two effects
  - The color of the light source (illumination + inter-reflections)
  - The color of the surface



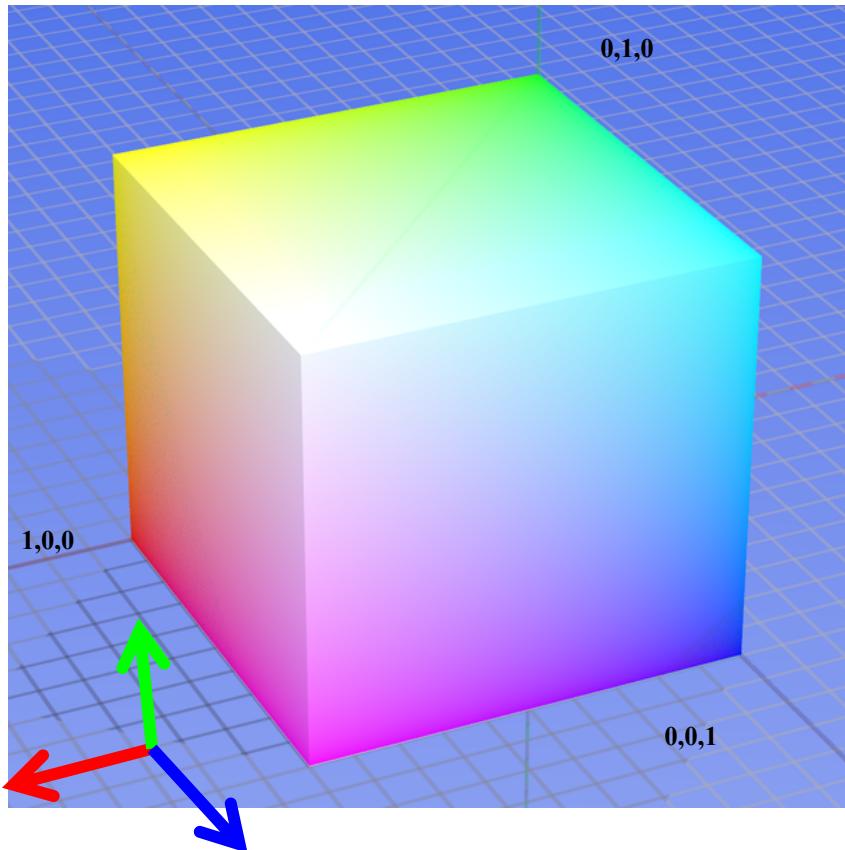
# Human color receptors



- Long (red), Medium (green), and Short (blue) cones, plus intensity rods
- Fun facts
  - “M” and “L” on the X-chromosome
    - That’s why men are more likely to be color blind
  - “L” has high variation, so some women are tetrachromatric
  - Some animals have 1 (night animals), 2 (e.g., dogs), 4 (fish, birds), 5 (pigeons, some reptiles/amphibians), or even 12 (mantis shrimp) types of cones

# Color spaces: RGB

Default color space



**R**  
 $(G=0, B=0)$



**G**  
 $(R=0, B=0)$



**B**  
 $(R=0, G=0)$

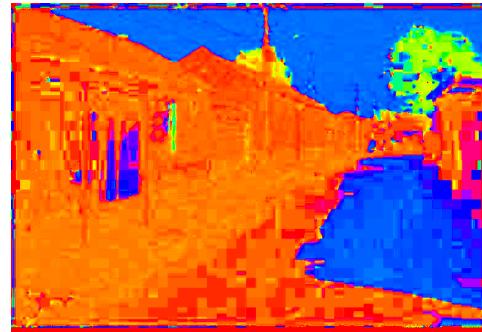
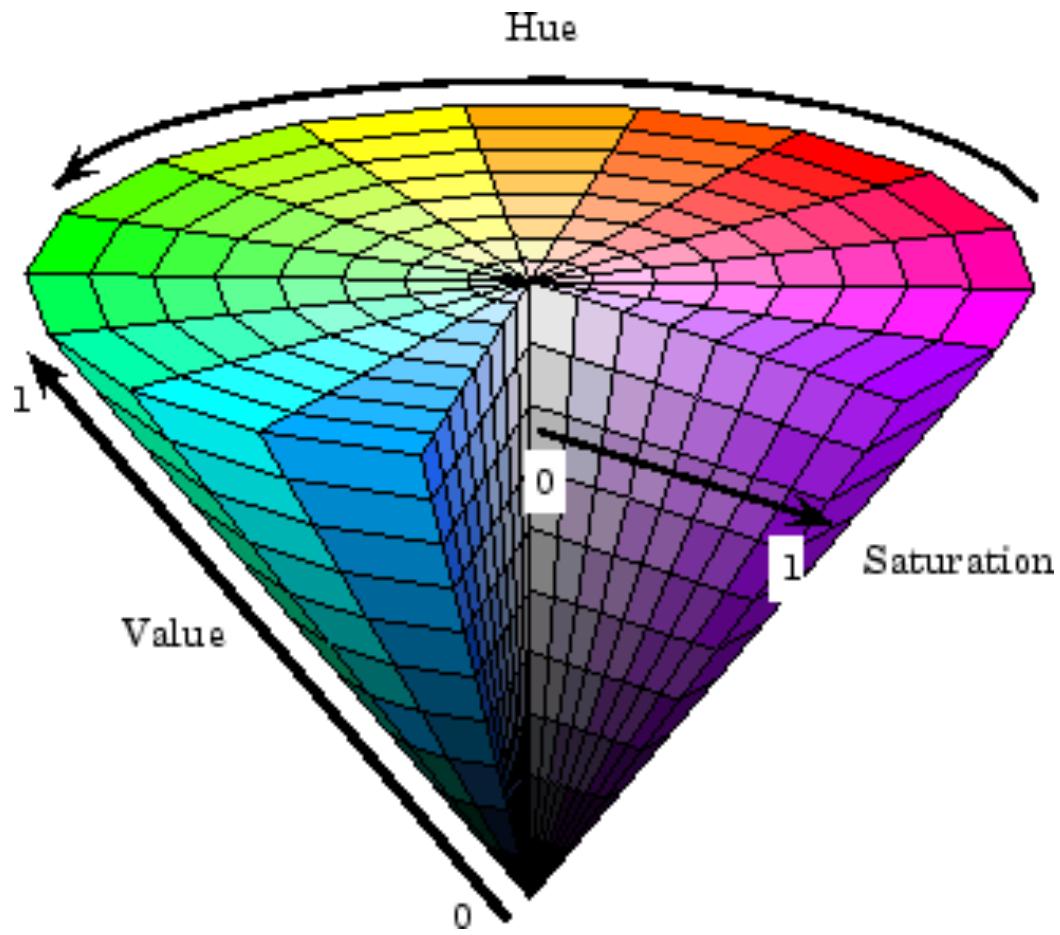
Some drawbacks

- Strongly correlated channels
- Non-perceptual

# Color spaces: HSV



Intuitive color space



H  
(S=1,V=1)



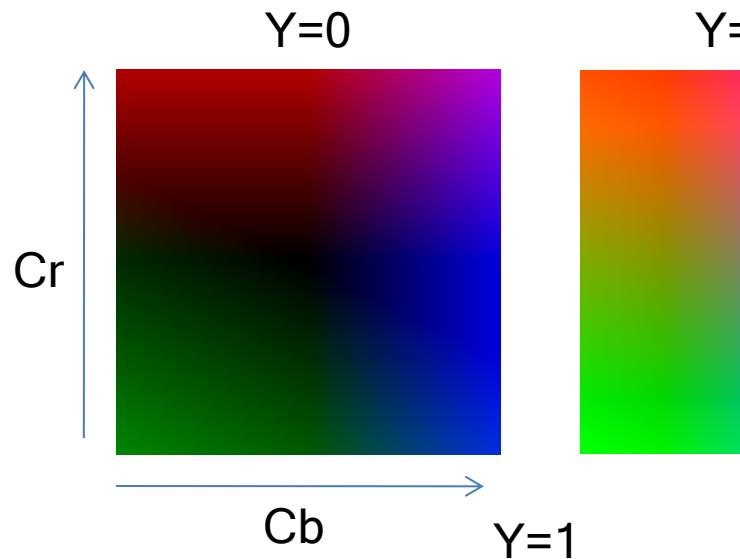
S  
(H=1,V=1)



V  
(H=1,S=0)

# Color spaces: YCbCr

Fast to compute, good for compression, used by TV



**Y**  
( $Cb=0.5, Cr=0.5$ )



**Cb**  
( $Y=0.5, Cr=0.5$ )



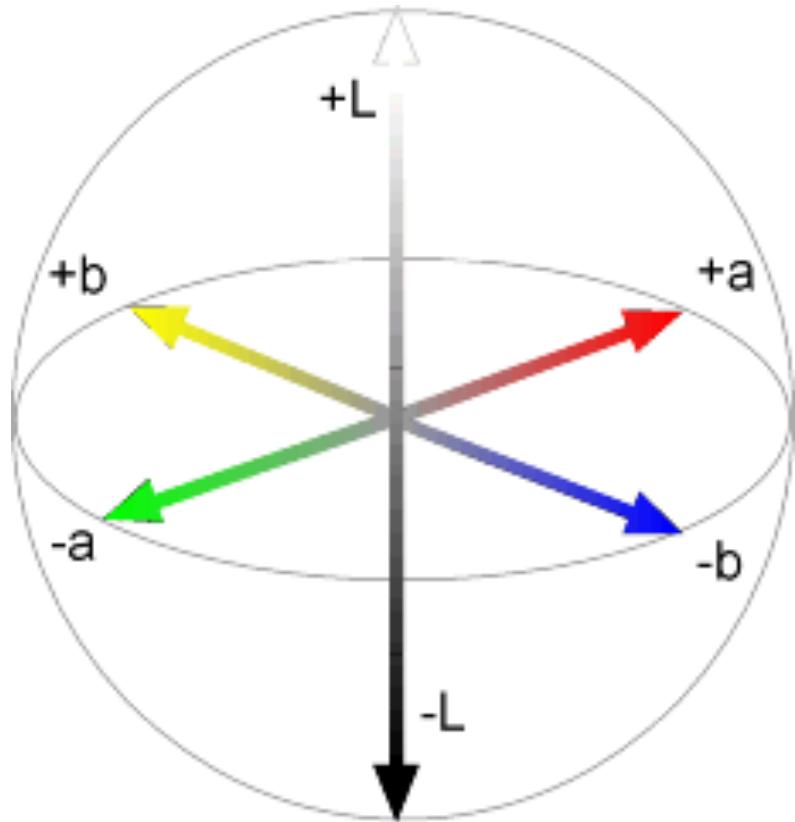
**Cr**  
( $Y=0.5, Cb=0.5$ )

$$Y' = 16 + \frac{65.738 \cdot R'_D}{256} + \frac{129.057 \cdot G'_D}{256} + \frac{25.064 \cdot B'_D}{256}$$
$$C_B = 128 + \frac{-37.945 \cdot R'_D}{256} - \frac{74.494 \cdot G'_D}{256} + \frac{112.439 \cdot B'_D}{256}$$
$$C_R = 128 + \frac{112.439 \cdot R'_D}{256} - \frac{94.154 \cdot G'_D}{256} - \frac{18.285 \cdot B'_D}{256}$$

# Color spaces: CIE L\*a\*b\*



“Perceptually uniform” color space



Luminance = brightness  
Chrominance = color



Which contains more information?

- (a) **intensity** (1 channel)
- (b) **chrominance** (2 channels)

# Most information in intensity



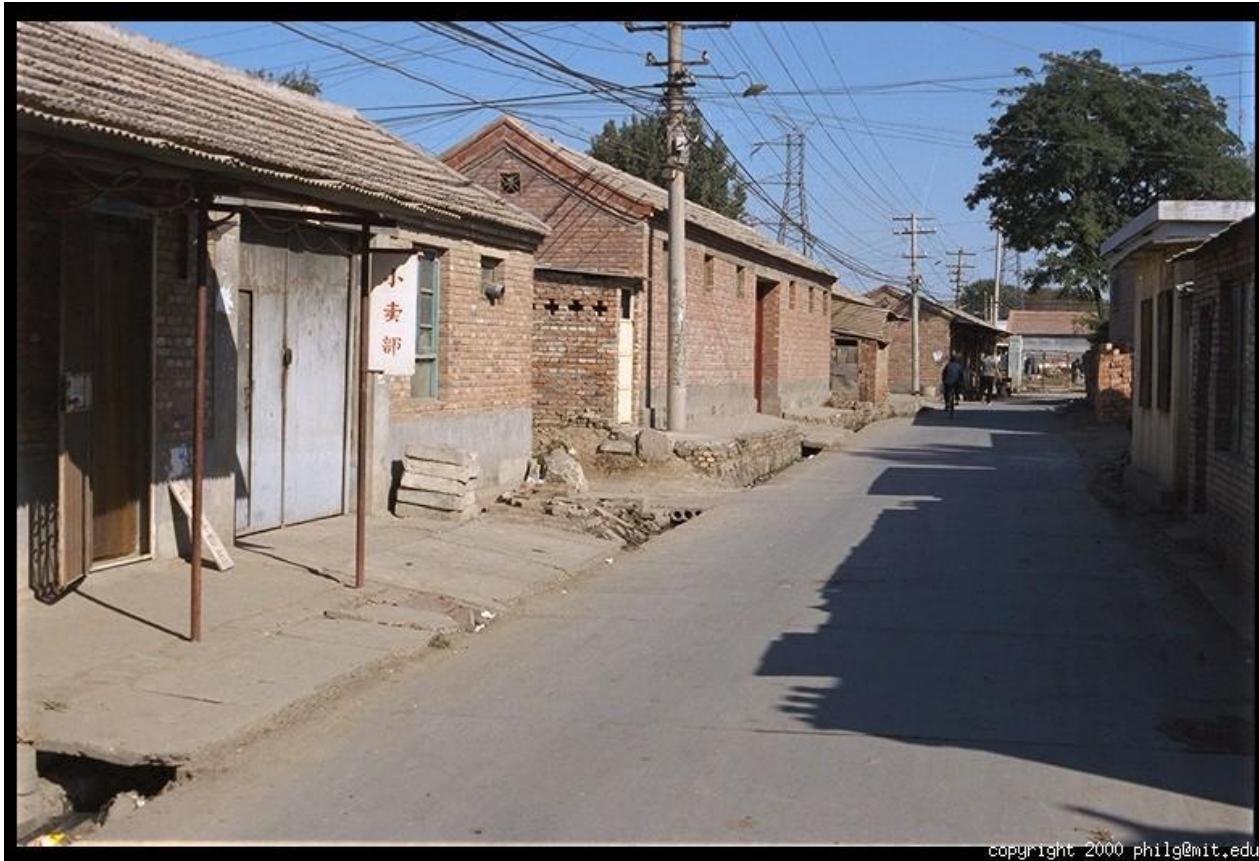
Only color shown – constant intensity

# Most information in intensity



Only intensity shown – constant color

# Most information in intensity

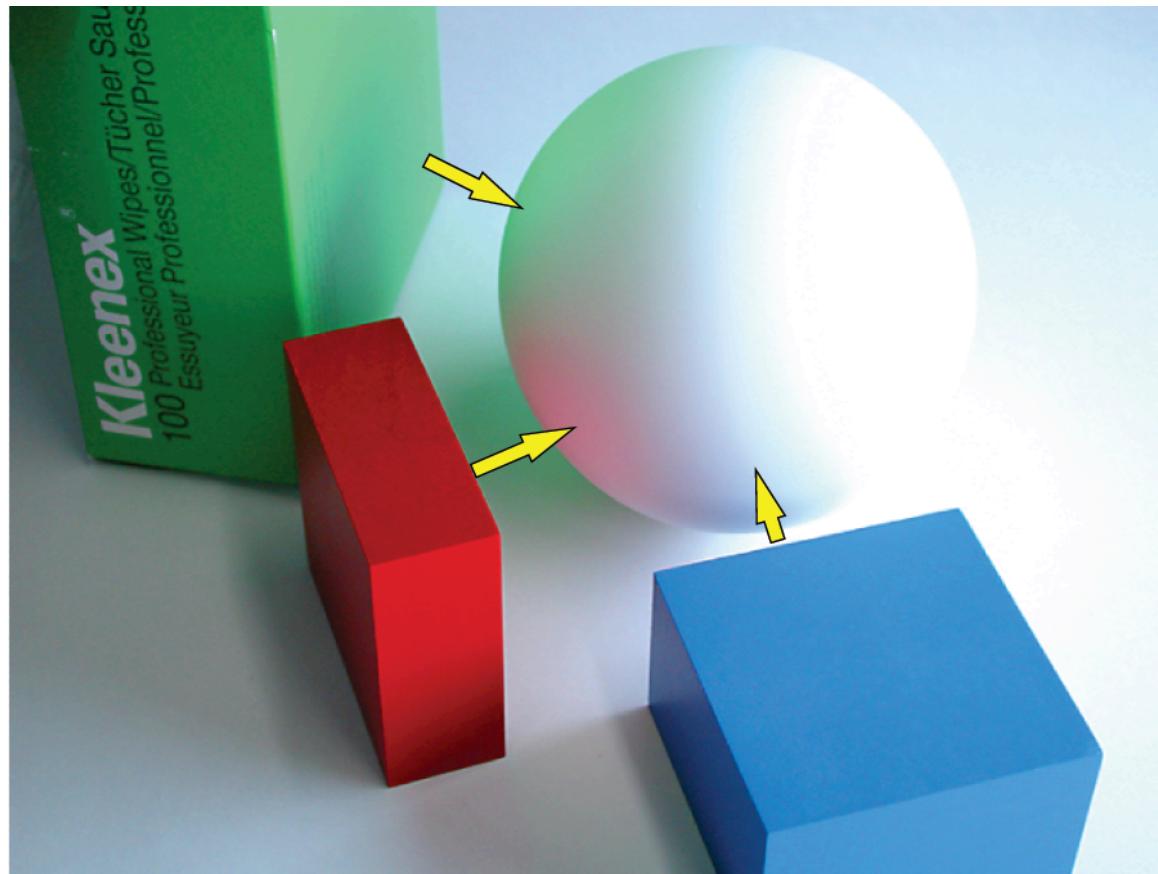


copyright 2000 philg@mit.edu

Original image

# So far: light $\rightarrow$ surface $\rightarrow$ camera

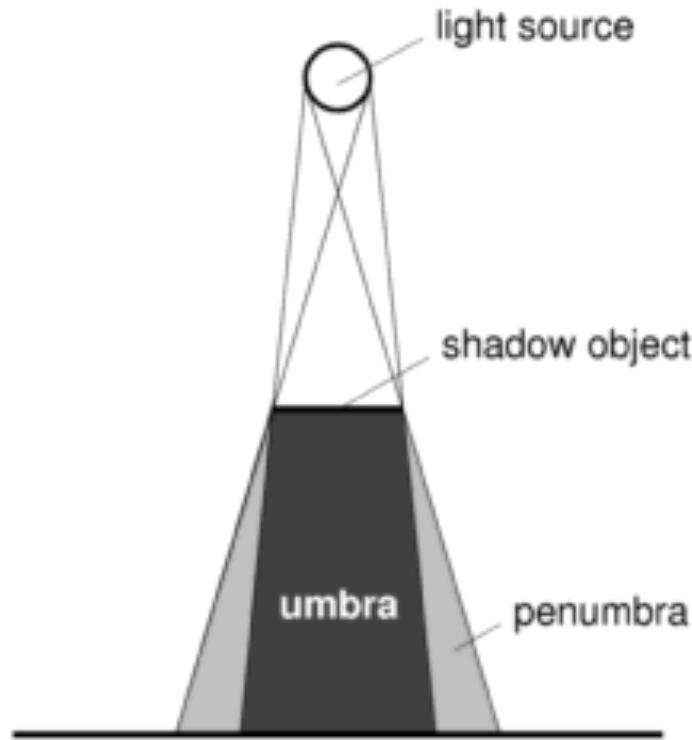
- Called a local illumination model
- But much light comes from surrounding surfaces



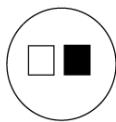
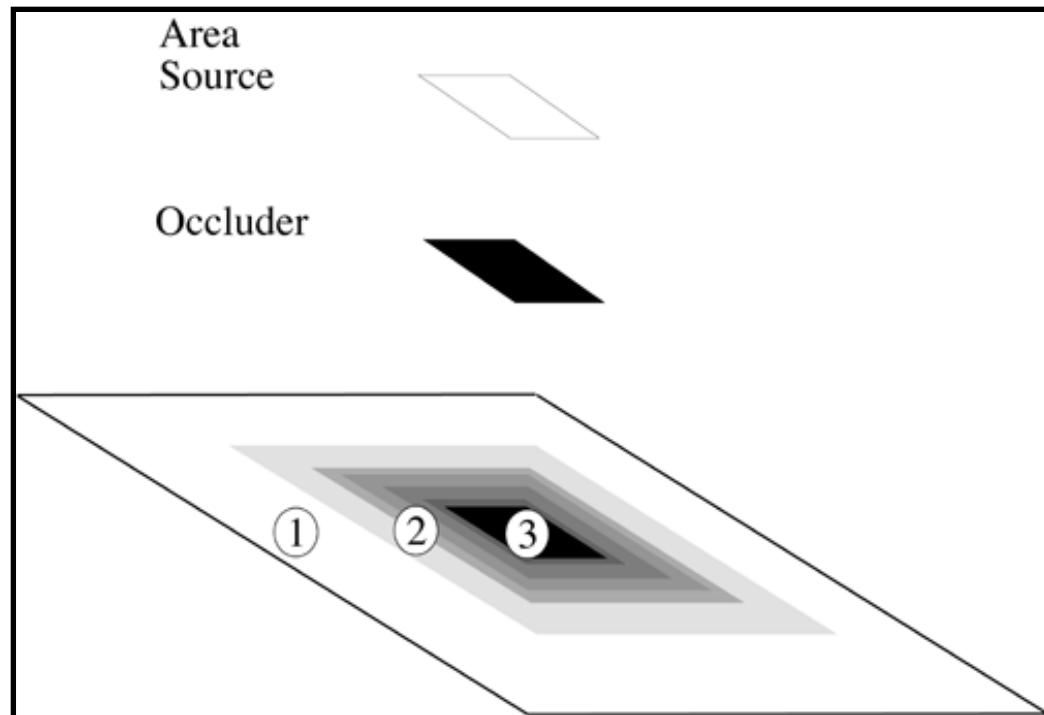
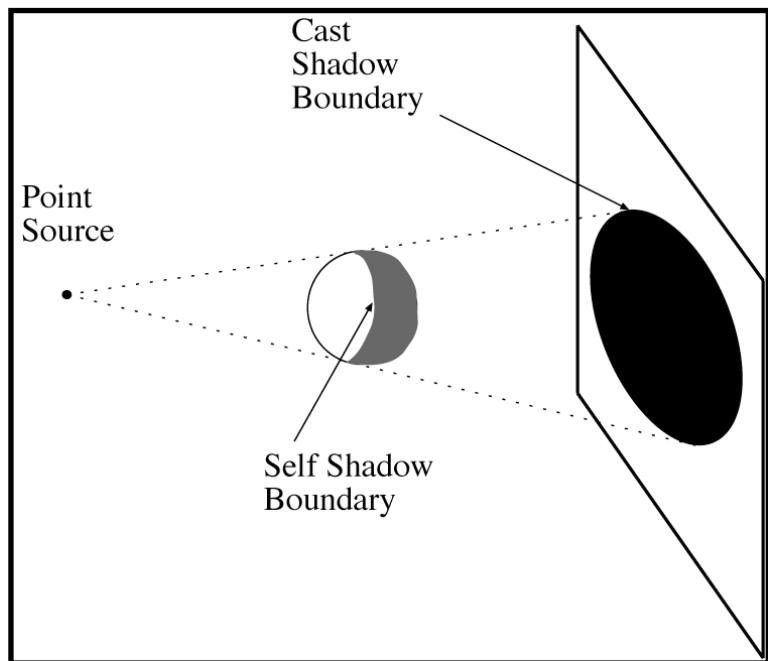
From Koenderink slides on image texture and the flow of light

# Scene surfaces also cause shadows

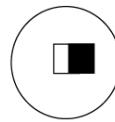
- Shadow: reduction in intensity due to a blocked source



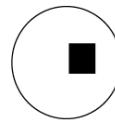
# Shadows



1



2



3

# Models of light sources

- Distant point source
  - One illumination direction
  - E.g., sun
- Area source
  - E.g., white walls, diffuser lamps, sky
- Ambient light
  - Substitute for dealing with interreflections
- Global illumination model
  - Account for interreflections in modeled scene

# The plight of the poor pixel

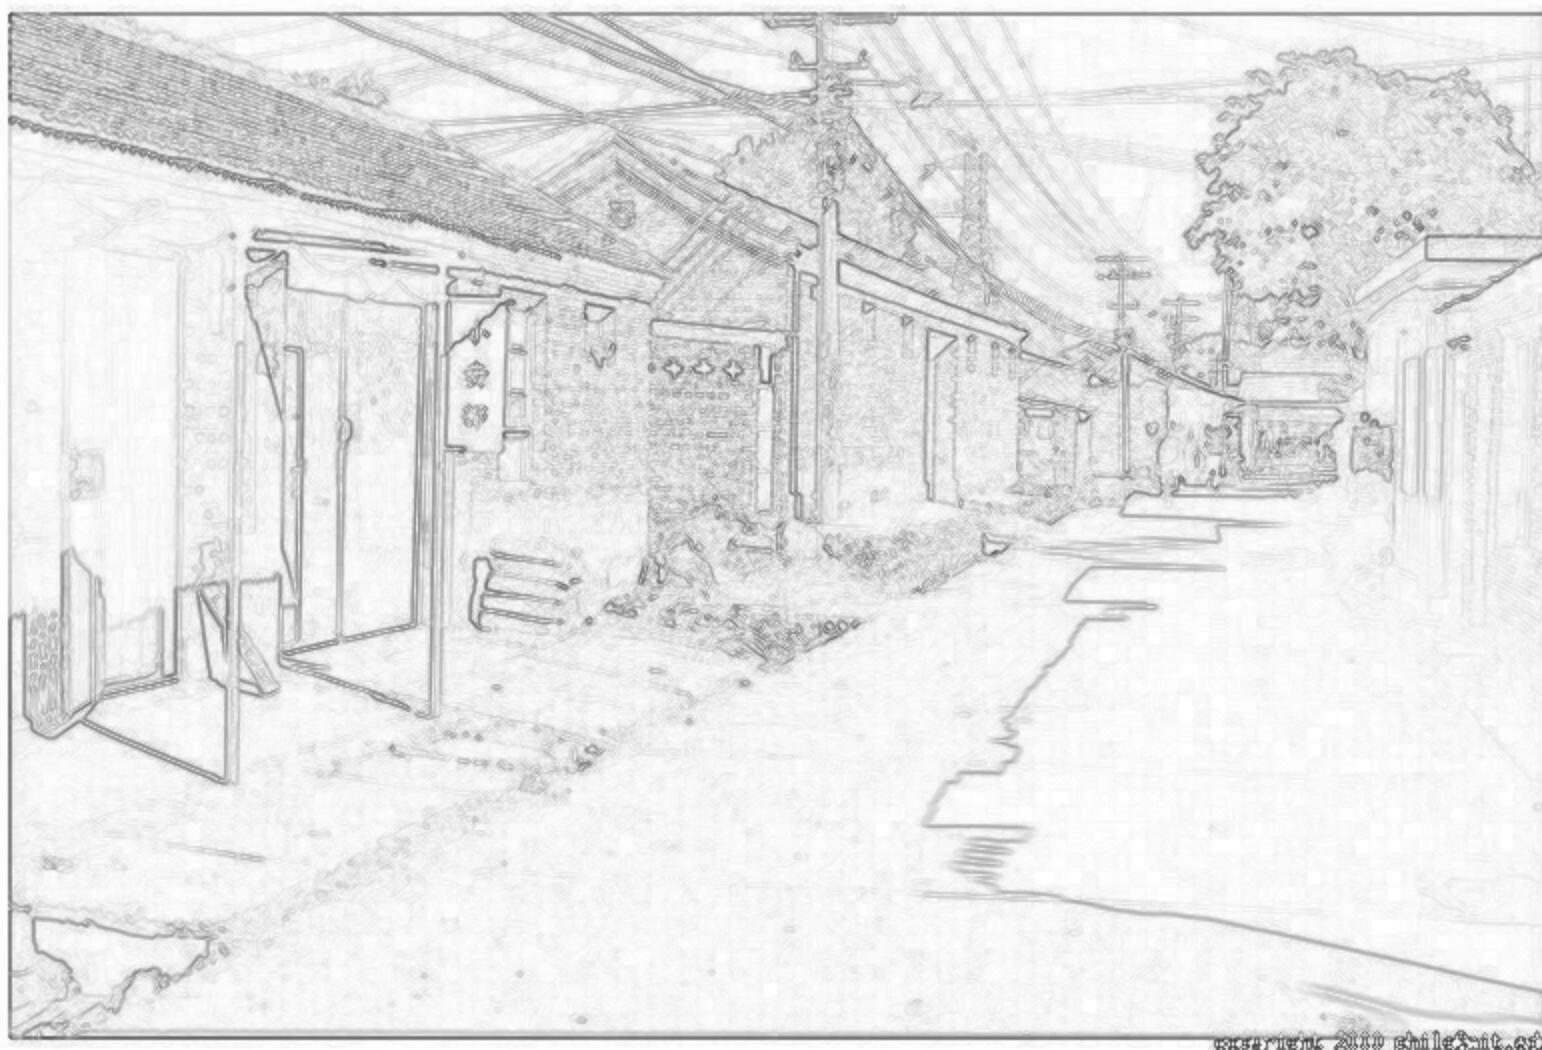
- A pixel's brightness is determined by
  - Light source (strength, direction, color)
  - Surface orientation
  - Surface material and albedo
  - Reflected light and shadows from surrounding surfaces
  - Gain on the sensor
- A pixel's brightness tells us nothing by itself

# And yet we can interpret images...



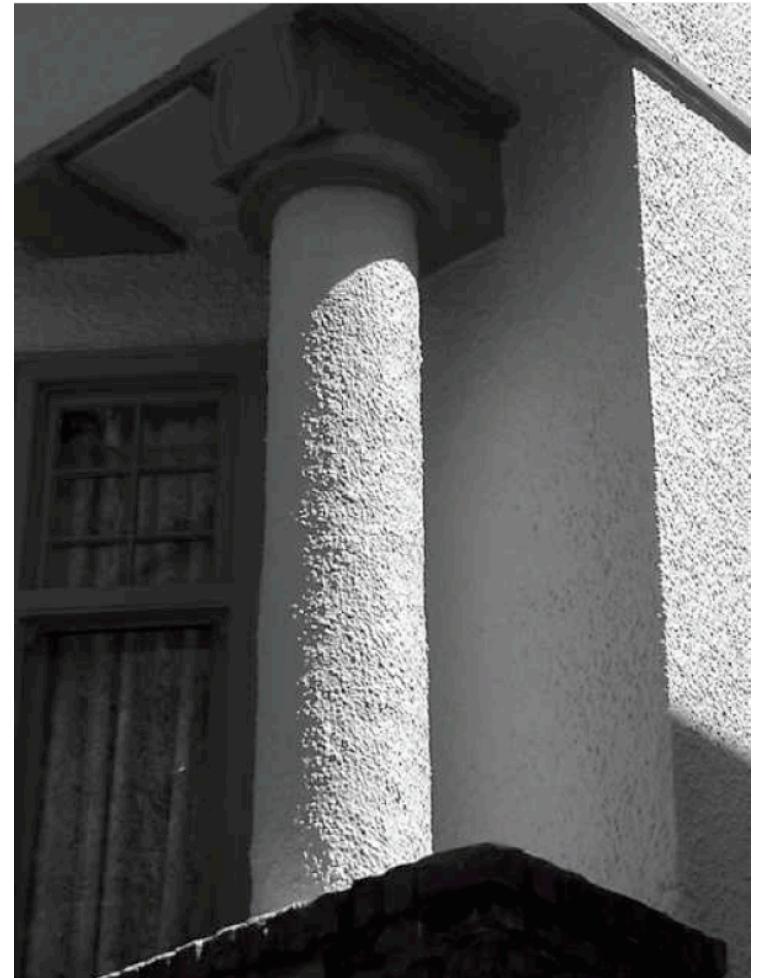
- Key idea: for nearby scene points, most factors do not change much
- The information is mainly contained in *local differences* of brightness

# Darkness = Large Difference in Neighboring Pixels



# What differences in intensity tell us about shape

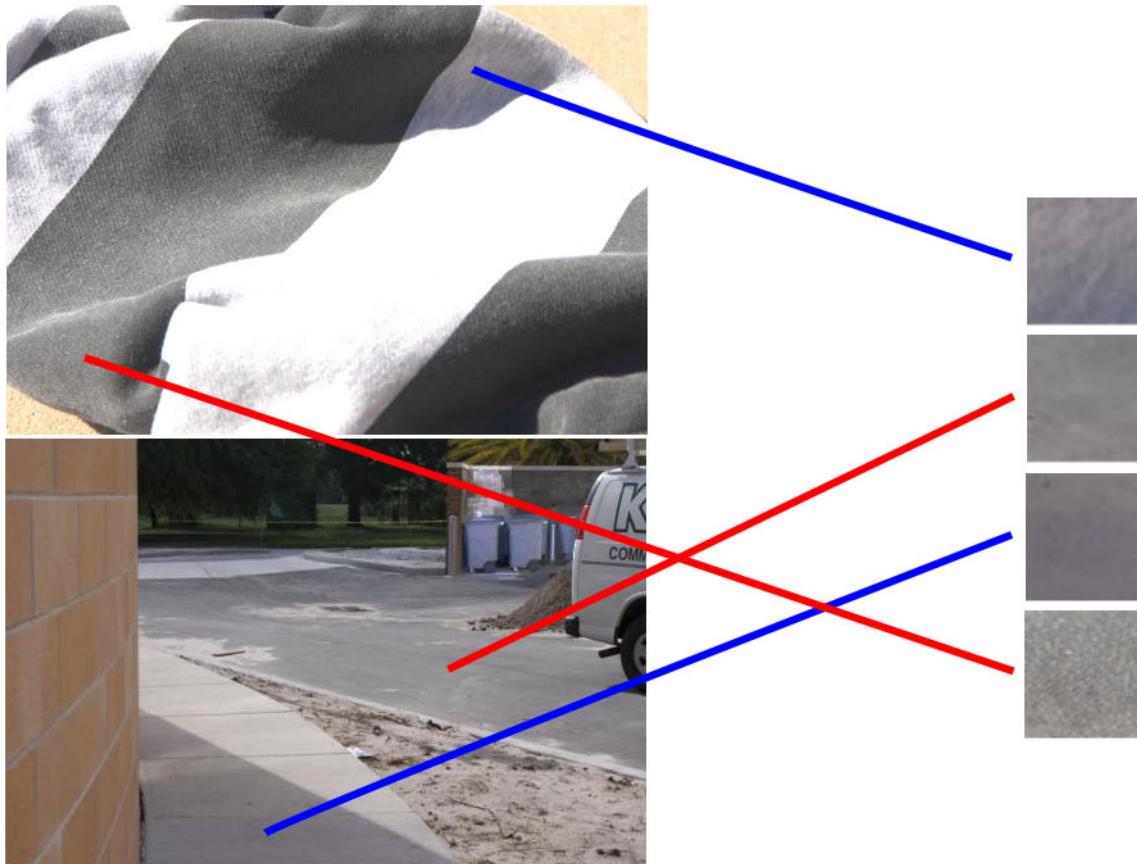
- Changes in surface normal
- Texture
- Proximity
- Indents and bumps
- Grooves and creases



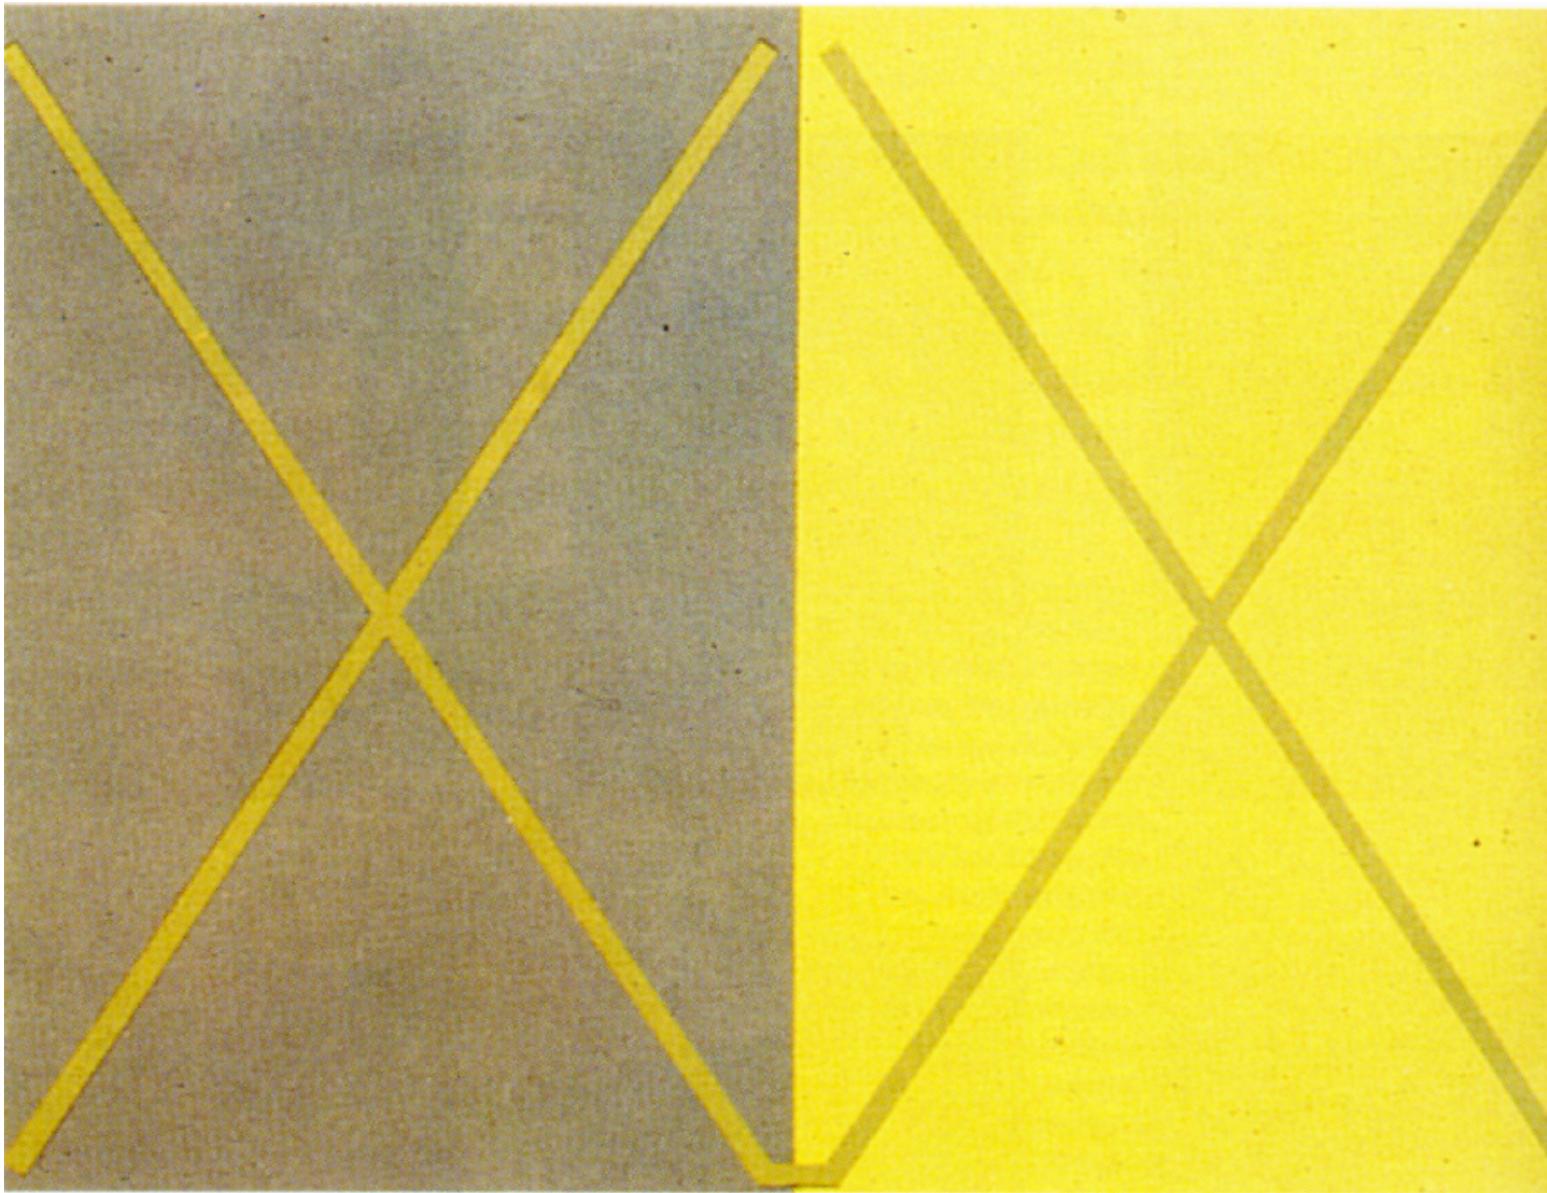
Photos Koenderink slides on image texture and the flow of light

# Color constancy

- Interpret surface in terms of albedo or “true color”, rather than observed intensity
  - Humans are good at it
  - Computers are not nearly as good

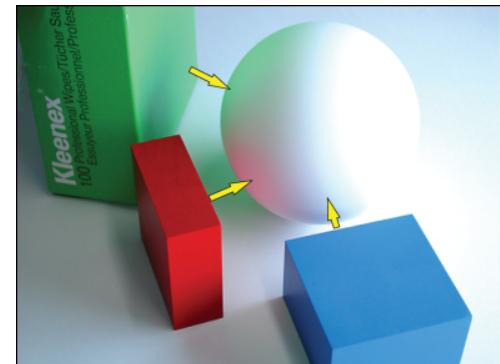
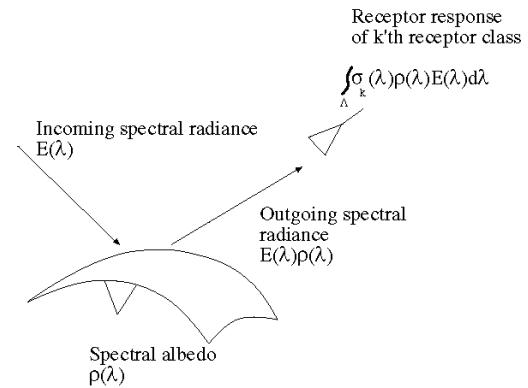


# One source of constancy: local comparisons



# Things to remember

- Important terms: diffuse/specular reflectance, albedo, umbra/penumbra
- Observed intensity depends on light sources, geometry/material of reflecting surface, surrounding objects, camera settings
- Objects cast light and shadows on each other
- Differences in intensity are primary cues for shape



# Pixels and Linear Filters

Slides by D. Hoiem

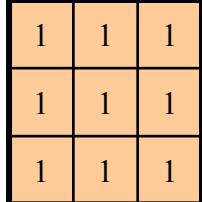
# Image filtering

- Image filtering: for each pixel, compute function of local neighborhood and output a new value
  - Same function applied at each position
  - Output and input image are typically the same size

# Image filtering

- Linear filtering: function is a weighted sum/difference of pixel values
- Really important
  - Enhance images
    - Denoise, smooth, increase contrast, etc.
  - Extract information from images
    - Texture, edges, distinctive points, etc.
  - Detect patterns
    - Template matching

# Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$


A 3x3 matrix where every element is 1/9. The matrix is enclosed in a black border.

$$f[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h[.,.]$$

	0	10	20	30	30				

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz

# Image filtering

$$g[\cdot, \cdot] \quad \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Credit: S. Seitz

# Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect  
(remove sharp features)

$$g[\cdot, \cdot]$$

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

# Basic gradient filters

Horizontal Gradient

0	0	0
-1	0	1
0	0	0

or

-1	0	1
----	---	---

Vertical Gradient

0	-1	0
0	0	0
0	1	0

or

-1
0
1

# Examples

Write as filtering operations, plus some pointwise operations: +, -, .\*, >

1. Sum of four adjacent neighbors plus 1

$$out(m, n) = 1 + \sum_{k,l \in \{-1,1\}} in(m+k, n+l)$$

2. Sum of squared values of 3x3 windows around each pixel:

$$out(m, n) = \sum_{k,l \in \{-1,0,1\}} in(m+k, n+l)^2$$

3. Center pixel value is larger than the average of the pixel values to the left and right:

$$out(m, n) = 1 \text{ if } in(m, n) > (in(m, n-1) + in(m, n+1)) / 2$$

$$out(m, n) = 0 \text{ if } in(m, n) \leq (in(m, n-1) + in(m, n+1)) / 2$$

# Filtering vs. Convolution

- 2d filtering
  - $h = \text{filter2}(g, f);$  or  
 $h = \text{imfilter}(f, g);$
  - $$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$
- 2d convolution
  - $h = \text{conv2}(g, f);$
  - $$h[m, n] = \sum_{k,l} g[k, l] f[m - k, n - l]$$

# Key properties of linear filters

## Linearity:

$$\text{filter}(f_1 + f_2) = \text{filter}(f_1) + \text{filter}(f_2)$$

**Shift invariance:** same behavior regardless of pixel location

$$\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$$

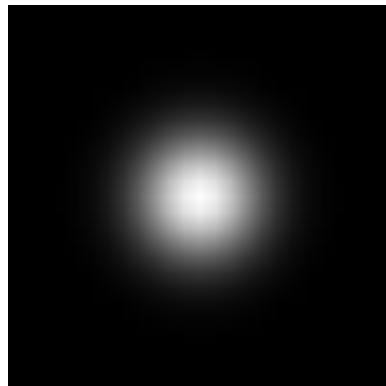
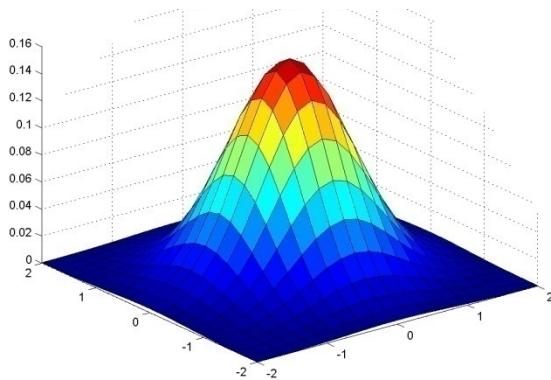
Any linear, shift-invariant operator can be represented as a convolution

# More properties

- Commutative:  $a * b = b * a$ 
  - Conceptually no difference between filter and signal
- Associative:  $a * (b * c) = (a * b) * c$ 
  - Often apply several filters one after another:  $((a * b_1) * b_2) * b_3$
  - This is equivalent to applying one filter:  $a * (b_1 * b_2 * b_3)$
- Distributes over addition:  $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out:  $ka * b = a * kb = k(a * b)$
- Identity: unit impulse  $e = [0, 0, 1, 0, 0]$ ,  $a * e = a$

# Important filter: Gaussian

- Spatially-weighted average



0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

$5 \times 5, \sigma = 1$

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

# Gaussian filters

- Remove “high-frequency” components from the image (low-pass filter)
  - Images become more smooth
- Convolution with self is another Gaussian
  - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
  - Convolving two times with Gaussian kernel of width  $\sigma$  is same as convolving once with kernel of width  $\sigma\sqrt{2}$
- *Separable* kernel
  - Factors into product of two 1D Gaussians

# Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

# Separability example

2D filtering  
(center location only)

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix}$$

The filter factors  
into a product of 1D  
filters:

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} = \begin{matrix} 1 \\ 2 \\ 1 \end{matrix} \times \begin{matrix} 1 & 2 & 1 \end{matrix}$$

Perform filtering  
along rows:

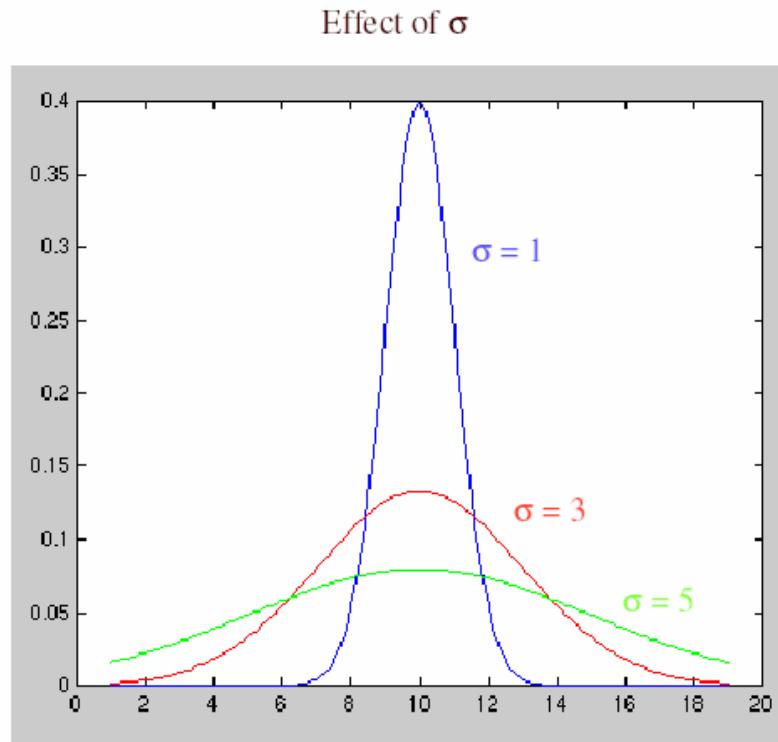
$$\begin{matrix} 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} = \begin{matrix} 11 \\ 18 \\ 18 \end{matrix}$$

Followed by filtering  
along the remaining column:

# Practical matters

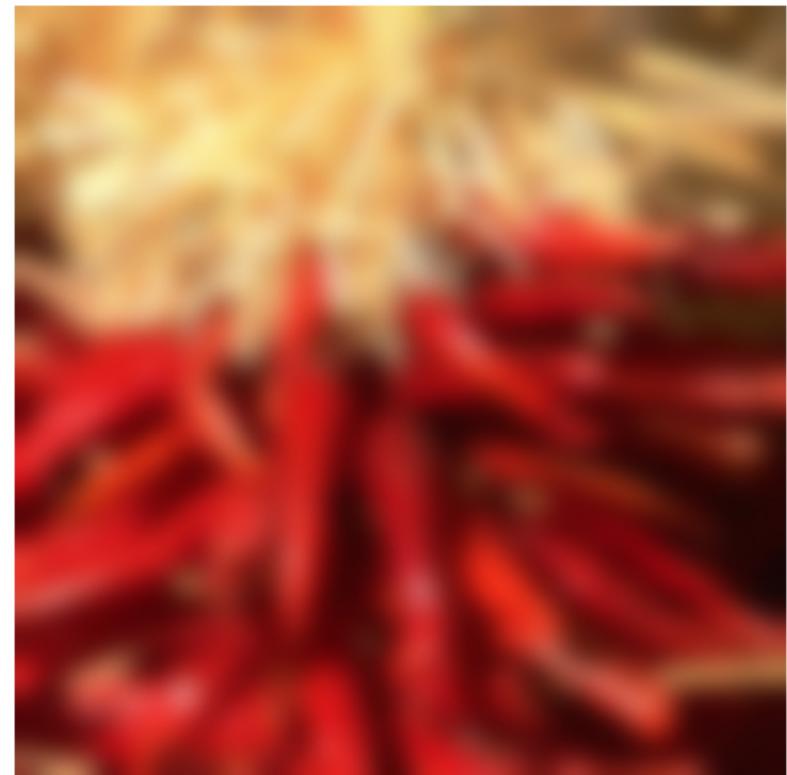
## How big should the filter be?

- Values at edges should be near zero  $\leftarrow$  important!
- Rule of thumb for Gaussian: set filter half-width to about  $3 \sigma$



# Practical matters

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - copy edge
    - reflect across edge



# Practical matters

- What is the size of the output?
- MATLAB: `filter2(g, f, shape)`
  - *shape* = ‘full’: output size is sum of sizes of *f* and *g*
  - *shape* = ‘same’: output size is same as *f*
  - *shape* = ‘valid’: output size is difference of sizes of *f* and *g*

