



# CS 558:

## Computer Vision

### 3rd Set of Notes

Instructor: Enrique Dunn

Webpage: [www.cs.stevens.edu/~edunn](http://www.cs.stevens.edu/~edunn)

E-mail: [edunn@stevens.edu](mailto:edunn@stevens.edu)

Office: North Bldg 219

# Overview

- Denoising
  - Based on slides by S. Lazebnik
- Edge detection
  - Based on slides by S. Lazebnik and D. Hoiem
- Feature extraction: Corners
  - Based on slides by S. Lazebnik
- Sampling images
  - Based on slides by D. Hoiem

# Image denoising

- How can we reduce noise in a photograph?



# Moving average

- Let's replace each pixel with a *weighted* average of its neighborhood
- The weights are called the *filter kernel*
- What are the weights for the average of a 3x3 neighborhood?

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

“box filter”

# Noise



Original



Salt and pepper noise



Impulse noise

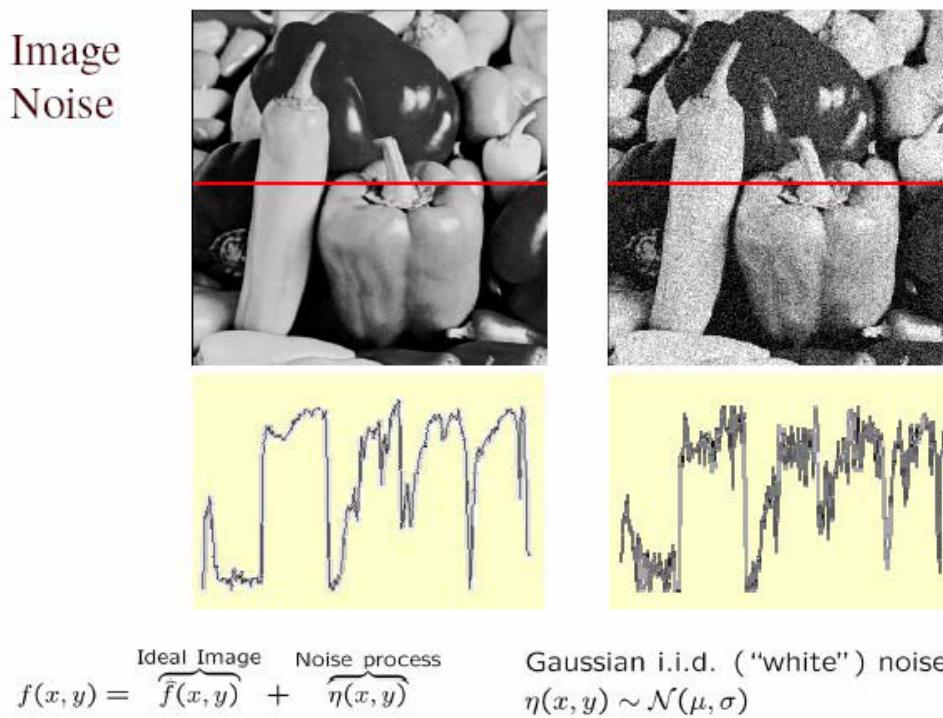


Gaussian noise

- **Salt and pepper noise:** contains random occurrences of black and white pixels
- **Impulse noise:** contains random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution

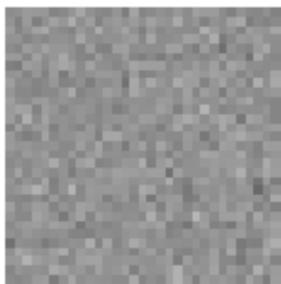
# Gaussian noise

- Mathematical model: sum of many independent factors
- Good for small standard deviations
- Assumption: independent, zero-mean noise

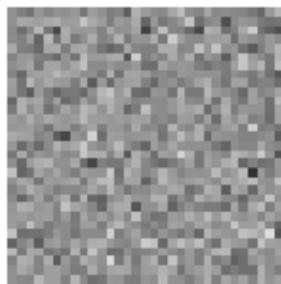


# Reducing Gaussian noise

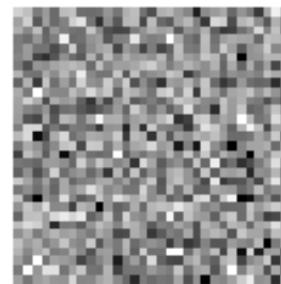
$\sigma=0.05$



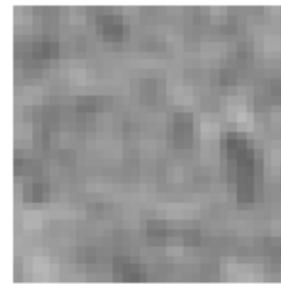
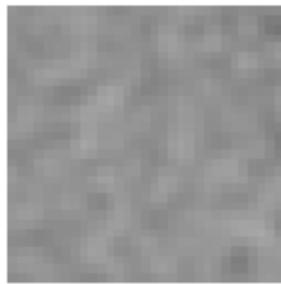
$\sigma=0.1$



$\sigma=0.2$



no  
smoothing



$\sigma=1$  pixel



$\sigma=2$  pixels



Smoothing with larger standard deviations suppresses noise, but also blurs the image

# Reducing salt-and-pepper noise

3x3



5x5



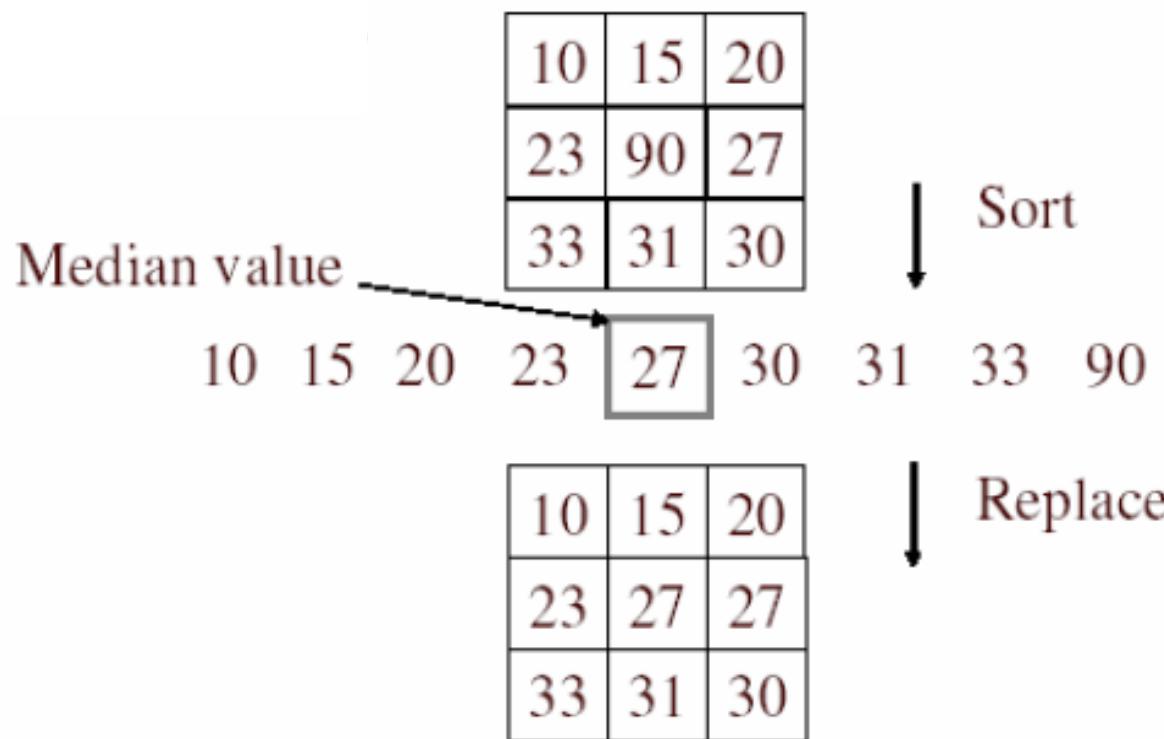
7x7



- What's wrong with the results?

# Alternative idea: Median filtering

- A **median filter** operates over a window by selecting the median intensity in the window

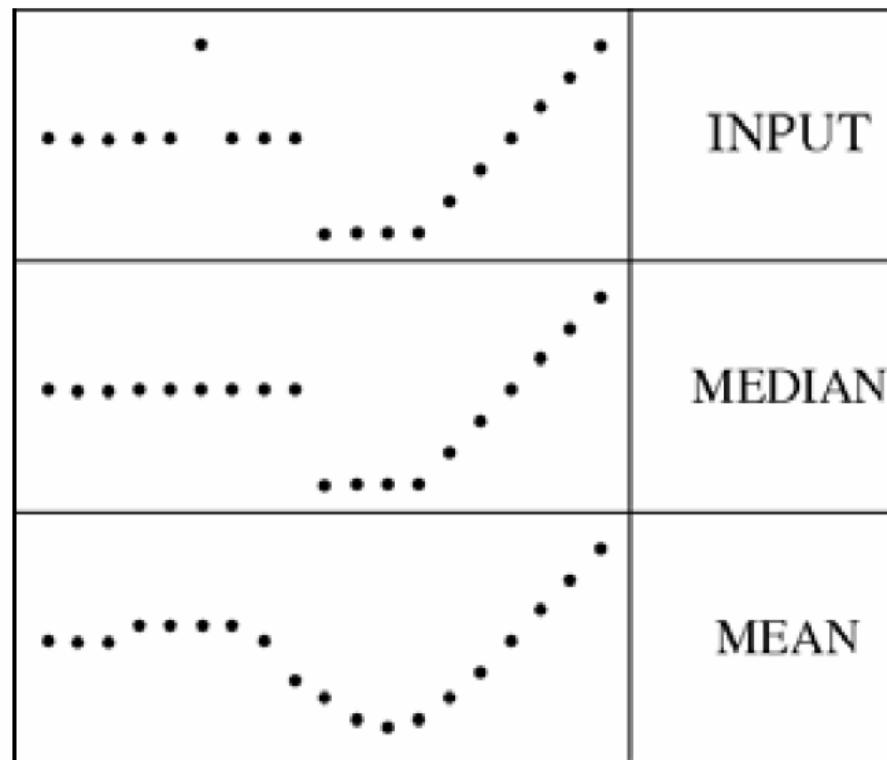


- Is median filtering linear?

# Median filter

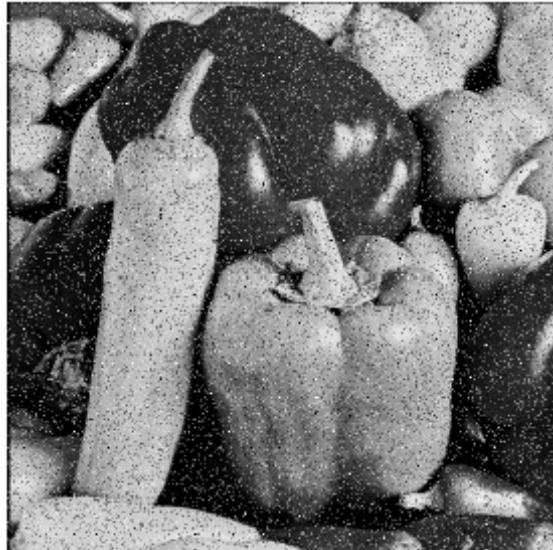
- What advantage does median filtering have over Gaussian filtering?
  - Robustness to outliers

filters have width 5 :

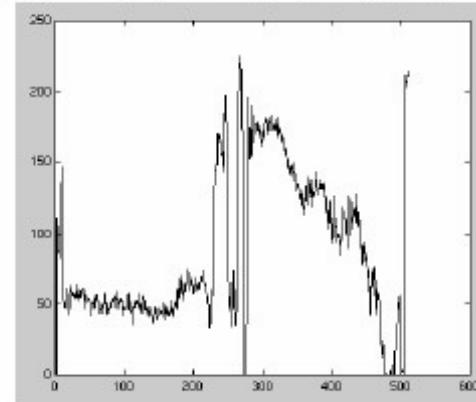
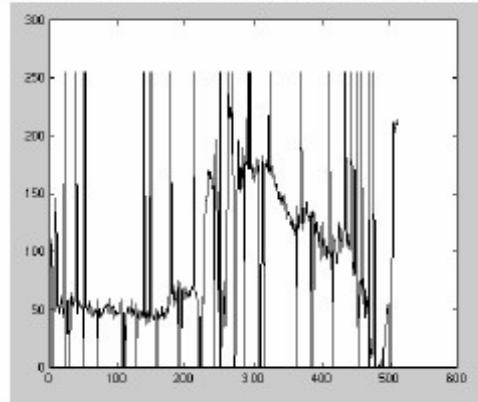


# Median filter

Salt-and-pepper noise



Median filtered



- MATLAB: `medfilt2(image, [h w])`

Source: M. Hebert

# Gaussian vs. median filtering

3x3



Gaussian

5x5



7x7



Median



# Sharpening filter



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

-

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



Original

## Sharpening filter

- Accentuates differences with local average

# Sharpening revisited

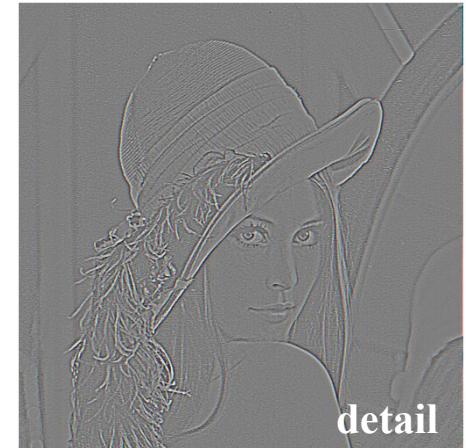
- What does blurring take away?



original



smoothed (5x5)



detail

-

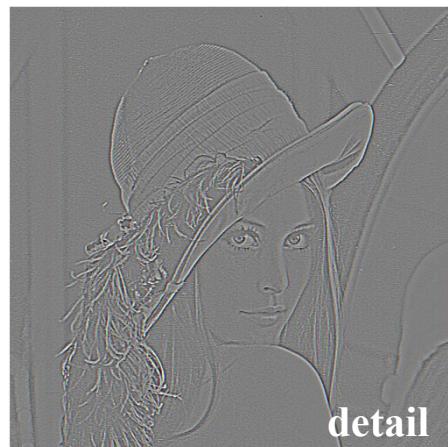
=

Let's add it back:



original

$+ \alpha$



detail

=



sharpened

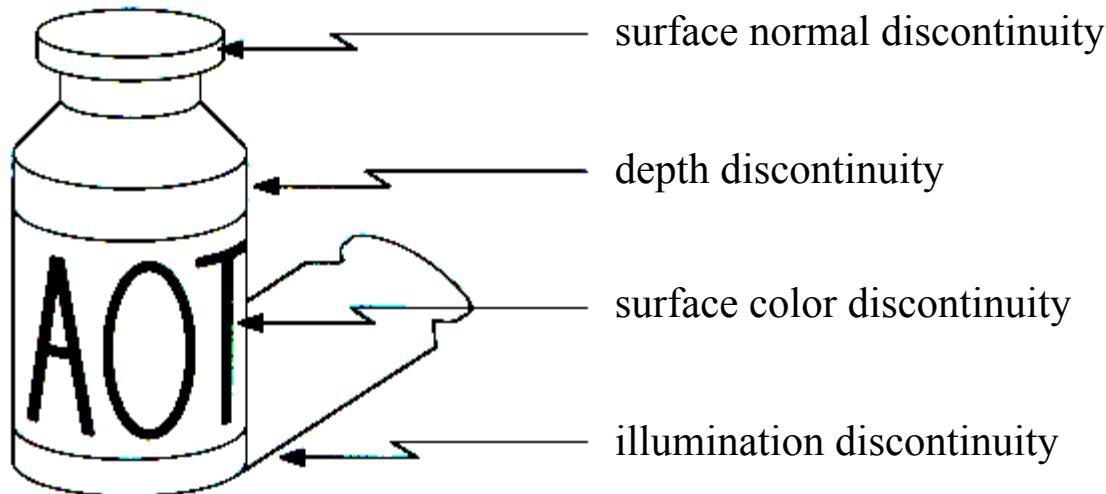
# Edge detection

- **Goal:** Identify sudden changes (discontinuities) in an image
  - Intuitively, most semantic and shape information from the image can be encoded in the edges
  - More compact than pixels
- **Ideal:** artist's line drawing (but artist is also using object-level knowledge)



# Origin of edges

- Edges are caused by a variety of factors:

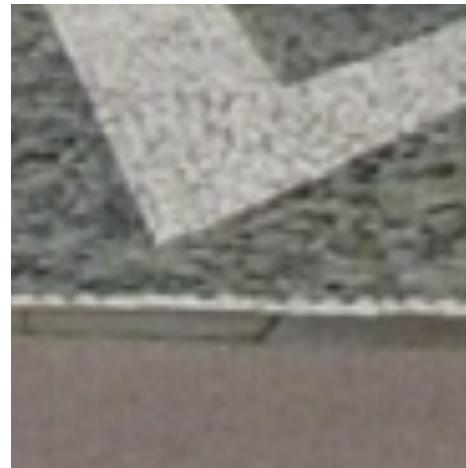


# Why finding edges is important

- Group pixels into objects or parts
- Cues for 3D shape
- Guiding interactive image editing

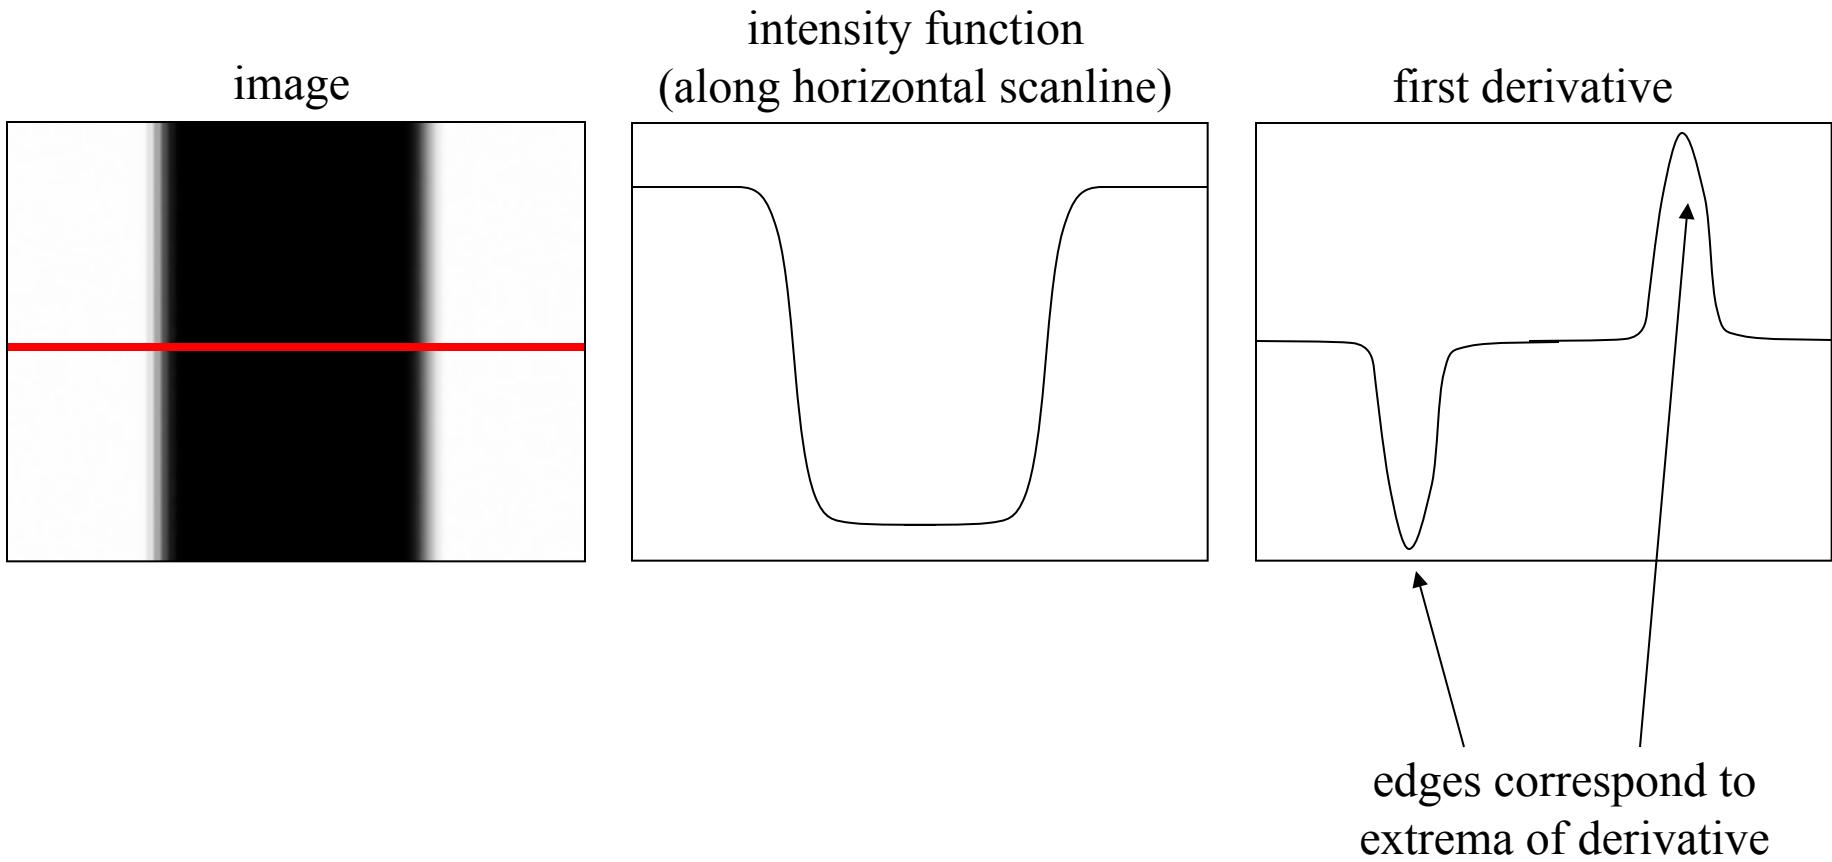


# Closeup of edges



# Edge detection

- An edge is a place of rapid change in the image intensity function



# Derivatives with convolution

For 2D function  $f(x,y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

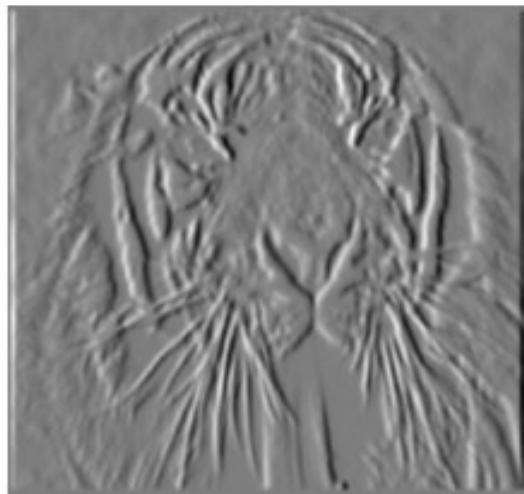
To implement the above as convolution, what would be the associated filter?

# Partial derivatives of an image



$$\frac{\partial f(x, y)}{\partial x}$$

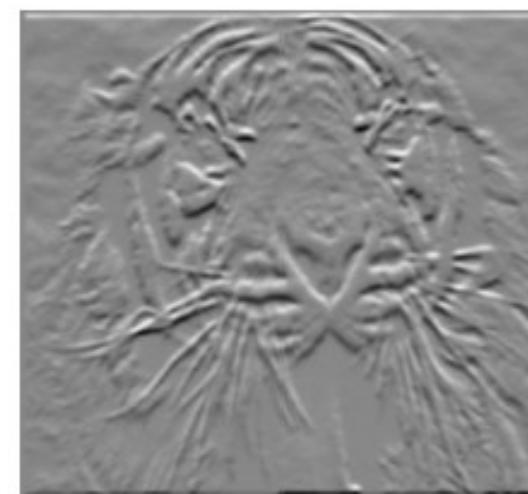
-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

-1  
or  
1  
1  
-1

-1	or	1
1		-1



Which shows changes with respect to x?

# Finite difference filters

- Other approximations of derivative filters exist:

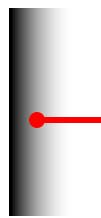
**Prewitt:**  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

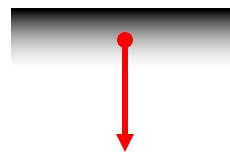
**Sobel:**  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

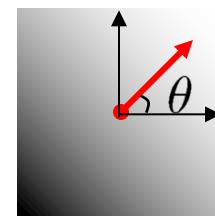
**Roberts:**  $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

# Image gradient

- The gradient of an image:  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$


$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid increase in intensity

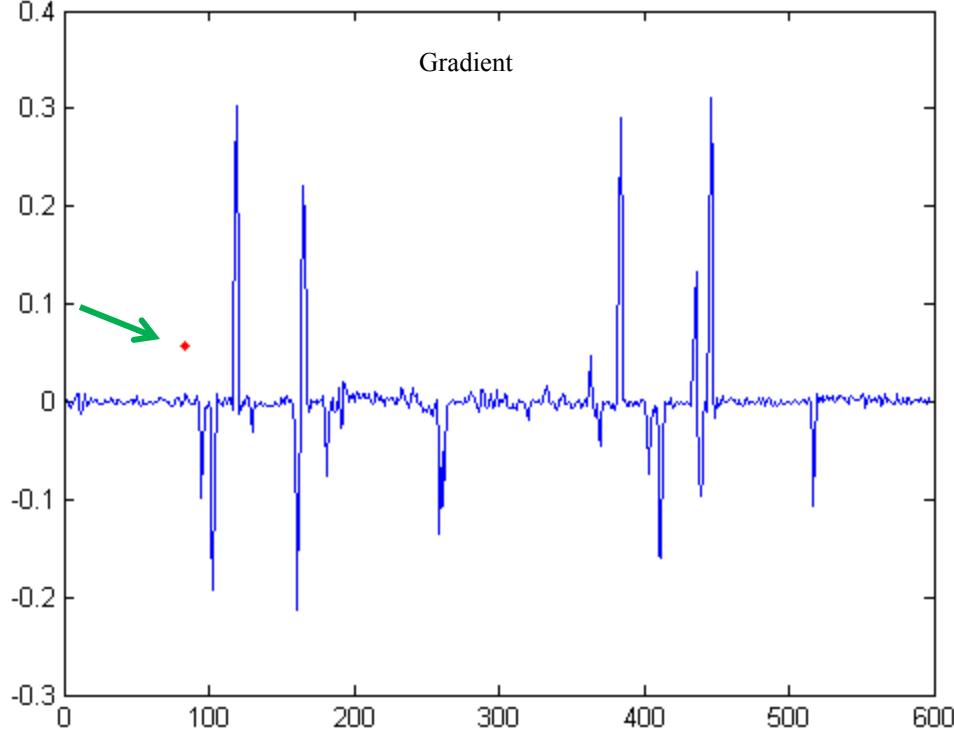
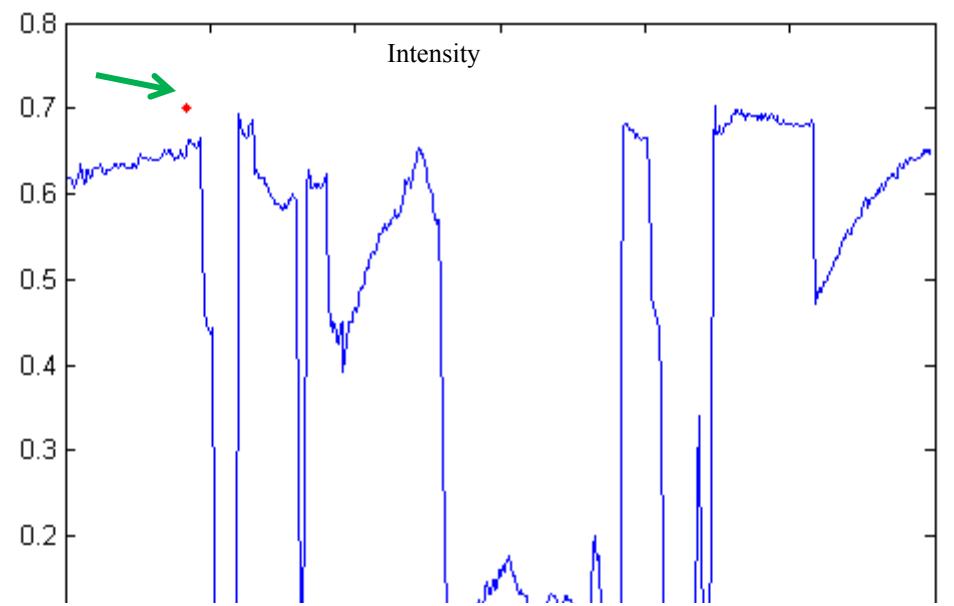
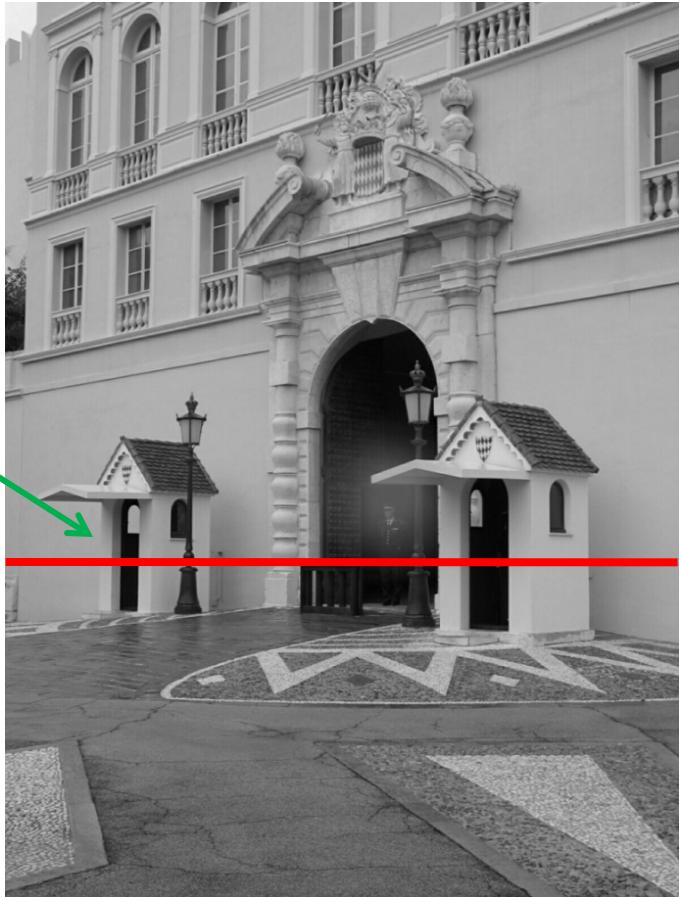
- How does this direction relate to the direction of the edge?

The gradient direction is given by  $\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

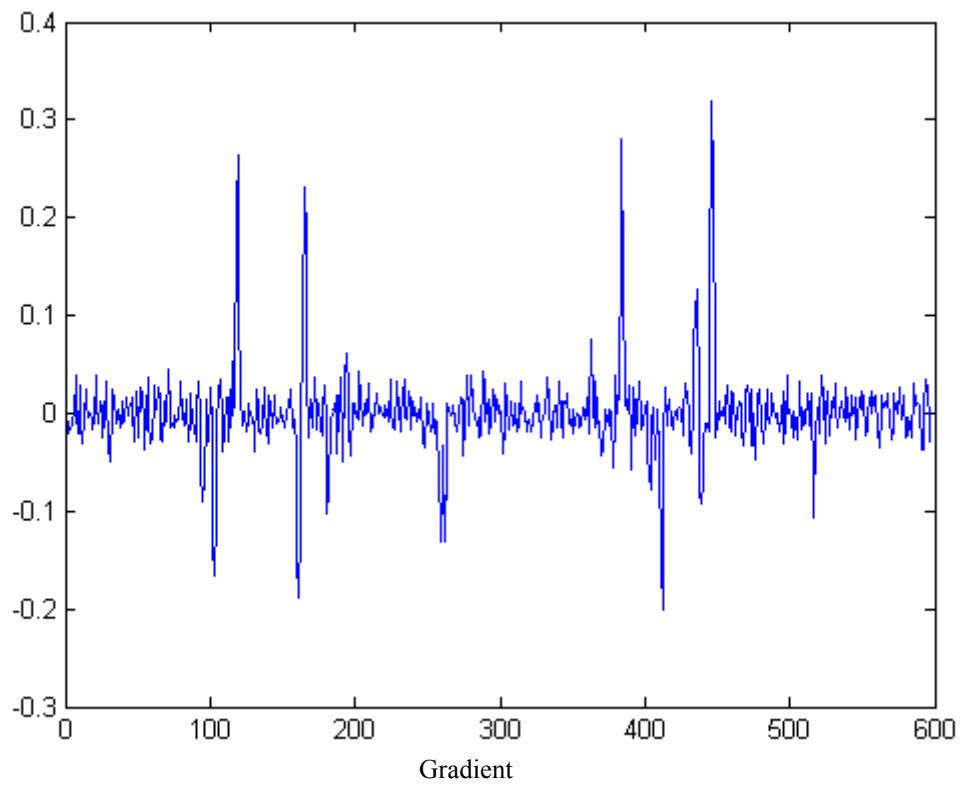
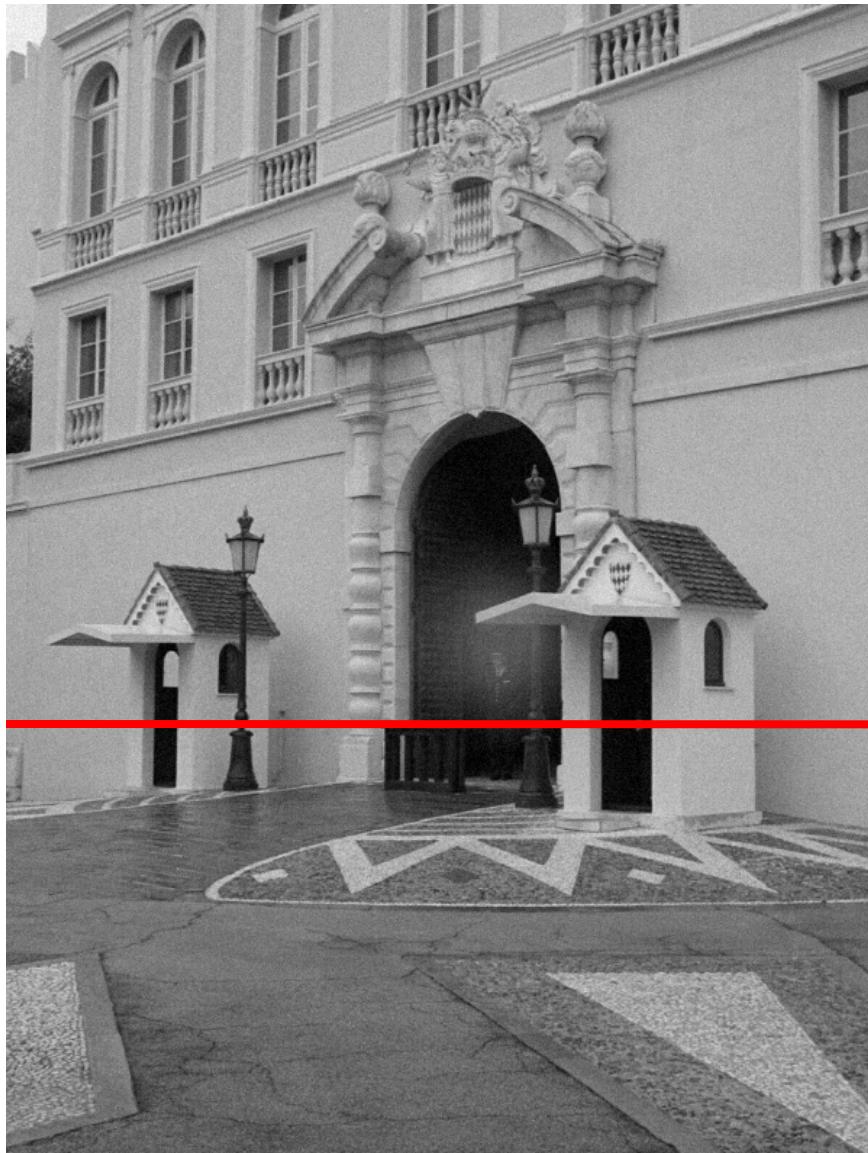
The edge strength is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

# Intensity profile

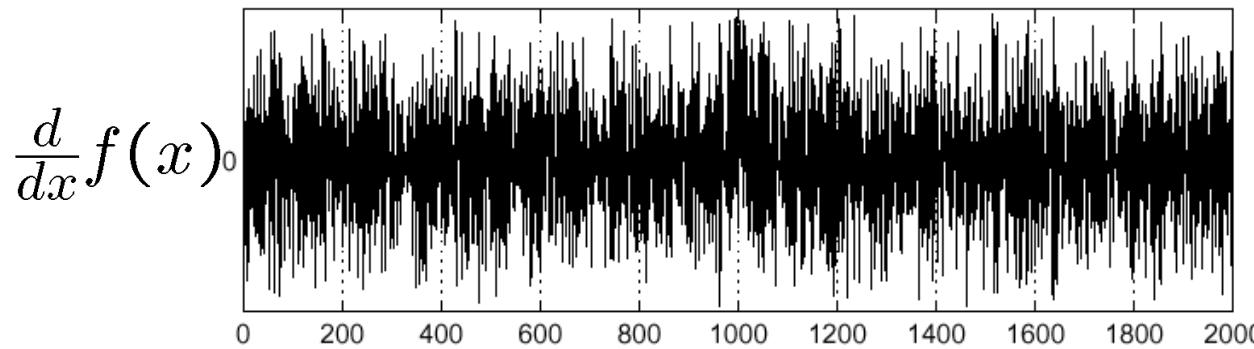
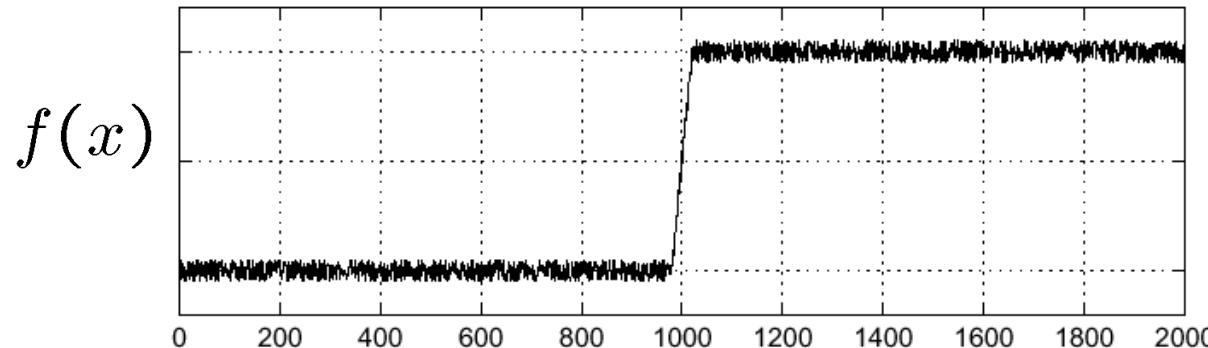


# With a little Gaussian noise



# Effects of noise

- Consider a single row or column of the image



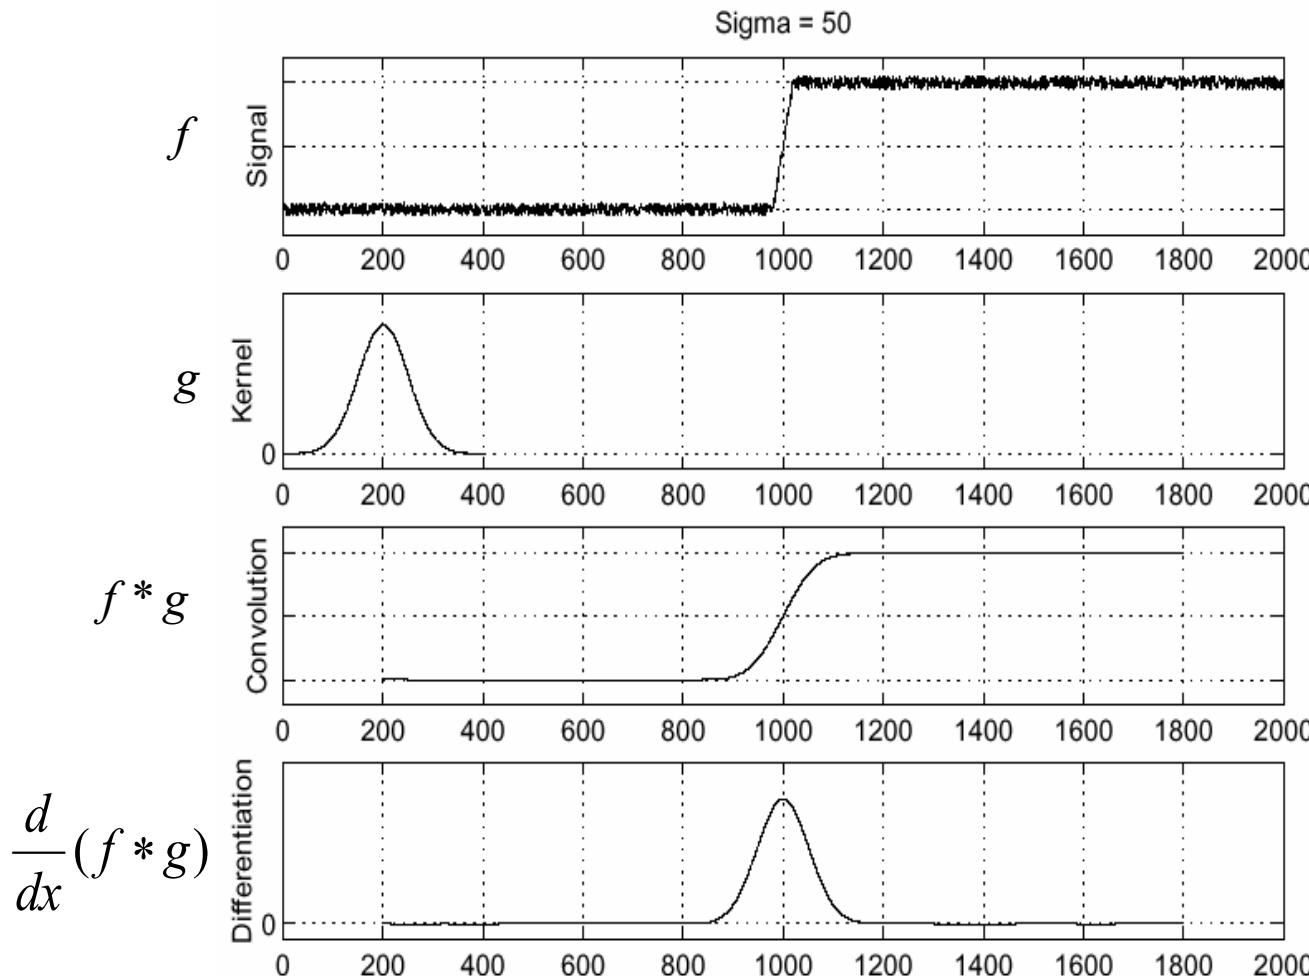
Where is the edge?

Source: S. Seitz

# Effects of noise

- Difference filters respond strongly to noise
  - Image noise results in pixels that look very different from their neighbors
  - Generally, the larger the noise the stronger the response
- What can we do about it?

# Solution: smooth first



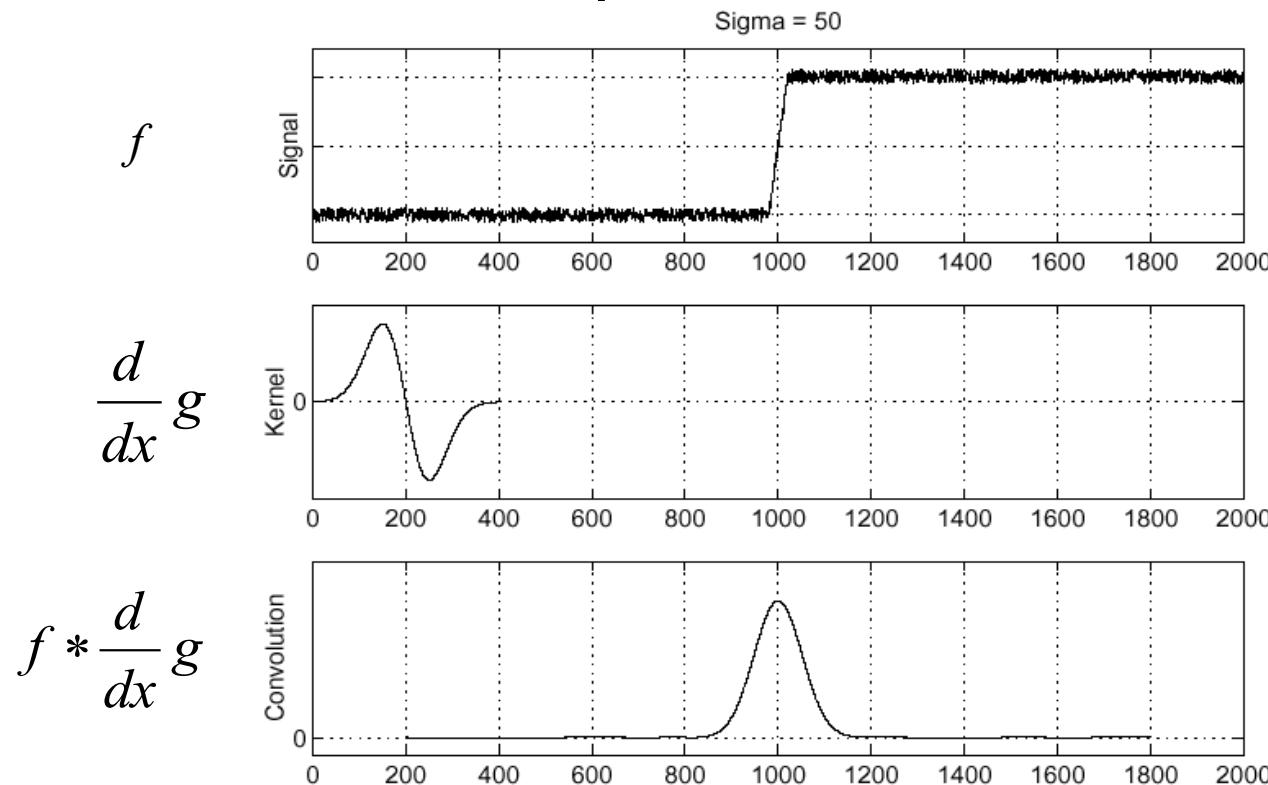
- To find edges, look for peaks in

$$\frac{d}{dx}(f * g)$$

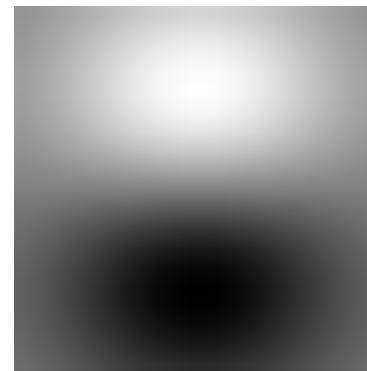
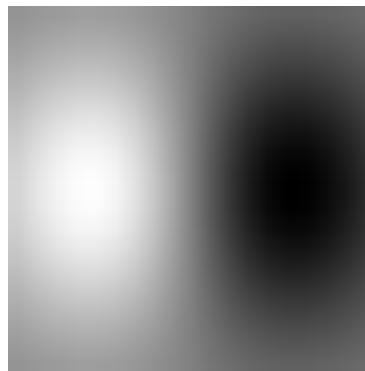
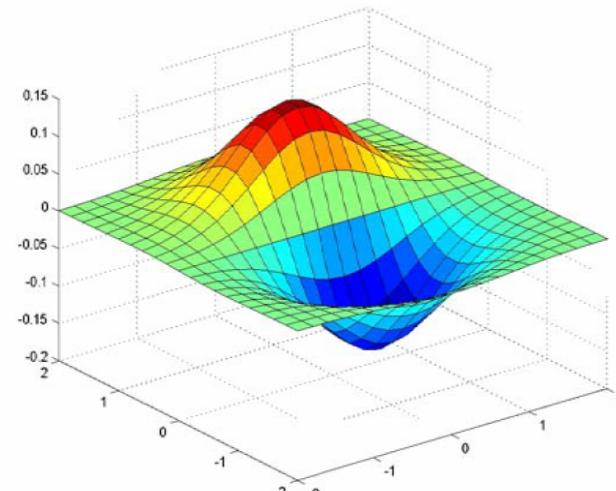
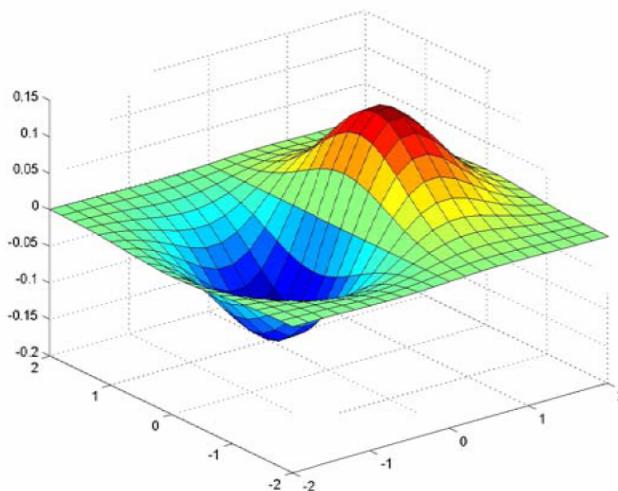
Source: S. Seitz

# Derivative theorem of convolution

- Differentiation is convolution, and convolution is associative:  
$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$
- This saves us one operation



# Derivative of Gaussian filters



- Which one finds horizontal/vertical edges?

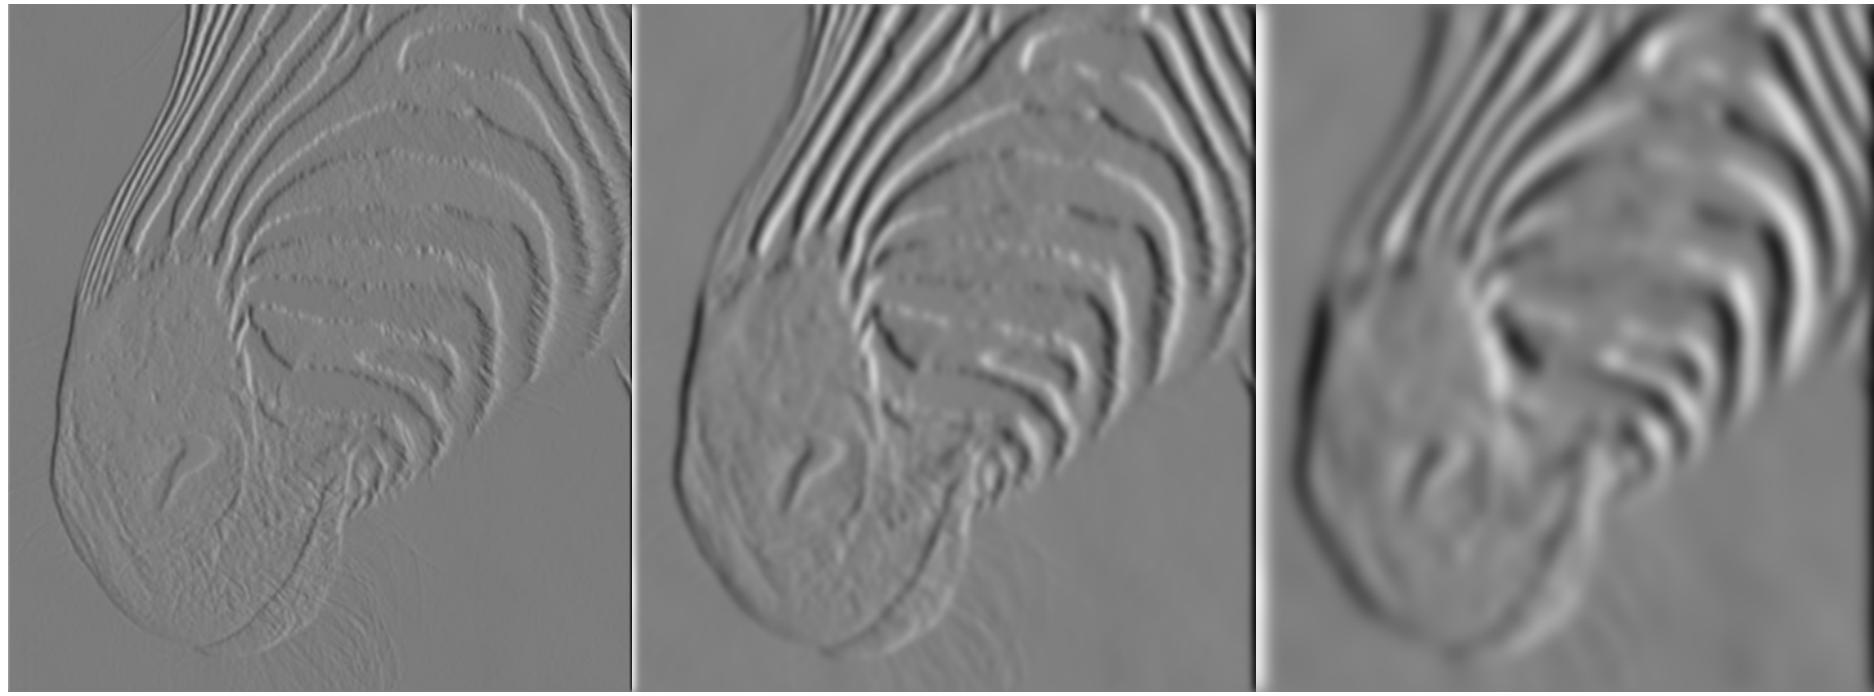
# Recall: Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

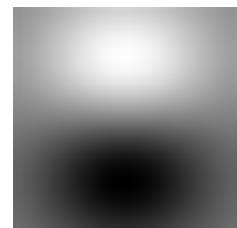
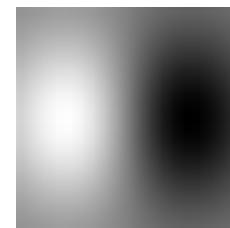
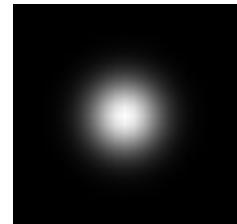
# Scale of Gaussian derivative filter



- Smoothed derivative removes noise, but blurs edge. Also finds edges at different “scales”

# Review: Smoothing vs. derivative filters

- Smoothing filters
  - Gaussian: remove “high-frequency” components; “low-pass” filter
  - Can the values of a smoothing filter be negative?
  - What should the values sum to?
    - One: constant regions are not affected by the filter
- Derivative filters
  - Derivatives of Gaussian
  - Can the values of a derivative filter be negative?
  - What should the values sum to?
    - Zero: no response in constant regions
    - High absolute value at points of high contrast



# The Canny edge detector

1. Filter image with derivative of Gaussian
2. Find magnitude and orientation of gradient
3. **Non-maximum suppression:**
  - Thin wide “ridges” down to single pixel width
4. **Linking and thresholding (hysteresis):**
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them

# The Canny edge detector



original image

# The Canny edge detector



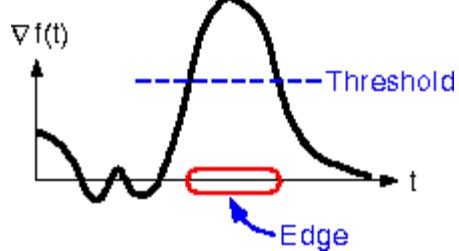
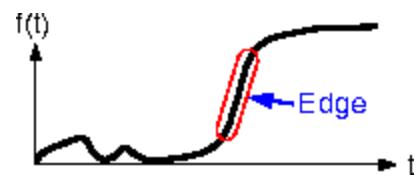
norm of the gradient

# The Canny edge detector



thresholding

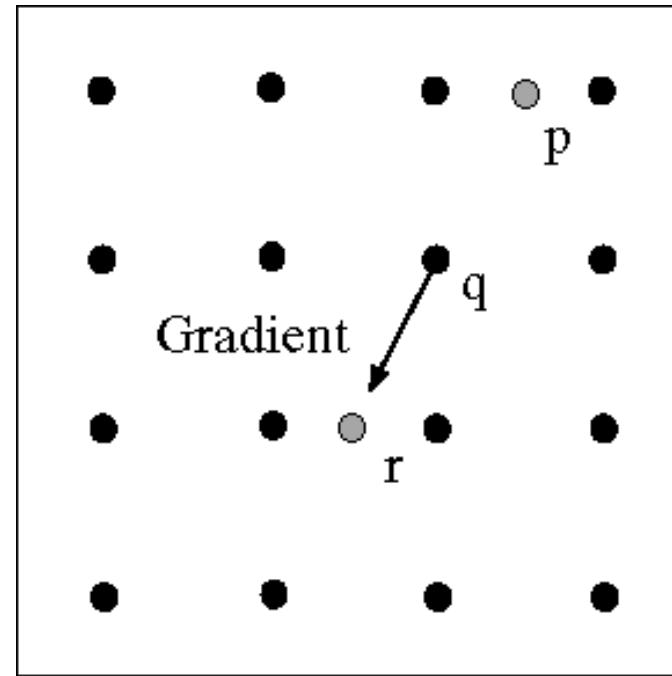
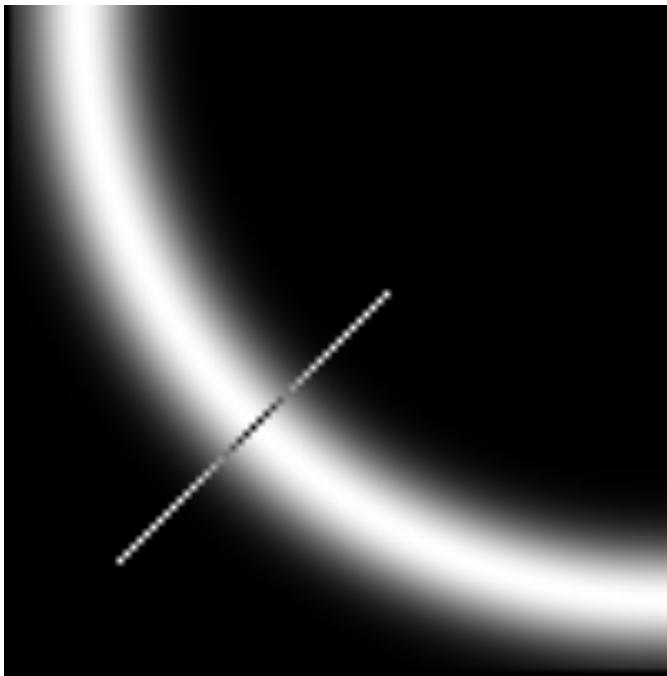
# The Canny edge detector



How to turn  
these thick  
regions of the  
gradient into  
curves?

thresholding

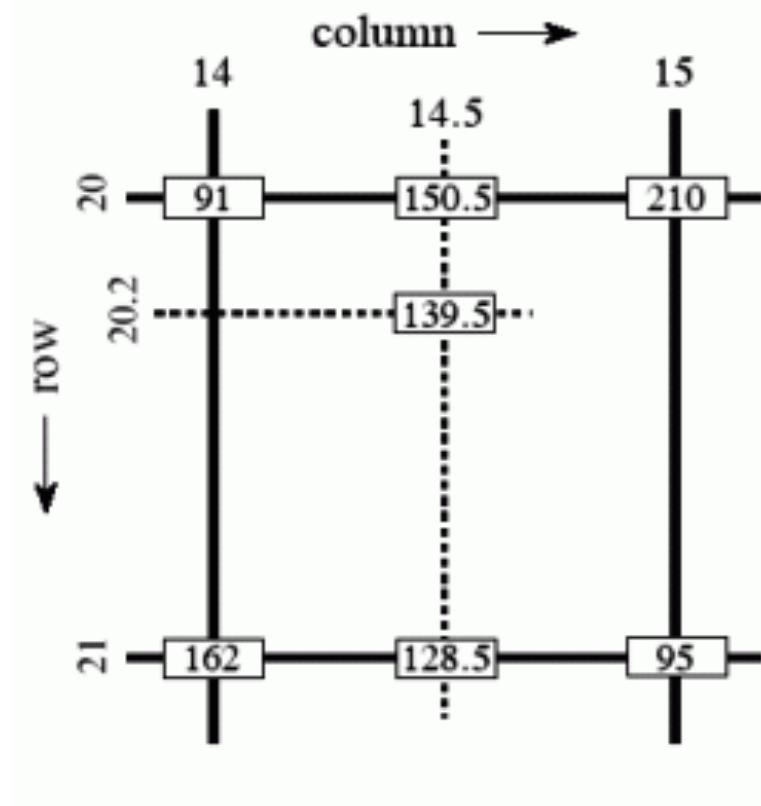
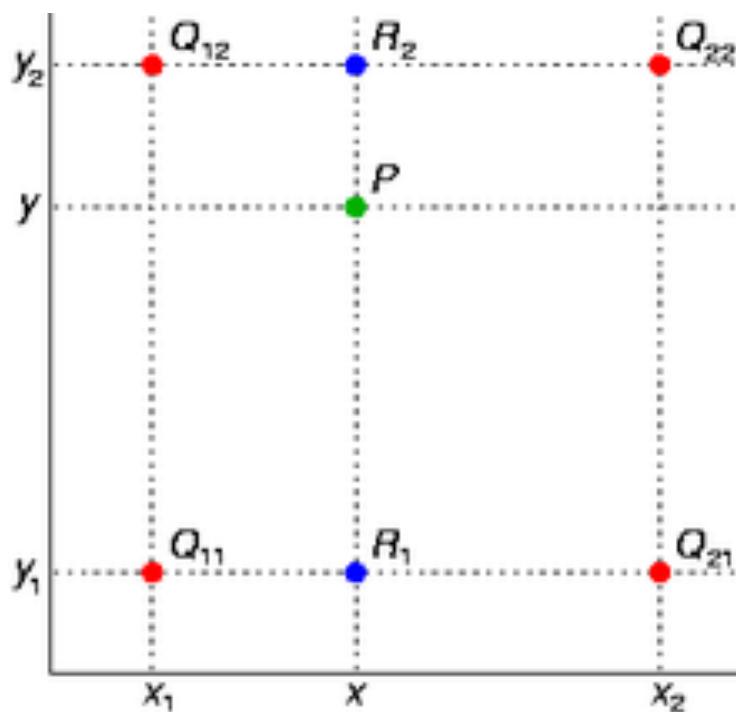
# Non-maximum suppression



Check if pixel is local maximum along gradient direction,  
select single max across width of the edge  
– requires checking interpolated pixels p and r

# Bilinear Interpolation

$$f(x, y) \approx [1 - x \quad x] \begin{bmatrix} f(0, 0) & f(0, 1) \\ f(1, 0) & f(1, 1) \end{bmatrix} \begin{bmatrix} 1 - y \\ y \end{bmatrix}.$$



# The Canny edge detector

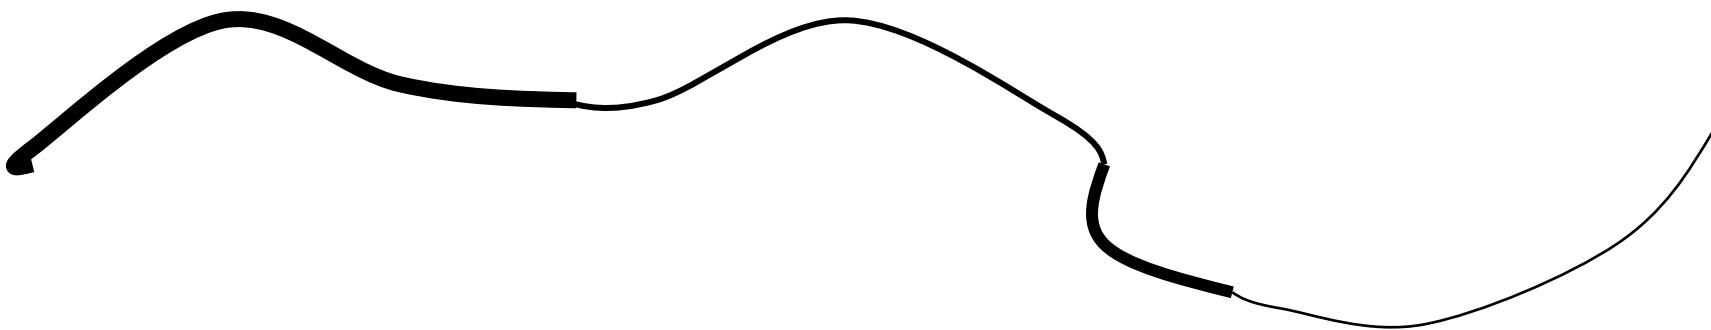


Thinning (non-maximum suppression)

Problem: pixels along this edge didn't survive the thresholding

# Hysteresis thresholding

- Use a high threshold to start edge curves, and a low threshold to continue them



# Hysteresis thresholding

- Threshold at low/high levels to get weak/strong edge pixels
- Trace connected components, starting from strong edge pixels



# Recap: Canny edge detector

1. Compute x and y gradient images
  2. Find magnitude and orientation of gradient
  3. **Non-maximum suppression:**
    - Thin wide “ridges” down to single pixel width
  4. **Linking and thresholding (hysteresis):**
    - Define two thresholds: low and high
    - Use the high threshold to start edge curves and the low threshold to continue them
- MATLAB: **edge (image, 'canny' ) ;**

J. Canny, *[A Computational Approach To Edge Detection](#)*, IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

# Effect of $\sigma$ (Gaussian kernel spread/size)



original



Canny with  $\sigma = 1$



Canny with  $\sigma = 2$

The choice of  $\sigma$  depends on desired behavior

- large  $\sigma$  detects large scale edges
- small  $\sigma$  detects fine features

# Feature extraction: Corners



# Why extract features?

- Motivation: panorama stitching
  - We have two images - how do we combine them?

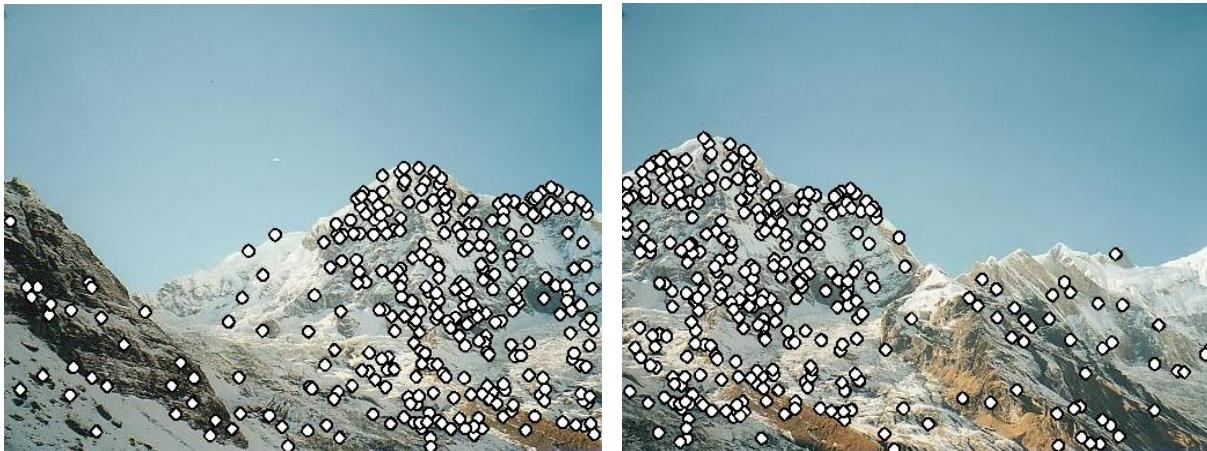


Step 1: extract features

Step 2: match features

Step 3: align images

# Characteristics of good features



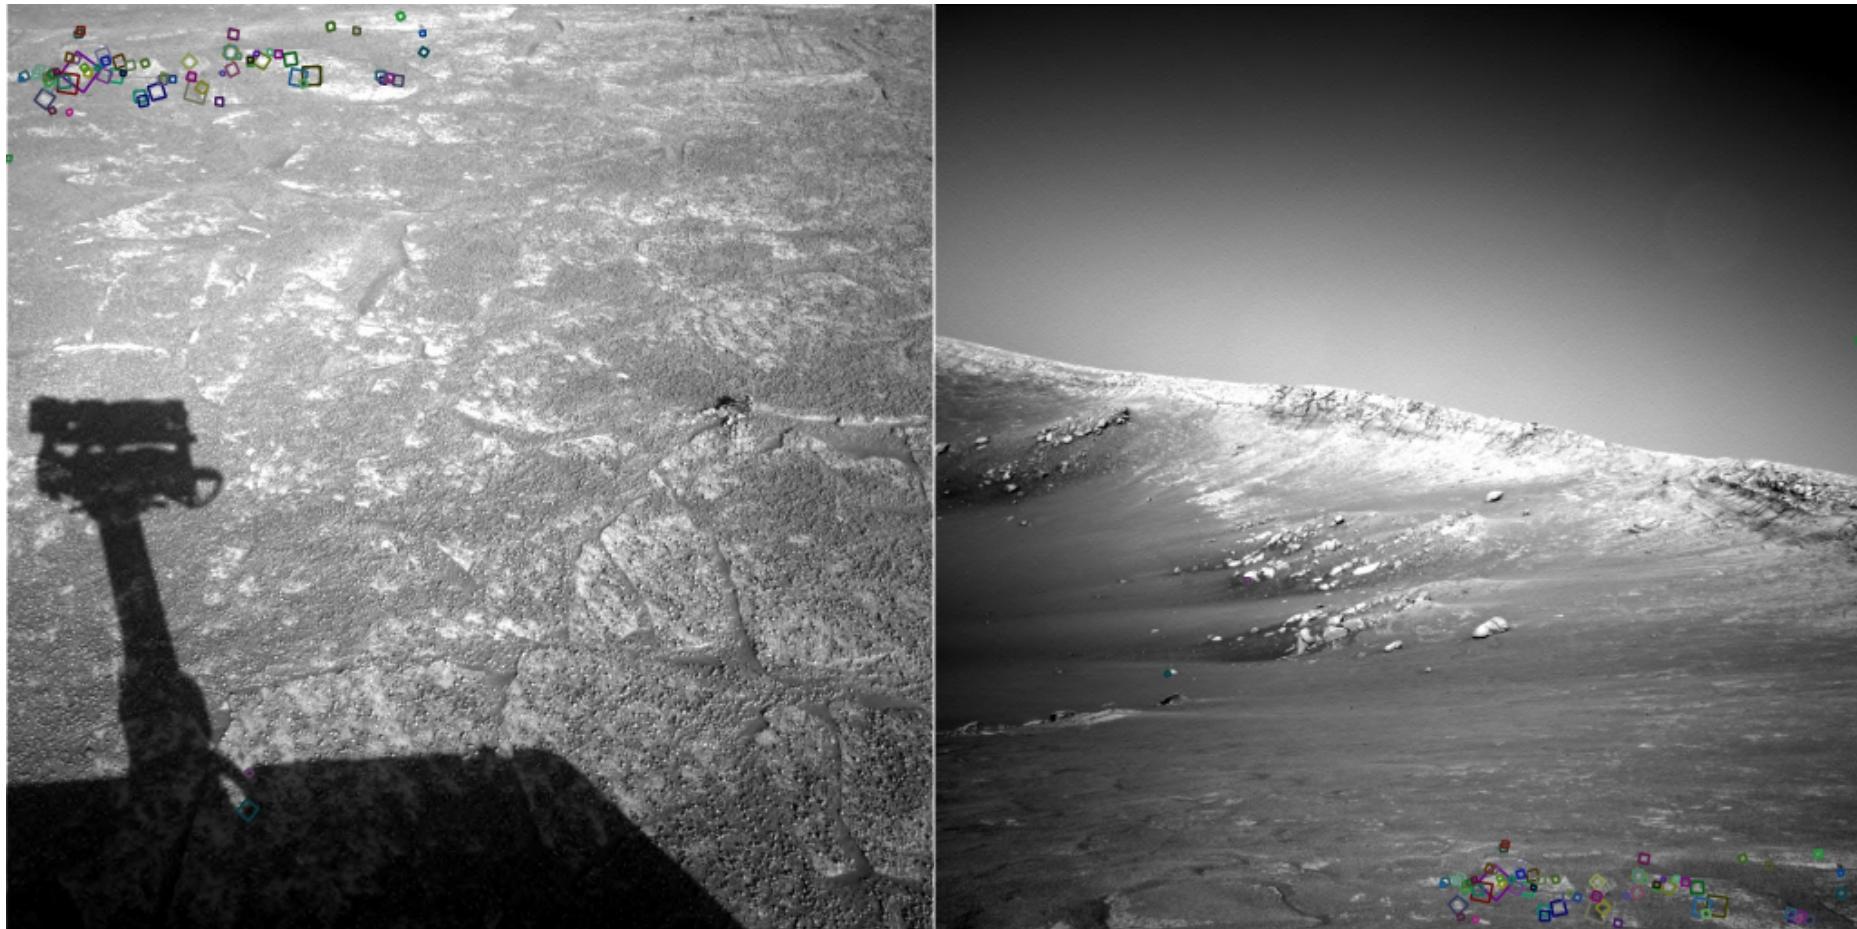
- **Repeatability**
  - The same feature can be found in several images despite geometric and photometric transformations
- **Saliency**
  - Each feature is distinctive
- **Compactness and efficiency**
  - Many fewer features than image pixels
- **Locality**
  - A feature occupies a relatively small area of the image; robust to clutter and occlusion

# Applications

- Feature points are used for:
  - Image alignment
  - 3D reconstruction
  - Motion tracking
  - Robot navigation
  - Indexing and database retrieval
  - Object recognition



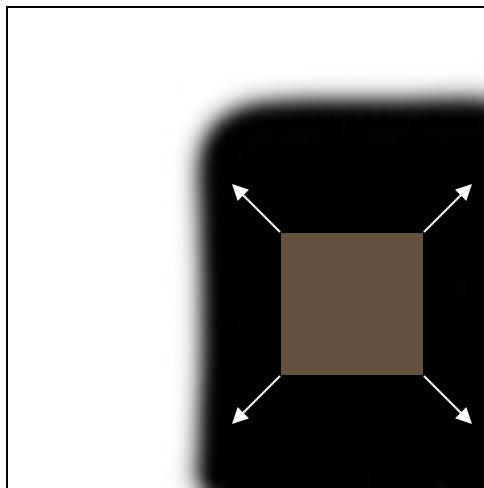
# Answer below (look for tiny colored squares...)



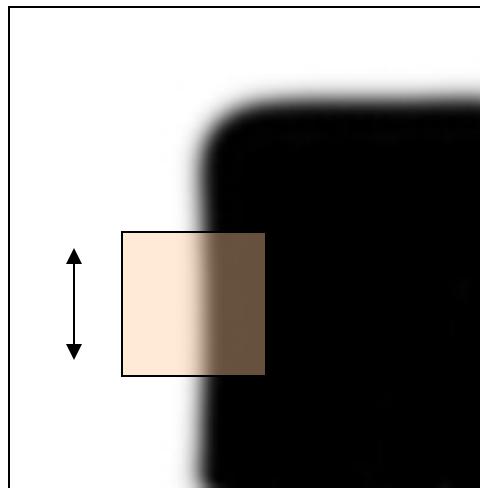
NASA Mars Rover images  
with SIFT feature matches  
Figure by Noah Snavely

# Corner Detection: Basic Idea

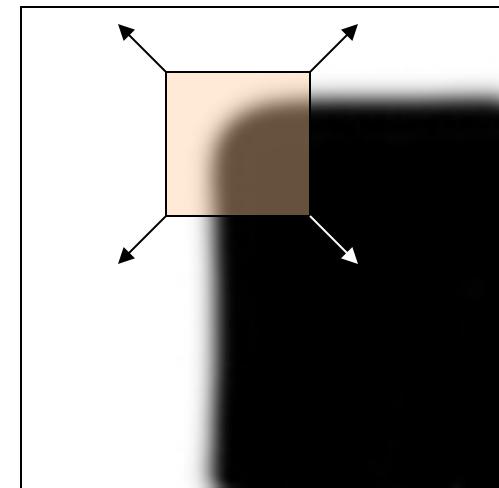
- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give *a large change* in intensity



“flat” region:  
no change in  
all directions



“edge”:  
no change  
along the edge  
direction

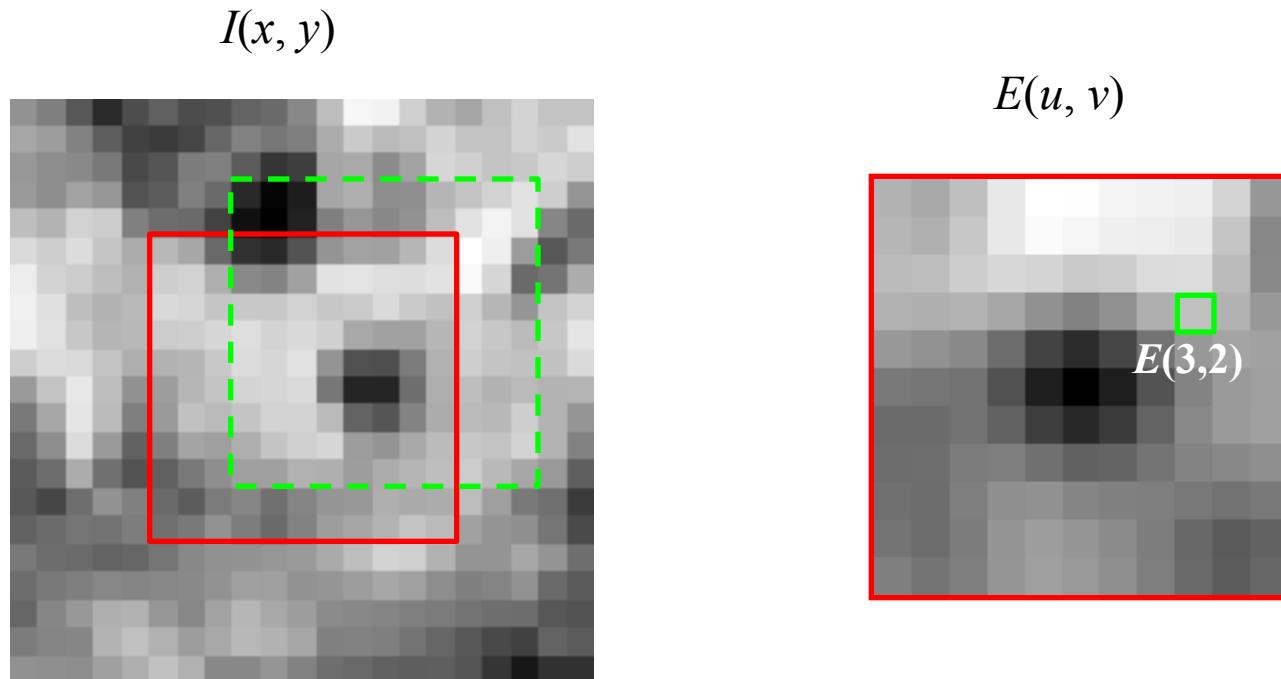


“corner”:  
significant  
change in all  
directions

# Corner Detection: Mathematics

Change in appearance of window  $W$  for the shift  $[u, v]$ :

$$E(u, v) = \sum_{(x, y) \in W} [I(x + u, y + v) - I(x, y)]^2$$



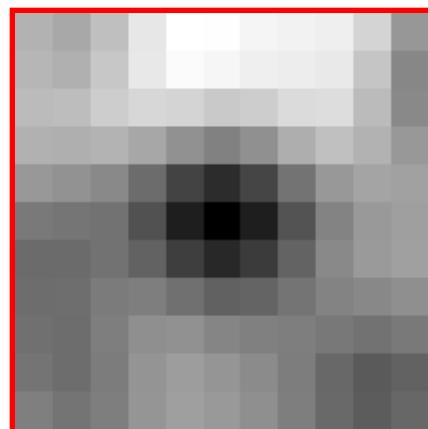
# Corner Detection: Mathematics

Change in appearance of window  $W$  for the shift  $[u, v]$ :

$$E(u, v) = \sum_{(x,y) \in W} [I(x+u, y+v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

$$E(u, v)$$



# Corner Detection: Mathematics

- First-order Taylor approximation for small motions  $[u, v]$ :

$$I(x + u, y + v) \approx I(x, y) + I_x u + I_y v$$

# Corner Detection: Mathematics

The quadratic approximation can be written as

$$E(u, v) \approx [u \quad v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where  $M$  is a *second moment matrix* computed from image derivatives:

$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix}$$

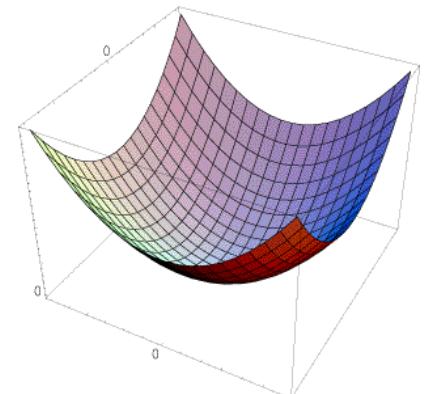
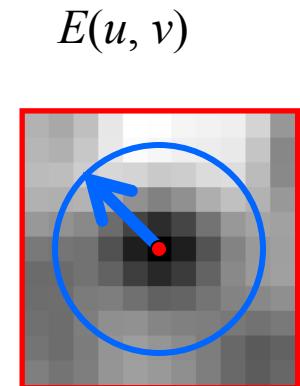
(the sums are over all the pixels in the window  $W$ )

# Interpreting the second moment matrix

- The surface  $E(u, v)$  is locally approximated by a quadratic form. Let's try to understand its shape.
  - Specifically, in which directions does it have the smallest/greatest change?

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix}$$



# Interpreting the second moment matrix

First, consider the axis-aligned case (gradients are either horizontal or vertical)

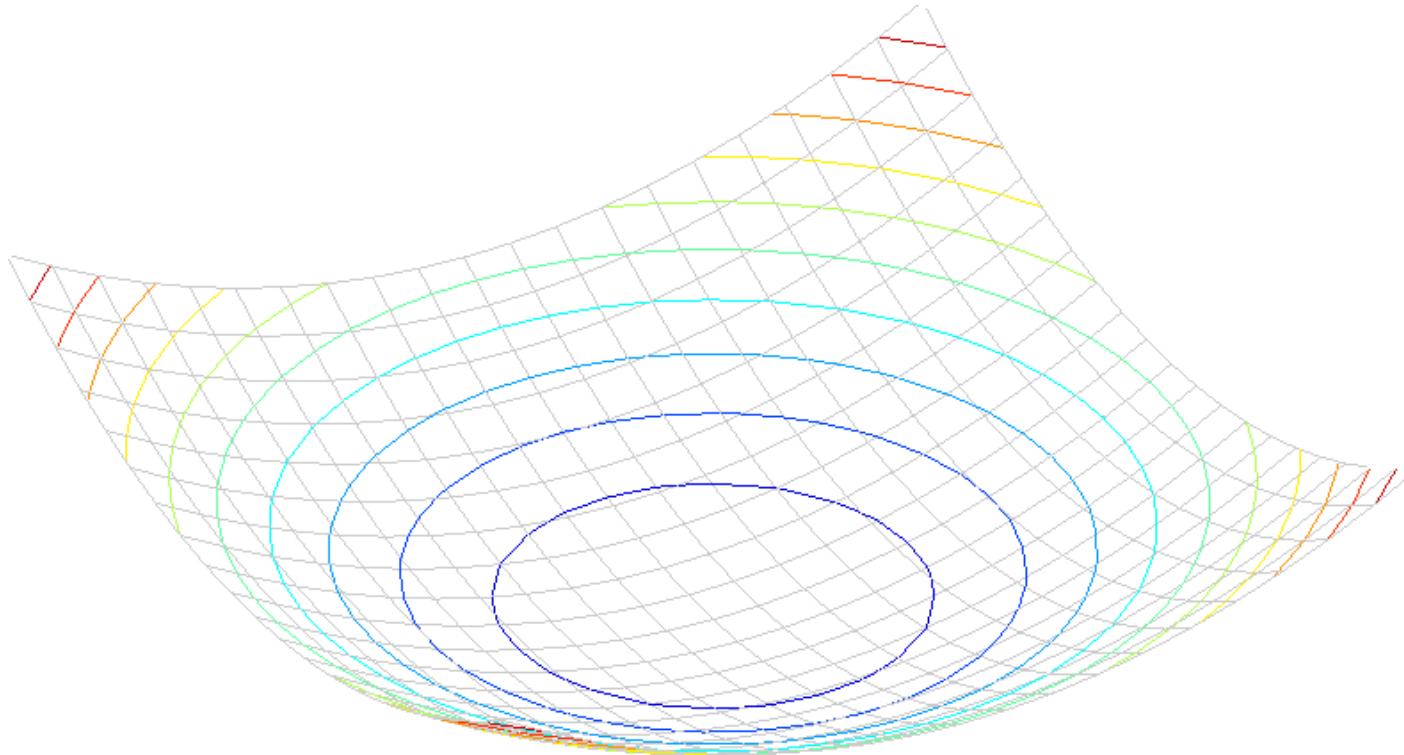
$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix} = \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$$

If either  $a$  or  $b$  is close to 0, then this is **not** a corner, so look for locations where both are large.

# Interpreting the second moment matrix

Consider a horizontal “slice” of  $E(u, v)$ :  $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$

This is the equation of an ellipse.



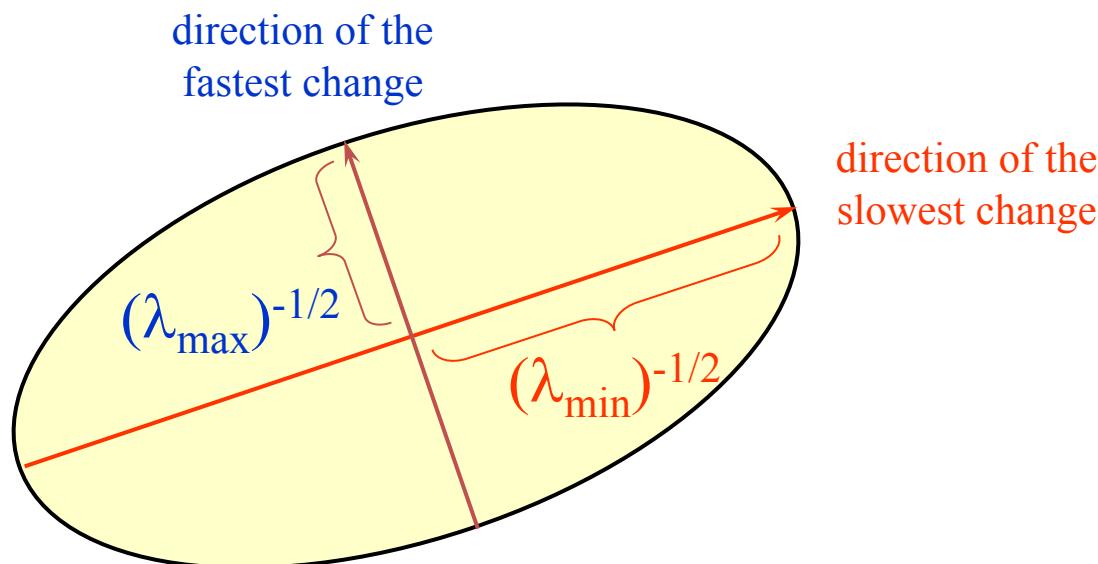
# Interpreting the second moment matrix

Consider a horizontal “slice” of  $E(u, v)$ :  $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$

This is the equation of an ellipse.

Diagonalization of  $M$ :  $M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$

The axis lengths of the ellipse are determined by the eigenvalues and the orientation is determined by  $R$



# Quick Eigenvalue/Eigenvector Review

The **eigenvectors** of a matrix  $A$  are the vectors  $x$  that satisfy:

$$Ax = \lambda x$$

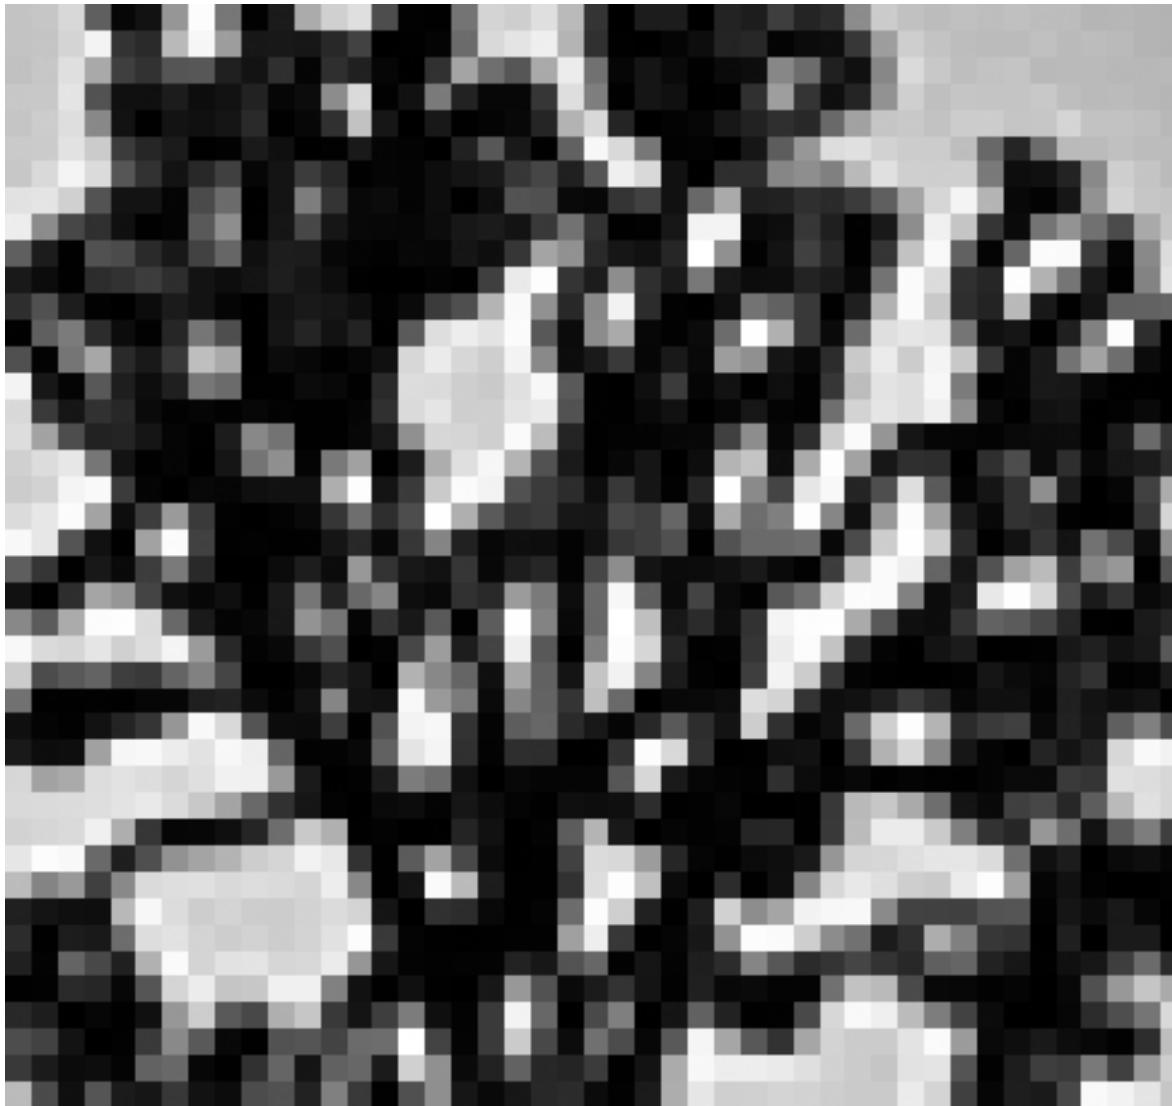
The scalar  $\lambda$  is the **eigenvalue** corresponding to  $x$

- The eigenvalues are found by solving:  $det(A - \lambda I) = 0$
- In our case,  $A = H$  is a  $2 \times 2$  matrix, so we have  $det \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} = 0$
- The solution:  $\lambda_{\pm} = \frac{1}{2} \left[ (h_{11} + h_{22}) \pm \sqrt{4h_{12}h_{21} + (h_{11} - h_{22})^2} \right]$

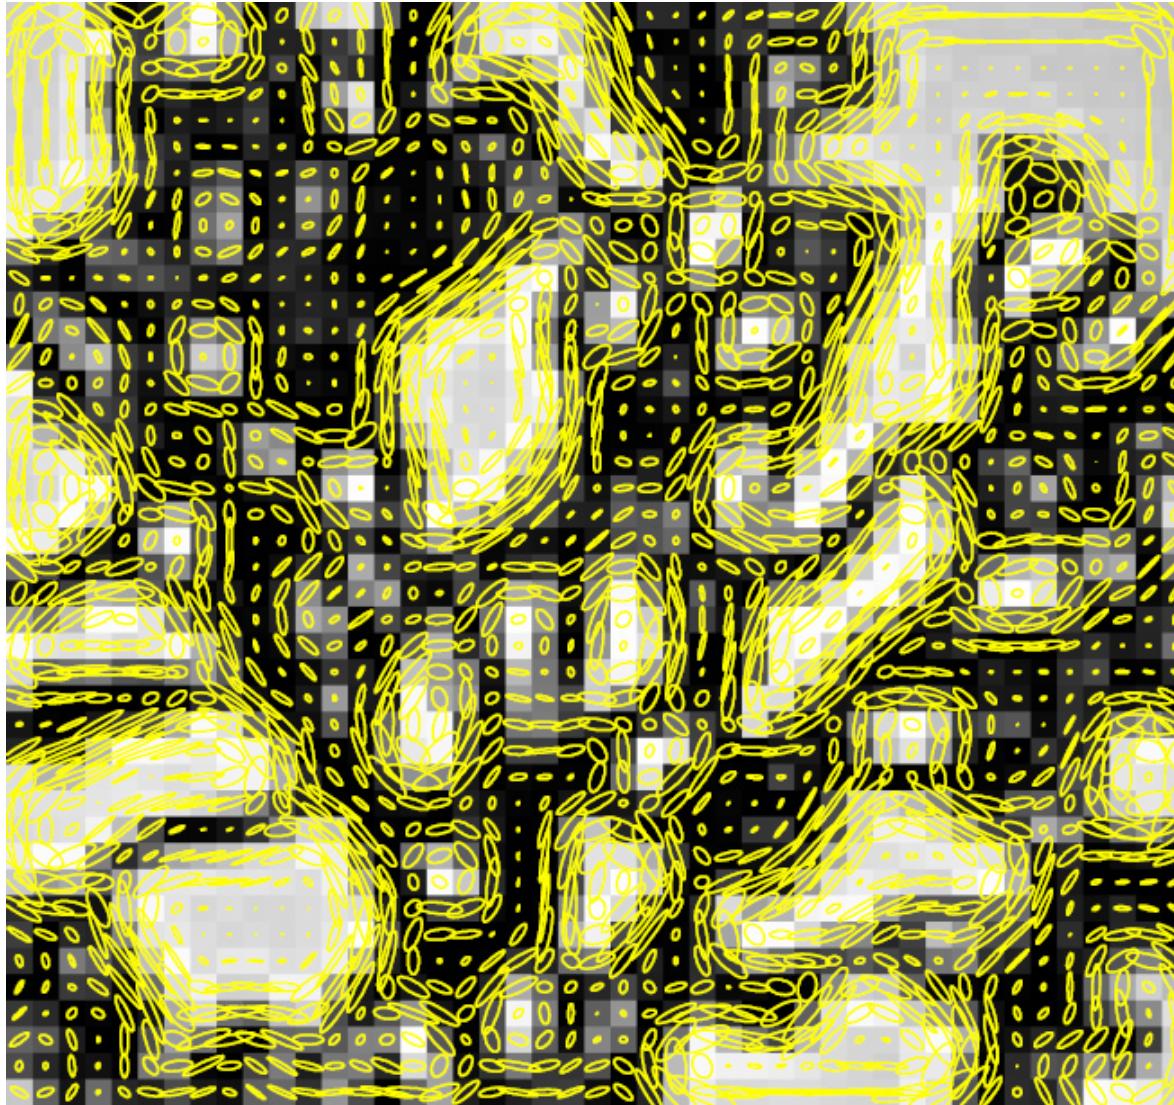
Once you know  $\lambda$ , you find  $x$  by solving

$$\begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = 0$$

# Visualization of second moment matrices

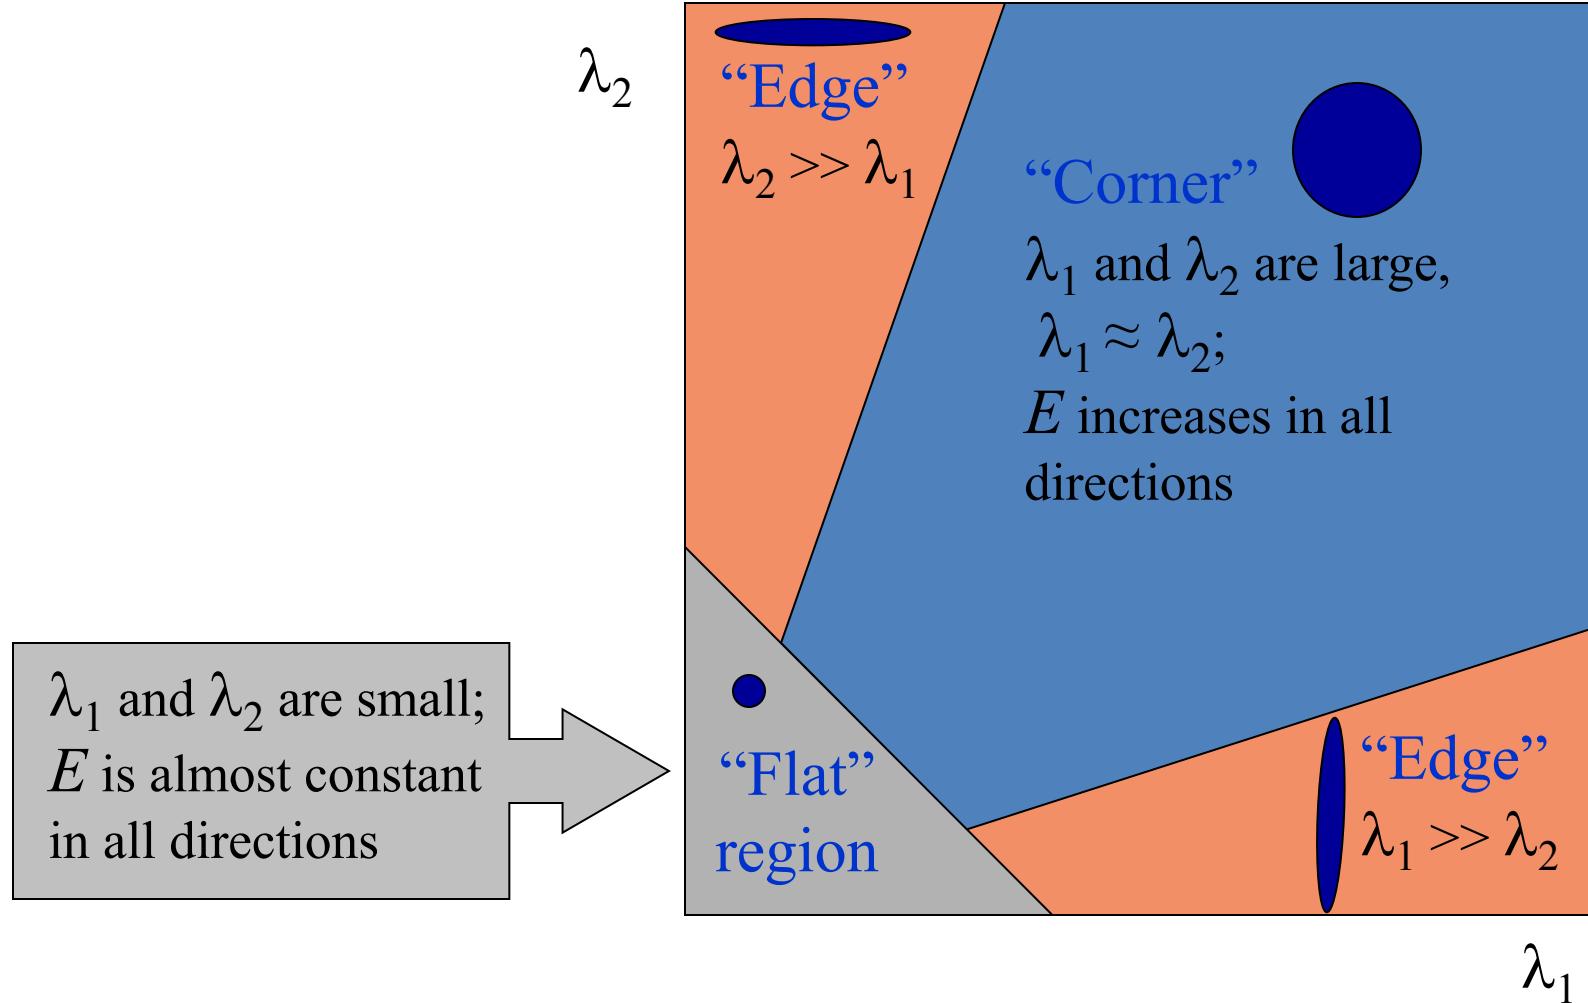


# Visualization of second moment matrices



# Interpreting the eigenvalues

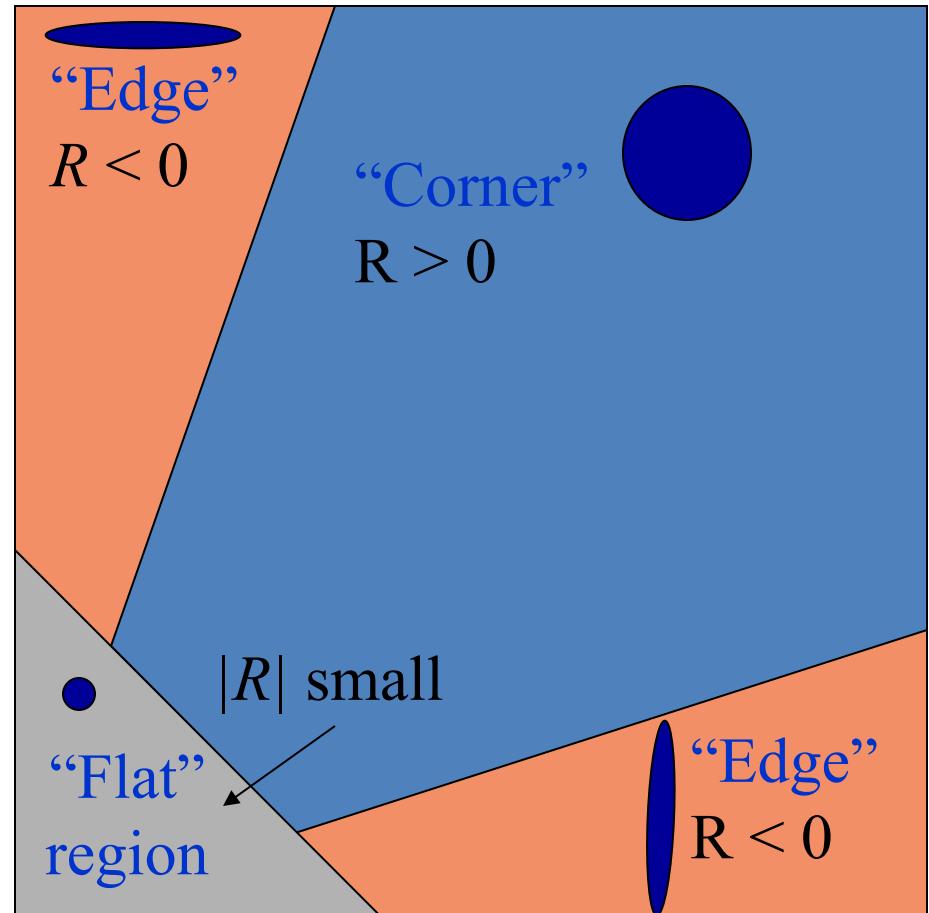
Classification of image points using eigenvalues of  $M$ :



# Corner response function

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

$\alpha$ : constant (0.04 to 0.06)



# The Harris corner detector

1. Compute partial derivatives at each pixel
2. Compute second moment matrix  $M$  in a Gaussian window around each pixel:

$$M = \begin{bmatrix} \sum_{x,y} w(x,y) I_x^2 & \sum_{x,y} w(x,y) I_x I_y \\ \sum_{x,y} w(x,y) I_x I_y & \sum_{x,y} w(x,y) I_y^2 \end{bmatrix}$$

C.Harris and M.Stephens. [“A Combined Corner and Edge Detector.”](#)  
*Proceedings of the 4th Alvey Vision Conference*: pages 147—151, 1988.

# The Harris corner detector

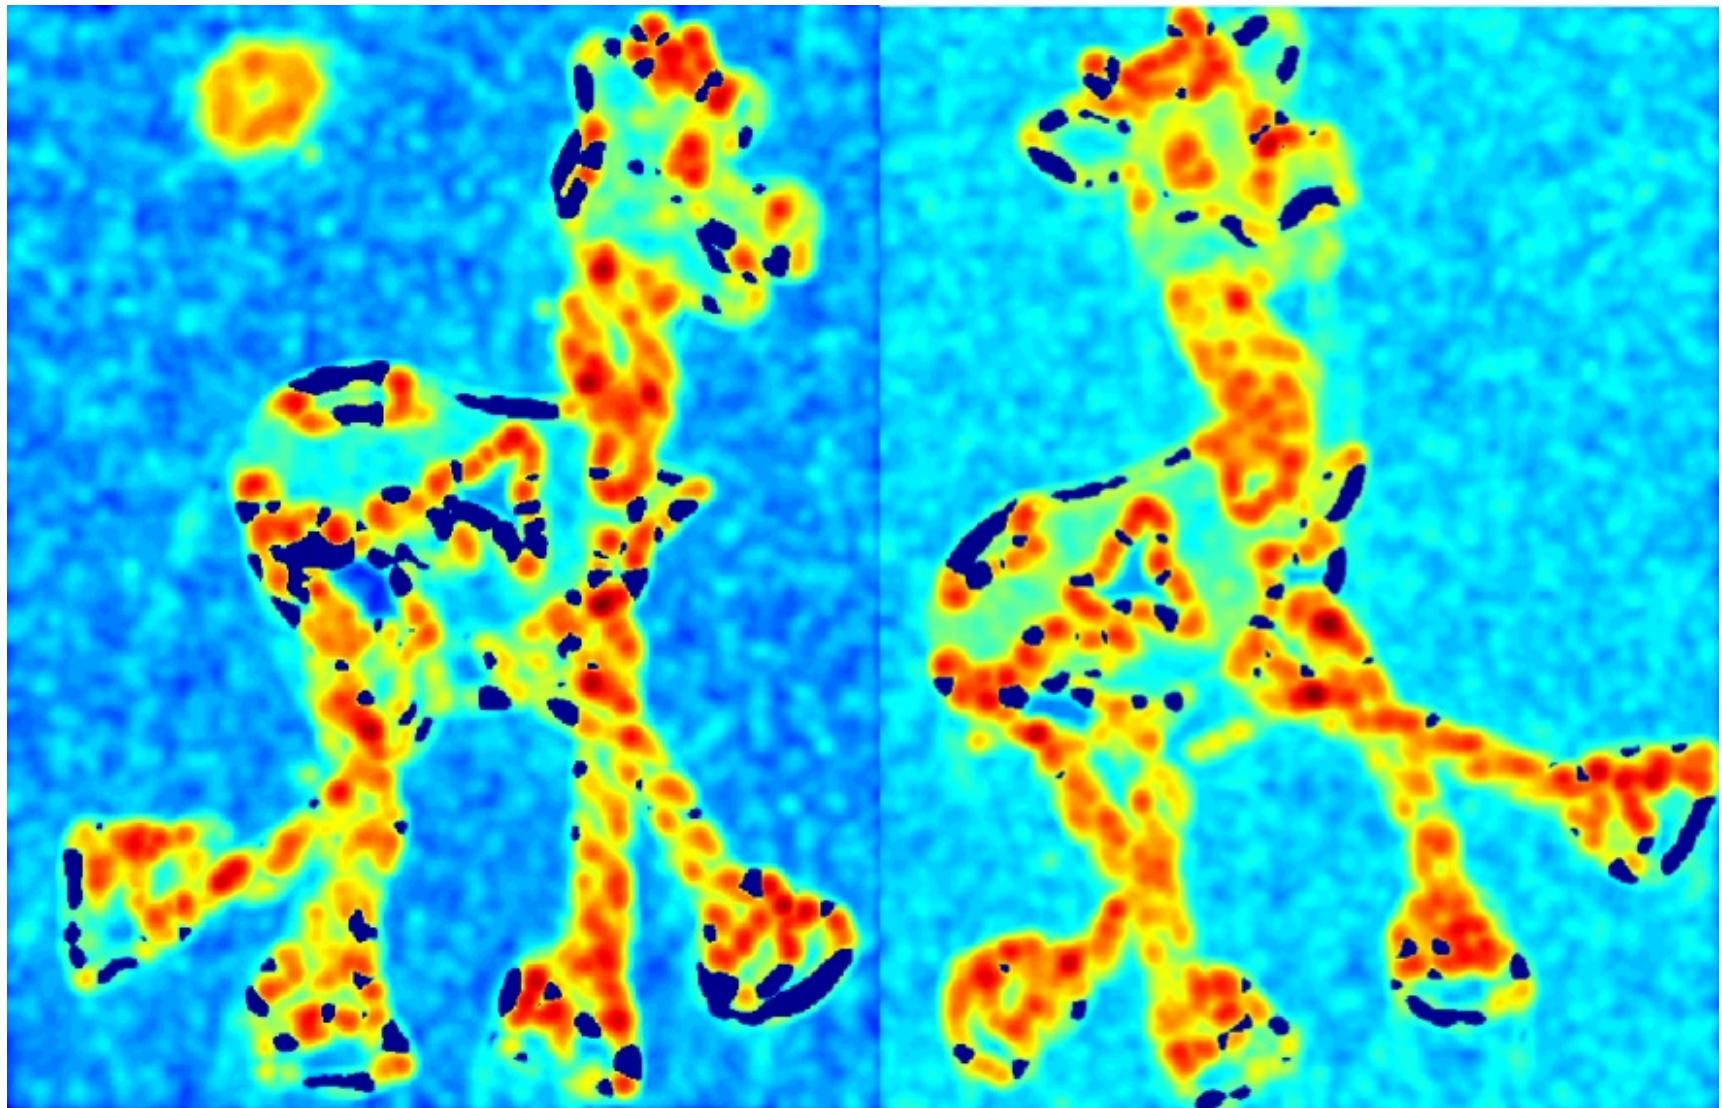
1. Compute partial derivatives at each pixel
2. Compute second moment matrix  $M$  in a Gaussian window around each pixel
3. Compute corner response function  $R$

# Harris Detector: Steps



# Harris Detector: Steps

Compute corner response  $R$

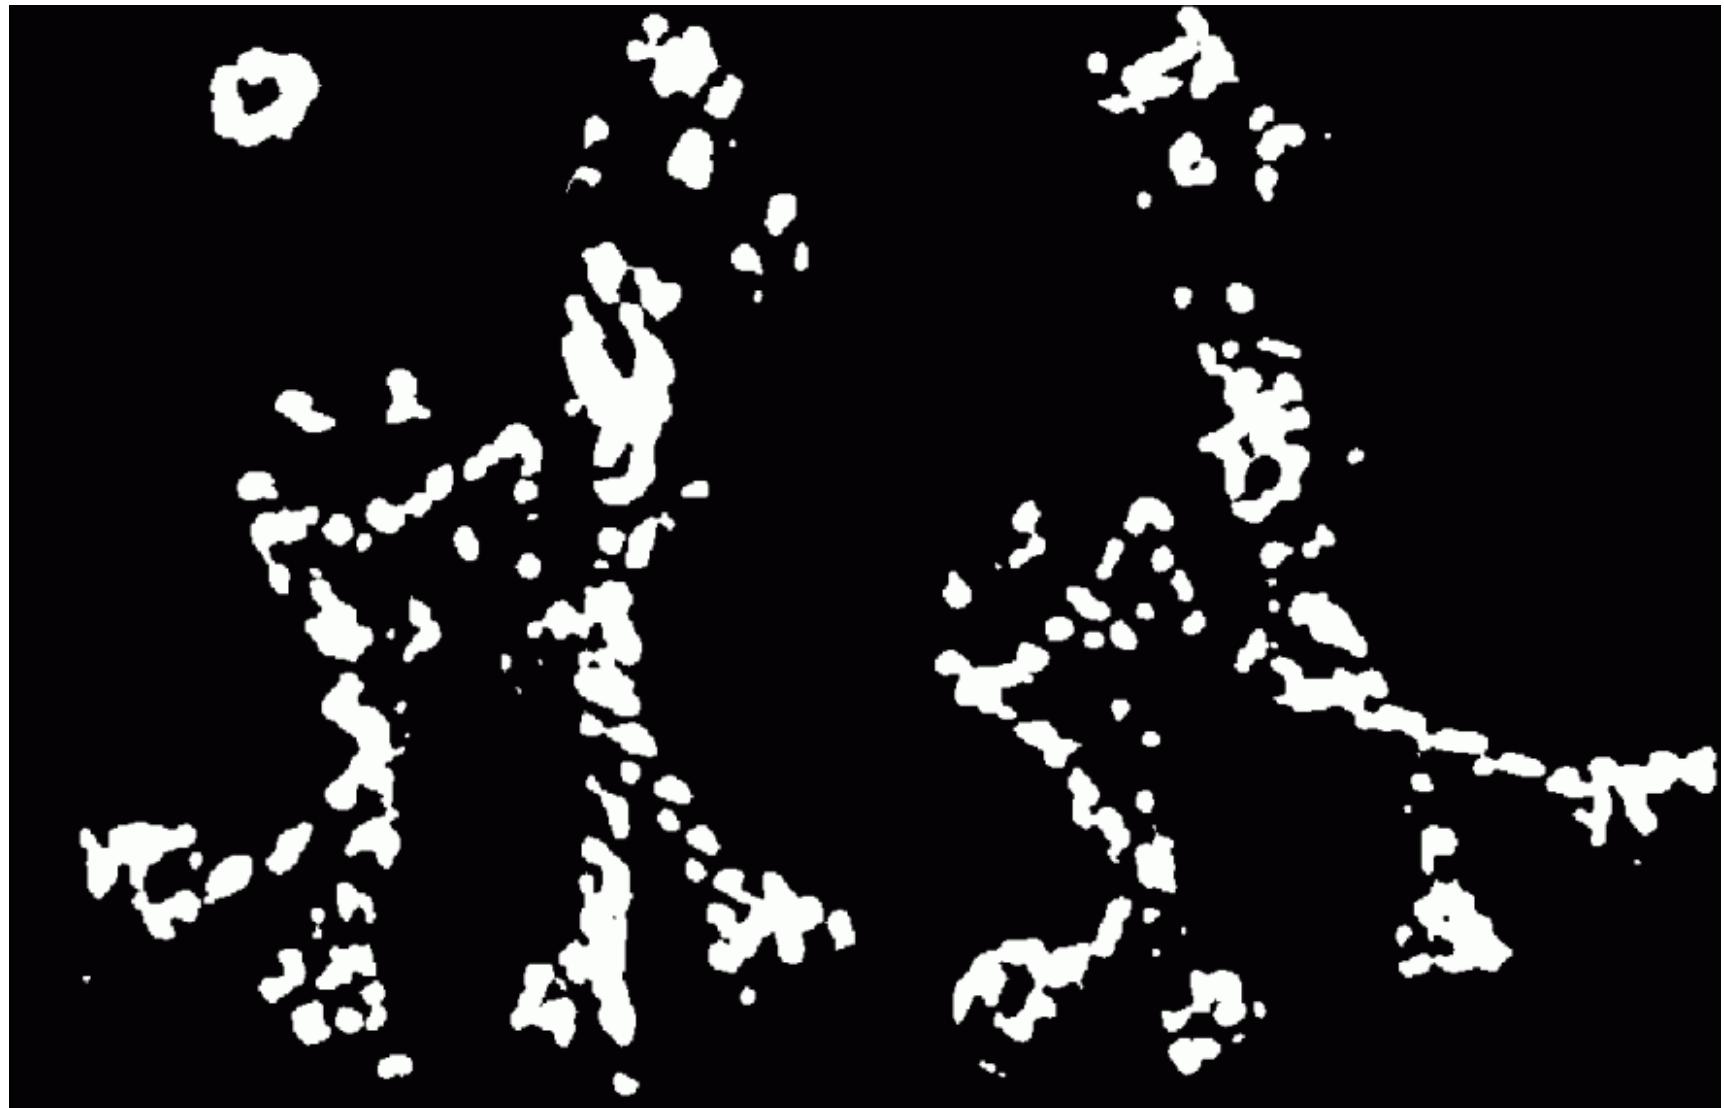


# The Harris corner detector

1. Compute partial derivatives at each pixel
2. Compute second moment matrix  $M$  in a Gaussian window around each pixel
3. Compute corner response function  $R$
4. Threshold  $R$
5. Find local maxima of response function  
(non-maximum suppression)

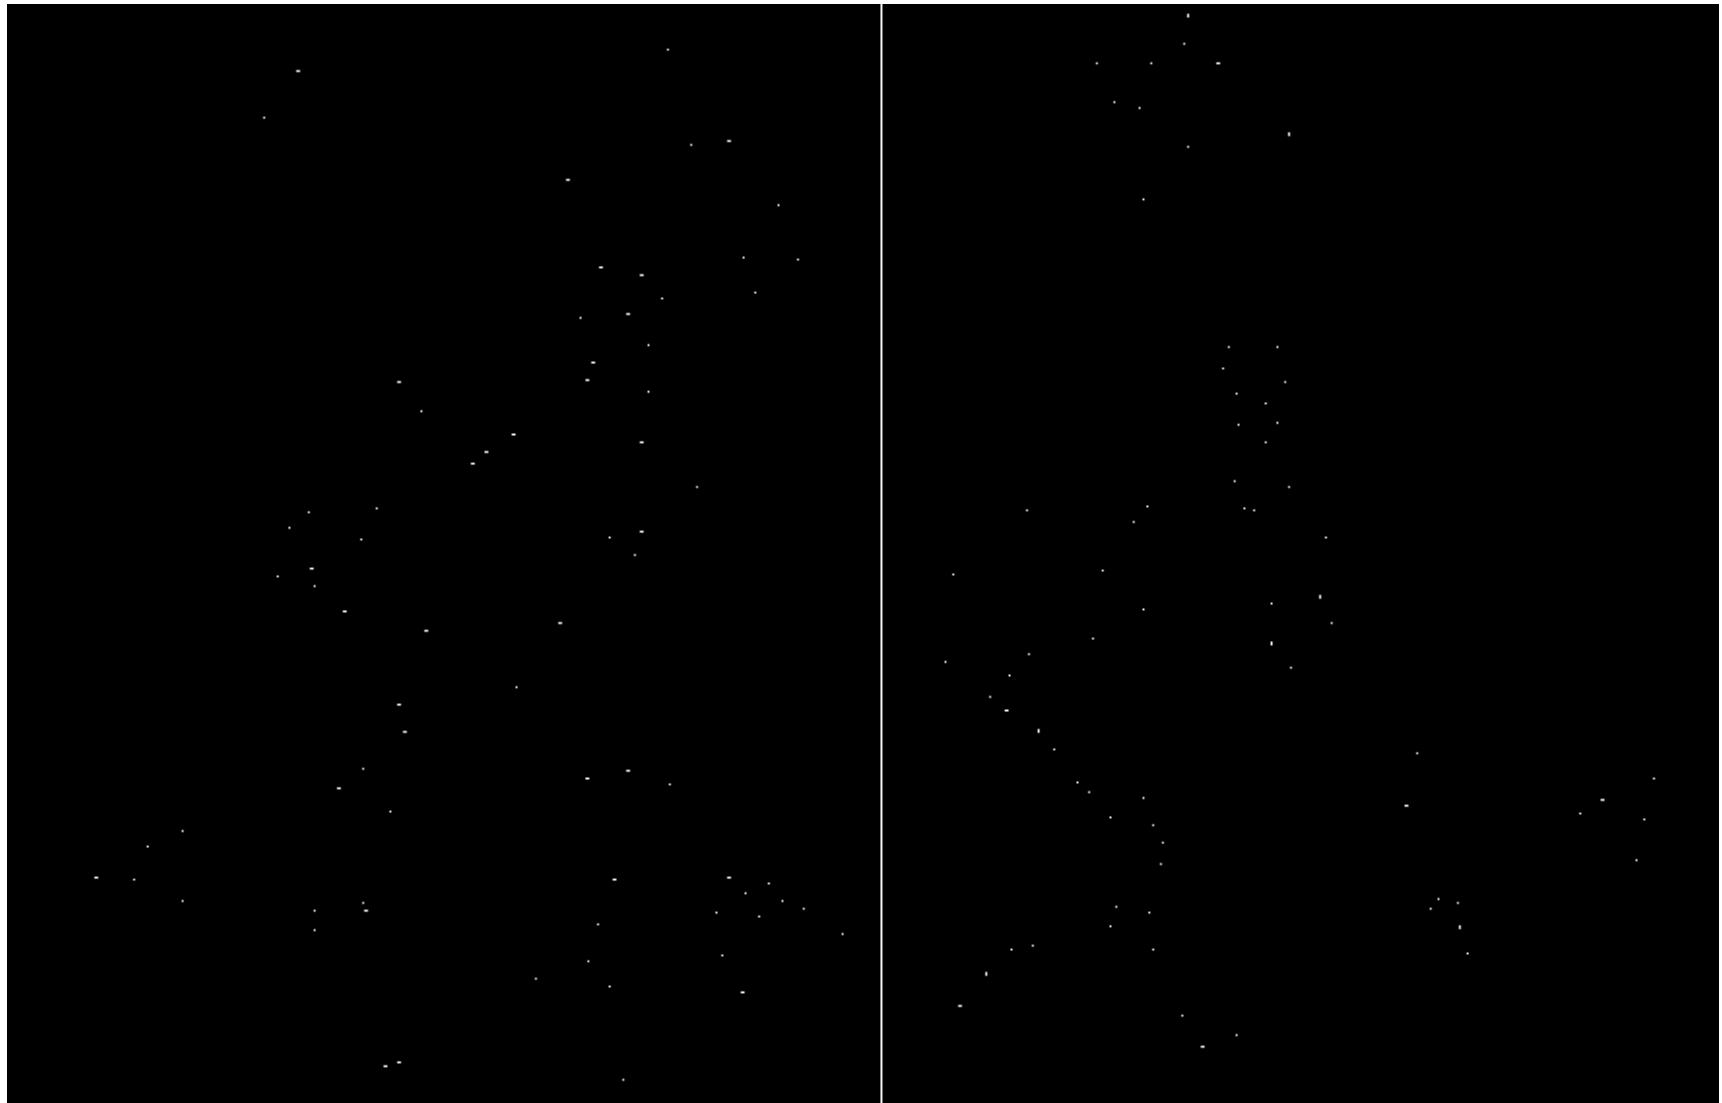
# Harris Detector: Steps

Find points with large corner response:  $R > \text{threshold}$



# Harris Detector: Steps

Take only the points of local maxima of  $R$



# Harris Detector: Steps

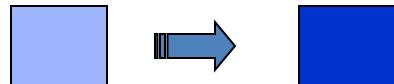


# Invariance and covariance

- We want corner locations to be *invariant* to photometric transformations and *covariant* to geometric transformations
  - **Invariance:** image is transformed and corner locations do not change
  - **Covariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations

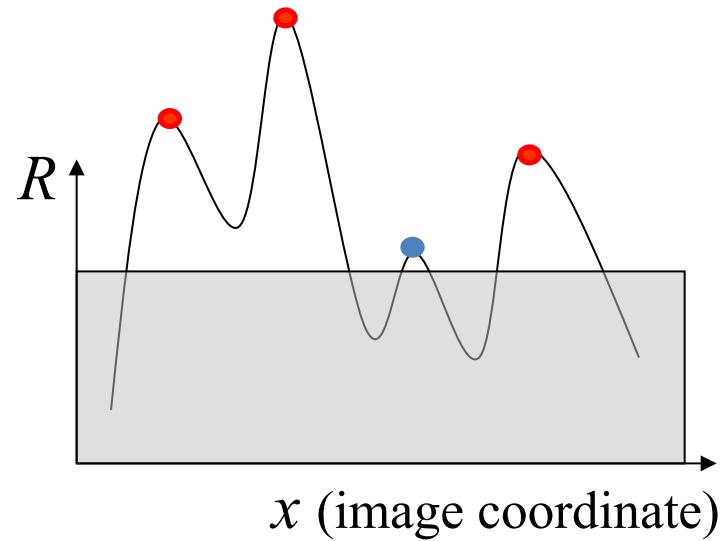
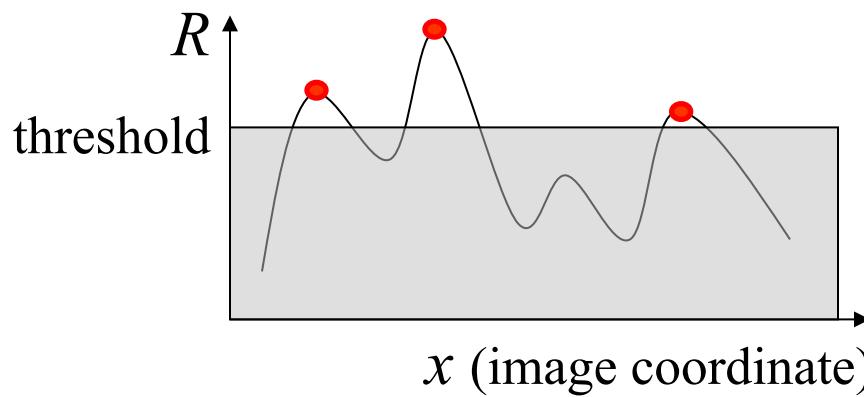


# Intensity changes



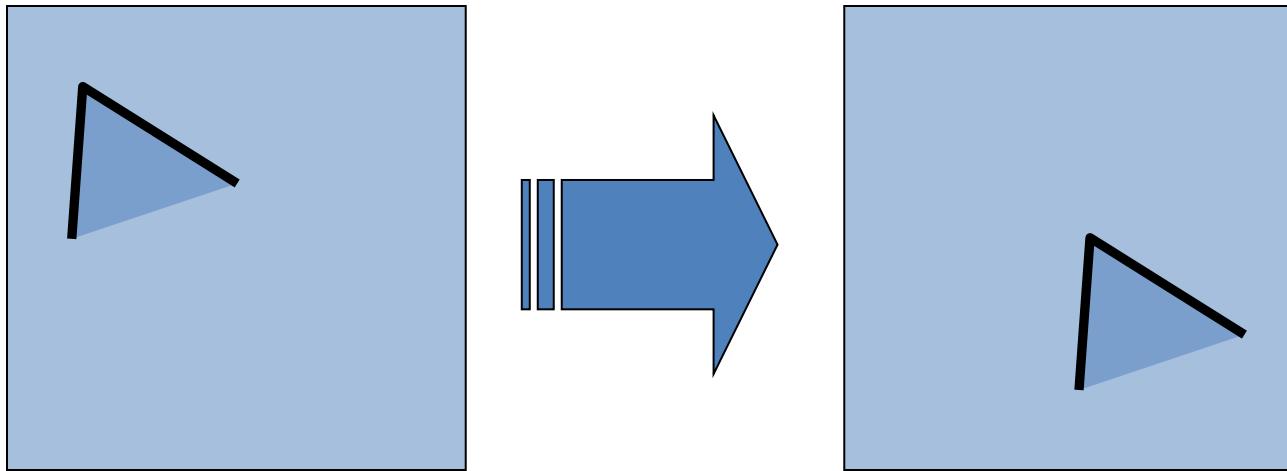
$$I \rightarrow a I + b$$

- Only derivatives are used => invariance to intensity shift  $I \rightarrow I + b$
- Intensity scaling:  $I \rightarrow a I$



*Partially invariant to affine intensity change*

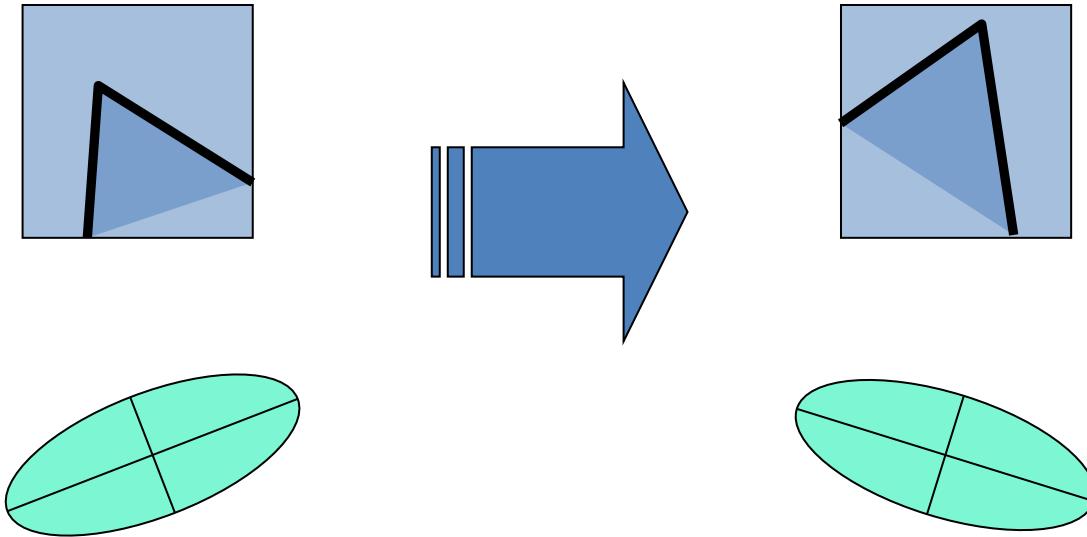
# Image translation



- Derivatives and window function are shift-invariant

Corner location is covariant w.r.t. translation

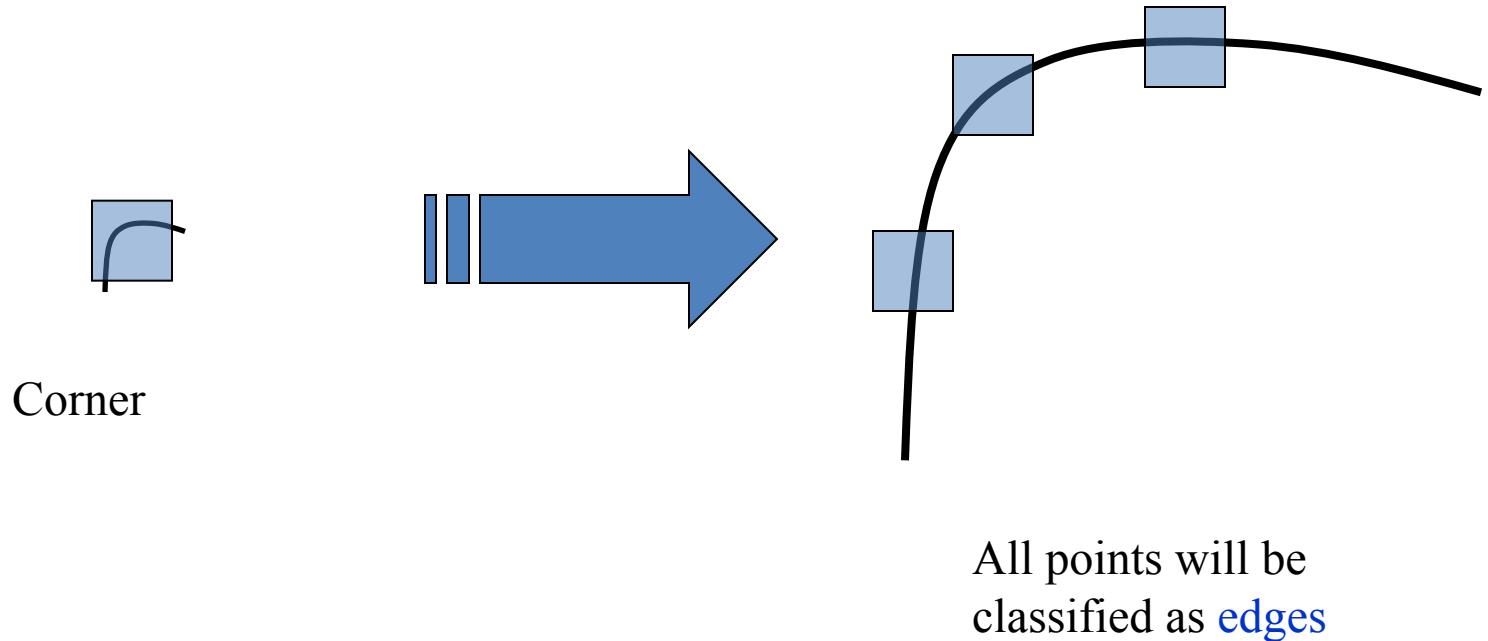
# Image rotation



Second moment ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner location is covariant w.r.t. rotation

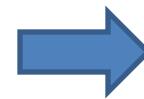
# Scaling



Corner location is not covariant to scaling!

# Sampling

**Why does a lower resolution image still make sense to us? What do we lose?**



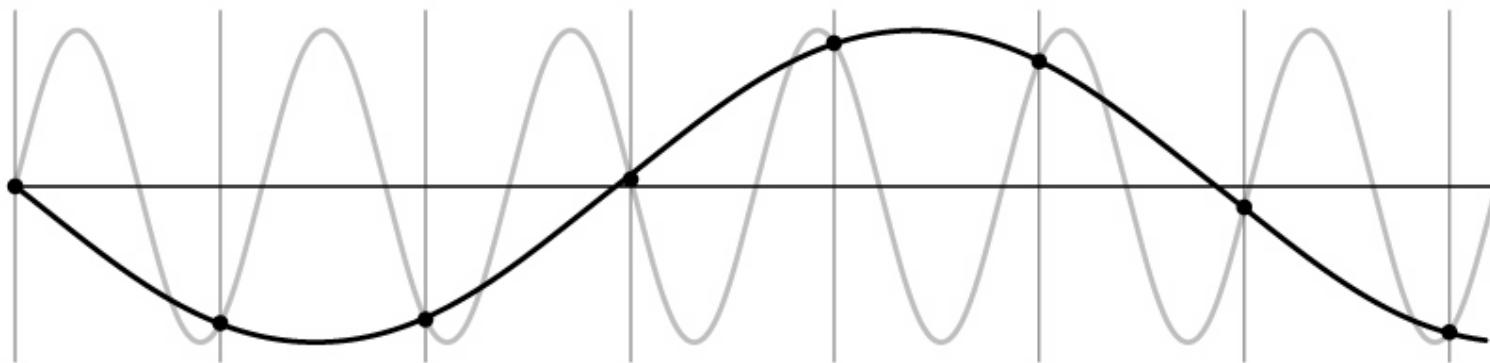
# Subsampling by a factor of 2



Throw away every other row and column  
to create a  $\frac{1}{2}$  size image

# Aliasing problem

- 1D example (sinewave):



# Aliasing problem

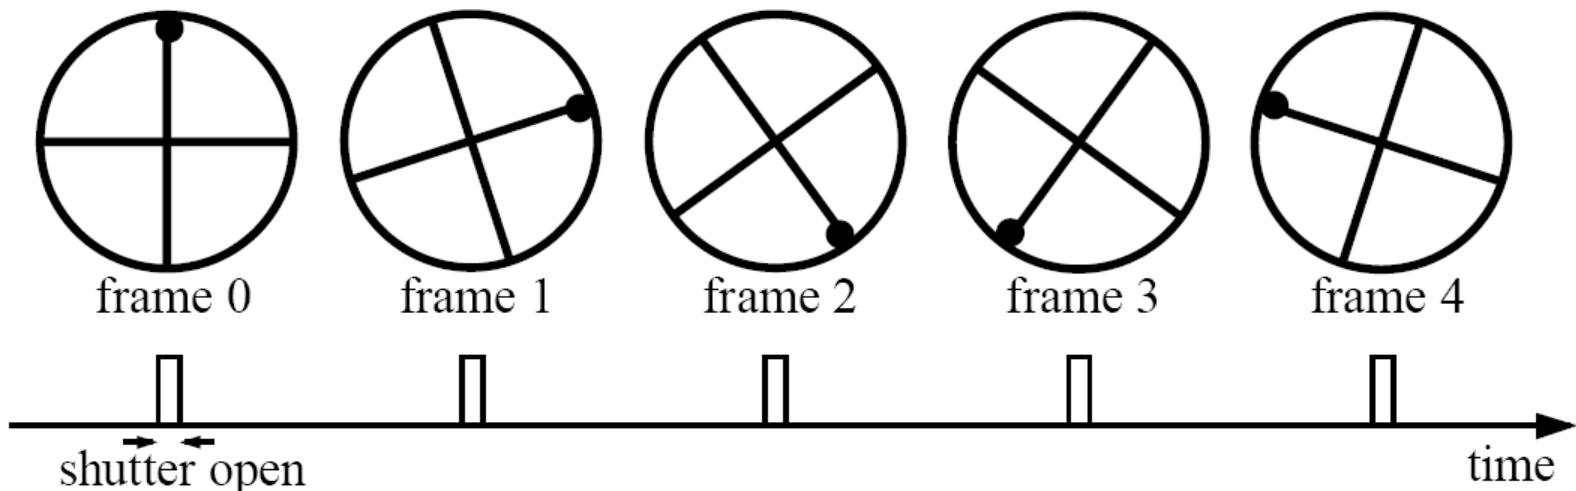
- Sub-sampling may be dangerous....
- Characteristic errors may appear:
  - “Wagon wheels rolling the wrong way in movies”
  - “Checkerboards disintegrate in ray tracing”
  - “Striped shirts look funny on color television”

# Aliasing in video

Imagine a spoked wheel moving to the right (rotating clockwise).

Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time =  $1/30$  sec. for video,  $1/24$  sec. for film):



Without dot, wheel appears to be rotating slowly backwards!  
(counterclockwise)

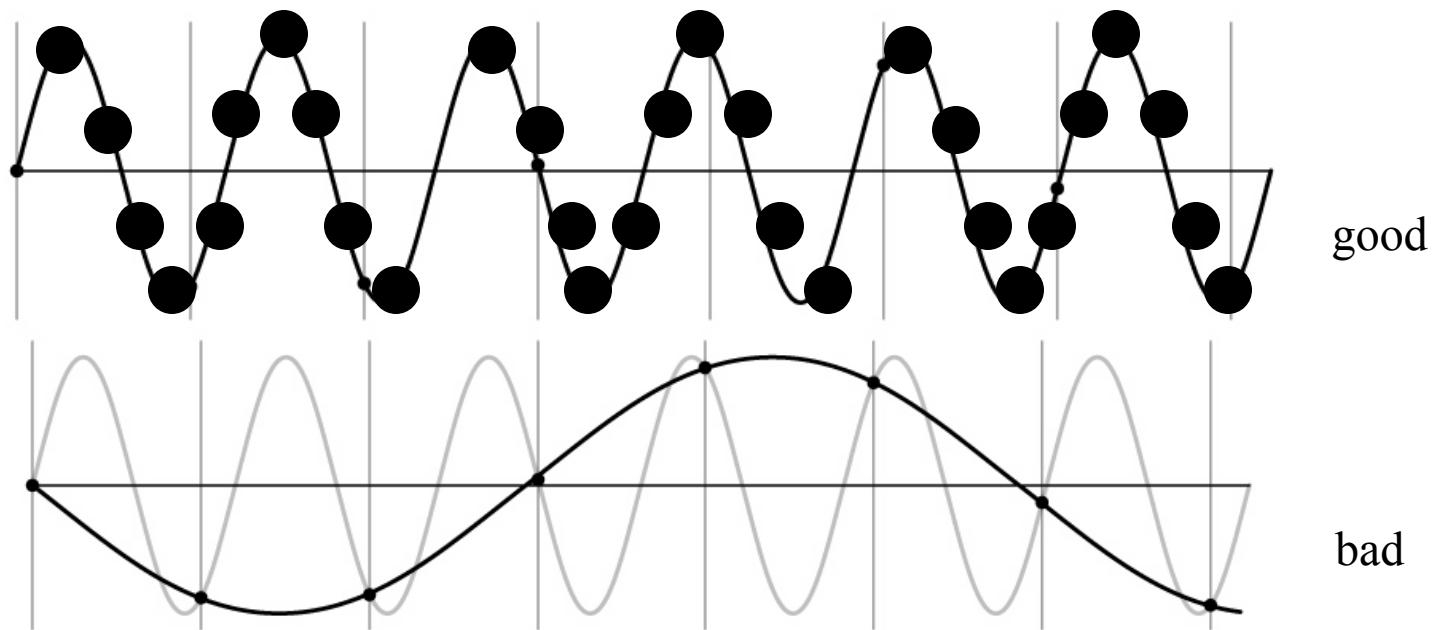
# Aliasing in graphics



**Disintegrating textures**

# Nyquist-Shannon Sampling Theorem

- When sampling a signal at discrete intervals, the sampling frequency must be  $\geq 2 \times f_{\max}$
- $f_{\max}$  = max frequency of the input signal
- This will allow to reconstruct the original perfectly from the sampled version



# Anti-aliasing

Solutions:

- Sample more often
- Get rid of all frequencies that are greater than half the new sampling frequency
  - Will lose information
  - But it's better than aliasing
  - Apply a smoothing filter

# Algorithm for downsampling by factor of 2

1. Start with  $\text{image}(h, w)$

2. Apply low-pass filter

```
im.blur = imfilter(image, fspecial('gaussian', 7, 1))
```

3. Sample every other pixel

```
im.small = im.blur(1:2:end, 1:2:end);
```

# Anti-aliasing

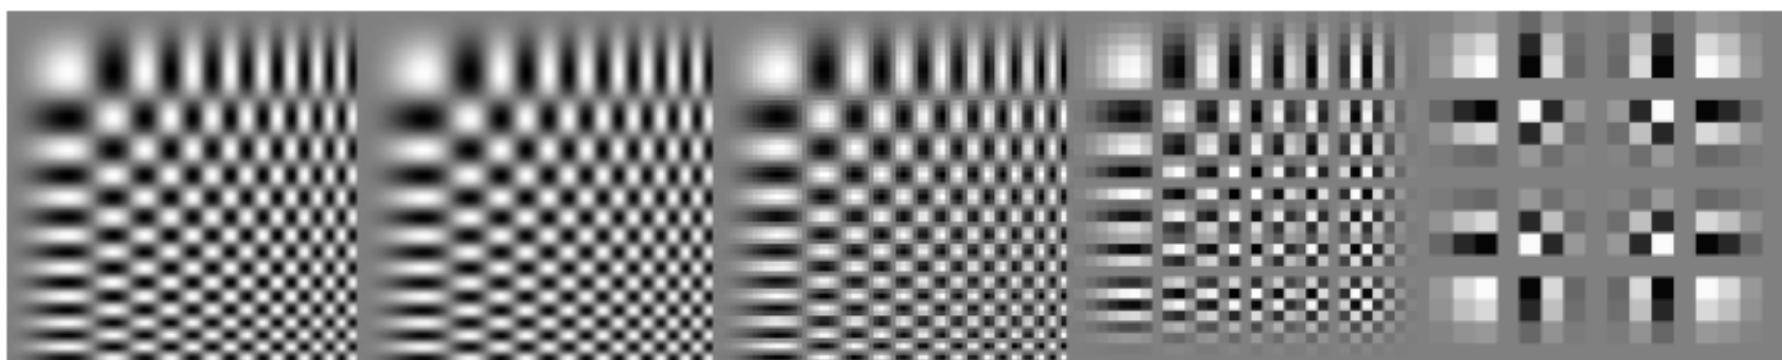
256x256

128x128

64x64

32x32

16x16



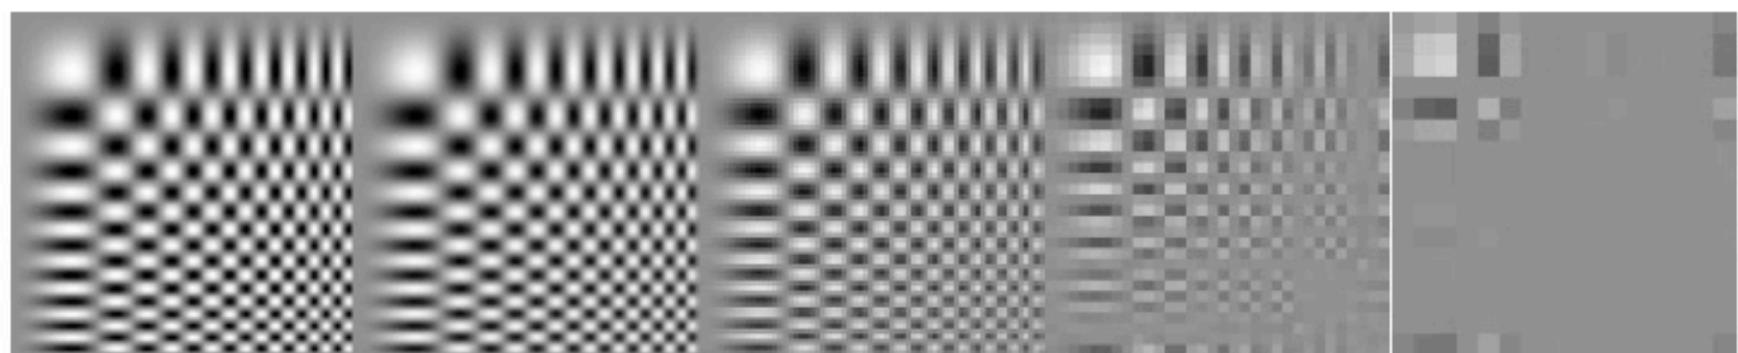
256x256

128x128

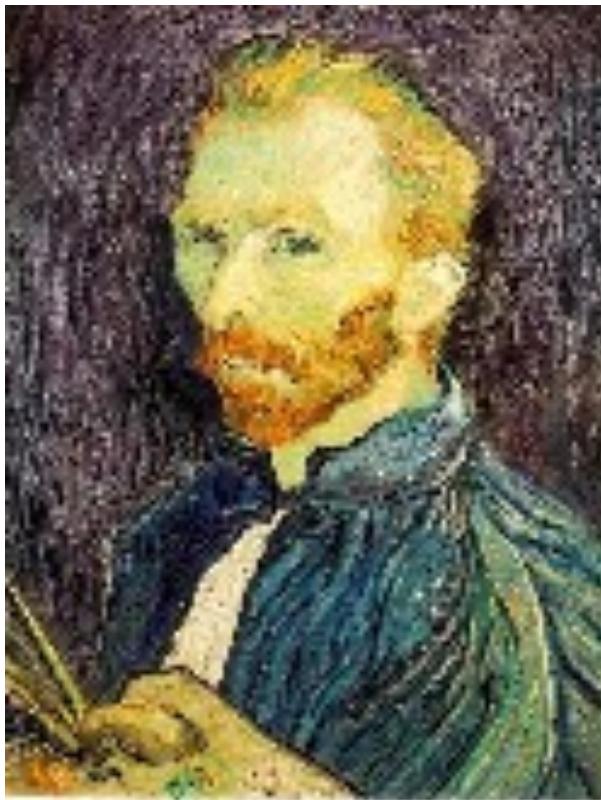
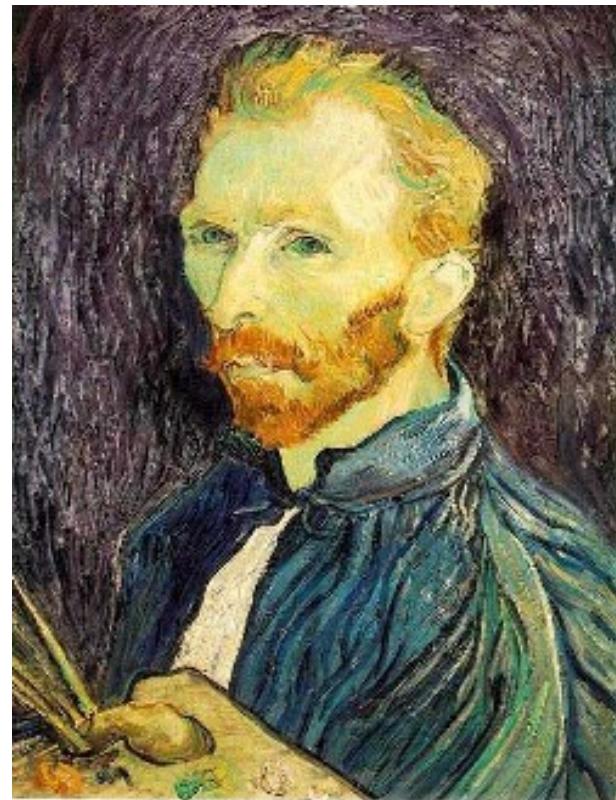
64x64

32x32

16x16



# Subsampling without pre-filtering

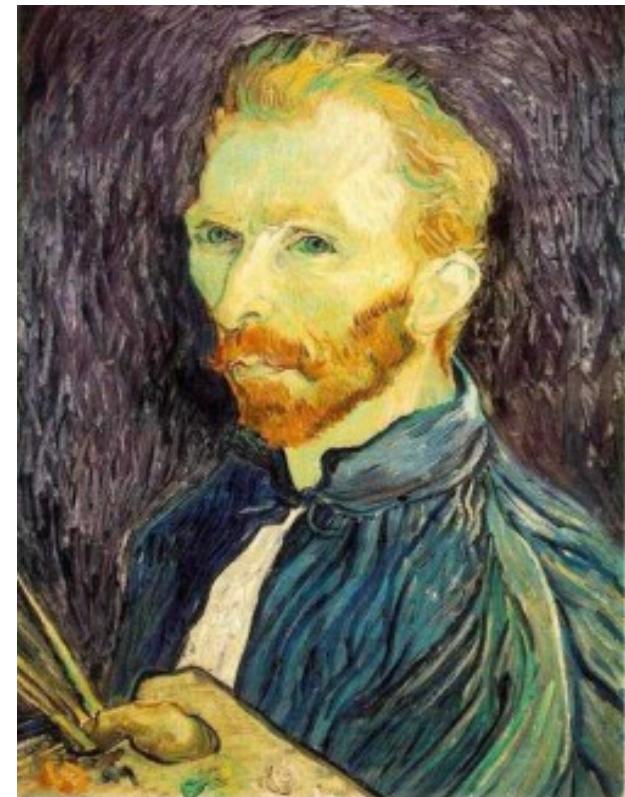


1/2

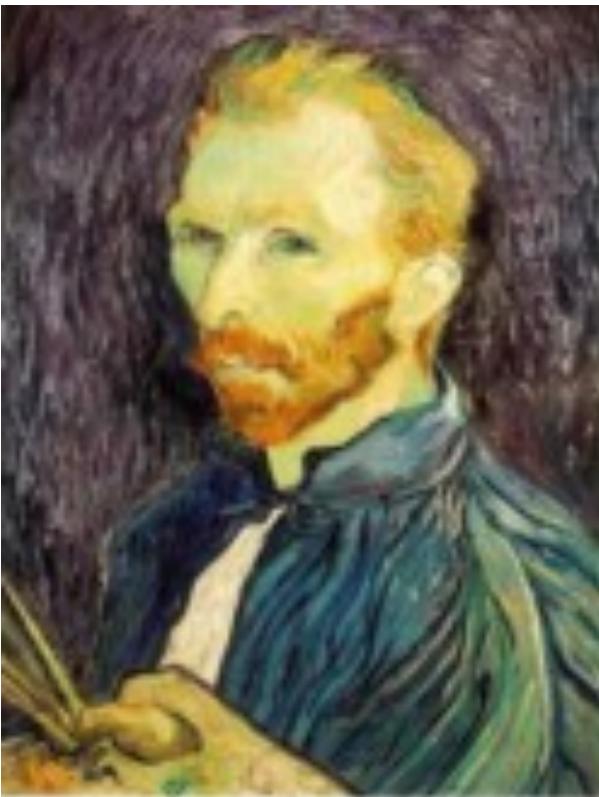
1/4 (2x zoom)

1/8 (4x zoom)

# Subsampling with Gaussian pre-filtering



Gaussian 1/2



G 1/4



G 1/8

# HW1 DISCUSSION